

RDMA for Apache Hadoop 2.x 0.9.9 User Guide

HIGH-PERFORMANCE BIG DATA TEAM

<http://hibd.cse.ohio-state.edu>

NETWORK-BASED COMPUTING LABORATORY
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
THE OHIO STATE UNIVERSITY

Copyright (c) 2011-2016
Network-Based Computing Laboratory,
headed by Dr. D. K. Panda.
All rights reserved.

Last revised: March 19, 2016

Contents

1	Overview of the RDMA for Apache Hadoop 2.x Project	1
2	Features	1
3	Setup Instructions	3
3.1	Prerequisites	3
3.2	Download	4
3.3	Installation steps	4
3.3.1	Installing intergrated RDMA for Apache Hadoop package	4
3.3.2	Installing RDMA-based plugin for Apache Hadoop, HDP, or CDH	4
3.4	Basic Configuration	5
3.4.1	HDFS with Heterogeneous Storage (HHH)	5
3.4.2	HDFS in Memory (HHH-M)	9
3.4.3	HDFS with Heterogeneous Storage and Lustre (HHH-L)	10
3.4.4	MapReduce over Lustre with Local Disks	11
3.4.5	MapReduce over Lustre without Local Disks	14
3.5	Advanced Configuration	17
3.5.1	Parallel Replication	18
3.5.2	Placement Policy Selection	18
3.5.3	Automatic Placement Policy Selection	18
3.5.4	Threshold of RAM Disk Usage	19
3.5.5	No. of In-Memory Replicas	19
3.5.6	Disk-assisted Shuffle	19
3.5.7	SSD Optimization for Intermediate Data	20
4	Basic Usage Instructions	20
4.1	Startup	20
4.1.1	MapReduce over HDFS	20
4.1.2	MapReduce over Lustre	20
4.2	Basic Commands	21
4.3	Shutdown	24
4.3.1	MapReduce over HDFS	24
4.3.2	MapReduce over Lustre	24
5	Running RDMA for Apache Hadoop with SLURM/PBS	25
5.1	Usage	25
5.1.1	Configure and Start a Hadoop Job	25
5.1.2	Shutdown Hadoop Cluster	26
5.2	Running MapReduce over HDFS with SLURM/PBS	27
5.2.1	Startup	27
5.2.2	Running Benchmarks	27
5.2.3	Shutdown	28
5.3	Running MapReduce over Lustre with SLURM/PBS	28
5.3.1	Startup	28
5.3.2	Running Benchmarks	28

5.3.3	Shutdown	29
6	Benchmarks	29
6.1	TestDFSIO	29
6.2	Sort	29
6.3	TeraSort	30
6.4	OHB Micro-benchmarks	31

1 Overview of the RDMA for Apache Hadoop 2.x Project

RDMA for Apache Hadoop 2.x is a high-performance design of Hadoop over RDMA-enabled Interconnects. This version of RDMA for Apache Hadoop 2.x 0.9.9 is based on Apache Hadoop 2.7.1 and is compliant with Apache Hadoop 2.7.1, Hortonworks Data Platform (HDP) 2.3.0.0, and Cloudera Distribution Including Apache Hadoop (CDH) 5.6.0 APIs and applications. This file is intended to guide users through the various steps involved in installing, configuring, and running RDMA for Apache Hadoop 2.x over InfiniBand.

Figure 1 presents a high-level architecture of RDMA for Apache Hadoop 2.x. In this package, many different modes have been included that can be enabled/disabled to obtain performance benefits for different kinds of applications in different Hadoop environments. This package can be configured to run MapReduce jobs on top of HDFS as well as Lustre.

Following are the different modes that are included in our package.

HHH: Heterogeneous storage devices with hybrid replication schemes are supported in this mode of operation to have better fault-tolerance as well as performance. This mode is enabled by **default** in the package.

HHH-M: A high-performance in-memory based setup has been introduced in this package that can be utilized to perform all I/O operations in-memory and obtain as much performance benefit as possible.

HHH-L: With parallel file systems integrated, HHH-L mode can take advantage of the Lustre available in the cluster.

MapReduce over Lustre, with/without local disks: Besides, HDFS based solutions, this package also provides support to run MapReduce jobs on top of Lustre alone. Here, two different modes are introduced: with local disks and without local disks.

Running with Slurm and PBS: Supports deploying RDMA for Apache Hadoop 2.x with Slurm and PBS in different running modes (HHH, HHH-M, HHH-L, and MapReduce over Lustre).

If there are any questions, comments or feedback regarding this software package, please post them to rdma-hadoop-discuss mailing list (rdma-hadoop-discuss@cse.ohio-state.edu).

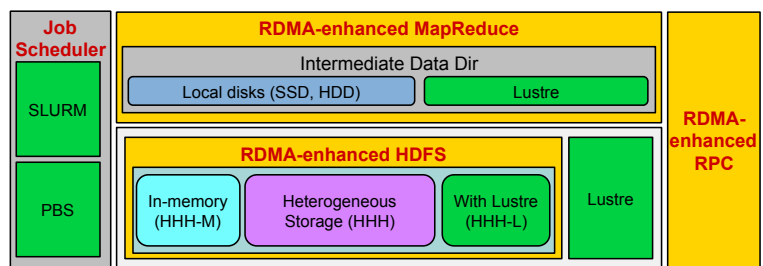


Figure 1: RDMA-based Hadoop architecture and its different modes

2 Features

High-level features of RDMA for Apache Hadoop 2.x 0.9.9 are listed below. New features and enhancements compared to 0.9.9 release are marked as **(NEW)**.

- Based on Apache Hadoop 2.7.1

- High performance design with native InfiniBand and RoCE support at the verbs level for HDFS, MapReduce, and RPC components
- **(NEW)** Compliant with Apache Hadoop 2.7.1, Hortonworks Data Platform (HDP) 2.3.0.0, and Cloudera Distribution Including Apache Hadoop (CDH) 5.6.0 APIs and applications
- Plugin-based architecture supporting RDMA-based designs for HDFS (HHH, HHH-M, HHH-L), MapReduce, MapReduce over Lustre, and RPC, etc.
 - Plugin for Apache Hadoop distribution (tested with 2.7.1)
 - Plugin for Hortonworks Data Platform (HDP) (tested with 2.3.0.0)
 - **(NEW)** Plugin for Cloudera Distribution Including Apache Hadoop (CDH) (tested with 5.6.0)
- Supports deploying Hadoop with Slurm and PBS in different running modes (HHH, HHH-M, HHH-L, and MapReduce over Lustre)
- Easily configurable for different running modes (HHH, HHH-M, HHH-L, and MapReduce over Lustre) and different protocols (native InfiniBand, RoCE, and IPoIB)
- On-demand connection setup
- HDFS over native InfiniBand and RoCE
 - RDMA-based write
 - RDMA-based replication
 - Parallel replication support
 - Overlapping in different stages of write and replication
 - Enhanced hybrid HDFS design with in-memory and heterogeneous storage (HHH)
 - * Supports three modes of operations
 - HHH (default) with I/O operations over RAM disk, SSD, and HDD
 - HHH-M (in-memory) with I/O operations in-memory
 - HHH-L (Lustre-integrated) with I/O operations in local storage and Lustre
 - * Policies to efficiently utilize heterogeneous storage devices (RAM Disk, SSD, HDD, and Lustre)
 - Greedy and Balanced policies support
 - Automatic policy selection based on available storage types
 - * Hybrid replication (in-memory and persistent storage) for HHH default mode
 - * Memory replication (in-memory only with lazy persistence) for HHH-M mode
 - * Lustre-based fault-tolerance for HHH-L mode
 - No HDFS replication
 - Reduced local storage space usage
- MapReduce over native InfiniBand and RoCE
 - RDMA-based shuffle
 - Prefetching and caching of map output

- In-memory merge
- Advanced optimization in overlapping
 - * map, shuffle, and merge
 - * shuffle, merge, and reduce
- Optional disk-assisted shuffle
- (NEW) Automatic Locality-aware Shuffle
- High performance design of MapReduce over Lustre
 - * Supports two shuffle approaches
 - Lustre read based shuffle
 - RDMA based shuffle
 - * Hybrid shuffle based on both shuffle approaches
 - Configurable distribution support
 - * In-memory merge and overlapping of different phases
- RPC over native InfiniBand and RoCE
 - JVM-bypassed buffer management
 - RDMA or send/recv based adaptive communication
 - Intelligent buffer allocation and adjustment for serialization
- Tested with
 - Mellanox InfiniBand adapters (DDR, QDR, and FDR)
 - RoCE support with Mellanox adapters
 - Various multi-core platforms
 - RAM Disks, SSDs, HDDs, and Lustre

3 Setup Instructions

3.1 Prerequisites

Prior to the installation of RDMA for Apache Hadoop 2.x, please ensure that you have the latest version of JDK installed on your system, and set the `JAVA_HOME` and `PATH` environment variables to point to the appropriate JDK installation. We recommend the use of JDK version 1.7 and later.

In order to use the RDMA-based features provided with RDMA for Apache Hadoop 2.x, install the latest version of the OFED distribution that can be obtained from <http://www.openfabrics.org>.

3.2 Download

Two kinds of tarballs are provided for downloading, intergrated package and plugin package.

Download the most recent distribution tarball of intergrated RDMA for Apache Hadoop package from <http://hibd.cse.ohio-state.edu/download/hibd/rdma-hadoop-2.x-0.9.9-bin.tar.gz>.

Download the most recent distribution tarball of RDMA for Apache Hadoop plugin package from <http://hibd.cse.ohio-state.edu/download/hibd/rdma-hadoop-2.x-0.9.9-plugin.tar.gz>.

3.3 Installation steps

The RDMA for Apache Hadoop can be installed using the intergrated distribution or as plugins on existing Apache Hadoop source in use.

3.3.1 Installing intergrated RDMA for Apache Hadoop package

Following steps can be used to install the integrated RDMA for Apache Hadoop package.

1. Unzip the intergrated RDMA for Apache Hadoop distribution tarball using the following command:

```
tar zxf rdma-hadoop-2.x-0.9.9-bin.tar.gz
```

2. Change directory to rdma-hadoop-2.x-0.9.9

```
cd rdma-hadoop-2.x-0.9.9
```

3.3.2 Installing RDMA-based plugin for Apache Hadoop, HDP, or CDH

Following steps can be used to build and install the RDMA-based plugin for Apache Hadoop, HDP, or CDH package.

1. Unzip the plugin tarball of the specific distribution using the following command:

```
tar zxf rdma-hadoop-2.x-0.9.9-plugin.tar.gz
```

2. Change directory to hadoop-plugin-2.7.1/rdma-plugins/

```
cd rdma-hadoop-2.x-0.9.9-plugin
```

3. Run the install script `install.sh`. This script applies the appropriate patch and re-builds the source from the path specified. This installation can be launched using the following command:

```
./install.sh HADOOP_SRC_DIR DISTRIBUTION
```

The installation script takes two arguments:

HADOOP_SRC_DIR

This mandatory argument is the directory location containing Hadoop source.

DISTRIBUTION

This mandatory argument specifies the Hadoop source distribution.

This is necessary to choose the appropriate patch to apply.

The current options are `apache`, `hdp`, or `cdh`.

3.4 Basic Configuration

The configuration files can be found in the directory `rdma-hadoop-2.x-0.9.9/etc/hadoop`, in the integrated RDMA for Apache Hadoop package. If the plugin-based Hadoop distribution is being used, the configuration files can be found at `HADOOP_SOURCE/hadoop-dist/target/etc/hadoop`, where `HADOOP_SOURCE` is the Hadoop source directory on which the plugin is being (or has been) applied.

3.4.1 HDFS with Heterogeneous Storage (HHH)

Steps to configure RDMA for Apache Hadoop 2.x include:

1. Configure `hadoop-env.sh` file.

```
export JAVA_HOME=/opt/java/1.7.0
```

2. Configure `core-site.xml` file. RDMA for Apache Hadoop 2.x 0.9.9 supports three different modes: IB, RoCE, and TCP/IP.

Configuration of the IB mode:

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://node001:9000</value>
    <description>NameNode URI.</description>
  </property>

  <property>
    <name>hadoop.ib.enabled</name>
    <value>true</value>
    <description>Enable the RDMA feature over IB. Default value of
      hadoop.ib.enabled is true.</description>
  </property>

  <property>
    <name>hadoop.roce.enabled</name>
    <value>false</value>
    <description>Disable the RDMA feature over RoCE. Default value
      of hadoop.roce.enabled is false.</description>
  </property>
</configuration>
```



```
</property>  
</configuration>
```

Configuration of the RoCE mode:

```
<configuration>  
  <property>  
    <name>fs.defaultFS</name>  
    <value>hdfs://node001:9000</value>  
    <description>NameNode URI.</description>  
  </property>  
  
  <property>  
    <name>hadoop.ib.enabled</name>  
    <value>false</value>  
    <description>Disable the RDMA feature over IB. Default value of  
      hadoop.ib.enabled is true.</description>  
  </property>  
  
  <property>  
    <name>hadoop.roce.enabled</name>  
    <value>true</value>  
    <description>Enable the RDMA feature over RoCE. Default value of  
      hadoop.roce.enabled is false.</description>  
  </property>  
</configuration>
```

Configuration of the TCP/IP mode:

```
<configuration>  
  <property>  
    <name>fs.defaultFS</name>  
    <value>hdfs://node001:9000</value>  
    <description>NameNode URI.</description>  
  </property>  
  
  <property>  
    <name>hadoop.ib.enabled</name>  
    <value>false</value>  
    <description>Disable the RDMA feature over IB. Default value of  
      hadoop.ib.enabled is true.</description>  
  </property>  
  
  <property>  
    <name>hadoop.roce.enabled</name>  
    <value>false</value>  
    <description>Disable the RDMA feature over RoCE. Default value  
      of hadoop.roce.enabled is false.</description>  
  </property>  
</configuration>
```

```
</property>
</configuration>
```

Note that we should not enable “hadoop.ib.enabled” and “hadoop.roce.enabled” at the same time. Also, for IB and RoCE mode, the speculative executions of map and reduce tasks are disabled by default. The Physical and Virtual memory monitoring is also disabled for these modes. As default, `org.apache.hadoop.mapred.HOMRShuffleHandler` is configured as the Shuffle Handler service and `org.apache.hadoop.mapreduce.task.reduce.HOMRShuffle` is configured as the default Shuffle plugin.

3. Configure `hdfs-site.xml` file.

```
<configuration>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///home/hadoop/rdma-hadoop-2.x-0.9.9/Name</value>
    <description>Path on the local filesystem where the NameNode
      stores the namespace and transactions logs
      persistently.</description>
  </property>

  <property>
    <name>dfs.datanode.data.dir</name>
    <value>[RAM_DISK]file:///data01, [SSD]file:///data02,
      [DISK]file:///data03</value>
    <description>Comma separated list of paths of the local storage
      devices with corresponding types on a DataNode where it
      should store its blocks.</description>
  </property>
</configuration>
```

4. Configure `yarn-site.xml` file.

```
<configuration>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>node001:8032</value>
    <description>ResourceManager host:port for clients to submit
      jobs.</description>
  </property>

  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>node001:8030</value>
    <description>ResourceManager host:port for ApplicationMasters to
      talk to Scheduler to obtain resources.</description>
  </property>
</configuration>
```

```

<property>
  <name>yarn.resourcemanager.resource-tracker.address</name>
  <value>node001:8031</value>
  <description>ResourceManager host:port for
    NodeManagers.</description>
</property>

<property>
  <name>yarn.resourcemanager.admin.address</name>
  <value>node001:8033</value>
  <description>ResourceManager host:port for administrative
    commands.</description>
</property>

<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
  <description>Shuffle service that needs to be set for MapReduce
    applications.</description>
</property>
</configuration>

```

5. Configure mapred-site.xml file.

```

<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
    <description>Execution framework set to Hadoop
      YARN.</description>
  </property>

  <property>
    <name>yarn.app.mapreduce.am.command-opts</name>
    <value>-Xmx1024m -Dhadoop.conf.dir=${HADOOP_CONF_DIR}</value>
    <description>Java opts for the MR App Master processes. The
      following symbol, if present, will be interpolated: @taskid@
      is replaced by current TaskID. Any other occurrences of '@'
      will go unchanged. For example, to enable verbose gc logging
      to a file named for the taskid in /tmp and to set the heap
      maximum to be a gigabyte, pass a 'value' of: -Xmx1024m
      -verbose:gc -Xloggc:/tmp/@taskid@.gc
  </property>
</configuration>

```

Usage of -Djava.library.path can cause programs to no longer function if hadoop native libraries are used. These values should instead be set as part of LD_LIBRARY_PATH in the map / reduce JVM env using the mapreduce.map.env and

```

    mapreduce.reduce.env config settings.
  </description>
</property>

<property>
  <name>yarn.app.mapreduce.am.env</name>
  <value>LD_LIBRARY_PATH=${HADOOP_HOME}/lib/native</value>
  <description>User added environment variables for the MR App
    Master
    processes. Example :
    1) A=foo This will set the env variable A to foo
    2) B=$B:c This is inherit tasktracker's B env variable.
  </description>
</property>
</configuration>

```

6. Configure slaves file. List all slave hostnames in this file, one per line.

```

node002
node003

```

We can also configure more specific items according to actual needs. For example, we can configure the item `dfs.blocksize` in `hdfs-site.xml` to change the HDFS block size. To get more detailed information, please visit <http://hadoop.apache.org>.

3.4.2 HDFS in Memory (HHH-M)

We can enable MapReduce over in-memory HDFS using the following configuration steps:

1. Configure `hdfs-site.xml` file.

```

<configuration>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///home/hadoop/rdma-hadoop-2.x-0.9.9/Name</value>
    <description>Path on the local filesystem where the NameNode
      stores the namespace and transactions logs
      persistently.</description>
  </property>

  <property>
    <name>dfs.datanode.data.dir</name>
    <value>[RAM_DISK]file:///data01, [SSD]file:///data02,
      [DISK]file:///data03</value>
  </property>
</configuration>

```

```

    <description>Comma separated list of paths of the local storage
        devices with corresponding types on a DataNode where it
        should store its blocks.</description>
</property>

<property>
    <name>dfs.rdma.hhh.mode</name>
    <value>In-Memory</value>
    <description>Select In-Memory mode (HHH-M). </description>
</property>
</configuration>

```

2. Configure the `yarn-site.xml` file as specified in Section 3.4.1.
3. Configure the `mapred-site.xml` file as specified in Section 3.4.1.
4. Configure the `core-site.xml` file as specified in Section 3.4.1.
5. Configure the `slaves` file as specified in Section 3.4.1.

3.4.3 HDFS with Heterogeneous Storage and Lustre (HHH-L)

We can enable MapReduce over Luster-integrated HDFS using the following configuration steps:

1. Configure `hdfs-site.xml` file.

```

<configuration>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///home/hadoop/rdma-hadoop-2.x-0.9.9/Name</value>
    <description>Path on the local filesystem where the NameNode
        stores the namespace and transactions logs
        persistently.</description>
  </property>

  <property>
    <name>dfs.datanode.data.dir</name>
    <value>[RAM_DISK]file:///data01, [SSD]file:///data02,
        [DISK]file:///data03</value>
    <description>Comma separated list of paths of the local storage
        devices with corresponding types on a DataNode where it
        should store its blocks.</description>
  </property>

  <property>

```

```
<name>dfs.rdma.hhh.mode</name>
<value>Lustre</value>
<description>Select Lustre-Integrated mode (HHH-L) .
</description>
</property>

<property>
  <name>dfs.replication</name>
  <value>1</value>
  <description>Lustre provides fault-tolerance; use replication
    factor of 1</description>
</property>

<property>
  <name>dfs.rdma.lustre.path</name>
  <value>/scratch/hadoop-lustre/</value>
  <description>The path of a directory in Lustre where HDFS will
    store the blocks</description>
</property>
</configuration>
```

2. Configure the `yarn-site.xml` file as specified in Section 3.4.1.
3. Configure the `mapred-site.xml` file as specified in Section 3.4.1.
4. Configure the `core-site.xml` file as specified in Section 3.4.1.
5. Configure the `slaves` file as specified in Section 3.4.1.

3.4.4 MapReduce over Lustre with Local Disks

In this mode, Hadoop MapReduce can be launched on top of Lustre file system without using HDFS-level APIs. We can enable MapReduce over Lustre with two different kinds of Hadoop configuration settings. In the first setting, we use the local disks of each node to hold the intermediate data generated in the runtime of job execution. To enable MapReduce over Lustre with this setting, we use the following configuration steps:

1. Create the data directory in the Lustre File System.

```
mkdir <lustre-path-data-dir>
```

2. Configure Lustre file striping. It is recommended to use a stripe size equal to the file system block size (“fs.local.block.size”) in “core-site.xml”. For our package, we recommend to use a stripe size value of 256 MB.

```
lfs setstripe -s 256M <lustre-path-data-dir>
```

3. Configure mapred-site.xml file.

```
<configuration>

  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
    <description>Execution framework set to Hadoop
      YARN.</description>
  </property>

  <property>
    <name>yarn.app.mapreduce.am.command-opts</name>
    <value>-Xmx1024m
      -Dhadoop.conf.dir=$HADOOP_HOME/etc/hadoop</value>
    <description>Java opts for the MR App Master processes. The
      following symbol, if present, will be interpolated: @taskid@
      is replaced by current TaskID. Any other occurrences of '@'
      will go unchanged. For example, to enable verbose gc logging
      to a file named for the taskid in /tmp and to set the heap
      maximum to be a gigabyte, pass a 'value' of: -Xmx1024m
      -verbose:gc -Xloggc:/tmp/@taskid@.gc
      Usage of -Djava.library.path can cause programs to no longer
      function if hadoop native libraries are used. These values
      should instead be set as part of LD_LIBRARY_PATH in the map /
      reduce JVM env using the mapreduce.map.env and
      mapreduce.reduce.env config settings.</description>
  </property>

  <property>
    <name>yarn.app.mapreduce.am.env</name>
    <value>LD_LIBRARY_PATH=${HADOOP_HOME}/lib/native</value>
    <description>User added environment variables for the MR App
      Master
      processes. Example :
      1) A=foo This will set the env variable A to foo
      2) B=$B:c This is inherit tasktracker's B env variable.
    </description>
  </property>

  <property>
    <name>mapreduce.jobtracker.system.dir</name>
    <value><lustre-path-data-dir>/mapred/system</value>
    <description>The directory where MapReduce stores control files.
    </description>
  </property>
```

```
<property>
  <name>mapreduce.jobtracker.staging.root.dir</name>
  <value><lustre-path-data-dir>/mapred/staging</value>
  <description>The root of the staging area for users' job files.
</description>
</property>

<property>
  <name>yarn.app.mapreduce.am.staging-dir</name>
  <value><lustre-path-data-dir>/yarn/staging</value>
  <description>The staging dir used while submitting jobs.
</description>
</property>

<property>
  <name>mapred.rdma.shuffle.lustre</name>
  <value>1</value>
  <description>This parameter enables mapreduce over lustre with
    all enhanced shuffle algorithms. For MapReduce over Lustre
    with local disks, this parameter value should be 1. The
    default value of this parameter is -1, which enables
    RDMA-based shuffle for MapReduce over HDFS.</description>
</property>
</configuration>
```

4. Configure the core-site.xml file.

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>file://<lustre-path-data-dir>/namenode</value>
    <description>The name of the default file system. A URI whose
      scheme and authority determine the FileSystem implementation.
      The uri's scheme determines the config property
      (fs.SCHEME.impl) naming the FileSystem implementation class.
      The uri's authority is used to determine the host, port, etc.
      for a filesystem.</description>
  </property>

  <property>
    <name>fs.local.block.size</name>
    <value>268435456</value>
    <description>This value should be equal to Lustre stripe
      size.</description>
  </property>
</configuration>
```



```
</property>

<property>
  <name>hadoop.ib.enabled</name>
  <value>true</value>
  <description>Enable/Disable the RDMA feature over IB. Default
    value of hadoop.ib.enabled is true.</description>
</property>

<property>
  <name>hadoop.roce.enabled</name>
  <value>false</value>
  <description>Enable/Disable the RDMA feature over RoCE. Default
    value of hadoop.roce.enabled is false.</description>
</property>

<property>
  <name>hadoop.tmp.dir</name>
  <value>/tmp/hadoop_local</value>
  <description>A base for other temporary
    directories.</description>
</property>
</configuration>
```

5. Configure the `yarn-site.xml` file as specified in Section 3.4.1.
6. Configure the `slaves` file as specified in Section 3.4.1.

3.4.5 MapReduce over Lustre without Local Disks

For MapReduce over Lustre without local disks, Lustre should provide the intermediate data directories for MapReduce jobs. However, to do this, Hadoop tmp directory (“`hadoop.tmp.dir`”) must have separate paths to Lustre for each NodeManager. This can be achieved by installing a separate Hadoop package on each node with difference in the configuration value for “`hadoop.tmp.dir`”. To enable MapReduce over Lustre with this setting, we use the following configuration steps:

1. Make the Lustre data directory and configure the stripe size as specified in Section 3.4.4.
2. Copy the hadoop installation on each of the slaves in a local dir of that node. Note that, this step is a requirement for this mode.
3. Configure `mapred-site.xml` file.

```
<configuration>

  <property>
```

```

<name>mapreduce.framework.name</name>
<value>yarn</value>
<description>Execution framework set to Hadoop
    YARN.</description>
</property>

<property>
<name>yarn.app.mapreduce.am.command-opts</name>
<value>-Xmx1024m
    -Dhadoop.conf.dir=$HADOOP_HOME/etc/hadoop</value>
<description>Java opts for the MR App Master processes. The
    following symbol, if present, will be interpolated: @taskid@
    is replaced by current TaskID. Any other occurrences of '@'
    will go unchanged. For example, to enable verbose gc logging
    to a file named for the taskid in /tmp and to set the heap
    maximum to be a gigabyte, pass a 'value' of: -Xmx1024m
    -verbose:gc -Xloggc:/tmp/@taskid@.gc
    Usage of -Djava.library.path can cause programs to no longer
    function if hadoop native libraries are used. These values
    should instead be set as part of LD_LIBRARY_PATH in the map /
    reduce JVM env using the mapreduce.map.env and
    mapreduce.reduce.env config settings.
</description>
</property>

<property>
<name>yarn.app.mapreduce.am.env</name>
<value>LD_LIBRARY_PATH=${HADOOP_HOME}/lib/native</value>
<description>User added environment variables for the MR App
    Master
    processes. Example :
    1) A=foo This will set the env variable A to foo
    2) B=$B:c This is inherit tasktracker's B env variable.
</description>
</property>

<property>
<name>mapreduce.jobtracker.system.dir</name>
<value><lustre-path-data-dir>/mapred/system</value>
<description>The directory where MapReduce stores control files.
</description>
</property>

<property>
<name>mapreduce.jobtracker.staging.root.dir</name>
<value><lustre-path-data-dir>/mapred/staging</value>

```

```

    <description>The root of the staging area for users' job files.
    </description>
  </property>

  <property>
    <name>yarn.app.mapreduce.am.staging-dir</name>
    <value><lustre-path-data-dir>/yarn/staging</value>
    <description>The staging dir used while submitting jobs.
    </description>
  </property>

  <property>
    <name>mapred.rdma.shuffle.lustre</name>
    <value>2</value>
    <description>This parameter enables Mapreduce over Lustre with
      all enhanced shuffle algorithms. For MapReduce over Lustre
      without local disks, it can have a value of 0, 1, 2, or any
      fractional value in between 0 and 1. A value of 1 indicates
      that the entire shuffle will go over RDMA; a value of 0
      indicates that the Lustre file system read operation will
      take place instead of data shuffle. Any value in between
      (e.g. 0.70) indicates some shuffle going over RDMA (70%);
      while the rest is Lustre Read (30%). A special value of 2
      enables a hybrid shuffle algorithm for MapReduce over Lustre
      with Lustre as intermediate data dir. The default value of
      this parameter is -1, which enables RDMA-based shuffle for
      MapReduce over HDFS.</description>
  </property>
</configuration>

```

4. Configure the `core-site.xml` file for each NodeManager separately.

```

<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>file://<lustre-path-data-dir>/namenode</value>
    <description>The name of the default file system. A URI whose
      scheme and authority determine the FileSystem implementation.
      The uri's scheme determines the config property
      (fs.SCHEME.impl) naming the FileSystem implementation class.
      The uri's authority is used to determine the host, port, etc.
      for a filesystem.</description>
  </property>

  <property>

```

```
<name>fs.local.block.size</name>
<value>268435456</value>
<description>This value should be equal to Lustre stripe
    size.</description>
</property>

<property>
  <name>hadoop.ib.enabled</name>
  <value>true</value>
  <description>Enable/Disable the RDMA feature over IB. Default
    value of hadoop.ib.enabled is true.</description>
</property>

<property>
  <name>hadoop.roce.enabled</name>
  <value>false</value>
  <description>Enable/Disable the RDMA feature over RoCE. Default
    value of hadoop.roce.enabled is false.</description>
</property>

<property>
  <name>hadoop.tmp.dir</name>
  <value><lustre-path-data-dir>/tmp_{$nodeID}</value>
  <description>A base for other temporary directories. For
    MapReduce over Lustre with Lustre providing the intermediate
    data storage, this parameter value should be different for
    each NodeManager. For example, the first NodeManager in the
    slaves file may have a value of <lustre-path-data-dir>/tmp_0,
    while the second NodeManager can be configured with
    <lustre-path-data-dir>/tmp_1, and so on. Any value can
    replace $nodeID in the value here as long as the value is
    different for each node.</description>
</property>
</configuration>
```

5. Configure the `yarn-site.xml` file as specified in Section 3.4.1.
6. Configure the `slaves` file as specified in Section 3.4.1.

3.5 Advanced Configuration

Some advanced features in RDMA for Apache Hadoop 2.x 0.9.9 can be manually enabled by users. Steps to configure these features in RDMA for Apache Hadoop 2.x 0.9.9 are discussed in this section.

3.5.1 Parallel Replication

Enable parallel replication in HDFS by configuring `hdfs-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 will choose the pipeline replication mechanism. This parameter is applicable for HHH and HHH-M modes as discussed in Section 3.4.1 and Section 3.4.2, respectively.

```
<property>
  <name>dfs.replication.parallel.enabled</name>
  <value>true</value>
  <description>Enable the parallel replication feature in HDFS.
    Default value of dfs.replication.parallel is false.
  </description>
</property>
```

3.5.2 Placement Policy Selection

Select specific placement policy (Greedy or Balanced) in HDFS by configuring `hdfs-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 selects policy automatically based on the storage types of the HDFS data directories. This parameter is applicable for HHH and HHH-L modes as discussed in Section 3.4.1 and Section 3.4.3, respectively.

```
<property>
  <name>dfs.rdma.placement.policy</name>
  <value>Greedy/Balanced</value>
  <description>Enable specific data placement policy. </description>
</property>
```

3.5.3 Automatic Placement Policy Selection

By default, RDMA for Apache Hadoop 2.x 0.9.9 selects policy automatically based on the storage types of the HDFS data directories. This parameter can be used if the user wants to disable automatic policy detection. This parameter is applicable for HHH mode as discussed in Section 3.4.1.

```
<property>
  <name>dfs.rdma.policy.autodetect</name>
  <value>false</value>
  <description>Disable automatic policy detection (default is
    true). </description>
</property>
```

In order to use the storage policies of default HDFS, users should not use the `dfs.rdma.placement.policy` parameter as discussed in Section 3.5.2 and disable policy auto detection.

3.5.4 Threshold of RAM Disk Usage

Select a threshold of RAM Disk usage in HDFS by configuring `hdfs-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 uses 70% of RAM Disk when RAM Disk is configured as a HDFS data directory. This parameter is applicable for HHH, HHH-M, and HHH-L modes as discussed in Section 3.4.1, Section 3.4.2, and Section 3.4.3, respectively.

```
<property>
  <name>dfs.rdma.memory.percentage</name>
  <value>0.5</value>
  <description>Select a threshold (default = 0.7) for RAM Disk usage.
  </description>
</property>
```

3.5.5 No. of In-Memory Replicas

Select the number of in-memory replicas in HHH mode by configuring `hdfs-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 writes two replicas to RAM Disk and one to persistent storage (replication factor = 3). The no. of in-memory replicas can be changed from one to no. of replication factor (all in-memory). This parameter is applicable for HHH mode as discussed in Section 3.4.1.

```
<property>
  <name>dfs.rdma.memory.replica</name>
  <value>3</value>
  <description>Select no. of in-memory replicas (default = 2).
  </description>
</property>
```

3.5.6 Disk-assisted Shuffle

Enable disk-assisted shuffle in MapReduce by configuring `mapred-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 assumes that disk-assisted shuffle is disabled. We encourage our users to enable this parameter if they feel that the local disk performance in their cluster is good. This parameter is applicable for all the modes in RDMA for Apache Hadoop 2.x 0.9.9.

```
<property>
  <name>mapred.disk-assisted.shuffle.enabled</name>
  <value>true</value>
  <description>Enable disk-assisted shuffle in MapReduce. Default
  value of mapred.disk-assisted.shuffle.enabled is false.
  </description>
</property>
```

3.5.7 SSD Optimization for Intermediate Data

Enable SSD-based shuffle in MapReduce by configuring `mapred-site.xml` file. By default, RDMA for Apache Hadoop 2.x 0.9.9 assumes that SSD is not used for intermediate data directories. We encourage our users to enable this parameter if they configure SSD for intermediate data directories. This parameter is applicable for all the modes in RDMA for Apache Hadoop 2.x 0.9.9.

```
<property>
  <name>mapred.rdma.local.ssd.used</name>
  <value>true</value>
  <description>Enable SSD-assisted shuffle in MapReduce. Default value
    of mapred.rdma.local.ssd.used is false. </description>
</property>
```

4 Basic Usage Instructions

RDMA for Apache Hadoop 2.x 0.9.9 has management operations similar to default Apache Hadoop 2.7.1. This section lists several of them for basic usage.

4.1 Startup

4.1.1 MapReduce over HDFS

To run MapReduce over HDFS with any of the modes (HHH/HHH-M/HHH-L), please follow these steps.

1. Use the following command to format the directory which stores the namespace and transactions logs for NameNode.

```
$ bin/hdfs namenode -format
```

2. Start HDFS with the following command:

```
$ sbin/start-dfs.sh
```

3. Start YARN with the following command:

```
$ sbin/start-yarn.sh
```

4.1.2 MapReduce over Lustre

To run MapReduce over Lustre, startup involves starting YARN only.

1. Start YARN with the following command:

```
$ sbin/start-yarn.sh
```

4.2 Basic Commands

1. Use the following command to manage HDFS:

```
$ bin/hdfs dfsadmin
Usage: java DFSAdmin
Note: Administrative commands can only be run as the HDFS
superuser.
    [-report]
    [-safemode enter | leave | get | wait]
    [-allowSnapshot <snapshotDir>]
    [-disallowSnapshot <snapshotDir>]
    [-saveNamespace]
    [-rollEdits]
    [-restoreFailedStorage true|false|check]
    [-refreshNodes]
    [-finalizeUpgrade]
    [-rollingUpgrade [<query|prepare|finalize>]]
    [-metasave filename]
    [-refreshServiceAcl]
    [-refreshUserToGroupsMappings]
    [-refreshSuperUserGroupsConfiguration]
    [-refreshCallQueue]
    [-printTopology]
    [-refreshNamenodes datanodehost:port]
    [-deleteBlockPool datanode-host:port blockpoolId [force]]
    [-setQuota <quota> <dirname>...<dirname>]
    [-clrQuota <dirname>...<dirname>]
    [-setSpaceQuota <quota> <dirname>...<dirname>]
    [-clrSpaceQuota <dirname>...<dirname>]
    [-setBalancerBandwidth <bandwidth in bytes per second>]
    [-fetchImage <local directory>]
    [-shutdownDatanode <datanode_host:ipc_port> [upgrade]]
    [-getDatanodeInfo <datanode_host:ipc_port>]
    [-help [cmd]]
```

Generic options supported are

```
-conf <configuration file> specify an application configuration
    file
-D <property=value> use value for given property
-fs <local|namenode:port> specify a namenode
-jt <local|jobtracker:port> specify a job tracker
-files <comma separated list of files> specify comma separated
    files to be copied to the map reduce cluster
-libjars <comma separated list of jars> specify comma separated jar
    files to include in the classpath.
```



```
-archives <comma separated list of archives> specify comma
separated archives to be unarchived on the compute machines.
```

For example, we often use the following command to show the status of HDFS:

```
$ bin/hdfs dfsadmin -report
```

2. Use the following command to manage files in HDFS:

```
$ bin/hdfs dfs
Usage: hadoop fs [generic options]
    [-appendToFile <localsrc> ... <dst>]
    [-cat [-ignoreCrc] <src> ...]
    [-checksum <src> ...]
    [-chgrp [-R] GROUP PATH...]
    [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
    [-chown [-R] [OWNER][:[GROUP]] PATH...]
    [-copyFromLocal [-f] [-p] <localsrc> ... <dst>]
    [-copyToLocal [-p] [-ignoreCrc] [-crc] <src> ...
    <localdst>]
    [-count [-q] <path> ...]
    [-cp [-f] [-p] <src> ... <dst>]
    [-createSnapshot <snapshotDir> [<snapshotName>]]
    [-deleteSnapshot <snapshotDir> <snapshotName>]
    [-df [-h] [<path> ...]]
    [-du [-s] [-h] <path> ...]
    [-expunge]
    [-get [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
    [-getfacl [-R] <path>]
    [-getmerge [-nl] <src> <localdst>]
    [-help [cmd ...]]
    [-ls [-d] [-h] [-R] [<path> ...]]
    [-mkdir [-p] <path> ...]
    [-moveFromLocal <localsrc> ... <dst>]
    [-moveToLocal <src> <localdst>]
    [-mv <src> ... <dst>]
    [-put [-f] [-p] <localsrc> ... <dst>]
    [-renameSnapshot <snapshotDir> <oldName> <newName>]
    [-rm [-f] [-r|-R] [-skipTrash] <src> ...]
    [-rmdir [--ignore-fail-on-non-empty] <dir> ...]
    [-setfacl [-R] [{-b|-k} {-m|-x <acl_spec>} <path>}] [--set
    <acl_spec> <path>]]
    [-setrep [-R] [-w] <rep> <path> ...]
    [-stat [format] <path> ...]
    [-tail [-f] <file>]
    [-test -[defsz] <path>]
    [-text [-ignoreCrc] <src> ...]
```

```
[-touchz <path> ...]
[-usage [cmd ...]]
```

Generic options supported are

```
-conf <configuration file> specify an application configuration
  file
-D <property=value> use value for given property
-fs <local|namenode:port> specify a namenode
-jt <local|jobtracker:port> specify a job tracker
-files <comma separated list of files> specify comma separated
  files to be copied to the map reduce cluster
-libjars <comma separated list of jars> specify comma separated jar
  files to include in the classpath.
-archives <comma separated list of archives> specify comma
  separated archives to be unarchived on the compute machines.
```

For example, we can use the following command to list directory contents of HDFS:

```
$ bin/hdfs dfs -ls /
```

3. Use the following command to interoperate with the MapReduce framework:

```
$ bin/mapred job
Usage: JobClient <command> <args>
  [-submit <job-file>]
  [-status <job-id>]
  [-counter <job-id> <group-name> <counter-name>]
  [-kill <job-id>]
  [-set-priority <job-id> <priority>]. Valid values for
    priorities are: VERY_HIGH HIGH NORMAL LOW VERY_LOW
  [-events <job-id> <from-event-#> <#-of-events>]
  [-history <jobOutputDir>]
  [-list [all]]
  [-list-active-trackers]
  [-list-blacklisted-trackers]
  [-list-attempt-ids <job-id> <task-type> <task-state>]

  [-kill-task <task-id>]
  [-fail-task <task-id>]
```

```
Usage: CLI <command> <args>
  [-submit <job-file>]
  [-status <job-id>]
  [-counter <job-id> <group-name> <counter-name>]
  [-kill <job-id>]
  [-set-priority <job-id> <priority>]. Valid values for
    priorities are: VERY_HIGH HIGH NORMAL LOW VERY_LOW
```

```

[-events <job-id> <from-event-#> <#-of-events>]
[-history <jobHistoryFile>]
[-list [all]]
[-list-active-trackers]
[-list-blacklisted-trackers]
[-list-attempt-ids <job-id> <task-type> <task-state>]. Valid
  values for <task-type> are REDUCE MAP. Valid values for
  <task-state> are running, completed
[-kill-task <task-attempt-id>]
[-fail-task <task-attempt-id>]
[-logs <job-id> <task-attempt-id>]

```

Generic options supported are

```

-conf <configuration file> specify an application configuration
  file
-D <property=value> use value for given property
-fs <local|namenode:port> specify a namenode
-jt <local|jobtracker:port> specify a job tracker
-files <comma separated list of files> specify comma separated
  files to be copied to the map reduce cluster
-libjars <comma separated list of jars> specify comma separated jar
  files to include in the classpath.
-archives <comma separated list of archives> specify comma
  separated archives to be unarchived on the compute machines.

```

For example, we can use the following command to list all active trackers of MapReduce:

```
$ bin/mapred job -list-active-trackers
```

4.3 Shutdown

4.3.1 MapReduce over HDFS

1. Stop HDFS with the following command:

```
$ sbin/stop-dfs.sh
```

2. Stop YARN with the following command:

```
$ sbin/stop-yarn.sh
```

4.3.2 MapReduce over Lustre

1. Stop YARN with the following command:

```
$ sbin/stop-yarn.sh
```

5 Running RDMA for Apache Hadoop with SLURM/PBS

To run RDMA for Apache Hadoop with SLURM/PBS, scripts in `HADOOP_HOME/bin/slurm_pbs/` directory can be used. These scripts can be used in interactive mode or batch mode. In the interactive mode, the user must allocate interactive nodes, and explicitly use the startup, run benchmark, and shutdown commands described Sections 5.2 or 5.3, in their interactive session. In the batch mode, the users must create and launch a SLURM/PBS batch script with the startup, run benchmark, and shutdown commands described Sections 5.2 or 5.3.

Detailed steps for using these scripts are described in Sections 5.1, 5.2 and 5.3.

5.1 Usage

This section gives an overview of the SLURM/PBS scripts and their usage for startup and shutdown for running a Hadoop cluster.

5.1.1 Configure and Start a Hadoop Job

For installing, configuring, and starting RDMA for Apache Hadoop with any particular mode of operation, `hibd_install_configure_start.sh` can be used. This script can configure and start Hadoop depending on the parameters provided by the user. Detailed parameter options available with this script are mentioned below:

```
$ ./hibd_install_configure_start.sh ?
Usage: hibd_install_configure_start.sh [options]
  -h <dir>
    specify location of hadoop installation a.k.a. hadoop home

  -m <hhh | hhh-m | hhh-l | mrlustre-local | mrlustre-lustre>
    specify the mode of operation (default: hhh). For more
    information, visit http://hibd.cse.ohio-state.edu/overview/

  -c <dir>
    specify the hadoop conf dir (default: ""). If user provides
    this directory, then the conf files are chosen from this
    directory. Otherwise, the conf files are generated automatically
    with/without user provided configuration with flag '-u'

  -j <dir>
    specify jdk installation or JAVA_HOME (default: ""). If user
    does not provide this, then java installation is searched in
    the environment.

  -u <file>
    specify a file containing all the configurations for hadoop
```

```

installation(default: n/a). Each line of this file must be
formatted as below:
"<C|H|M|Y>\t<parameter_name>\t<parameter_value>"
C = core-site.xml, H = hdfs-site.xml, M = mapred-site.xml,
Y = yarn-site.xml

-l <dir>
specify the Lustre path to use for hhh-l, mrlustre-local,
and mrlustre-lustre modes (default: "")

-r <dir>
specify the ram disk path to use for hhh and hhh-m modes
(default: /dev/shm)

-s
specify to start hadoop after installation and configuration

-?
show this help message

```

5.1.2 Shutdown Hadoop Cluster

After running a benchmark with the script indicated in Section 5.1.1, stopping the Hadoop cluster with cleanup of all the directories can be achieved by using the script `hibd_stop_cleanup.sh`. Similar to the startup script, this cleanup script can makes different parameters available to the user. Detailed parameter options available with this script are mentioned below:

```

$ ./hibd_stop_cleanup.sh ?
Usage: hibd_stop_cleanup.sh [options]
-h <dir>
specify location of hadoop installation a.k.a. hadoop home

-m <hhh | hhh-m | hhh-l | mrlustre-local | mrlustre-lustre>
specify the mode of operation (default: hhh).
For more information, visit
http://hibd.cse.ohio-state.edu/overview/

-c <dir>
specify the hadoop conf dir (default: "").

-l <dir>
specify the Lustre path to use for hhh-l, mrlustre-local,
and mrlustre-lustre modes (default: "")

-r <dir>

```

```

    specify the ram disk path to use for hhh and hhh-m modes
    (default: /dev/shm)

-d
    specify to delete logs and data after hadoop stops

-?
    show this help message

```

Details of the usage of the above-mentioned scripts can also be found in `slurm-script.sh`

5.2 Running MapReduce over HDFS with SLURM/PBS

The user can run MapReduce over HDFS in one of the three modes: HHH, HHH-M or HHH-L. Based on the parameters supplied to the `hibd_install_configure_start.sh` script, Hadoop cluster will start with the requested mode of operation configuration setting.

5.2.1 Startup

To start Hadoop in HHH mode, the following command can be used:

```
$ hibd_install_configure_start.sh -s -m hhh-default -r /dev/shm -h
  $HADOOP_HOME -j $JAVA_HOME
```

To start Hadoop in HHH-M mode, the following command can be used:

```
$ hibd_install_configure_start.sh -s -m hhh-m -r /dev/shm -h
  $HADOOP_HOME -j $JAVA_HOME
```

To start Hadoop in HHH-L mode, the following command can be used:

```
$ hibd_install_configure_start.sh -s -m hhh-l -r /dev/shm -h
  $HADOOP_HOME -j $JAVA_HOME -l <lustre_path>
```

5.2.2 Running Benchmarks

User can launch a benchmark after successful start of the Hadoop cluster. While running a benchmark, user should provide the Hadoop config directory using `--config` flag with `hadoop` script. If user does not have any pre-configured files, the default config directory will be created in the present working directory named as `conf` concatenated with job id.

```
$ $HADOOP_HOME/bin/hadoop --config ./conf_<job_id> jar
  $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  randomwriter -Dmapreduce.randomwriter.mapsperhost=4
  -Dmapreduce.randomwriter.bytespermap=67108864 rand_in
```

```
$ $HADOOP_HOME/bin/hadoop --config ./conf_<job_id> jar
  $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
sort rand_in rand_out
```

5.2.3 Shutdown

In order to stop Hadoop and cleanup the corresponding logs, the following command can be used:

```
$ hibd_stop_cleanup.sh -d -l <lustre_path> -h $HADOOP_HOME
```

5.3 Running MapReduce over Lustre with SLURM/PBS

5.3.1 Startup

To start Hadoop in MapReduce over Lustre with local disks mode, the following command can be used:

```
$ hibd_install_configure_start.sh -s -m mrlustre-local -l <lustre_path>
  -j $JAVA_HOME -h $HADOOP_HOME
```

To start Hadoop in MapReduce over Lustre without local disks mode, the following command can be used:

```
$ hibd_install_configure_start.sh -s -m mrlustre-lustre -l
  <lustre_path> -j $JAVA_HOME -h $HADOOP_HOME
```

5.3.2 Running Benchmarks

Running benchmark for MapReduce over Lustre with local disks follows the same guidelines as shown above. For benchmarks running on MapReduce over Lustre without local disks, the following command should be used.

```
$ /tmp/hadoop_install_<job_id>/bin/hadoop jar
  /tmp/hadoop_install_<job_id>/share/hadoop/mapreduce/
  hadoop-mapreduce-examples-*.jar randomwriter
  -Dmapreduce.randomwriter.mapsperhost=4
  -Dmapreduce.randomwriter.bytespermap=67108864
  file://<lustre_path>/hibd_data_<job_id>/rand_in

$ /tmp/hadoop_install_<job_id>/bin/hadoop jar
  /tmp/hadoop_install_<job_id>/share/hadoop/mapreduce/
  hadoop-mapreduce-examples-*.jar sort
  file://<lustre_path>/hibd_data_<job_id>/rand_in
  file://<lustre_path>/hibd_data_<job_id>/rand_out
```

5.3.3 Shutdown

For stopping Hadoop and clean the used directories, the same command as shown above can be used.

6 Benchmarks

6.1 TestDFSIO

The TestDFSIO benchmark is used to measure I/O performance of the underlying file system. It does this by using a MapReduce job to read or write files in parallel. Each file is read or written in a separate map task and the benchmark reports the average read/write throughput per map.

On a client node, the TestDFSIO write experiment can be run using the following command:

```
$ bin/hadoop jar
  share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-*-tests.jar
  TestDFSIO -write -nrFiles <nfiles> -fileSize <fsize>
```

This command writes ‘nfiles’ files and ‘fsize’ MB each.

To run the same job with MapReduce over Lustre, one additional config parameter must be added in `mapred-site.xml`.

```
<property>
  <name>test.build.data</name>
  <value><lustre-path-data-dir>/benchmarks/TestDFSIO</value>
</property>
```

After adding this config parameter, the same command as mentioned above can be used to launch TestDFSIO experiment on top of Lustre.

6.2 Sort

The Sort benchmark uses the MapReduce framework to sort the input directory into the output directory. The inputs and outputs are sequence files where the keys and values are `BytesWritable`. Before running the Sort benchmark, we can use `RandomWriter` to generate the input data. `RandomWriter` writes random data to HDFS using the MapReduce framework. Each map takes a single file name as input and writes random `BytesWritable` keys and values to the HDFS sequence file.

On a client node, the `RandomWriter` experiment can be run using the following command:

```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  randomwriter -Dmapreduce.randomwriter.bytespermap=<nbytes>
  -Dmapreduce.randomwriter.mapsperhost=<nmaps> <out-dir>
```

This command launches ‘nmaps’ maps per node, and each map writes ‘nbytes’ data to ‘out-dir’.

On a client node, the Sort experiment can be run using the following command:


```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  sort -r <nreds> <in-dir> <out-dir>
```

This command launches ‘nreds’ reduces to sort data from ‘in-dir’ to ‘out-dir’.

The input directory of Sort can be the output directory of RandomWriter.

To run the same job with MapReduce over Lustre, the following commands can be used.

```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  randomwriter -Dmapreduce.randomwriter.bytespermap=<nbytes>
  -Dmapreduce.randomwriter.mapsperhost=<nmaps>
  file:///<lustre-path-data-dir>/<out-dir>
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  sort -r <nreds> file:///<lustre-path-data-dir>/<in-dir>
  file:///<lustre-path-data-dir>/<out-dir>
```

6.3 TeraSort

TeraSort is probably the most well-known Hadoop benchmark. It is a benchmark that combines testing the HDFS and MapReduce layers of a Hadoop cluster. The input data for TeraSort can be generated by the TeraGen tool, which writes the desired number of rows of data in the input directory. By default, the key and value size is fixed for this benchmark at 100 bytes. TeraSort takes the data from the input directory and sorts it to another directory. The output of TeraSort can be validated by the TeraValidate tool.

Before running the TeraSort benchmark, we can use TeraGen to generate the input data as follows:

```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  teragen <nrows> <out-dir>
```

This command writes ‘nrows’ of 100-byte rows to ‘out-dir’.

On a client node, the TeraSort experiment can be run using the following command:

```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  terasort <in-dir> <out-dir>
```

This command sorts data from ‘in-dir’ to ‘out-dir’.

The input directory of TeraSort can be the output directory of TeraGen.

To run the same job with MapReduce over Lustre, the following commands can be used.

```
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  teragen <nrows> file:///<lustre-path-data-dir>/<out-dir>
$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar
  terasort file:///<lustre-path-data-dir>/<in-dir>
  file:///<lustre-path-data-dir>/<out-dir>
```

6.4 OHB Micro-benchmarks

The OHB Micro-benchmarks support standalone evaluations of Hadoop Distributed File System (HDFS) and Memcached (See [here](#)). These benchmarks help fine-tune each component avoiding the impact of others.

OSU HiBD-Benchmarks (OHB-0.8) have HDFS benchmarks for Sequential Write Latency (SWL), Sequential Read Latency (SRL), Random Read Latency (RRL), Sequential Write Throughput (SWT), Sequential Read Throughput (SRT).

The source code can be downloaded from <http://hibd.cse.ohio-state.edu/download/hibd/osu-hibd-benchmarks-0.8.tar.gz>. The source can be compiled with the help of the Makefile provided. More details on building and running the OHB Micro-benchmark are provided in the README.

A brief description of the benchmark is provided below:

Sequential Write Latency (SWL): This benchmark measures the latency of sequential write to HDFS. The benchmark takes five parameters: file name (`-fileName`), file size (`-fileSize`), block size (`-bSize`), replication factor (`-rep`), and buffer size (`-bufSize`). The mandatory parameters are file name and size (in MB). The output of the benchmark is the time taken to write the file to HDFS. The buffer size indicates the size of the write buffer. HDFS block size and replication factor can also be tuned through this benchmarks. The benchmark also prints the important configuration parameters of HDFS.

Sequential Read Latency (SRL): This benchmark measures the latency of sequential read from HDFS. The benchmark takes two parameters: file name and buffer size. The mandatory parameter is file name. The output of the benchmark is the time taken to read the file from HDFS. The buffer size indicates the size of the read buffer. The benchmark also prints the important configuration parameters of HDFS.

Random Read Latency (RRL): This benchmark measures the latency of random read from HDFS. The benchmark takes four parameters: file name (`-fileName`), file size (`-fileSize`), skip size (`-skipSize`) and buffer size (`-bufSize`). The mandatory parameters are file name and file size. The benchmark first creates a file 2x the file (read) size and then randomly reads from it with a default skip size of 10. The output of the benchmark is the time taken to read the file from HDFS. The buffer size indicates the size of the read buffer. The benchmark also prints the important configuration parameters of HDFS.

Sequential Write Throughput (SWT): This benchmark measures the throughput of sequential write to HDFS. The benchmark takes five parameters: file size (`-fileSize`), block size (`-bSize`), replication factor (`-rep`), buffer size (`-bufSize`), and an output directory (`-outDir`) for the output files. The mandatory parameters are file size (in MB) and the output directory. Linux xargs command is used to launch multiple concurrent writers. File size indicates the write size per writer. A hostfile contains the hostnames where the write processes are launched. The benchmark outputs the total write throughput in MBps. The buffer size indicates the size of the write buffer. HDFS block size and replication factor can also be tuned through this benchmarks.

Sequential Read Throughput (SRT): This benchmark measures the throughput of sequential read from HDFS. The benchmark takes three parameters: file size (`-fileSize`), buffer size (`-bufSize`), and an output directory (`-outDir`) for the output files. The mandatory parameters are file size (in MB) and the output directory. Linux xargs command is used to launch multiple concurrent readers. File size indicates the write size per reader. A hostfile contains the hostnames where the write processes are launched. The

benchmark outputs the total read throughput in MBps. The buffer size indicates the size of the read buffer.