

Today

Distributed Subsystems (Lesson 7)

- ✓ * Global memory system
- ✓ * Distributed Shared Memory
- ⇒ * Distributed file Systems

Monday

Failures + recovery (Lesson 8)

* Please watch LRVM lecture *

LRC with Multi-Writer Coherence Protocol

P1

LOCK(L);

$x \leftarrow \begin{cases} x, y, z \text{ pages} \\ y \leftarrow \text{modified in C.S} \end{cases} \Rightarrow x_d, y_d, z_d \} \text{ diff's}$

Unlock(L);

LRC With Multi-Writer Coherence Protocol

P1

LOCK(L);

$x \leftarrow \begin{cases} x, y, z \text{ pages} \\ y \leftarrow \text{modified in C.S} \end{cases} \Rightarrow x_d, y_d, z_d \} \text{ diff's}$

Unlock(L);

P2

invalidate \rightarrow LOCK(L) ;

x, y, z at

lock acquisition

(CLRC)

LRC With Multi-Writer Coherence Protocol

P1

LOCK(L);

$x \leftarrow \begin{cases} x, y, z \text{ pages} \\ y \leftarrow \text{modified in C.S} \end{cases}$ $\Rightarrow x_d, y_d, z_d \} \text{ diff's}$

Unlock(L);

P2

invalidate \rightarrow LOCK(L);

x, y, z at
lock acquisition
(CLRC)

$\leftarrow x \leftarrow$
fetch at
point of
access

LRC with Multi-Writer Coherence Protocol

P1

LOCK(L);

$x \leftarrow \begin{cases} x, y, z \text{ pages} \\ y \leftarrow \text{modified in C.S.} \\ z \leftarrow \end{cases}$ $\Rightarrow \{x_d, y_d, z_d\}$ diff's

Unlock(L);

P3

LOCK(L);

$x \leftarrow \Rightarrow x'_d$

unlock(L);

invalidate \rightarrow LOCK(L);
 x, y, z at
lock acquisition
(CLRC)

P2

LOCK(L);

$\leftarrow x$ \leftarrow
fetch at
point of
access

LRC with Multi-Writer Coherence Protocol

P1

LOCK(L);

$x \leftarrow \begin{cases} x, y, z \text{ pages} \\ y \leftarrow \text{modified in C.S.} \\ z \leftarrow \end{cases}$ $\Rightarrow \{x_d, y_d, z_d\}$ diff's

Unlock(L);

P3

LOCK(L);

$x \leftarrow \Rightarrow x'_d$

unlock(L);

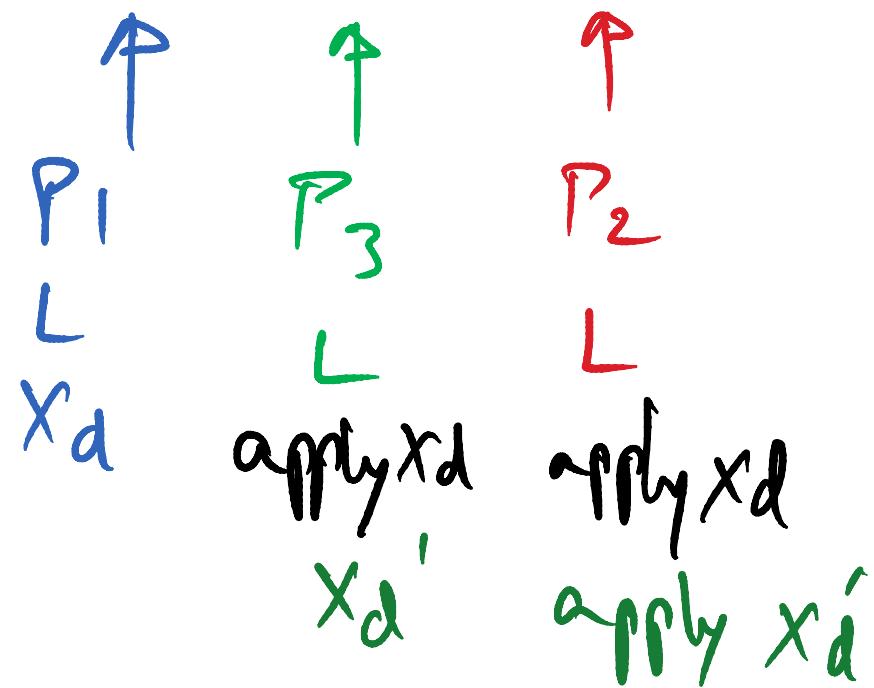
invalidate \rightarrow LOCK(L);

x, y, z at

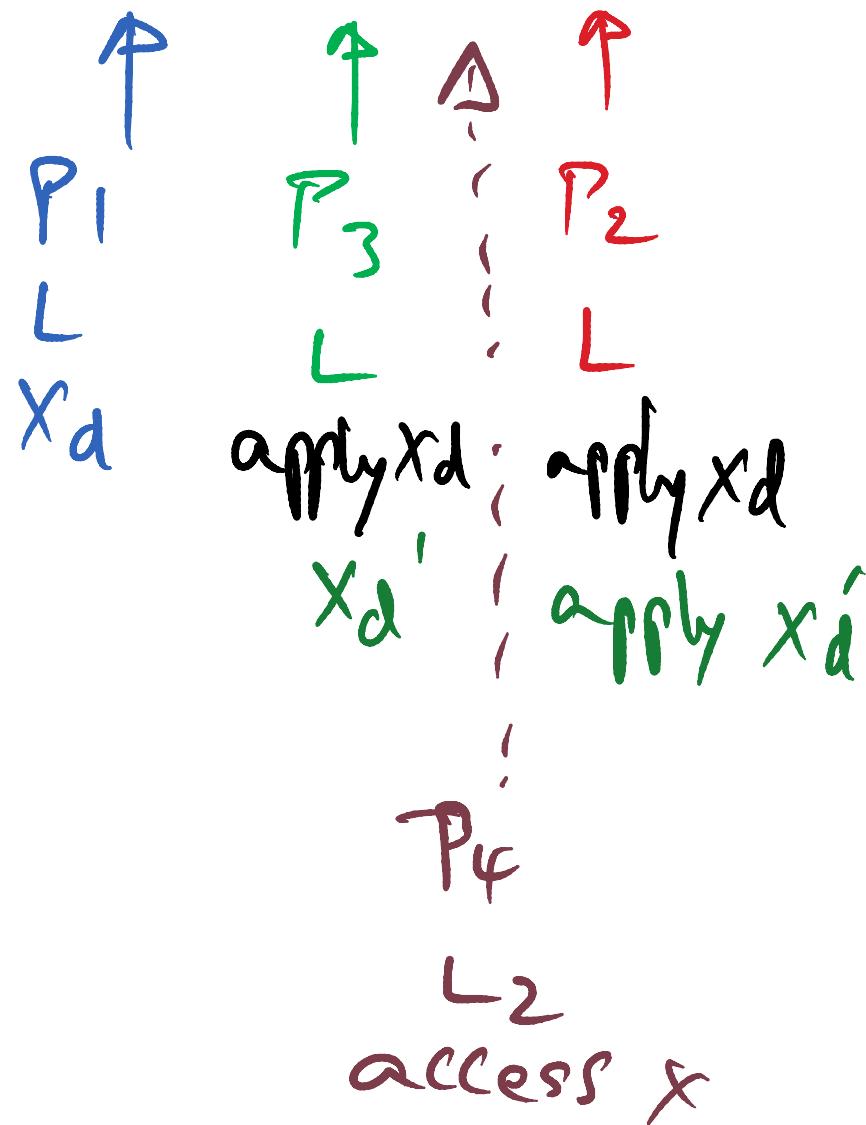
lock acquisition
(CLRC)

$x \leftarrow$ $Q \leftarrow$ fetch at
point of access

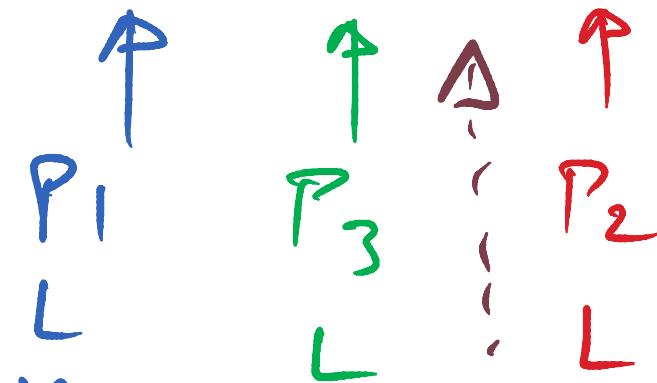
\Rightarrow true



\Rightarrow true



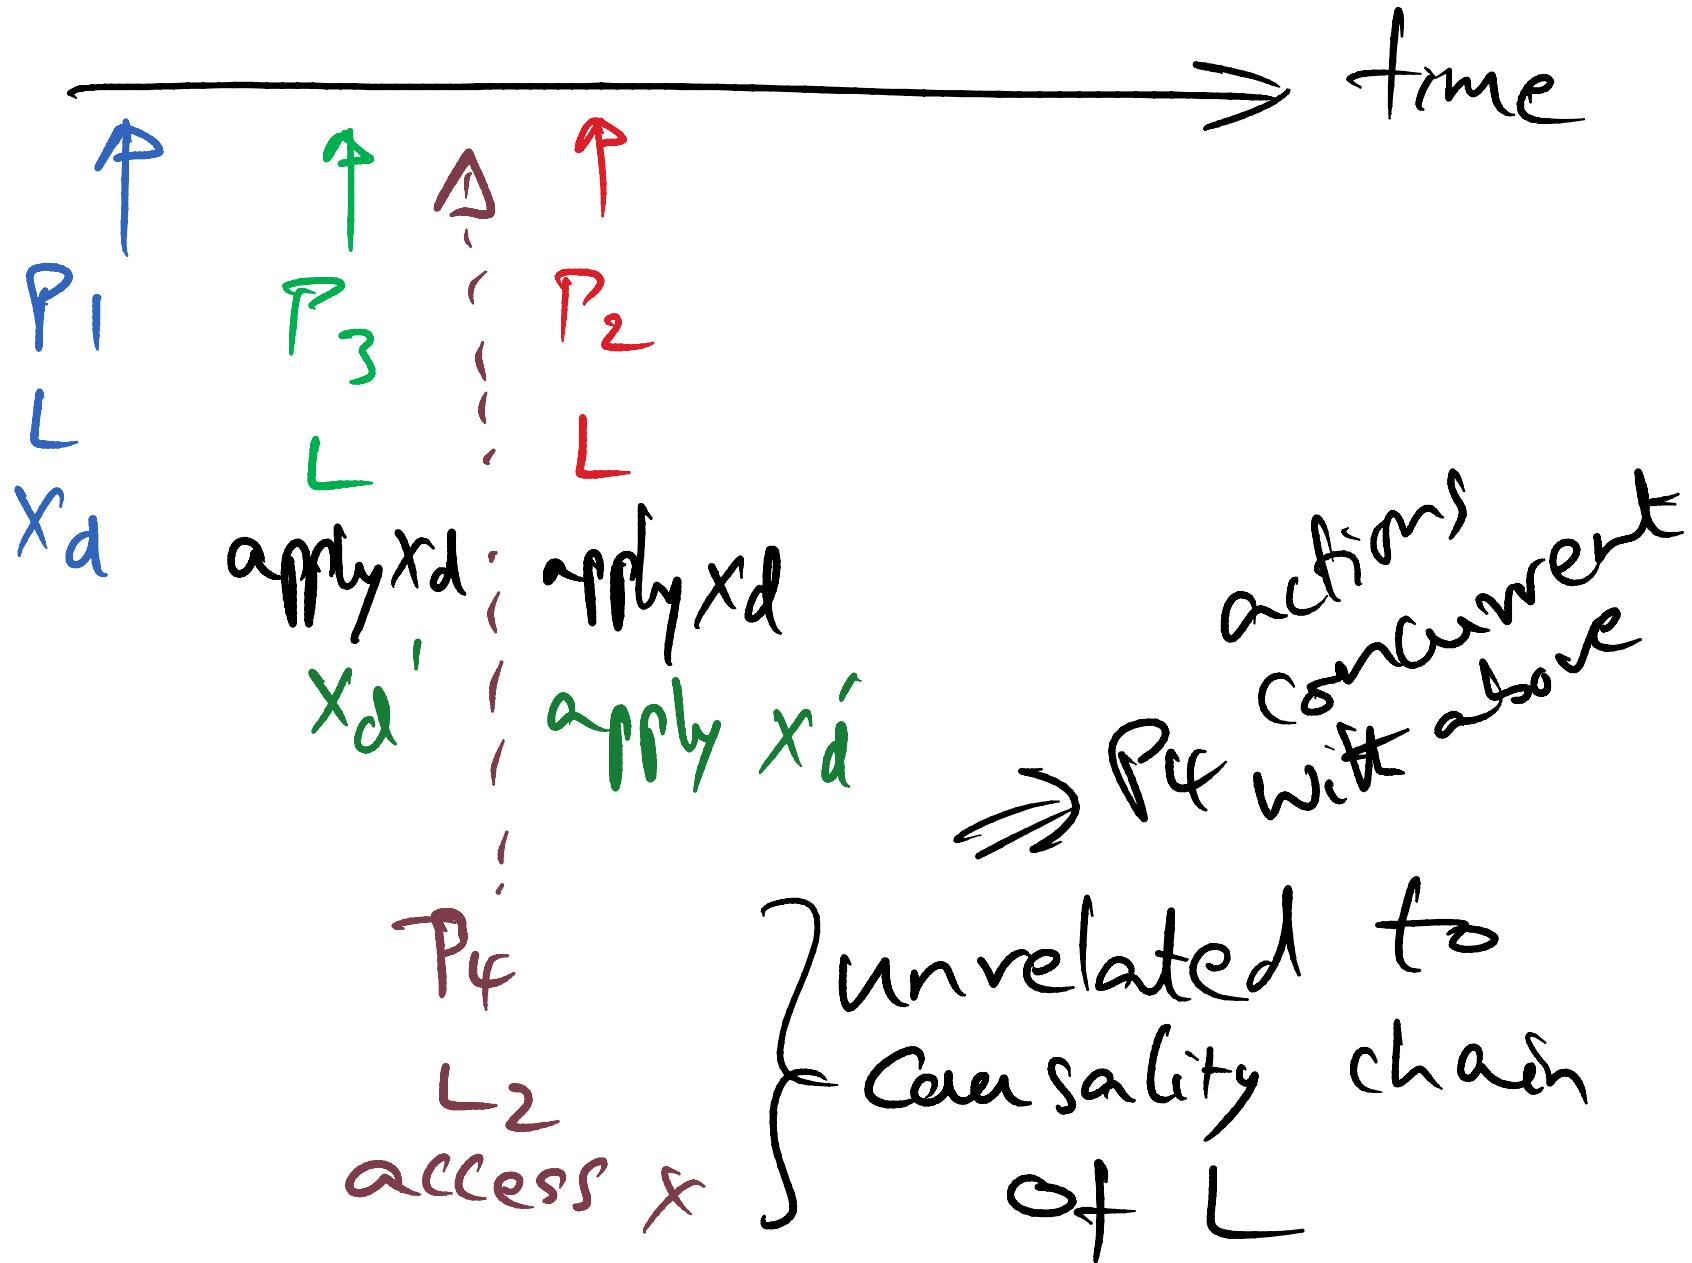
\Rightarrow time

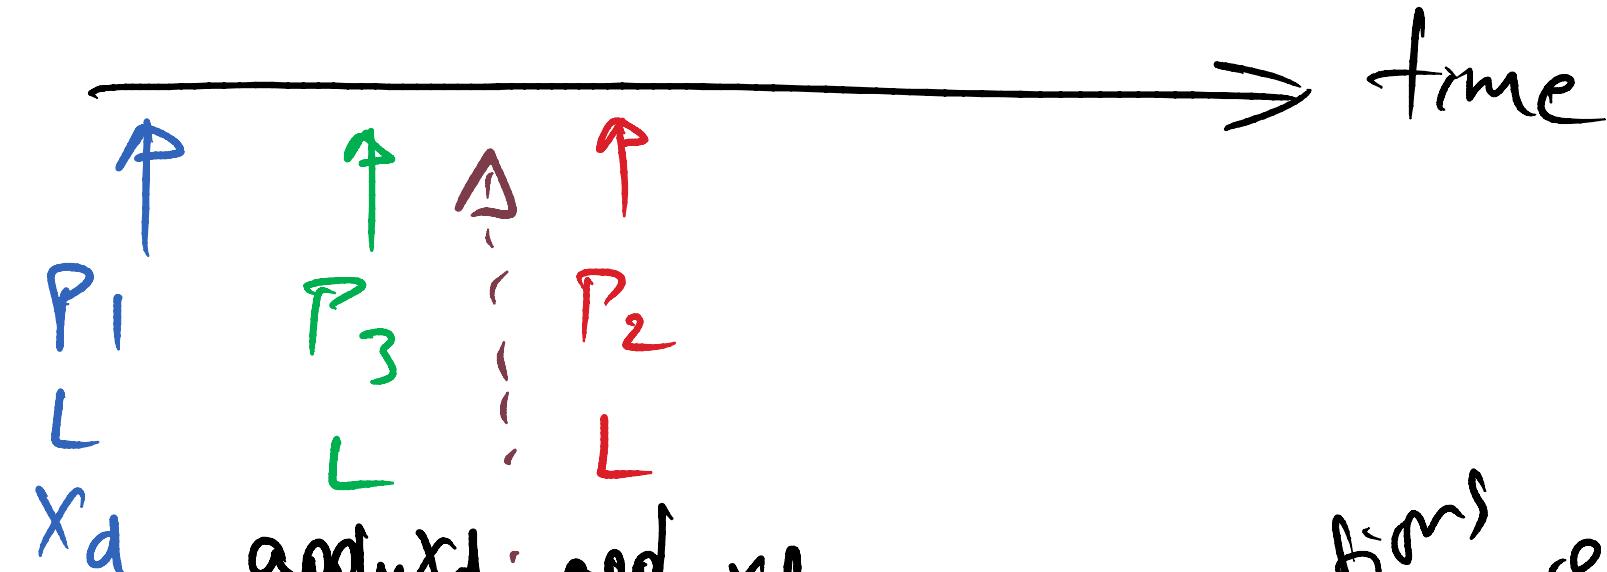


apply X_d; apply X_d
X_{d'}; apply X_{d'}

P₄
L₂
access x

unrelated to
causality chain
of L





Multi-writer Coherence Protocol

P_4

L_2

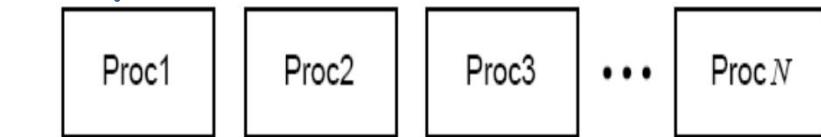
access x

unrelated to
Causality chain
of L

actions concurrent with above

Software DSM

Application



Pagefault

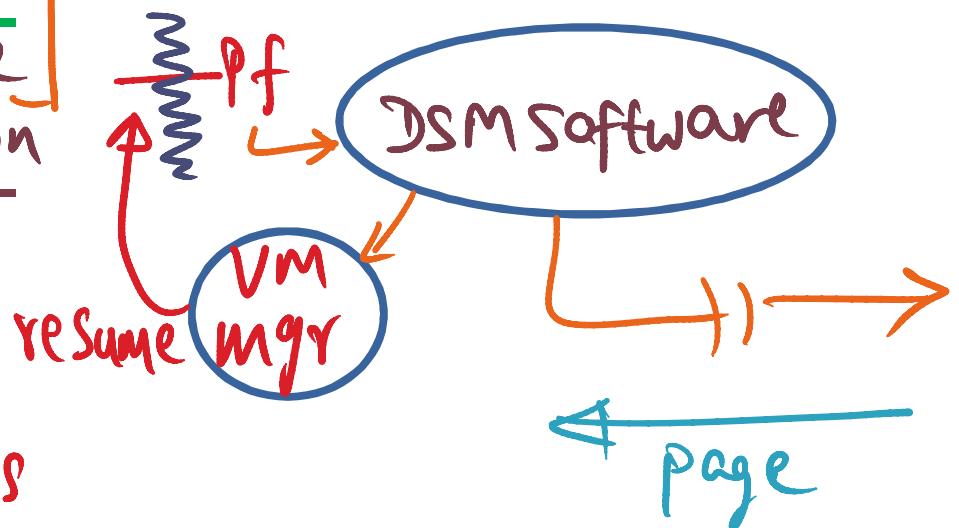
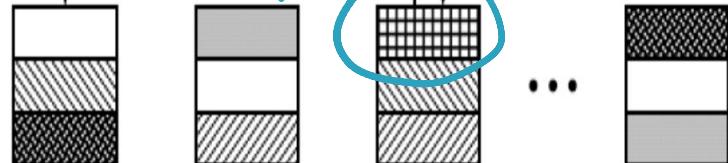
Global Virtual
memory
abstraction

Address space partitioned
Address equivalence
Distributed ownership

Contact owner of page
to get current copy

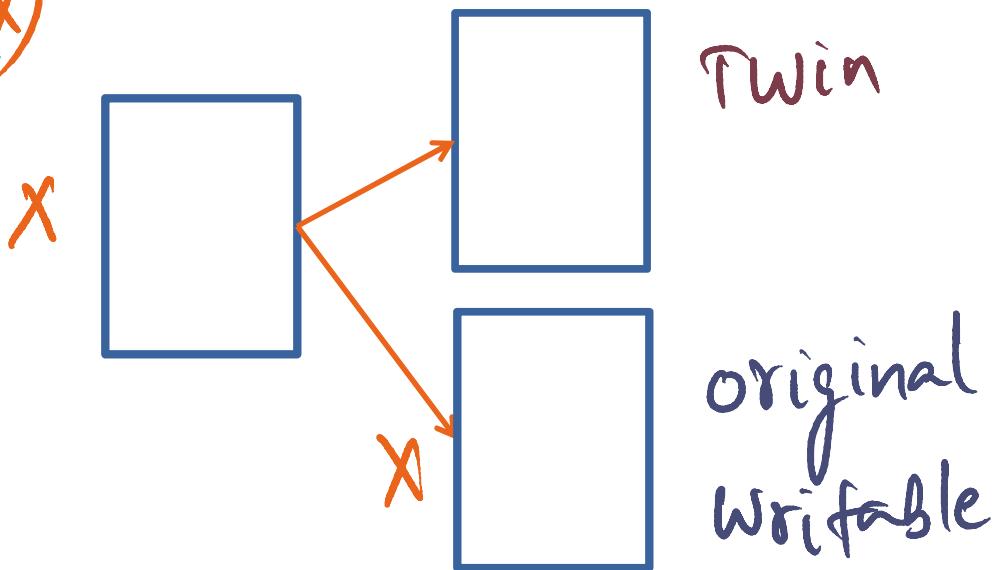
DSM Software
implementation

data xfer



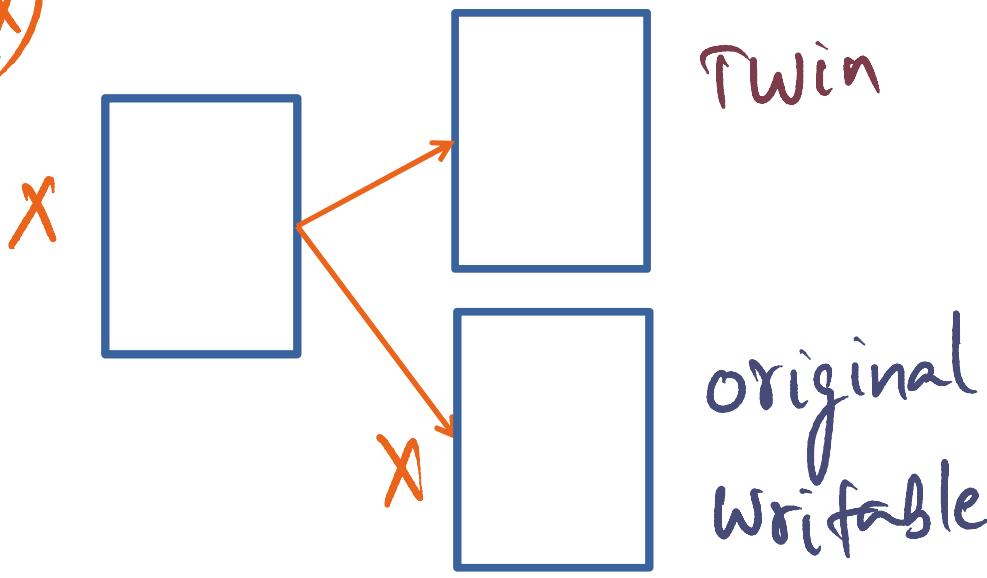
Implementation

write(\times)

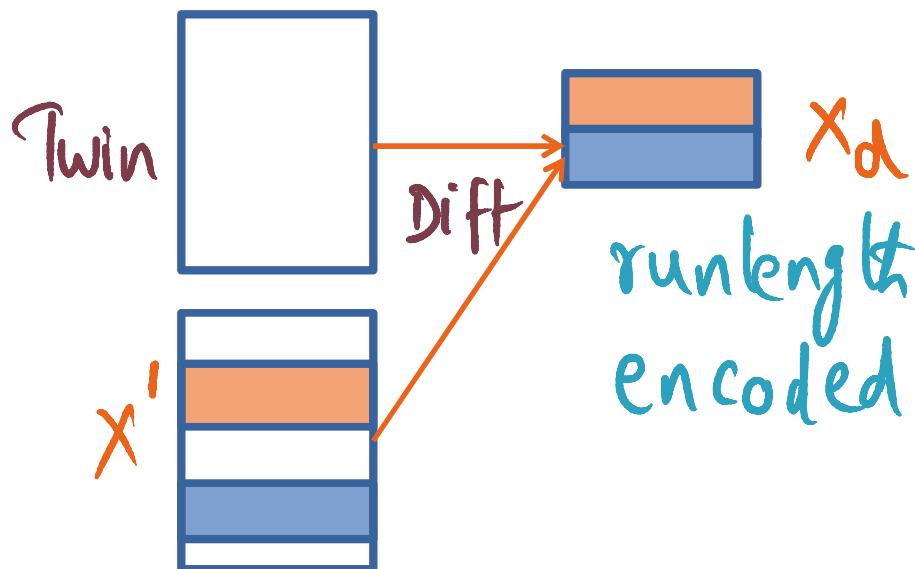


Implementation

write(x)

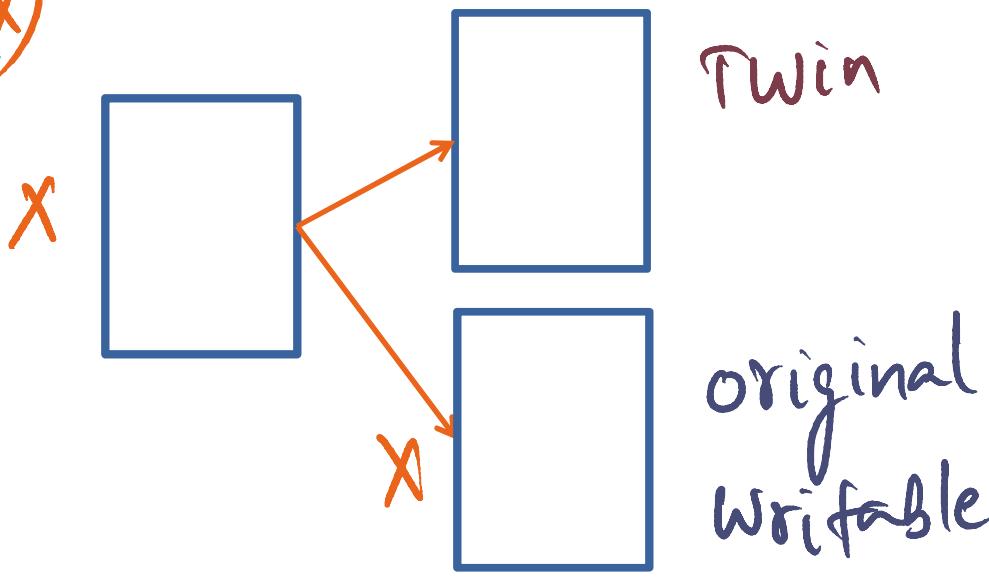


release

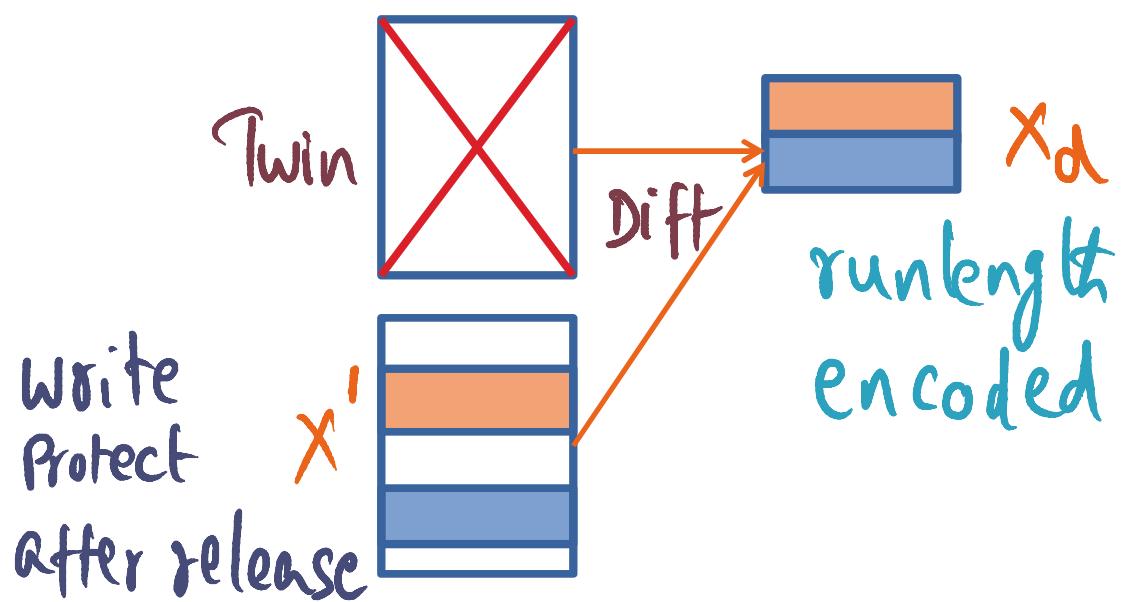


Implementation

write(x)

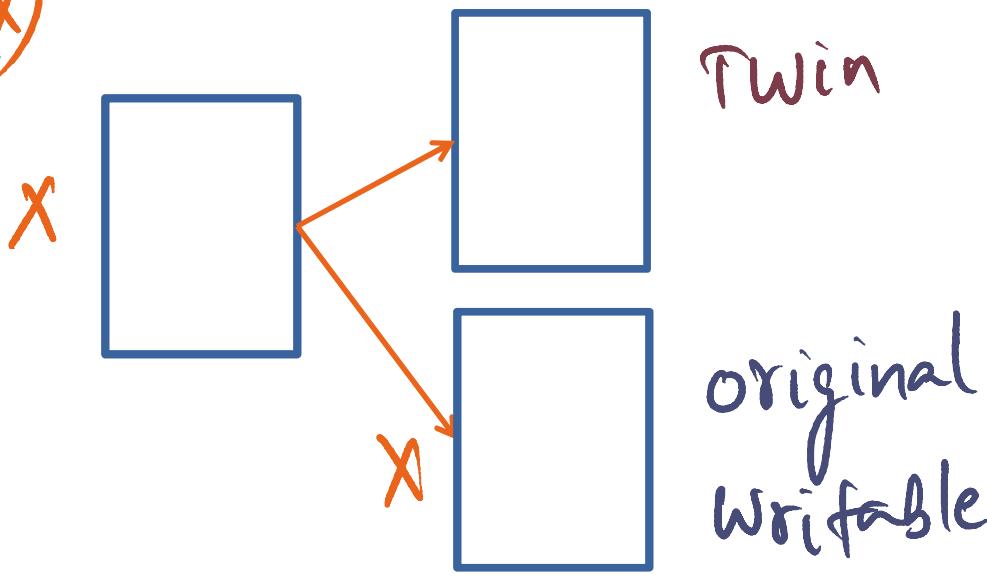


release

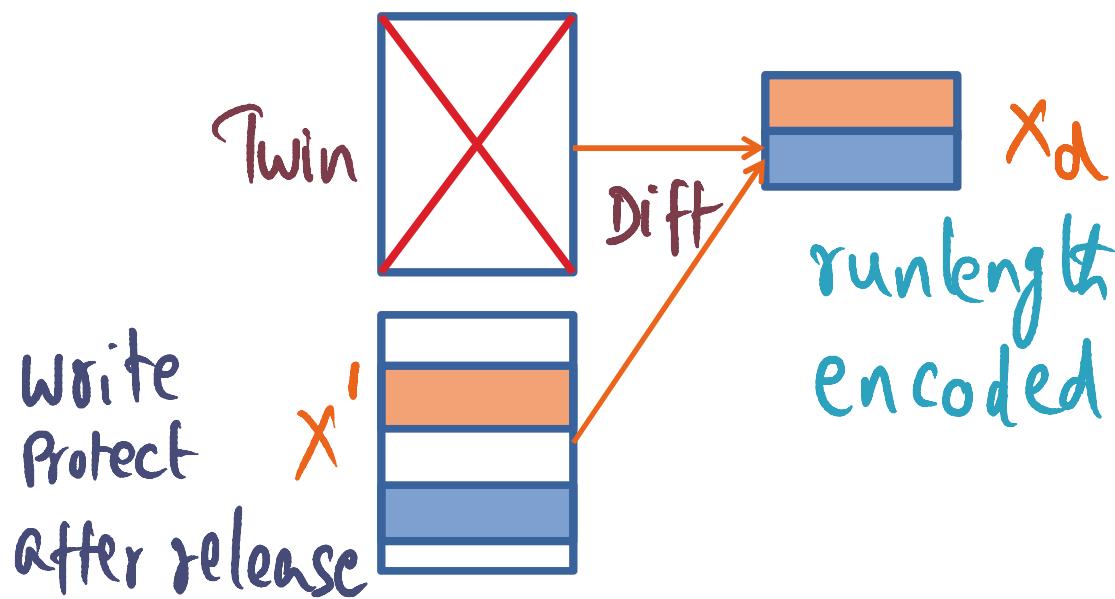


Implementation

write(x)



release



writes to same
portion of a
page under
different locks

⇒ Data race



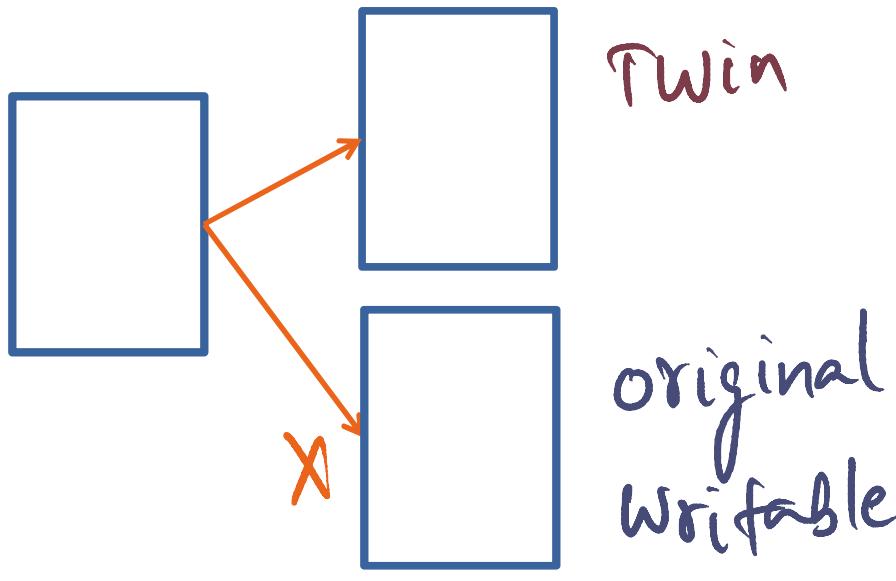
application
problem

Implementation

write(χ)

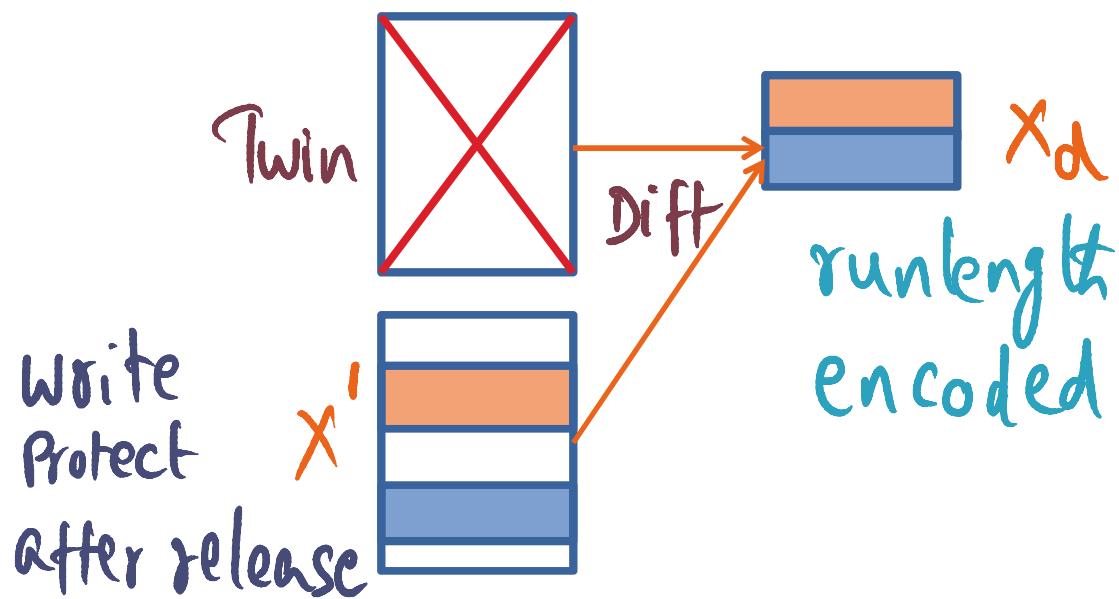
DSM
+
OS coop

release



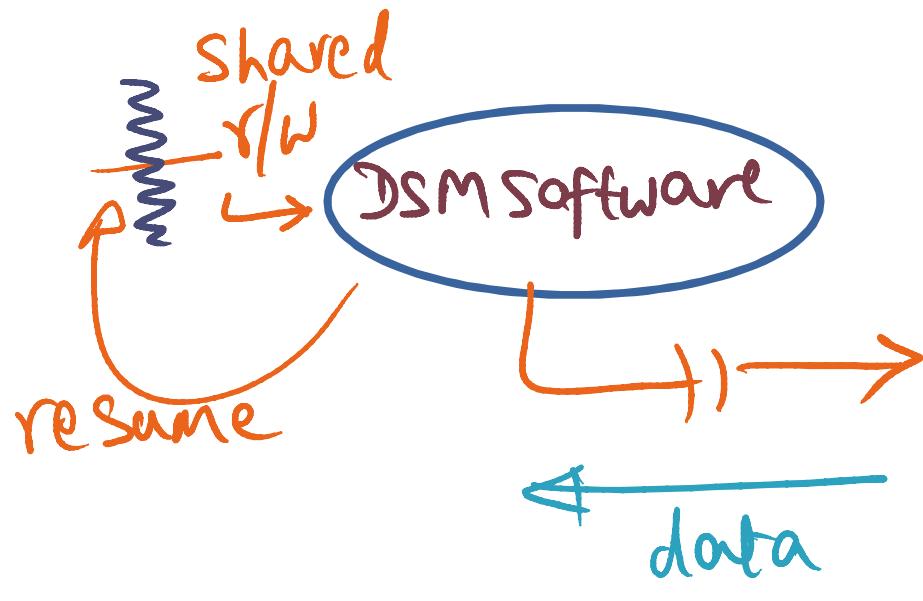
writes to same
portion of a
page under
different locks

⇒ Data race



application
problem

Non-page-based DSM

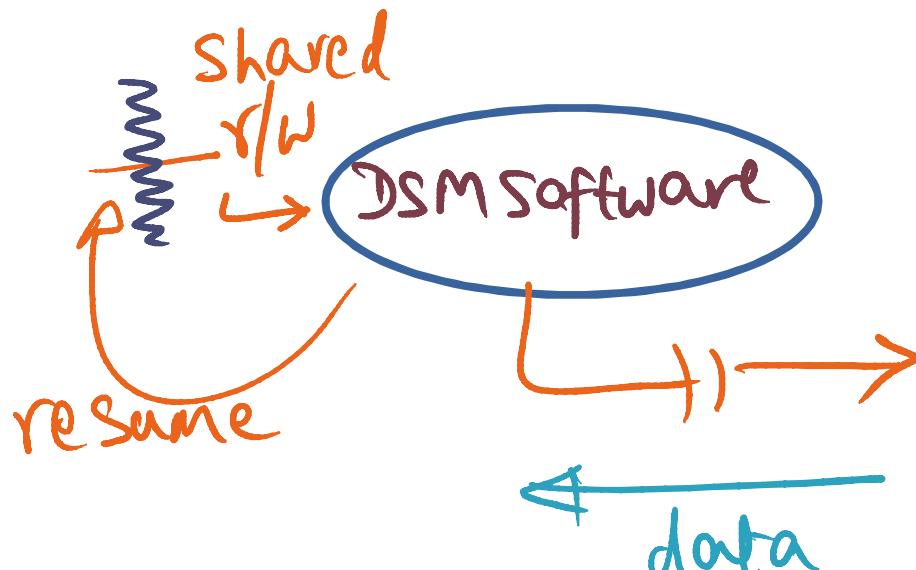


Library-based

- annotate shared variables

- Coherence actions inserted at point of access

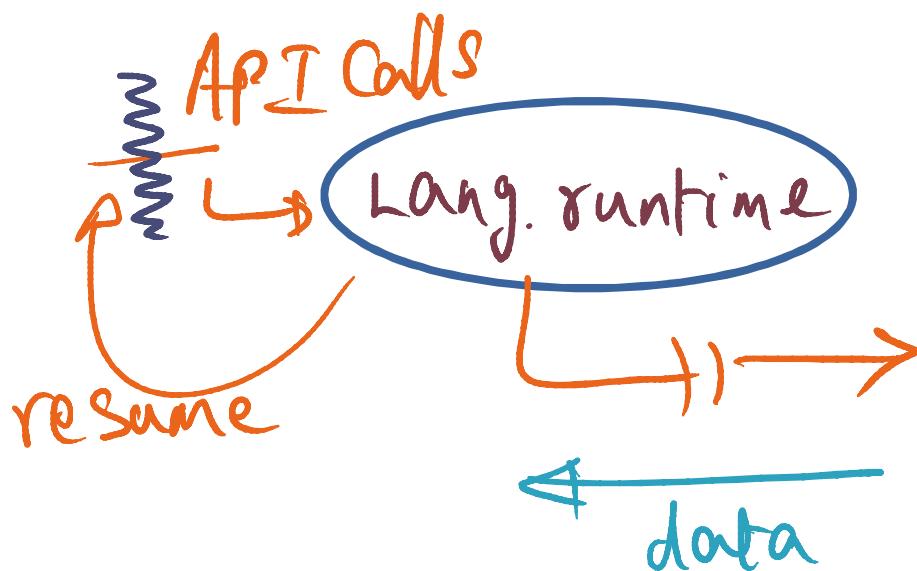
Non-page-based DSM



Library-based

- annotate shared variables

- Coherence actions inserted at point of access



Structured DSM

- API for structs
- Coherence actions on API calls

Scalability

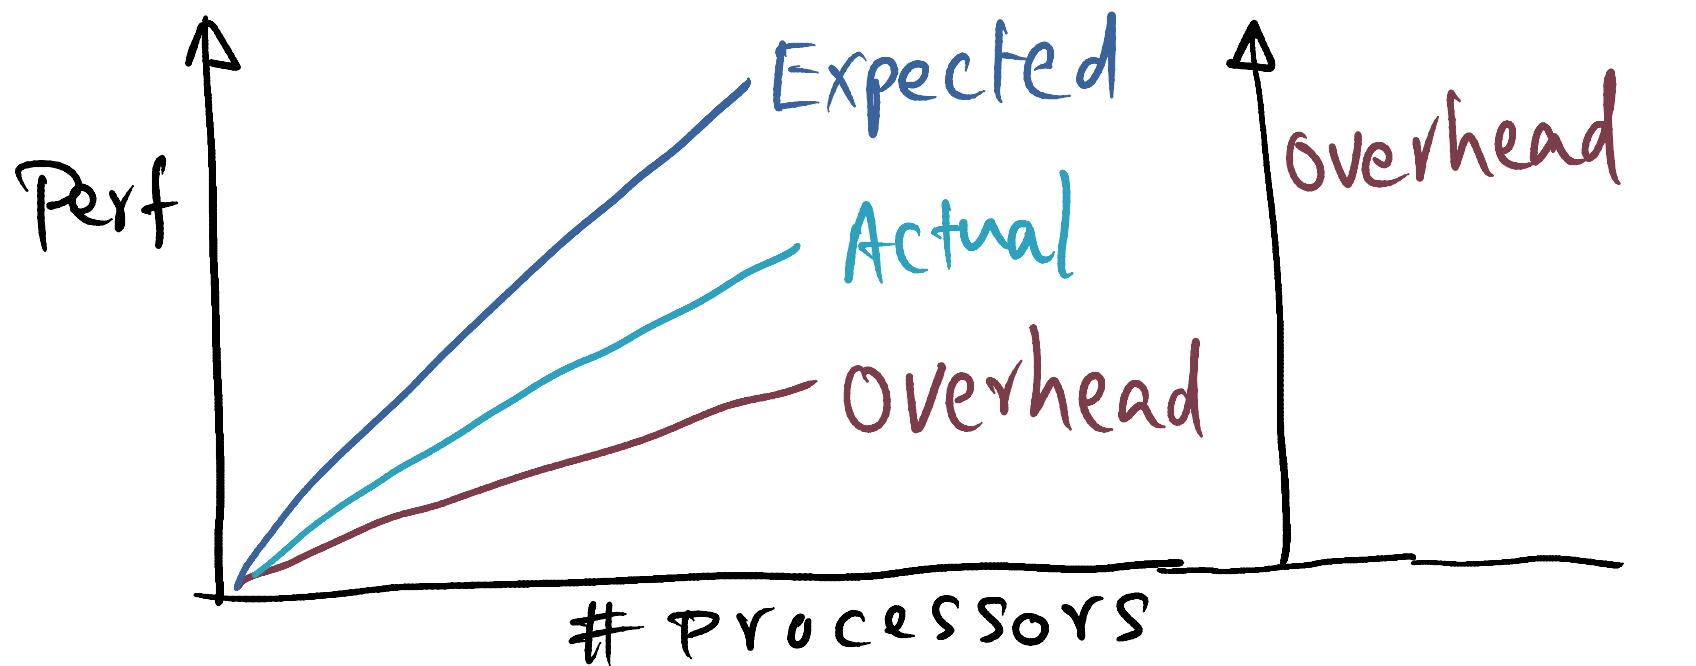
Expectation with more Processors?

Pro:

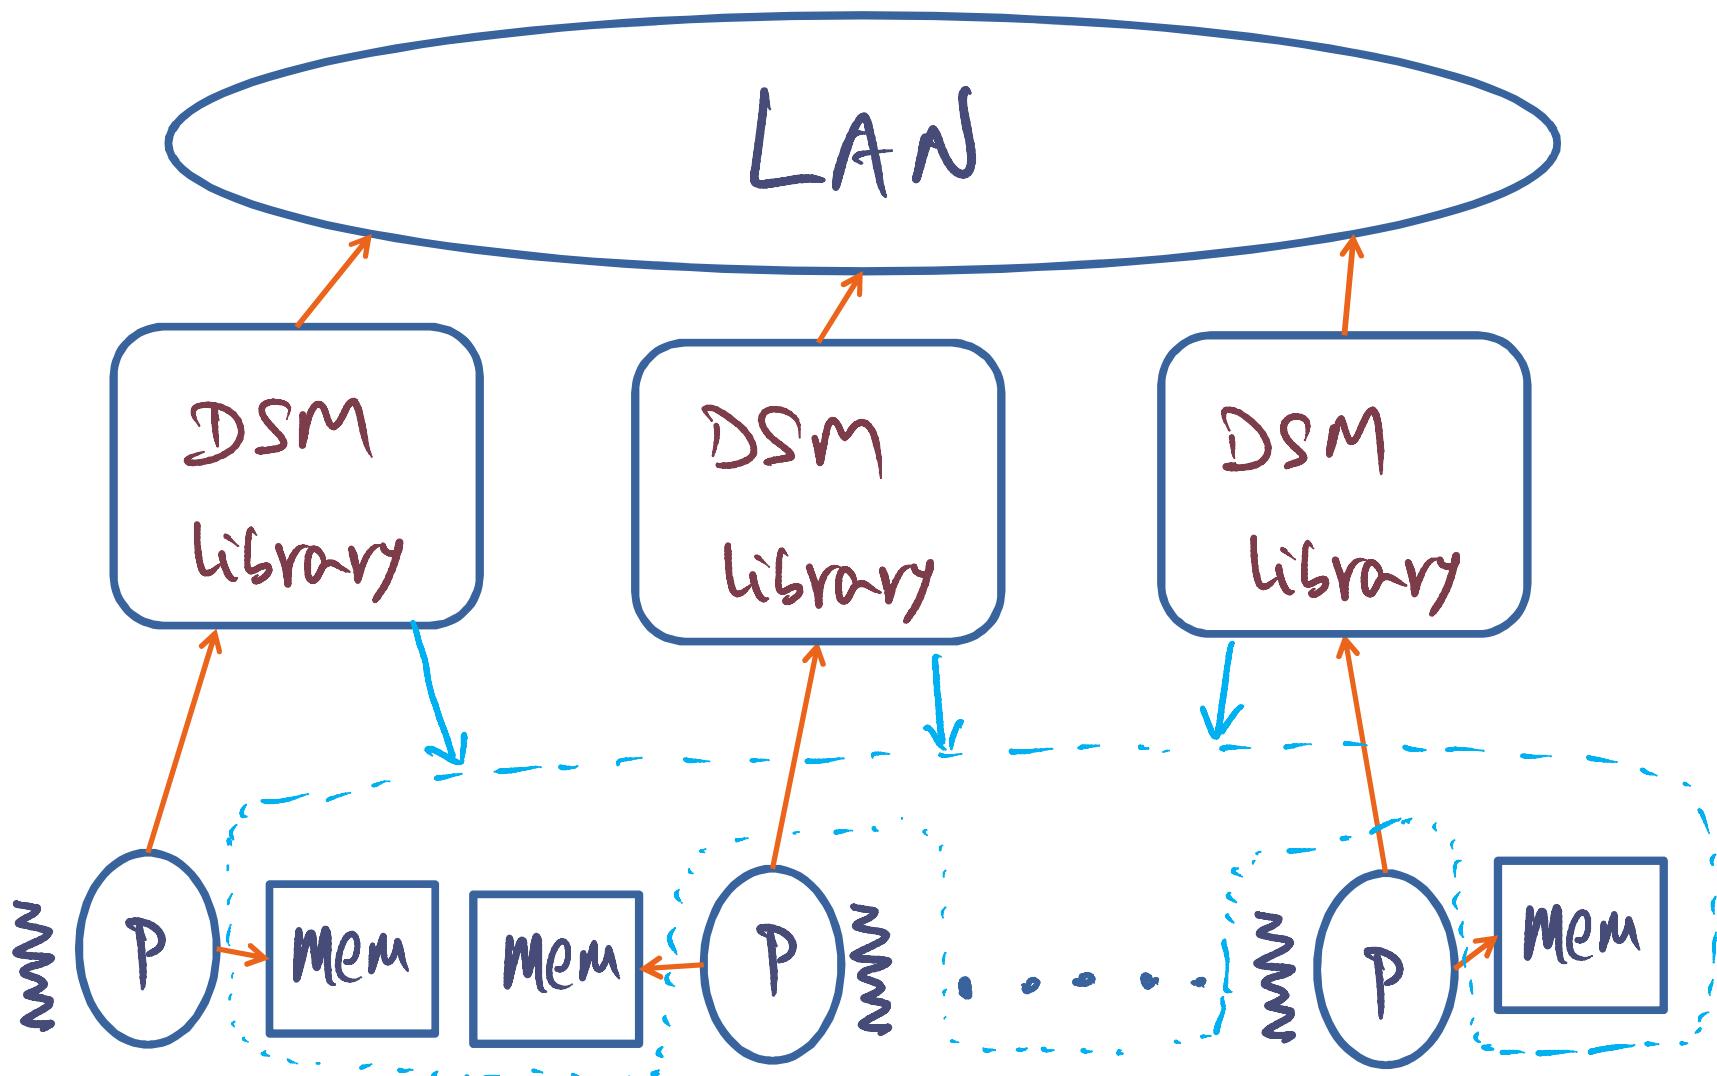
Exploit parallelism

Con:

Increased Overhead



DSM and Speedup



- Speedup is not automatic
 - If sharing is too fine grain no hope of speedup on DSM systems
 - Illusion of shared memory via software, not really physically shared memory => you have to be very careful how much memory you actually share between the threads that are running on different nodes of the cluster

Recall Thacker's quote regarding sharing memory!

- Basic principle
 - Comp:comm ratio has to be high for any hope of speedup
 - What does this mean for shared memory codes?
 - Dynamic data structures can lead to a lot of "implicit" communication across the LAN

Some classes of scientific applications with coarse-grained sharing may be able to achieve speedup with increasing number of processors if carefully constructed.

Key takeaways

- DSM as originally envisioned, i.e., as a “threads package” for a cluster is dead.
- Structured DSM, namely, providing higher level data abstractions for sharing among threads executing on different nodes of the cluster is attractive to reduce the programming pain for the developers of distributed applications on a cluster.

We will discuss one such system called “Persistent Temporal Streams” in a later lesson.