

Two Issues with Content distribution

- * meta server overload
(storing $\langle \text{key}, \text{value} \rangle$)
- * Origin Server overload
(Storing content)

Origin Server overload

Two Solutions

- Web proxy
- not good enough for slashdot effect

Origin Server overload

Two Solutions

- Web proxy
 - not good enough for slashdot effect
- CDN,
 - content mirrored stored geographically
 - user request dynamically re-routed to geo-local mirror

Origin Server overload

Two Solutions

- Web proxy
 - not good enough for slashdot effect
 - CDN,
 - content mirrored stored geographically
 - user request dynamically re-routed to geo-local mirror
- ⇒ pay CDN providers to avoid origin overload

Origin Server overload

Two Solutions

- Web proxy
 - not good enough for slashdot effect
 - CDN,
 - content mirrored stored geographically
 - user request dynamically re-routed to geo-local mirror
- ⇒ pay CDN providers to avoid origin overload

How to democratize content distribution?

Origin Server overload

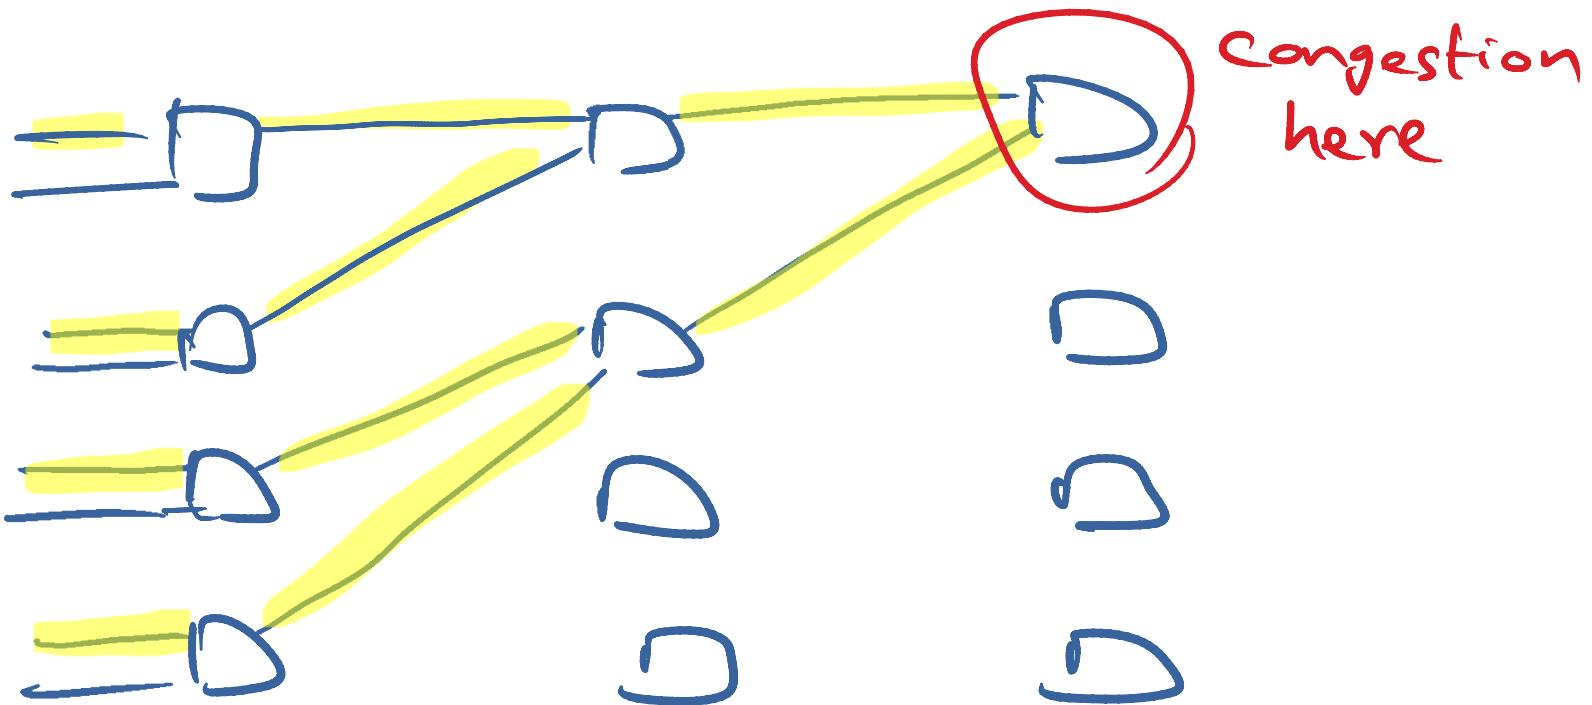
Two Solutions

- Web proxy
 - not good enough for slashdot effect
 - CDN,
 - content mirrored stored geographically
 - user request dynamically re-routed to geo-local mirror
- ⇒ pay CDN providers to avoid origin overload

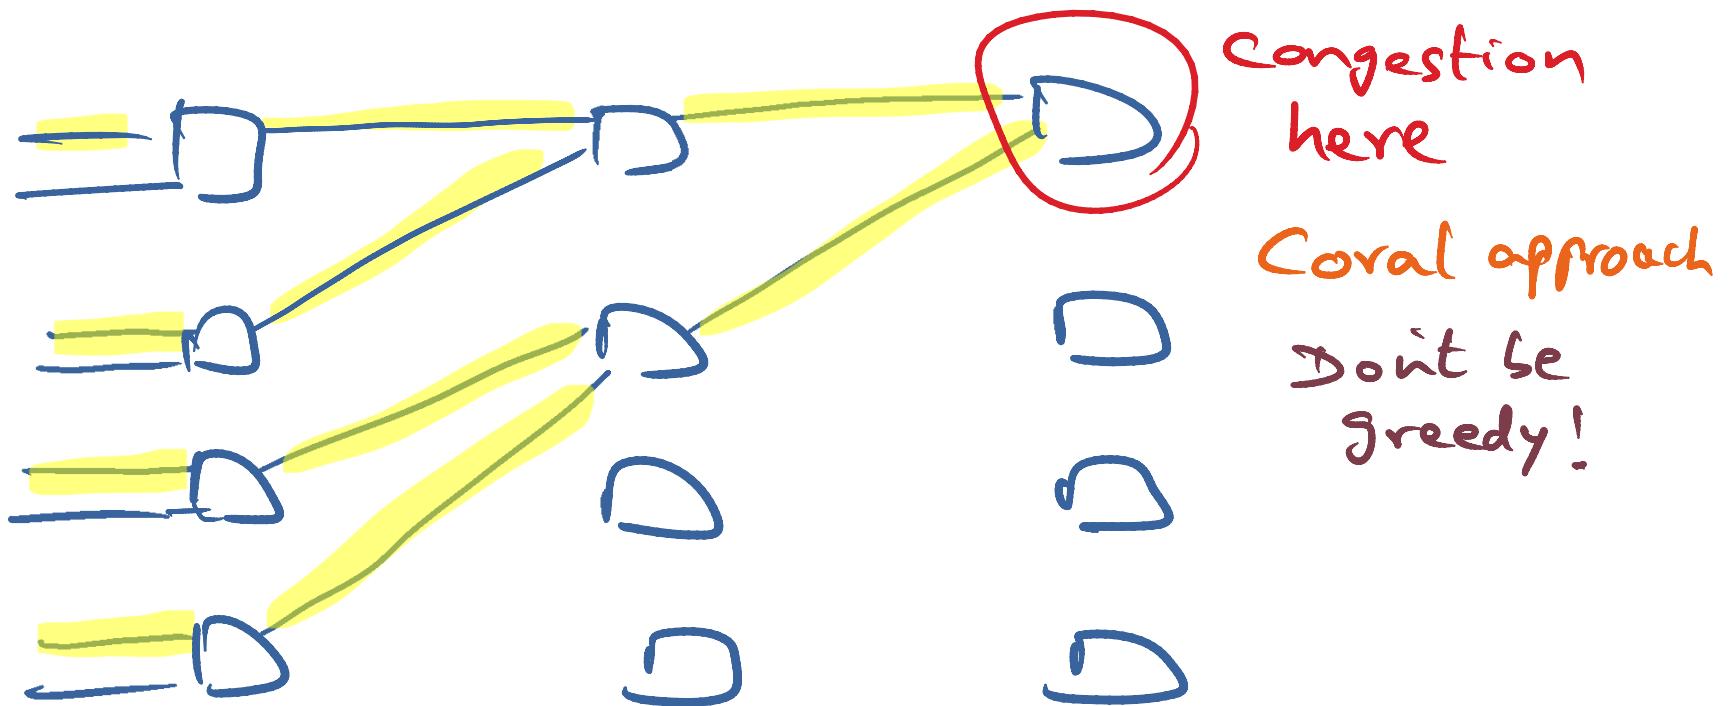
How to democratize content distribution?

Enter Coral System

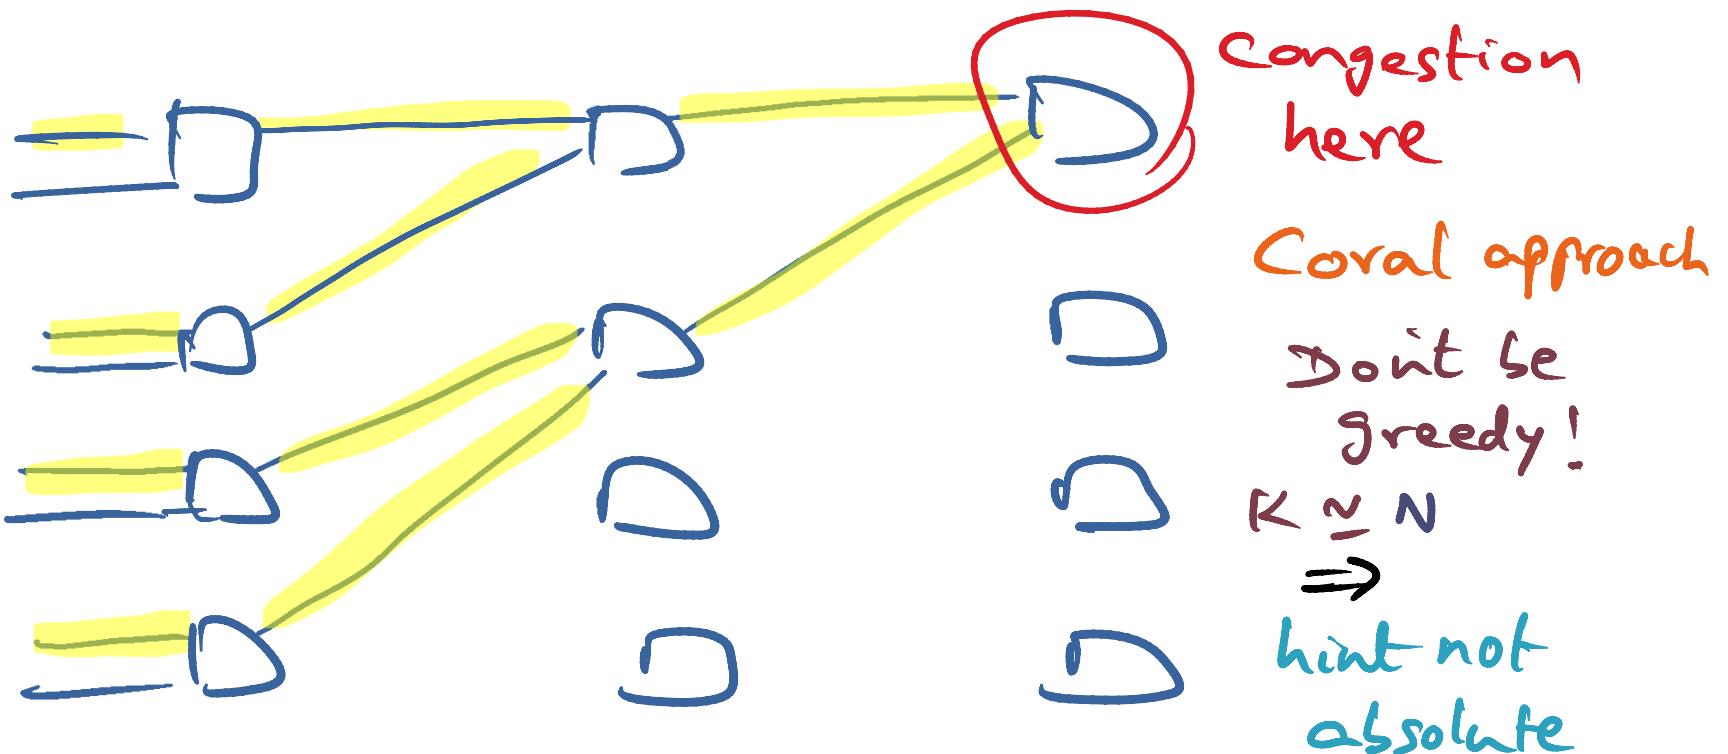
Greedy Approach leads to "tree saturation"



Greedy Approach leads to "tree saturation"



Greedy Approach leads to "tree saturation"



Coral DHT

- Get/put satisfied by nodes different from $\langle N \cong \text{key} \rangle \Rightarrow$ Sloppy DHT

Rationale

- avoid tree saturation
- spread metadata overload

Coral DHT

- Get/put satisfied by nodes different from $\langle N \cong \text{key} \rangle \Rightarrow$ Sloppy DHT

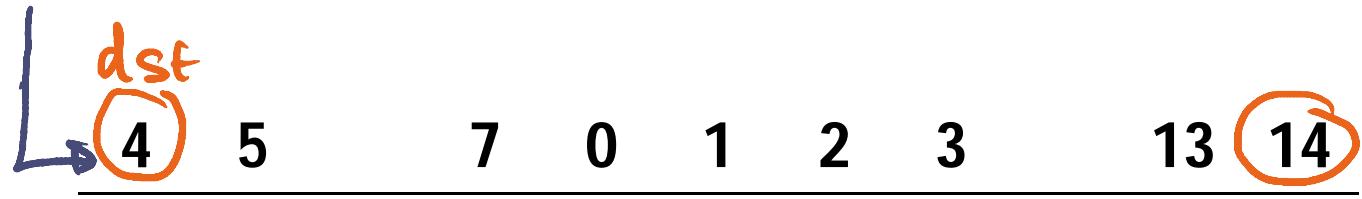
Rationale

- avoid tree saturation
- spread metadata overload

How?

- novel key-based routing
 - basis
 - XOR distance between src & dst
 $N_{src} \oplus N_{dst}$ (e.g., if $xor f = 10$)

node-ids



	V			V		V	V		V	V
--	---	--	--	---	--	---	---	--	---	---

node-ids

dst
→ 4

Key Based Routing

State of routing table @ Src

7 0 1 2 3

13 14

entries \Rightarrow reachable

	V			V		V	V		V	V
--	---	--	--	---	--	---	---	--	---	---

node-ids



Key Based Routing

State of routing table @ src

7 0 1 2 3

13 14

entries \Rightarrow reachable

	V			V		V	V		V	V
--	---	--	--	---	--	---	---	--	---	---

Intuition in greedy approach

- get as close to the desired dst in node-id namespace

node-ids



entries \Rightarrow reachable

	V			V		V	V		V	V					
--	---	--	--	---	--	---	---	--	---	---	--	--	--	--	--

Intuition in greedy approach

- get as close to the desired dst in node-id namespace
- he may know route to get me to desired dst

node-ids



entries \Rightarrow reachable

	V			V		V	V		V	V
--	---	--	--	---	--	---	---	--	---	---

Intuition in greedy approach

- get as close to the desired dst in node-id namespace
- he may know route to get me to desired dst

Objective?

- reach dst with fewest number of hops

node-ids



entries \Rightarrow reachable

	V			V		V	V		V	V
--	---	--	--	---	--	---	---	--	---	---

Intuition in greedy approach

- get as close to the desired dst in node-id namespace
- he may know route to get me to desired dst

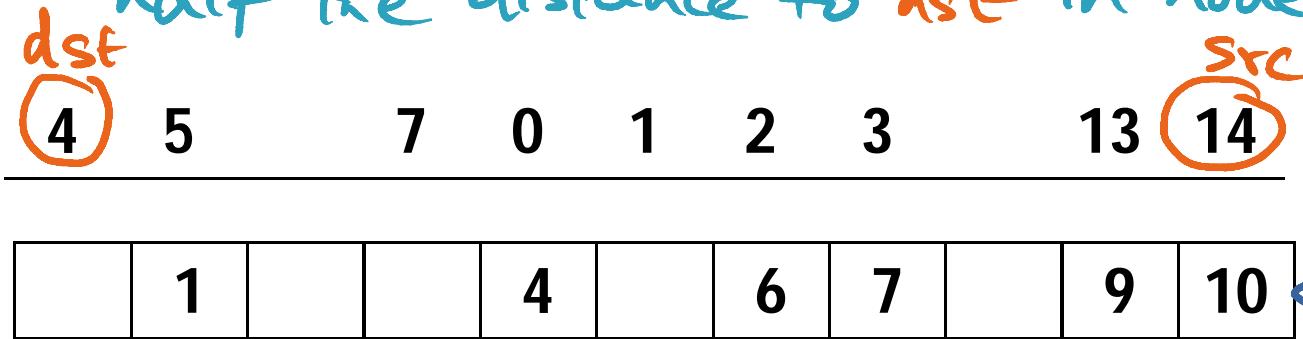
Objective?

- reach dst with fewest number of hops
- \Rightarrow "me first" approach leads to congestion

Conal Key-based routing

- Each hop go to some node that is

half the distance to dst in node-id namespace



node-id XOR 4

XOR distance
from dst

Conal Key-based routing

- Each hop go to some node that is

half the distance to dst in node-id namespace



node-id XOR 4

XOR distance
from dst

XOR distance from src to dst = 10

Conal Key-based routing

- Each hop go to some node that is

half the distance to dst in node-id namespace



node-id XOR 4

XOR distance
from dst

XOR distance from src to dst = 10

1st hop go to $\frac{10}{2} \Rightarrow 5$ distant

Conal Key-based routing

- Each hop go to some node that is

half the distance to dst in node-id namespace



node-id XOR 4

↓
XOR distance
from dst

XOR distance from src to dst = 10

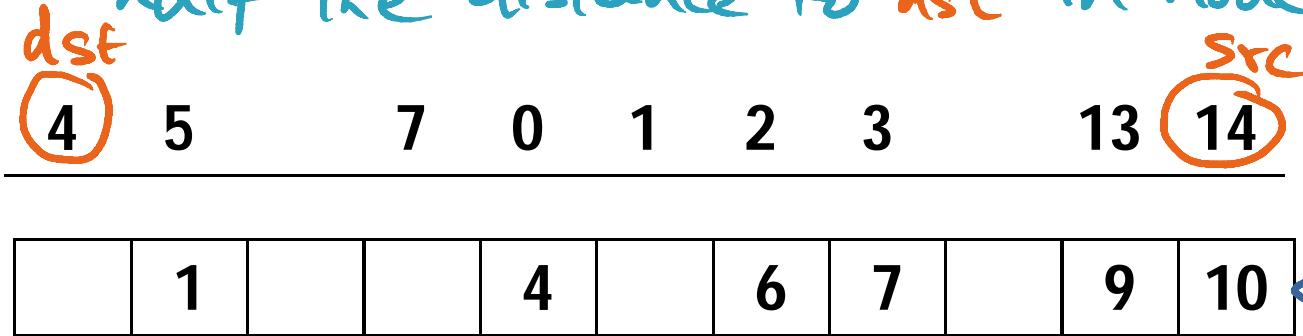
1st hop go to $\frac{10}{2} \Rightarrow 5$ distant

2nd hop go to $\frac{5}{2} \Rightarrow 2$ distant

Coral Key-based routing

- Each hop go to some node that is

half the distance to dst in node-id namespace



node-id XOR 4

↓
XOR distance
from dst

XOR distance from src to dst = 10

1st hop go to $\frac{10}{2} \Rightarrow 5$ distant

2nd hop go to $\frac{5}{2} \Rightarrow 2$ distant

3rd hop go to $\frac{2}{2} \Rightarrow 1$ distant
... finally home !!

node-ids

Key-based routing in coral



entries \Rightarrow reachable

	1			4		6	7		9	10
--	---	--	--	---	--	---	---	--	---	----

node-id XOR 4

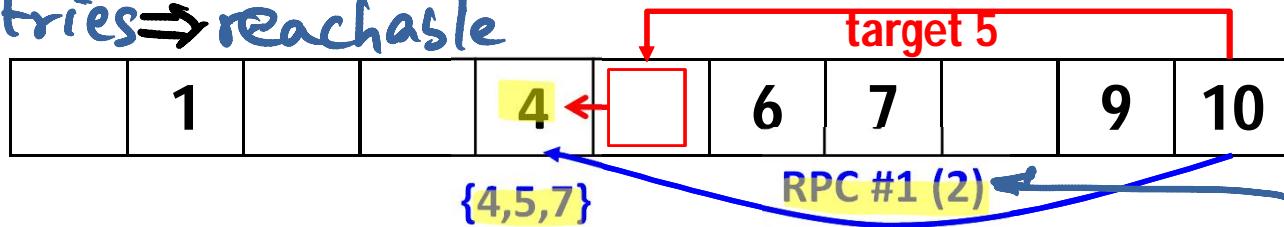
XOR distance
from dst

node-ids

Key-based routing in coral



entries \Rightarrow reachable



XOR distance
from dst

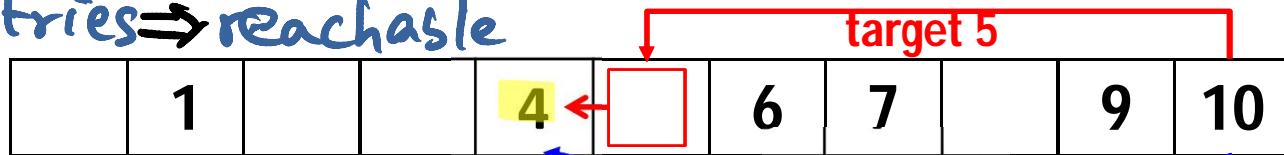
looking for
distance 2

node-ids

Key-based routing in coral



entries \Rightarrow reachable



XOR distance
from dst

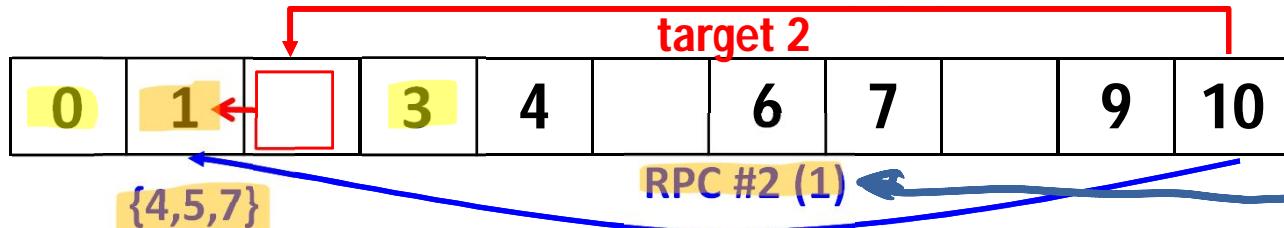
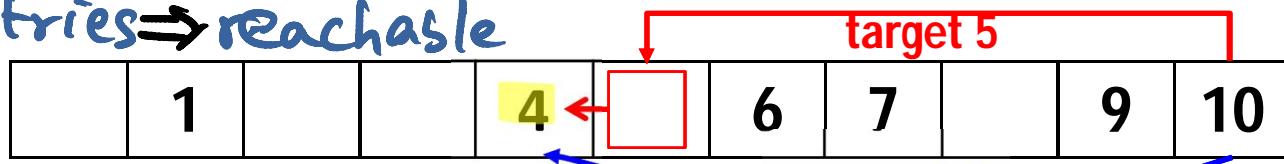
looking for
distance 2

node-ids

Key-based routing in coral



entries \Rightarrow reachable



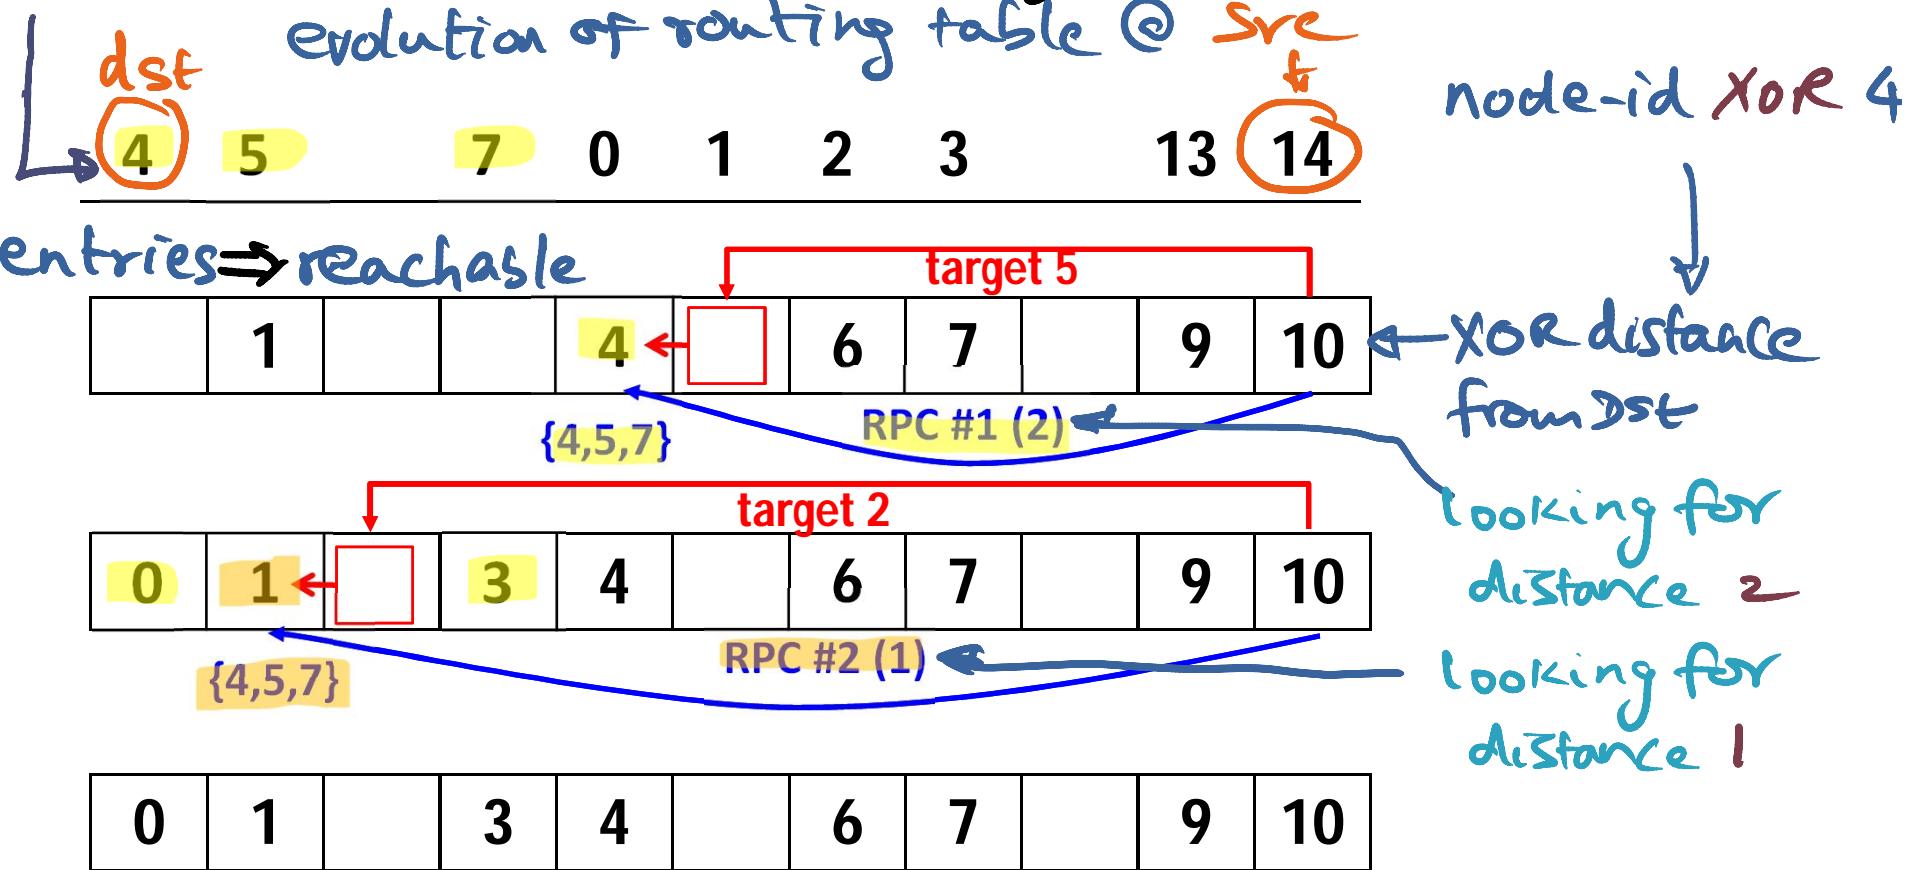
XOR distance
from dst

looking for
distance 2

looking for
distance 1

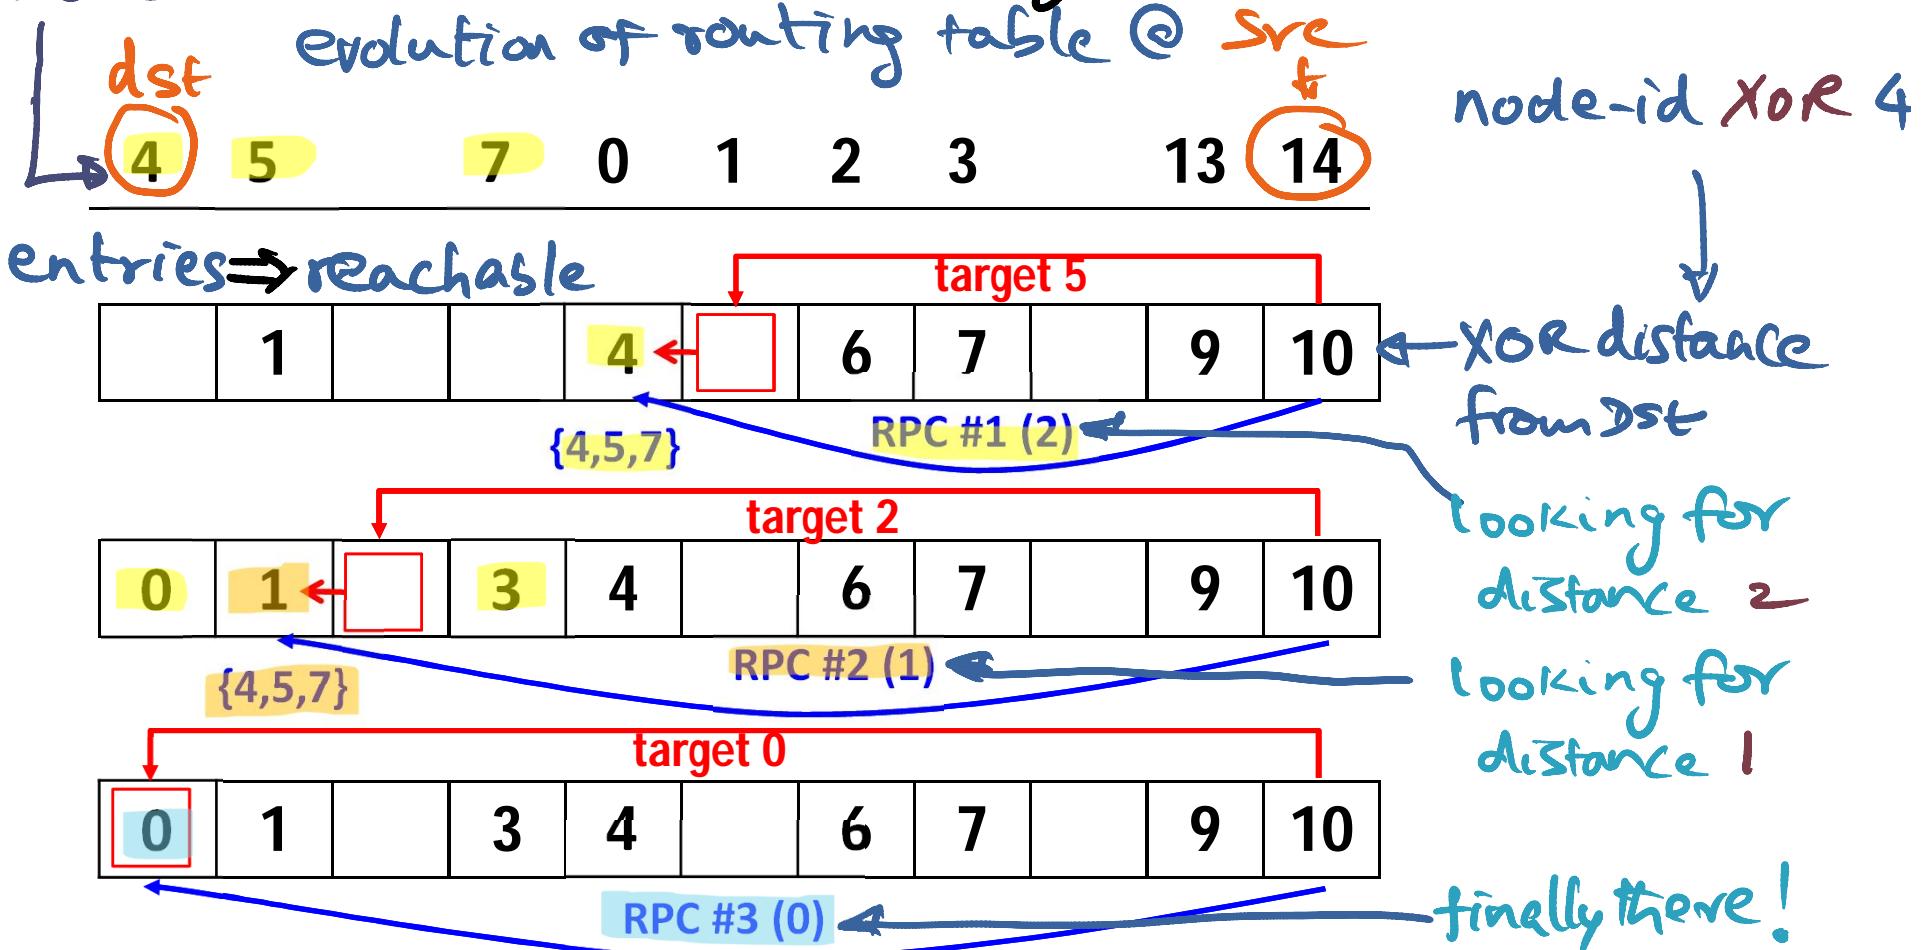
node-ids

Key-based routing in coral



node-ids

Key-based routing in coral



Coral Sloppy DHT

Coral Primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

Coral Sloppy DHT

Coral Primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

* put <key,value> in an "appropriate" node

Coral Sloppy DHT

Coral primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

* put <key,value> in an "appropriate" node



Coral Sloppy DHT

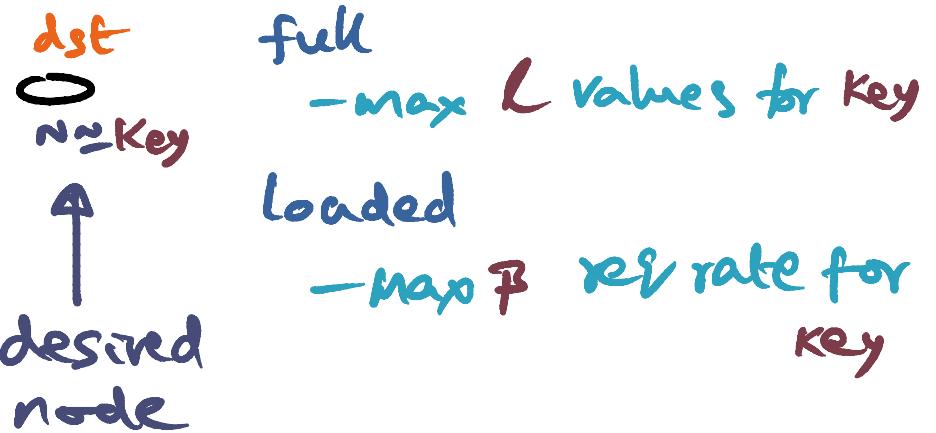
Coral primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

* put <key,value> in an "appropriate" node



Coral Sloppy DHT

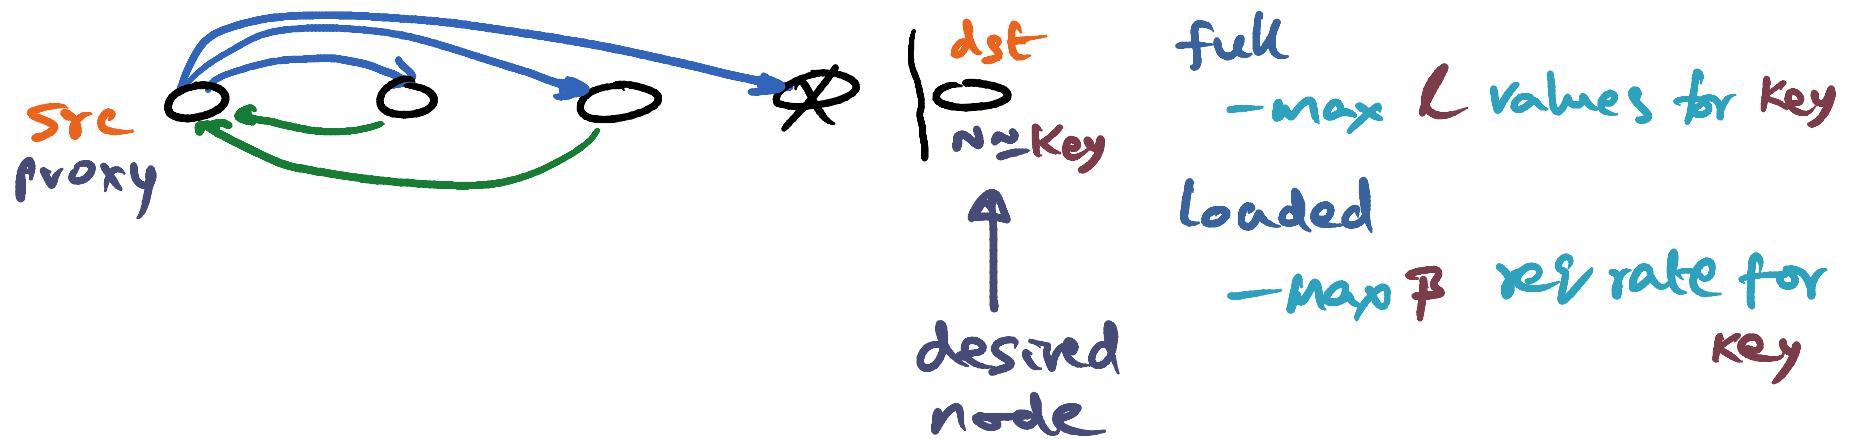
Coral primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

* put <key,value> in an "appropriate" node



Coral Sloppy DHT

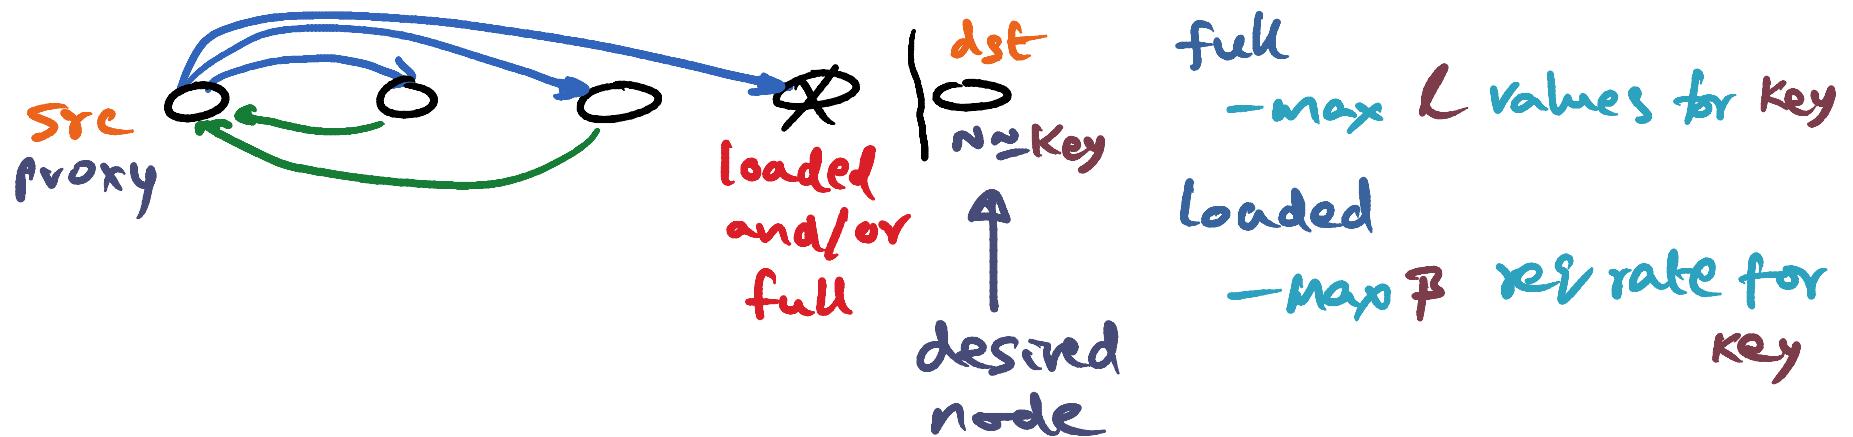
Coral primitives

- put (key, value)

↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

* put <key,value> in an "appropriate" node



Coral Sloppy DHT

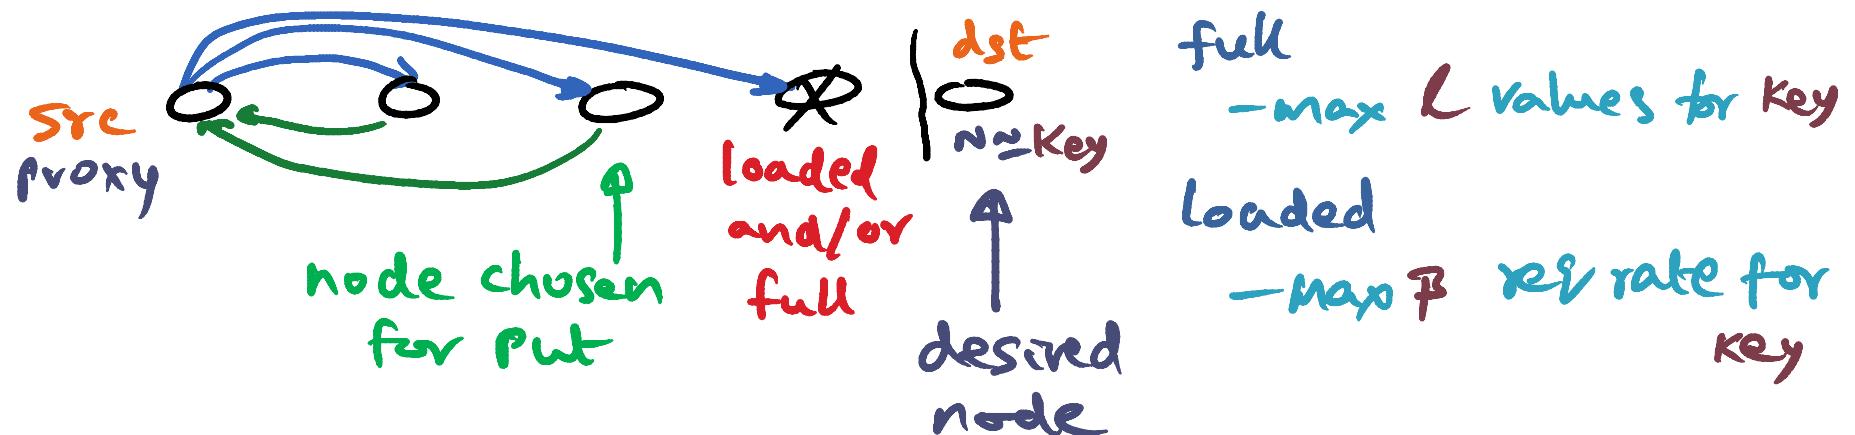
Coral primitives

- put (*key, value*)

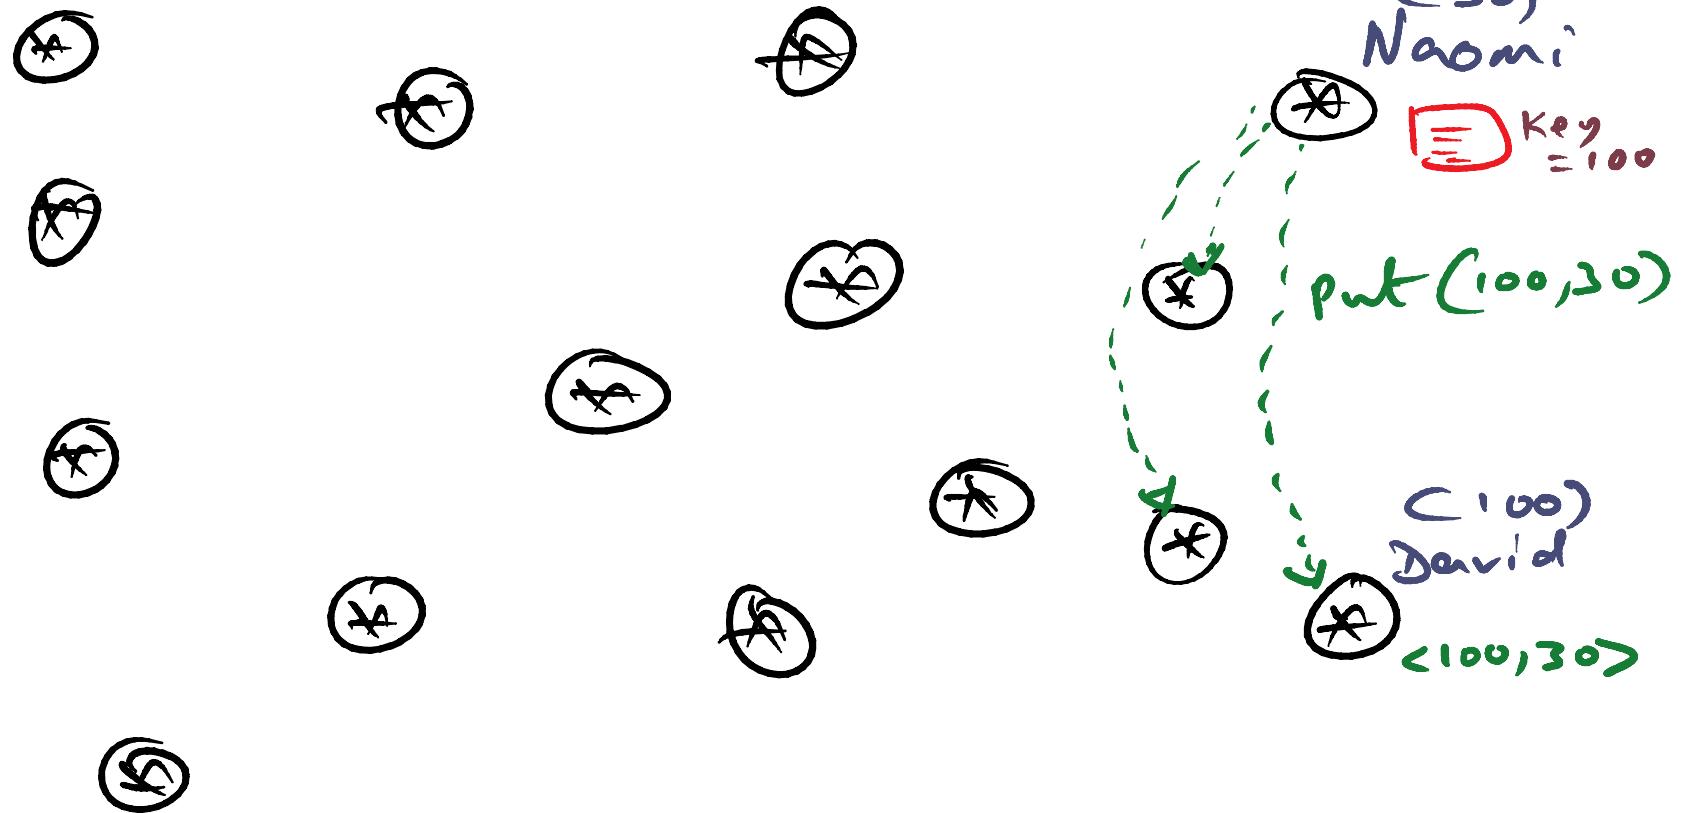
↑
node-id of proxy with content for key

* announcing willingness to serve as proxy

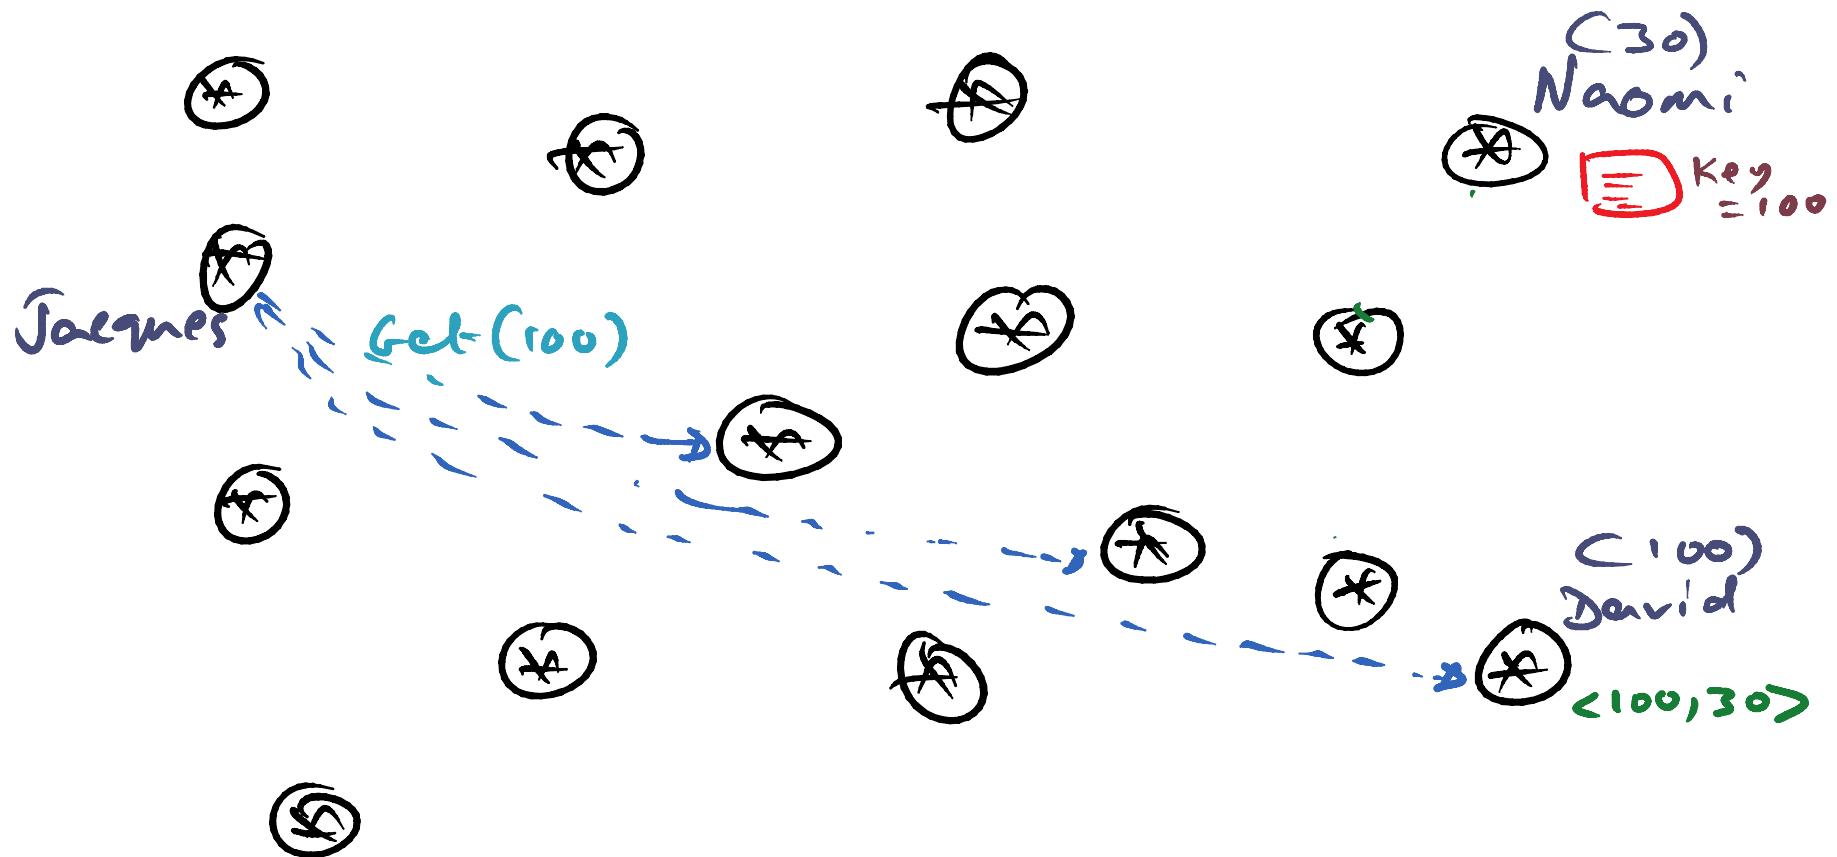
* put <key,value> in an "appropriate" node



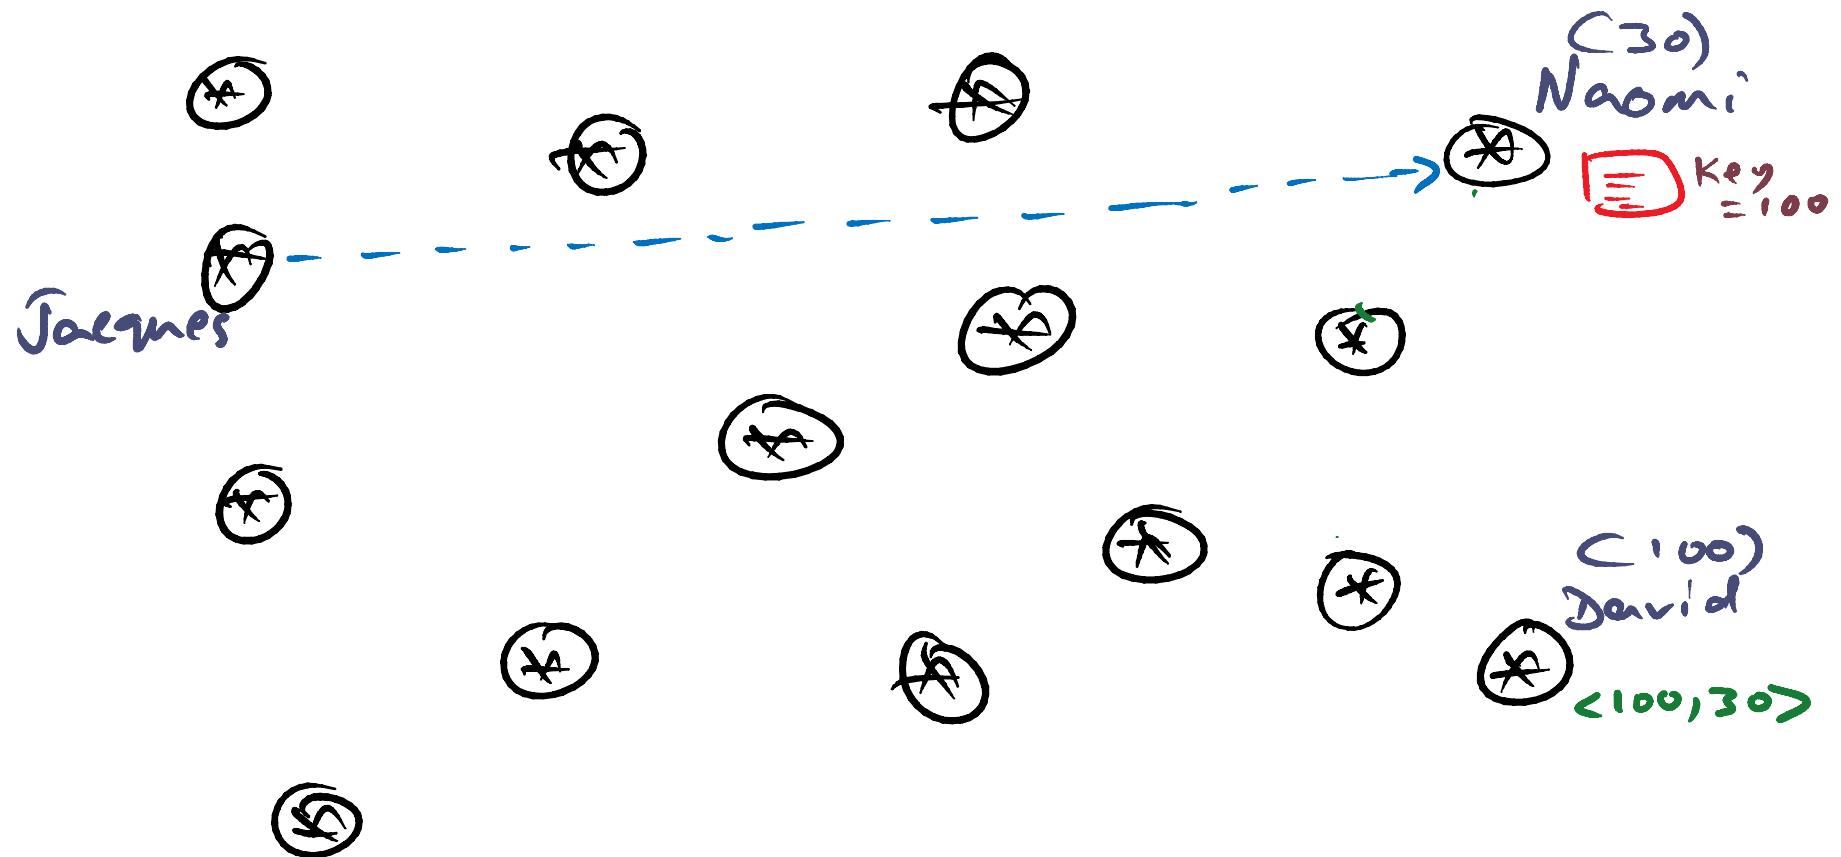
Coral in action



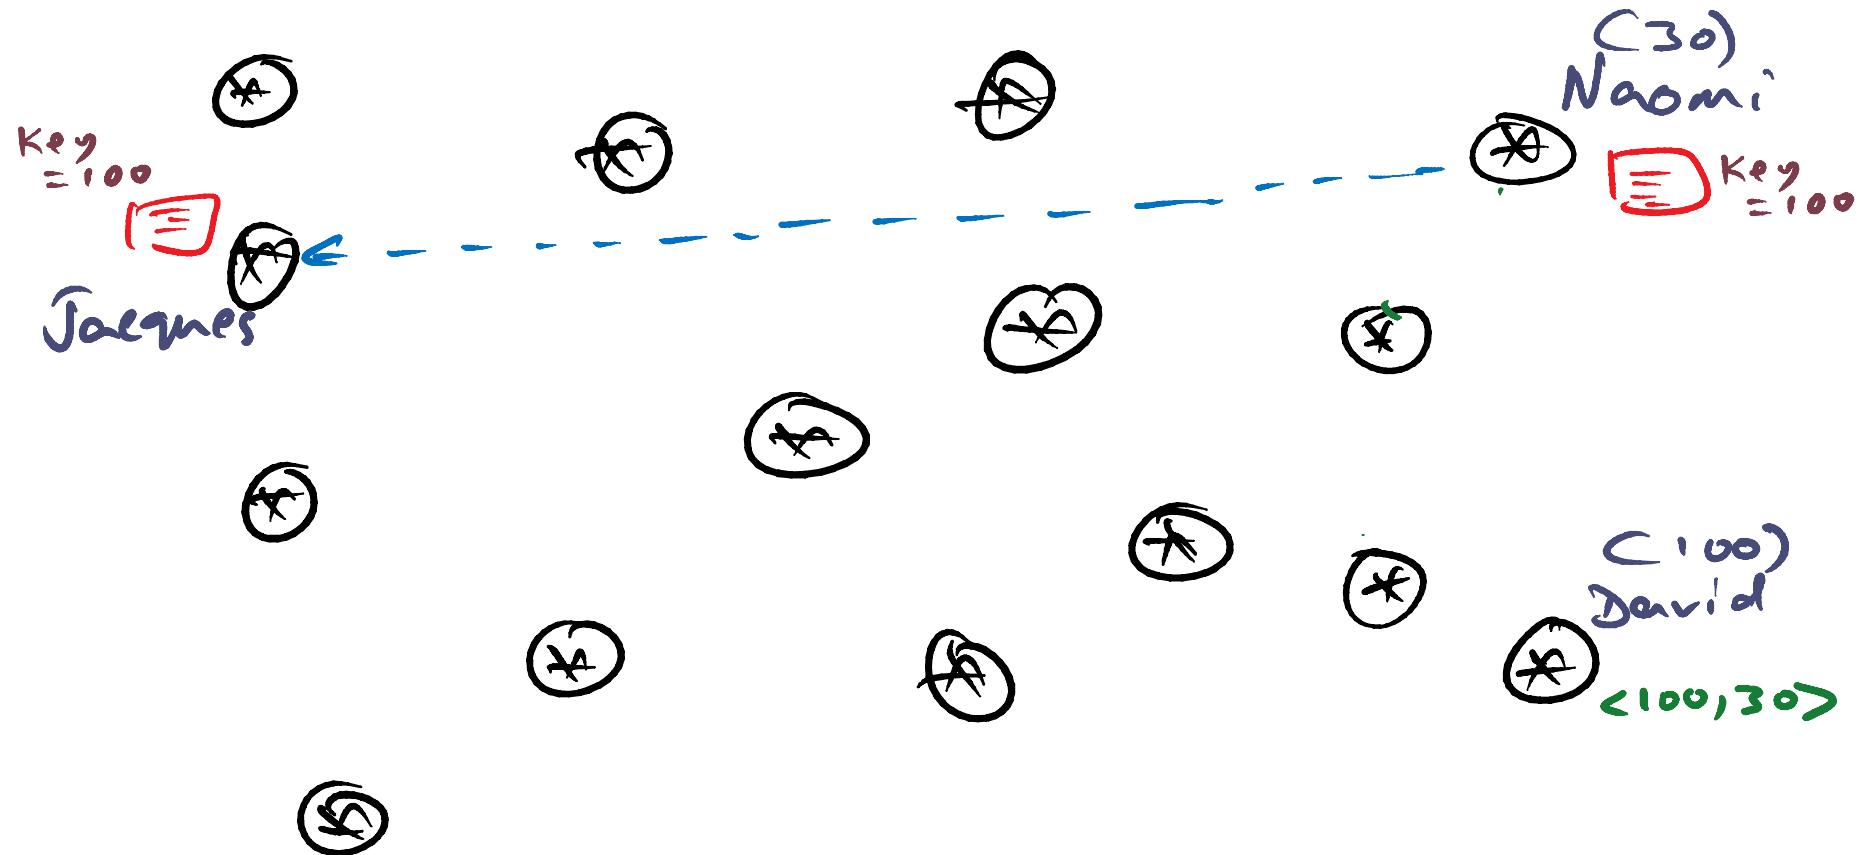
Coral in action



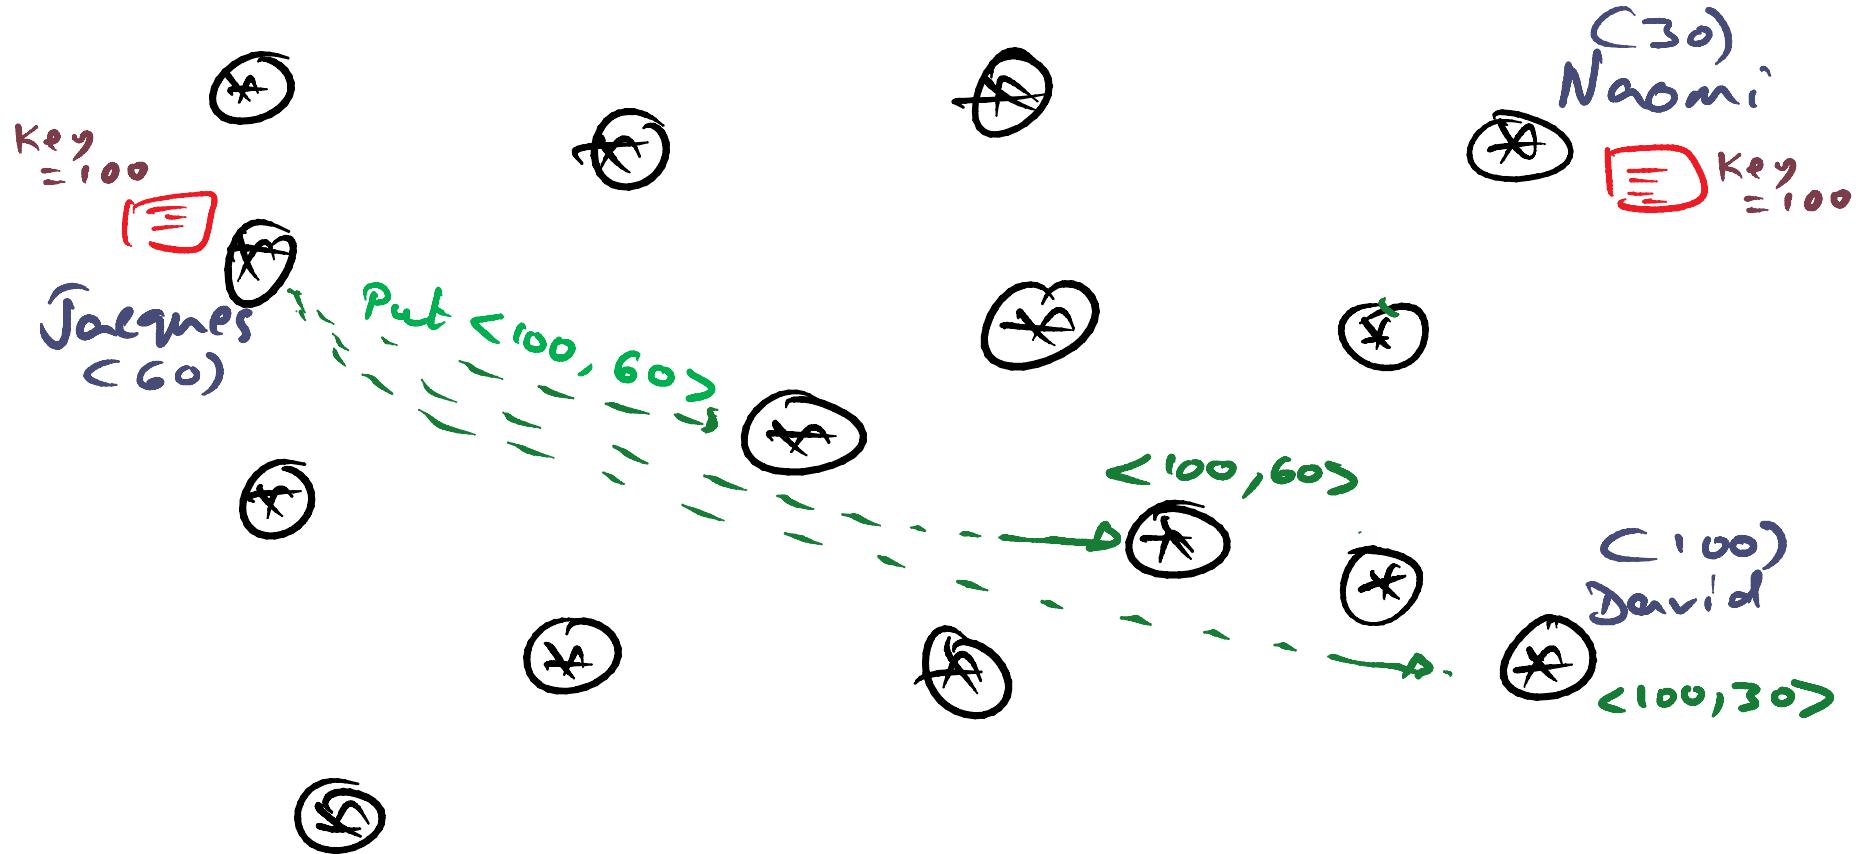
Coral in action



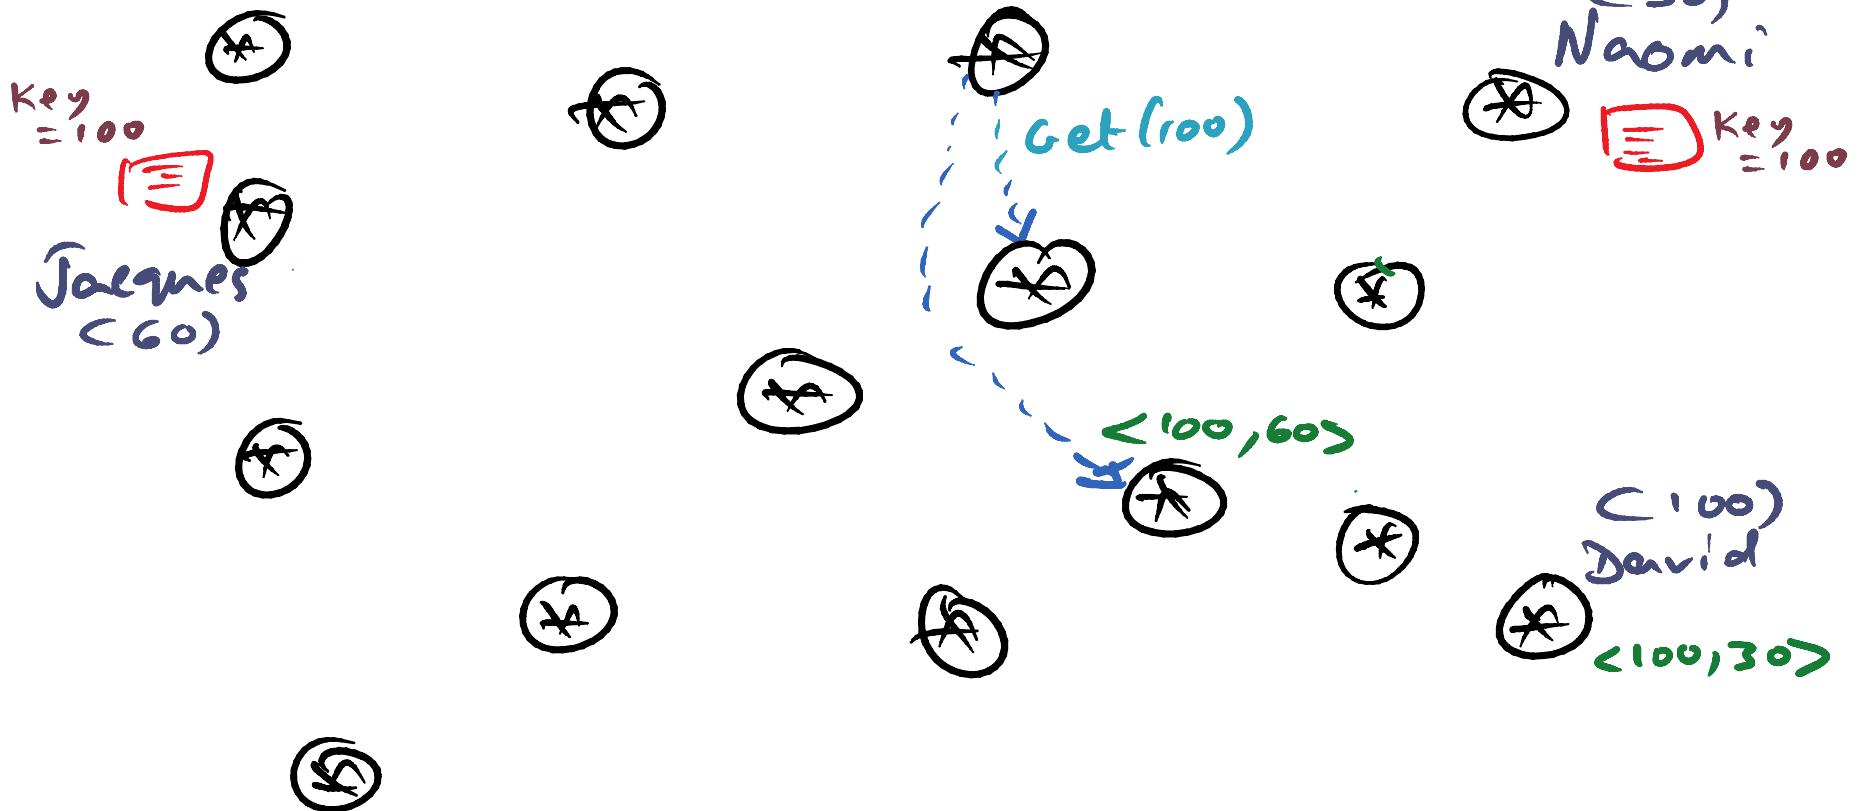
Coral in action



Coral in action



Coral in action



Coral in action



Coral in action



Key takeaways

- How to avoid overloading the meta-server and origin server
- Coral system democratize content generation, storage, and distribution
- Of course, commercial CDNs like Akamai do not operate this way. They are in it for the money!