

Today

Distributed System Basics

✓ Lamport's clock

✓ Efficient Communication Software

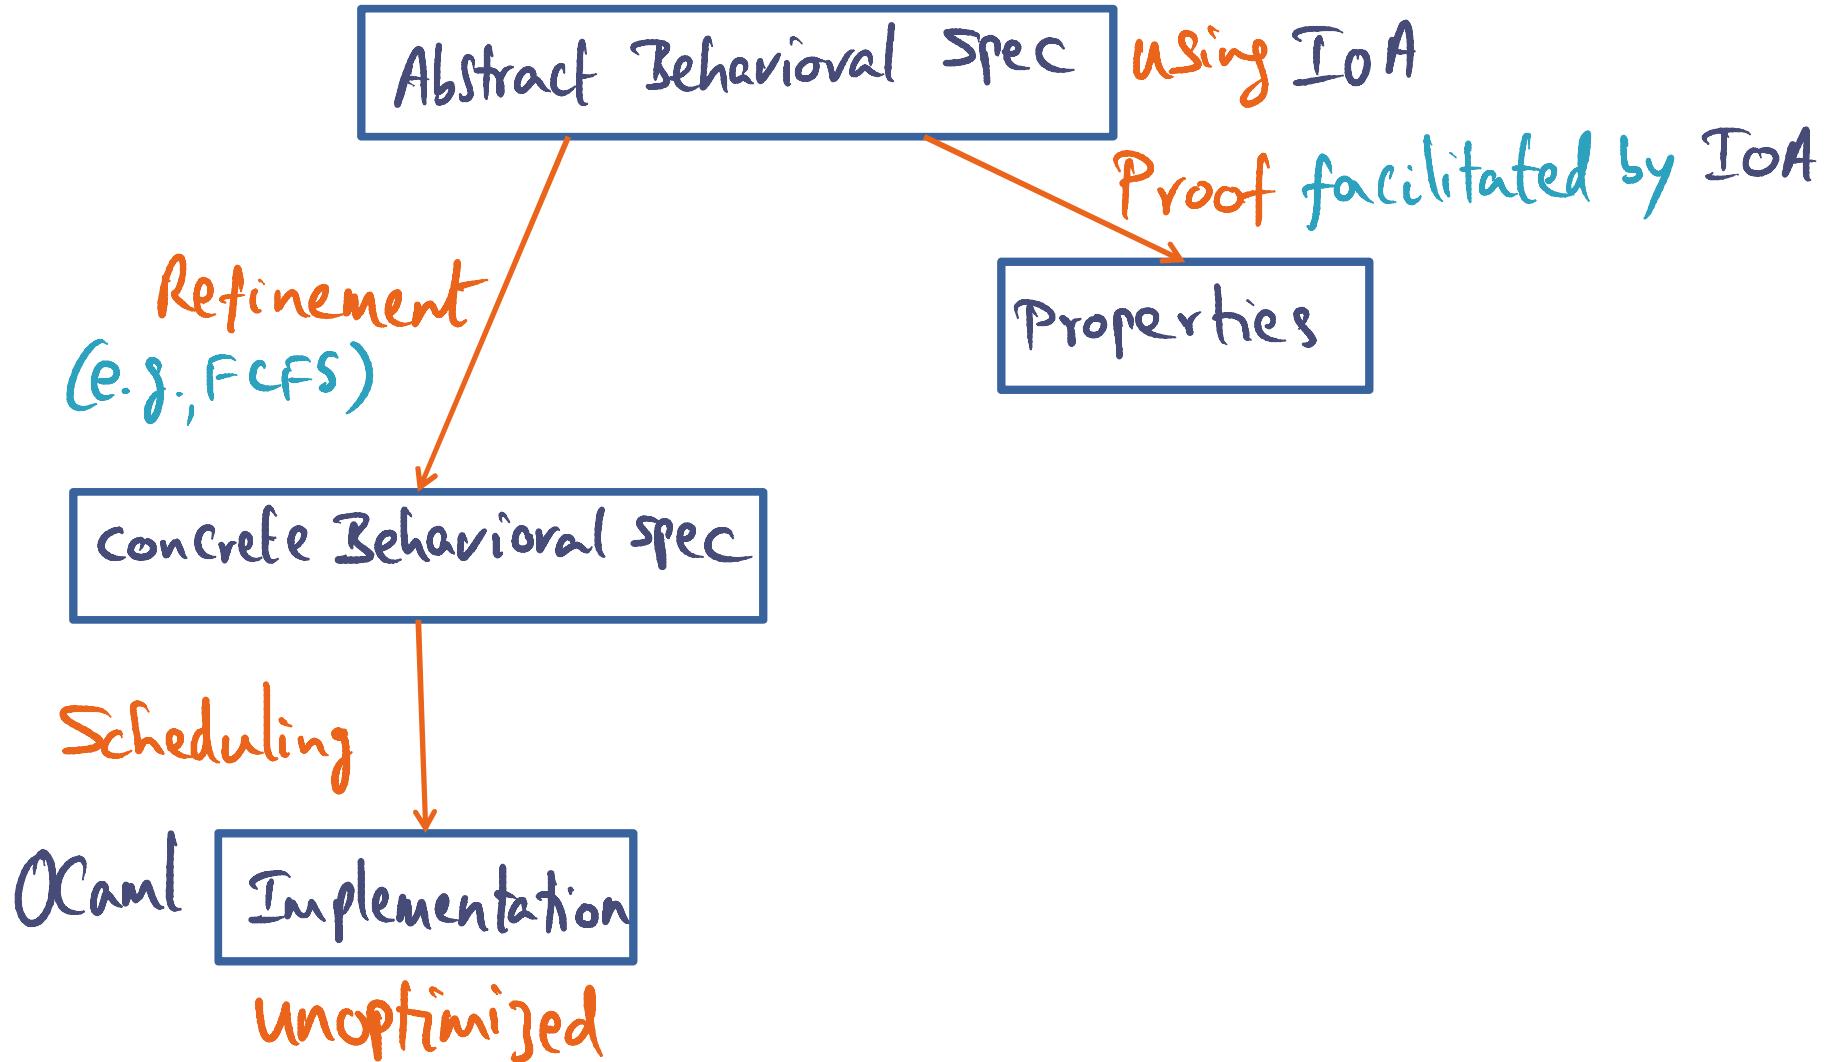
The kernel  Application interface to the Kernel
Levy

~~we'll~~ { End-to-End QoS via Active Networks

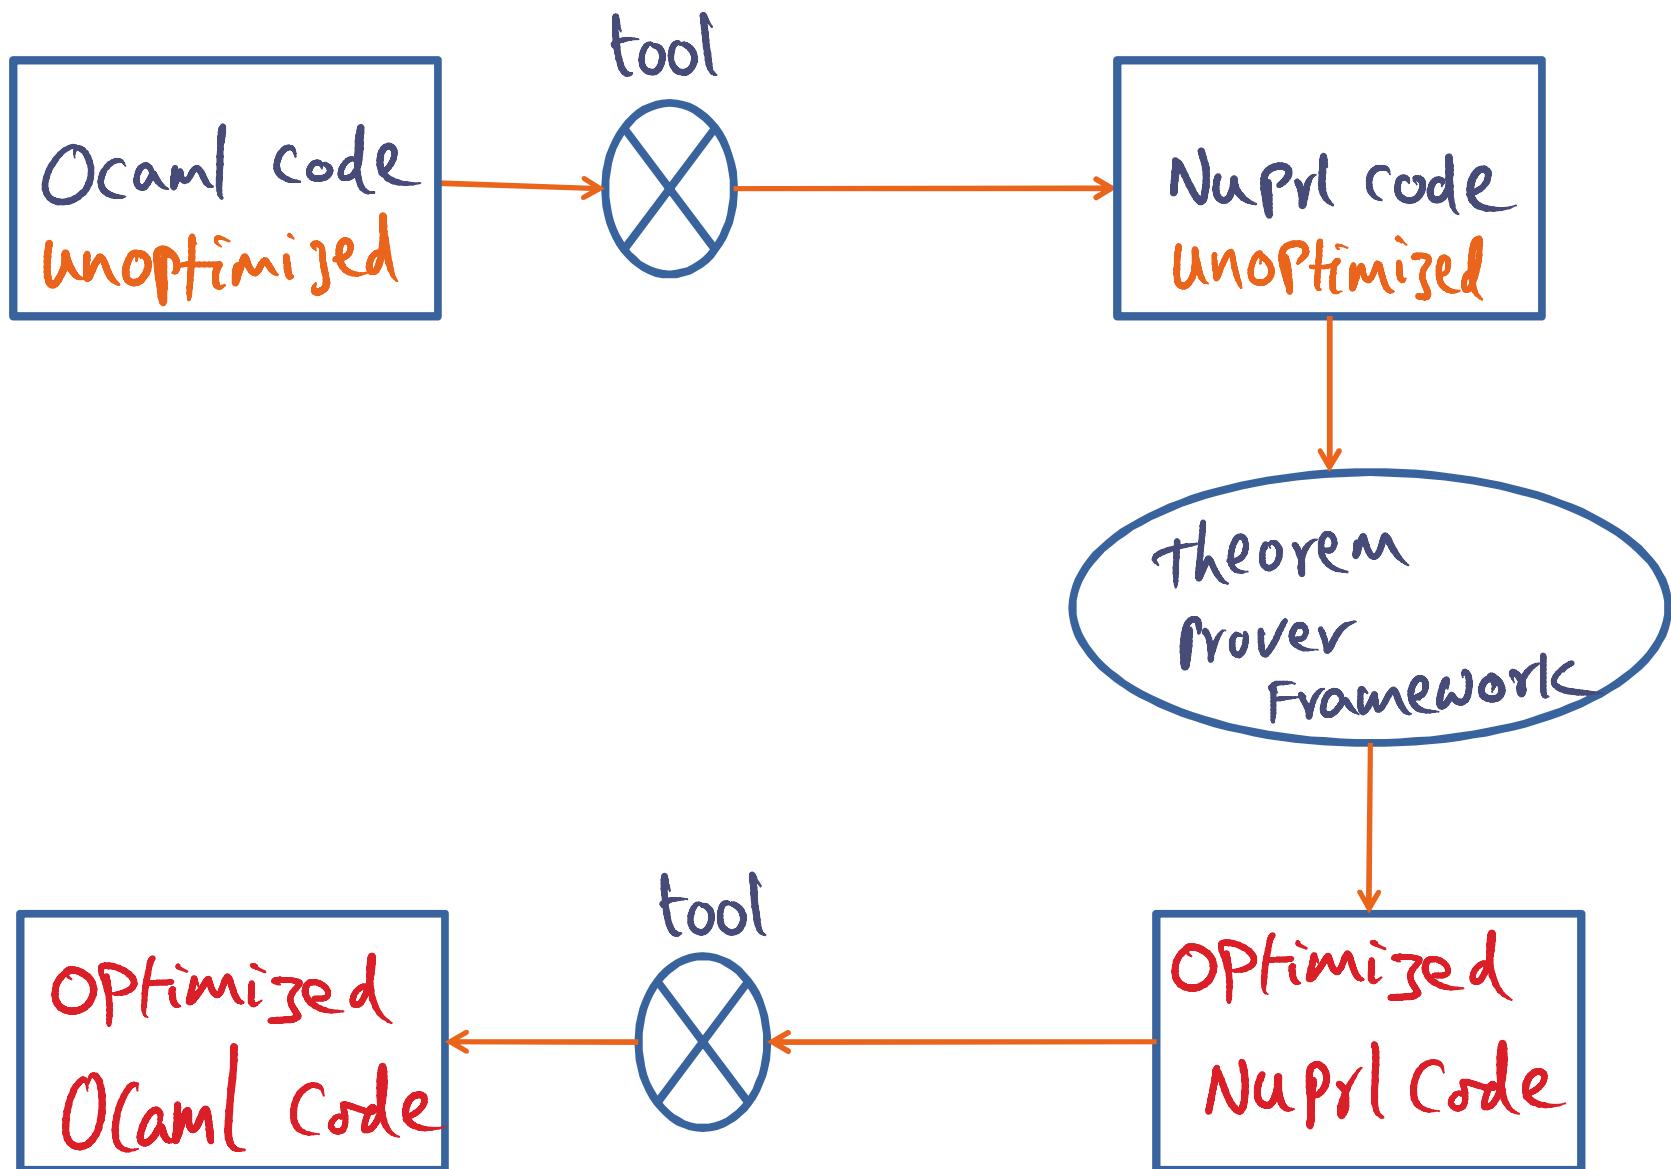
⇒ Synthesizing Network Protocol Stacks

- Synthesizing Real-time
- Spring Kernel (Disp. Objects)
Lesson 6

Digging Deeper – From Spec to Implementation

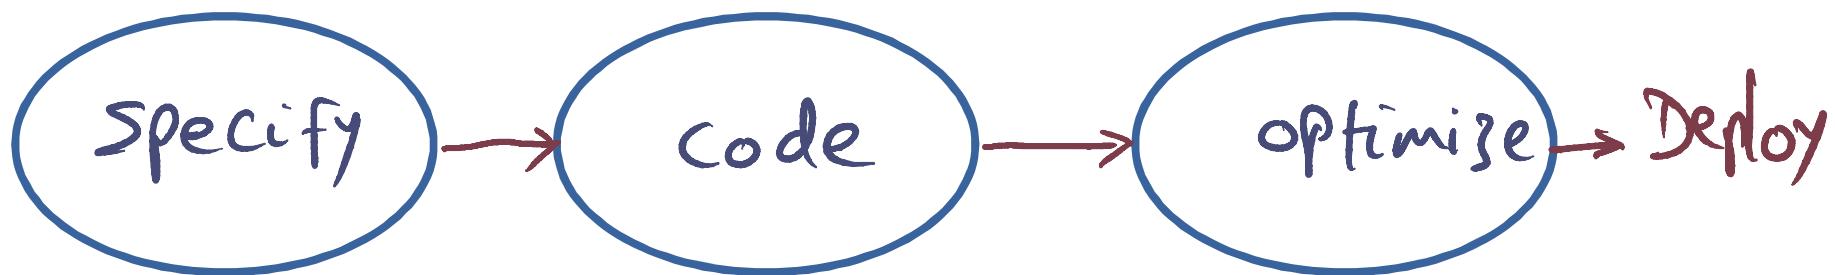


Digging Deeper – From Implementation to optimization



Ensemble - Big Picture

Design Cycle



IOA

- C-like syntax
- composition operator

Ocaml

- object oriented
- efficient code
similar to C
- nice complement
to IOA

Nuprl

- optimization of
Ocaml code
- output verified to be
functionally equivalent

Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

How to synthesize the stack from concrete Spec?

Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

How to synthesize the stack from concrete Spec?

Getting to an unoptimized Ocaml implementation

- Ensemble Suite of micro protocols

Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

How to synthesize the stack from concrete Spec?

Getting to an unoptimized Ocaml implementation

- Ensemble Suite of micro protocols
 - * flow control, sliding window, encryption, scatter/gather, etc.

Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

How to synthesize the stack from concrete Spec?

Getting to an unoptimized Ocaml implementation

- Ensemble Suite of micro protocols
 - * flow control, sliding window, encryption, scatter/gather, etc.
 - * well defined interfaces allowing composition
 - * facilitates component based design



Putting this methodology to work

Start with IOA Spec

- Abstract spec
- Concrete Spec

How to synthesize the stack from concrete Spec?

Getting to an unoptimized Ocaml implementation

- Ensemble Suite of micro protocols
 - * flow control, sliding window, encryption, scatter/gather, etc.
 - * well defined interfaces allowing composition
 - * facilitates component based design



Recall original goal: mimic VLSI design

How to optimize the Protocol stack ?

How to optimize the Protocol stack ?

Layering Could lead to inefficiencies

How to optimize the Protocol stack ?

Layering Could lead to inefficiencies

- Analogy to VLSI component-based design breaks down for software

How to optimize the Protocol stack ?

Layering Could lead to inefficiencies

- Analogy to VLSI component-based design breaks down for software

Several sources of optimization possible

- explicit memory management instead of implicit GC
- Avoid marshaling/unmarshaling across layers
- Buffering in parallel with transmission
- Header compression
- locality enhancement for common code sequences

How to optimize the Protocol stack ?

Layering Could lead to inefficiencies

- Analogy to VLSI component-based design breaks down for software

Several sources of optimization possible

- explicit memory management instead of implicit GC
- Avoid marshaling/unmarshaling across layers
- Buffering in parallel with transmission
- Header compression
- locality enhancement for common code sequences

But how to automate the process ?

NuPrl to the rescue

Static opt

- NuPrl expert
- plus OCaml expert
- layer by layer

NuPrl to the rescue

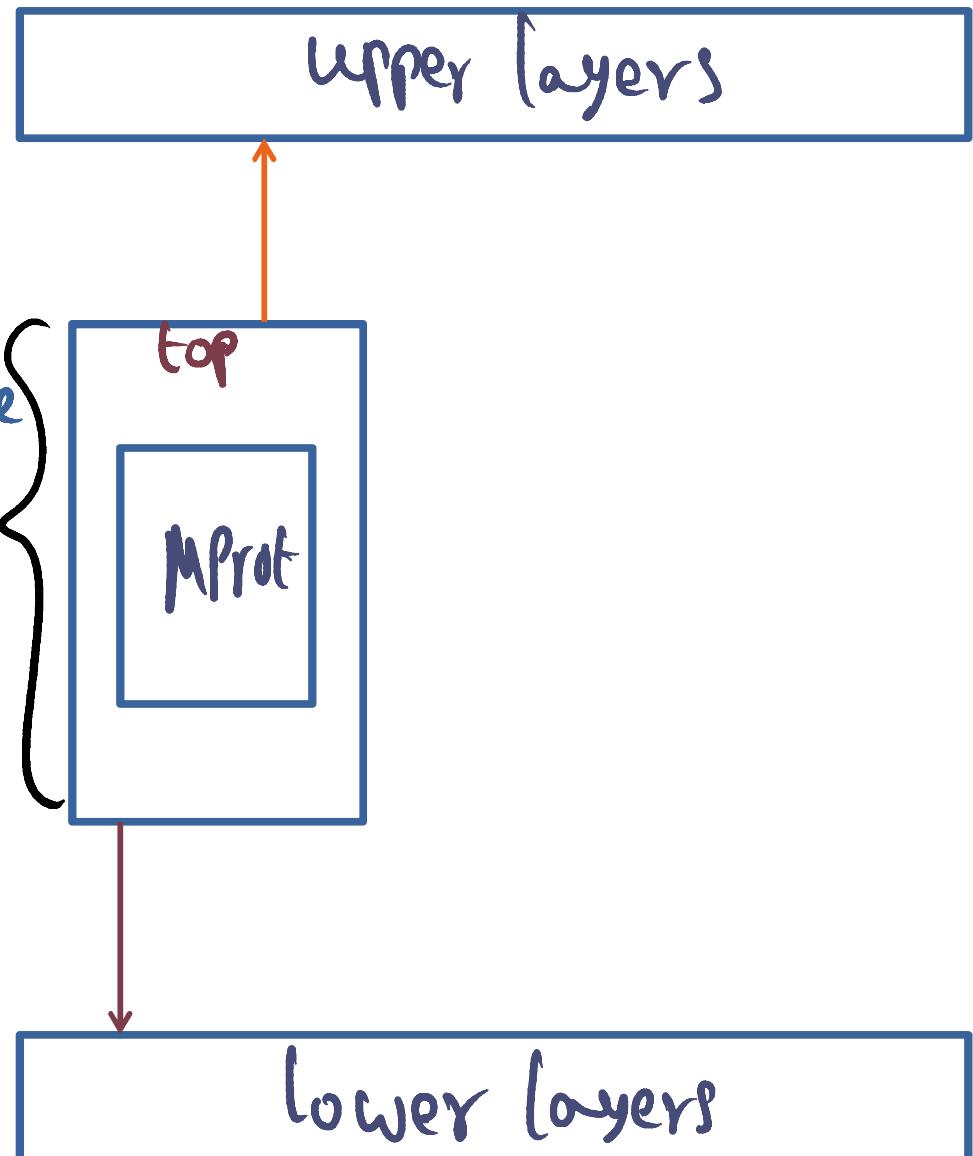
Static opt

- NuPrl expert

plus OCaml expert

- layer by layer

multiple
layers



NuPrl to the rescue

Static opt

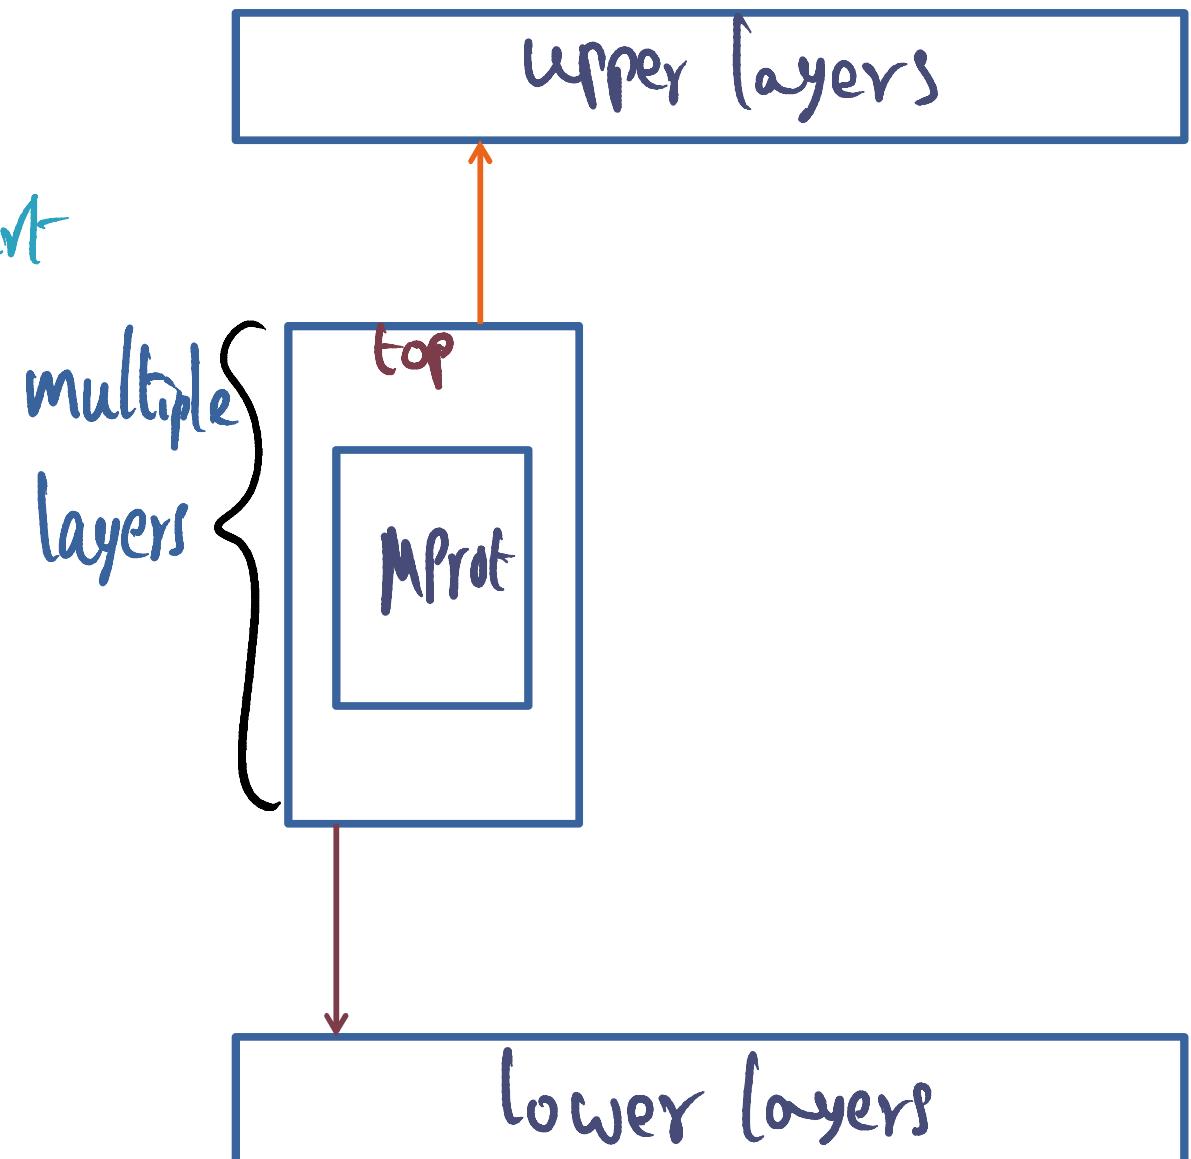
- NuPrl expert

plus OCaml expert

- layer by layer

Dynamic opt

- Collapse layers



NuPrl to the rescue

Static opt

- NuPrl expert

plus OCaml expert

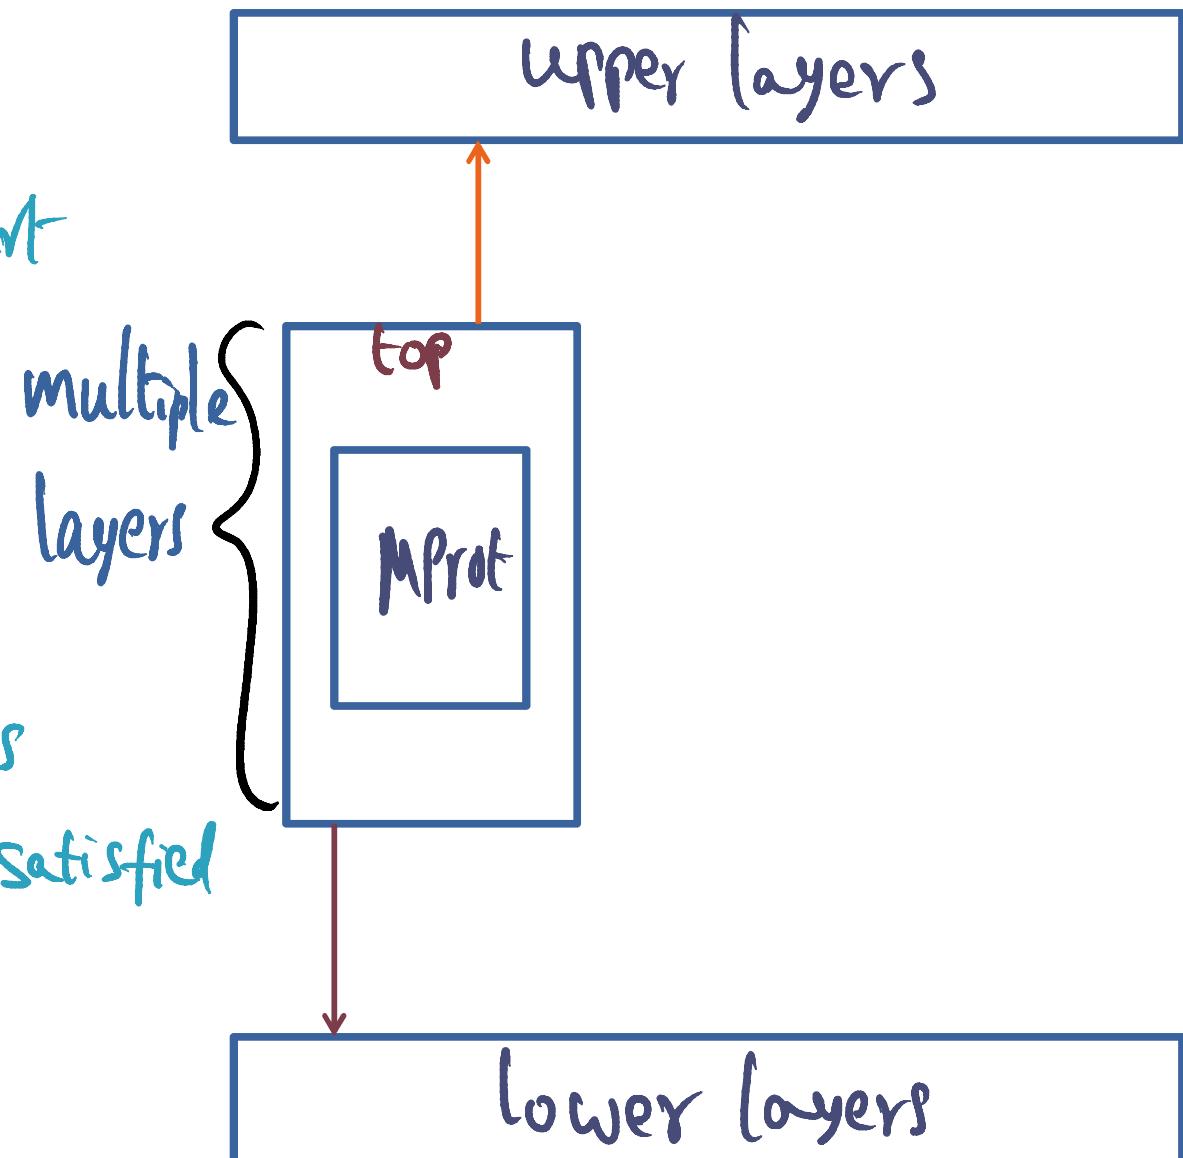
- layer by layer

Dynamic opt

- Collapse layers

- generate bypass

code if CCP satisfied



NuPrl to the rescue

Static opt

- NuPrl expert

plus OCaml expert

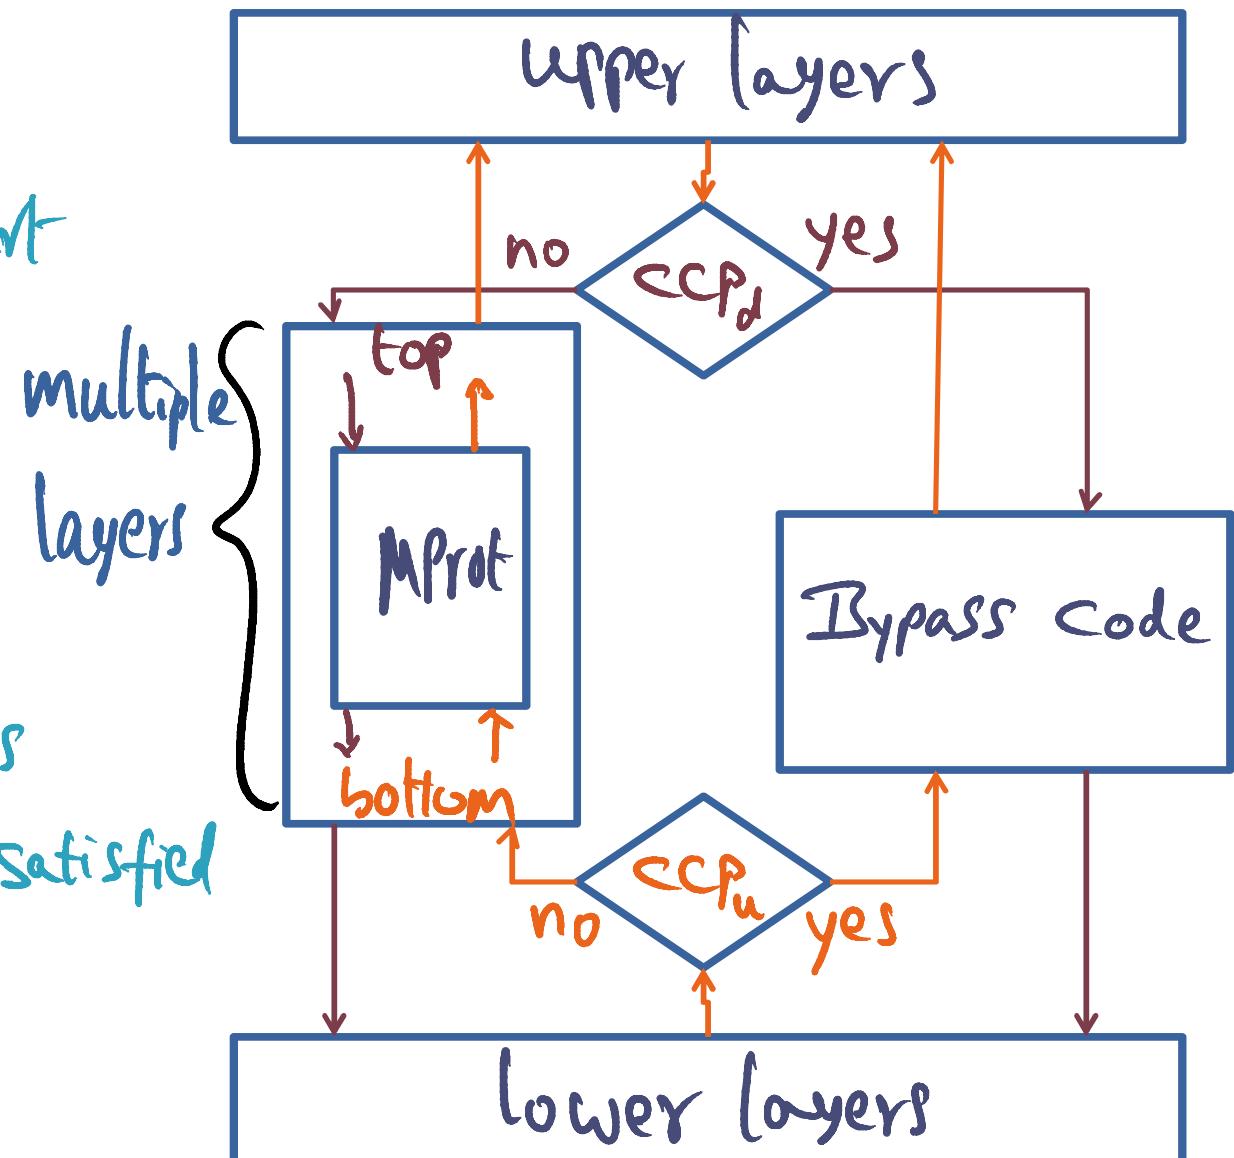
- layer by layer

Dynamic opt

- Collapse layers

- generate bypass

code if CCP satisfied



NuPrl to the rescue

Static opt

- NuPrl expert

plus Ocaml expert

- layer by layer

Dynamic opt

- Collapse layers

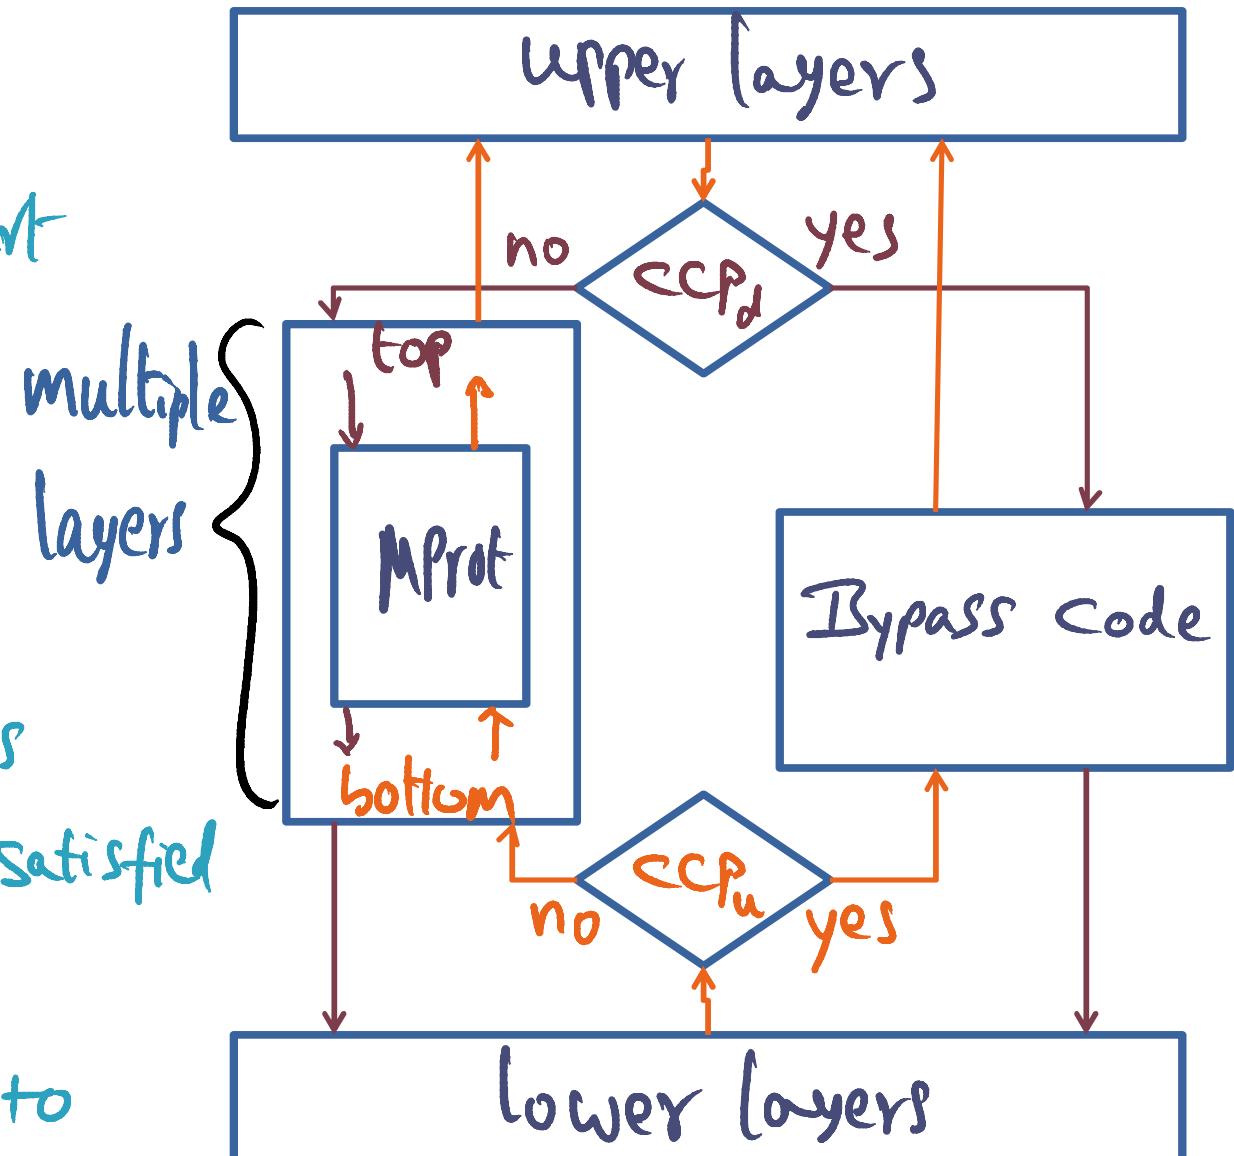
- generate bypass

code if CCP satisfied

Final step

- Convert back to

Ocaml



Key Takeaways

- Do we lose out on performance for this convenience of component-based design?
 - Recall this is the same question that came up when we wanted to go for a micro-kernel based design of an OS
- This experiment concerns synthesizing individual subsystems of an OS from modular components
- Key to the success of this experiment is putting theory with practice
- Ensemble is one concrete example of putting this methodology to work for designing protocol stacks
 - Proof again by construction similar to Liedtke
- Specialization is not new (SPIN, Exokernel, Thekkath & Levy)
- Cornell experiment is achieving this at a subsystem level systematically