

Today

## Lesson 8 Videos

→ — LRVM

Friday

- RioVista
  - Hg
- ( \* please  
watch  
videos \* )

LRVM

in Software

- persistent memory layer in support of system services

Why ?

## Persistence

Why?

- need of OS Subsystems

How?

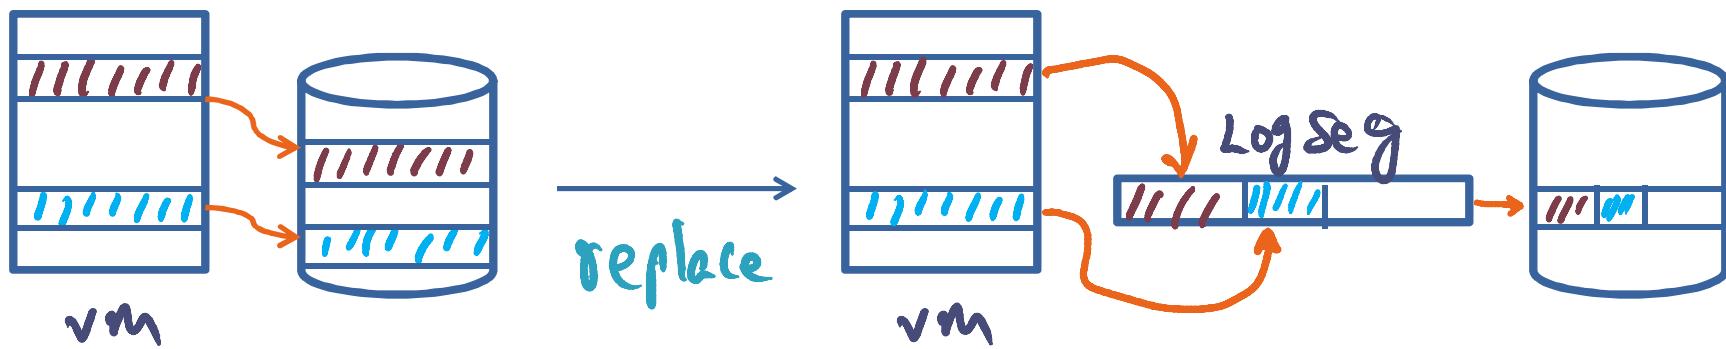
- make Virtual memory persistent

Who will use it?

- Subsystem designers if it is **performant**

How to make it efficient?

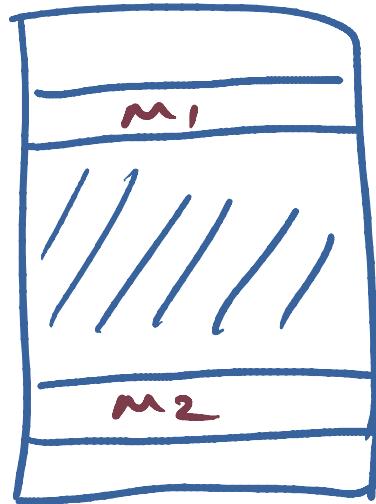
- use **persistent logs** to record changes to VM



# Server Design

- persistent metadata  $m_1, m_2, \dots, m_n$

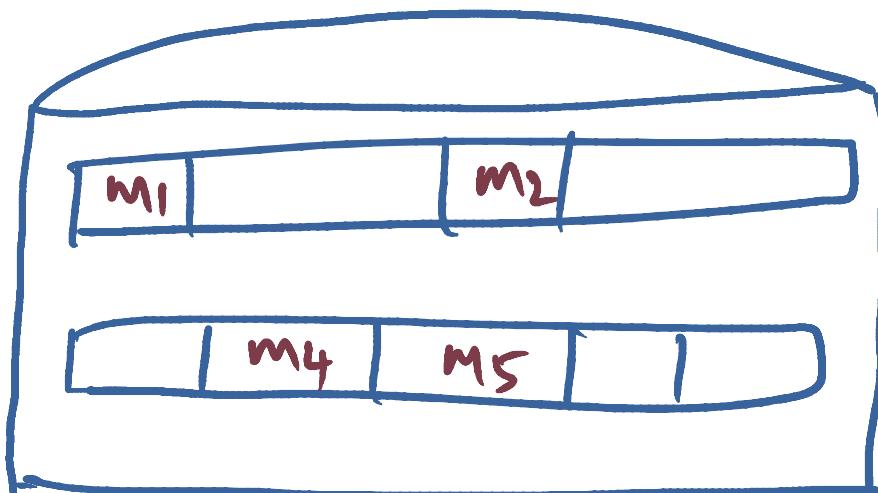
- normal data structures + code →



Virtual  
address  
space of  
server

Create external data segments to back persistent data structures

- apps manage their persistence needs



Designer's choice  
to use single or  
multiple data  
segments

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

### Body of server code

- begin-xact(tid, restore-mode)
- set-range(tid, addr, size)
- end-xact(tid, commit-mode)
- abort-xact(tid)

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

### Gc to reduce log space

- flush()
- truncate()

### Body of server code

- begin-xact(tid, restore-mode)
- set-range(tid, addr, size)
- end-xact(tid, commit-mode)
- abort-xact(tid)

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

### Body of server code

- begin-xact(tid, restore-mode)
- set-range(tid, addr, size)
- end-xact(tid, commit-mode)
- abort-xact(tid)

### GC to reduce log space

- flush()
  - truncate()
- ↓
- done by  
RVM  
automatically

Provided for  
app flexibility

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

### Gc to reduce log space

- flush()
  - truncate()
- ↓
- done by  
RVM  
automatically

Provided for  
app flexibility

### Body of server code

- begin-xact(tid, restore-mode)
- set-range(tid, addr, size)
- end-xact(tid, commit-mode)
- abort-xact(tid)

### Miscellaneous

- query-options(region)
- set-options(options)
- create-log(options, len, mode)

## RVM Primitives

### Initialization

- initialize(options)
- map(region, options)
- unmap(region)

### Gc to reduce log space

- flush()
  - truncate()
- ↓
- done by  
RVM  
automatically

Provided for  
app flexibility

### Body of server code

- begin-xact(tid, restore-mode)
- set-range(tid, addr, size)
- end-xact(tid, commit-mode)
- abort-xact(tid)

### Miscellaneous

- query-options(region)
- set-options(options)
- create-log(options, len, mode)

⇒ Simplicity - small set of primitives

## How the Server uses the primitives

Initialize address space from Ext Segs

:  
begin-xact(tid, mode);

Set range (tid, base-addr, #bytes);

:  
write metadata m<sub>1</sub>

:  
write metadata m<sub>2</sub>

end-xact(tid, mode);

; // contained in range

; // contained in range

// or this can be abort

## How the Server uses the primitives

Initialize address space from Ext Segs

:

begin-xact(tid, mode);

Set range (tid, base-addr, #bytes);

;

write metadata m<sub>1</sub>

;

write metadata m<sub>2</sub>

end-xact(tid, mode);

LRVM creates  
undo record

;

// contained in range

;

// contained in range

// or this can be abort

## How the Server uses the primitives

Initialize address space from Ext Segs

:

begin-xact(tid, mode);

Set range (tid, base-addr, #bytes);

no action  
by  
LRVM

{ write metadata m<sub>1</sub>

:

{ write metadata m<sub>2</sub>

end-xact (tid, mode);

LRVM creates  
undo record

;

; // contained in range

;

; // contained in range

|| or this can be abort

## How the Server uses the primitives

Initialize address space from Ext Segs

⋮

begin-xact(tid, mode);

Set range (tid, base-addr, #bytes);

no action

by  
LRVM

{ write metadata m<sub>1</sub>

⋮

{ write metadata m<sub>2</sub>

end-xact(tid, mode);



LRVM creates redo log in memory

LRVM creates  
undo record

; // contained in range

; // contained in range

// or this can be abort

## How the Server uses the primitives

Initialize address space from Ext Segs

⋮

begin-xact(tid, mode);

Set range (tid, base-addr, #bytes);

no action

by  
LRVM

{ write metadata m<sub>1</sub>

⋮

{ write metadata m<sub>2</sub>

end-xact(tid, mode);

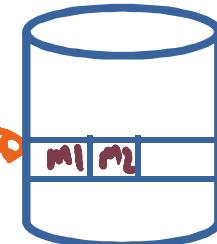


LRVM creates redo log in memory

- flush to disk sync or later depending on Mode



redo log



redo log on disk

LRVM creates  
undo record

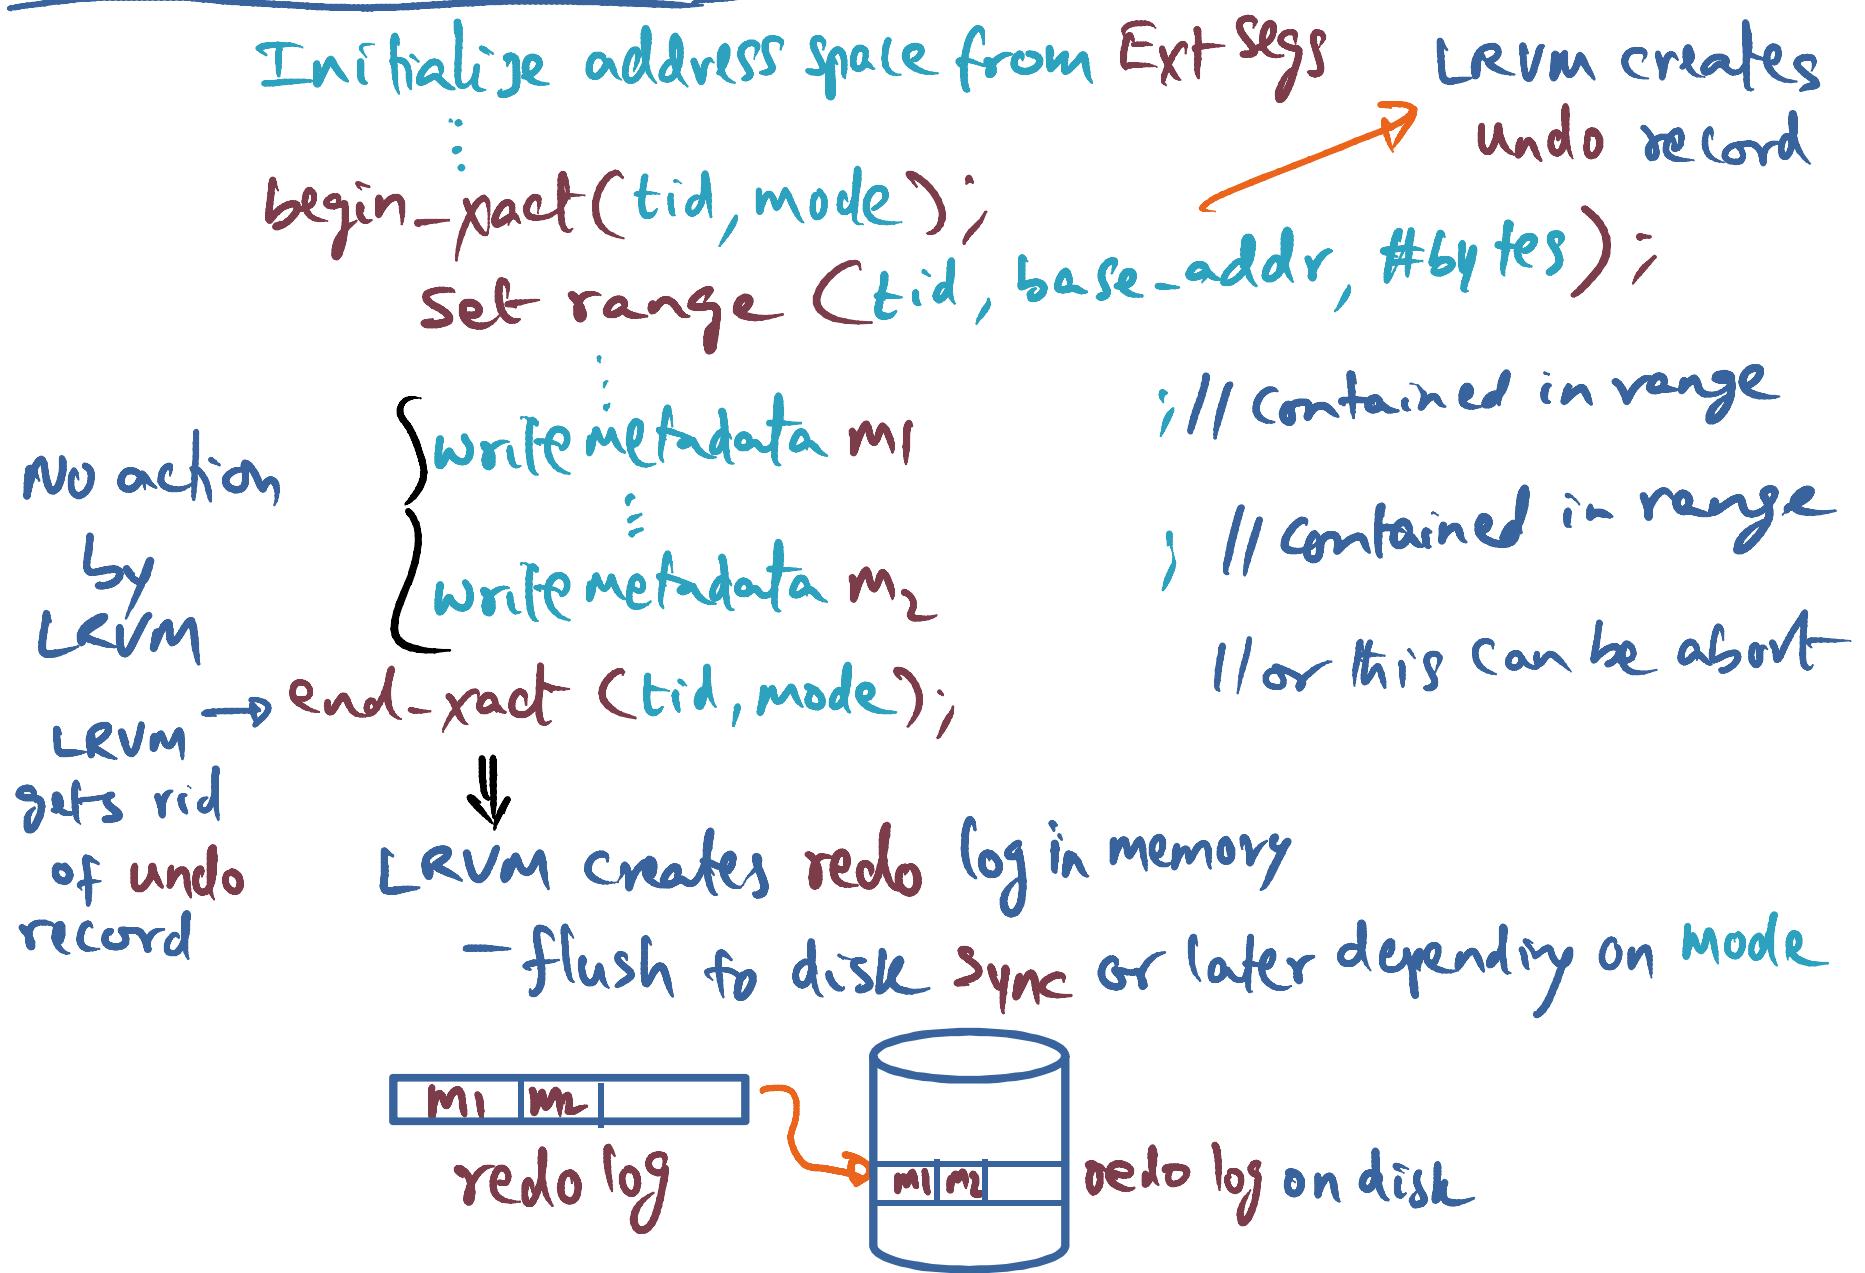
⋮

; // contained in range

; // contained in range

// or this can be abort

## How the Server uses the primitives



## How the Server uses the primitives

Initialize address space from Ext Segs  
...  
begin-xact(tid, mode);  
Set range (tid, base-addr, #bytes);  
LRVM creates undo record

no action by LRVM

LRVM gets rid of undo record

LRVM creates metadata m<sub>1</sub>  
LRVM creates metadata m<sub>2</sub>  
end-xact(tid, mode);  
; // contained in range  
; // contained in range  
// or this can be abort  
↓  
LRVM creates redo log in memory  
- flush to disk sync or later depending on Mode

