

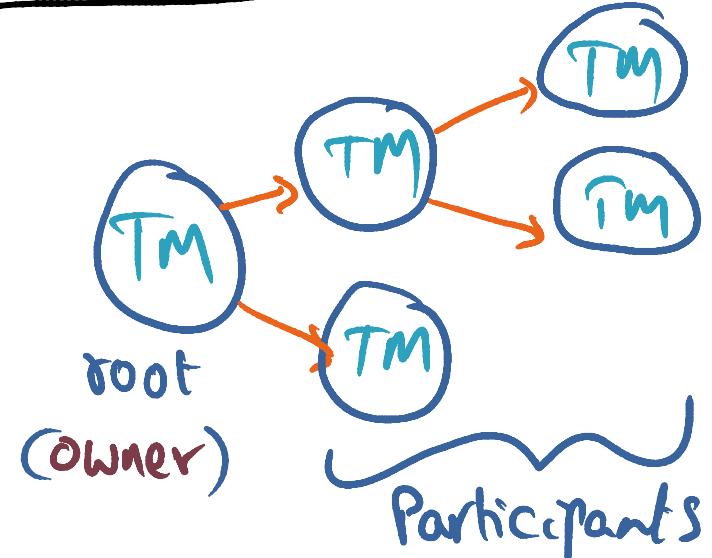
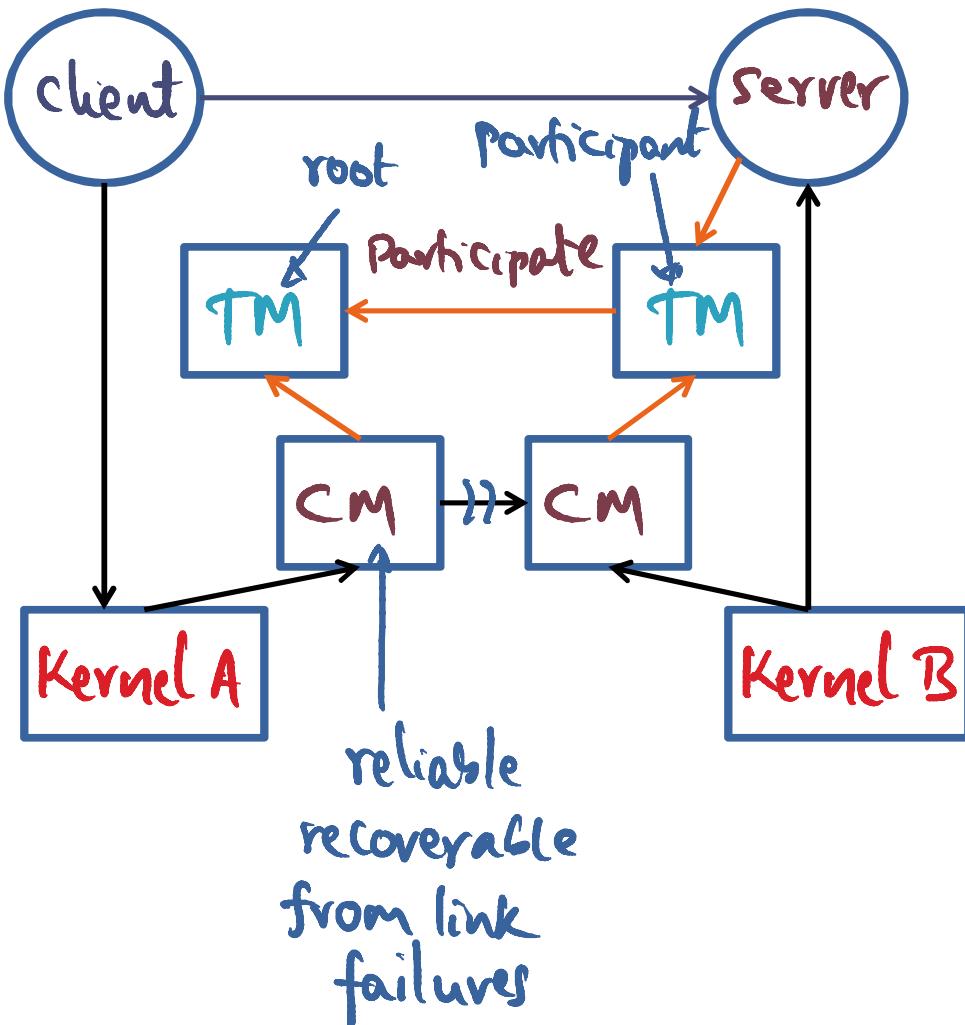
# Today

- \* Quicksilver  
Internet scale computing  
(Lesson 9)
- ⇒ \* giant scale services
  - \* Map-reduce programming model
  - \* CDN (Coral)

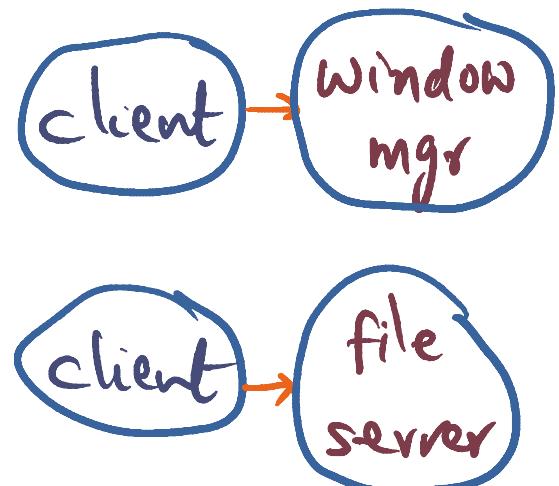
\* Please watch videos before class \*

# Bundling Distributed IPC & Actions

Transaction  $\Rightarrow$  Secret Share for recovery management



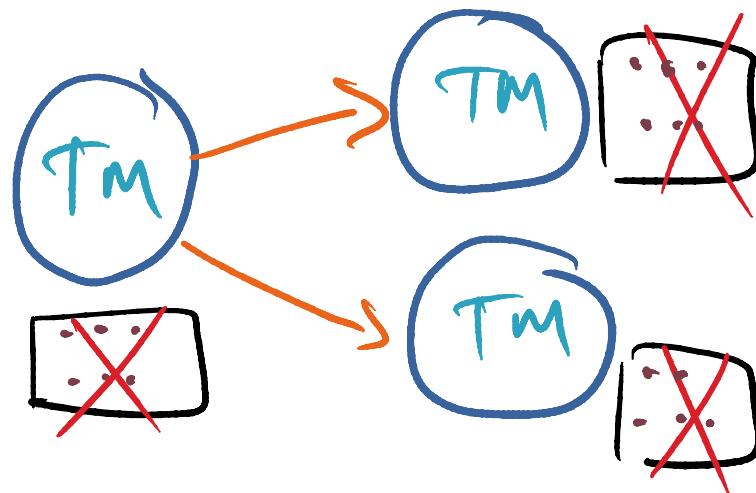
## Examples



## upshot of bundling IPC + Recovery

### Reclaim Resources

- breadcrumbs left behind by failed clients/servers

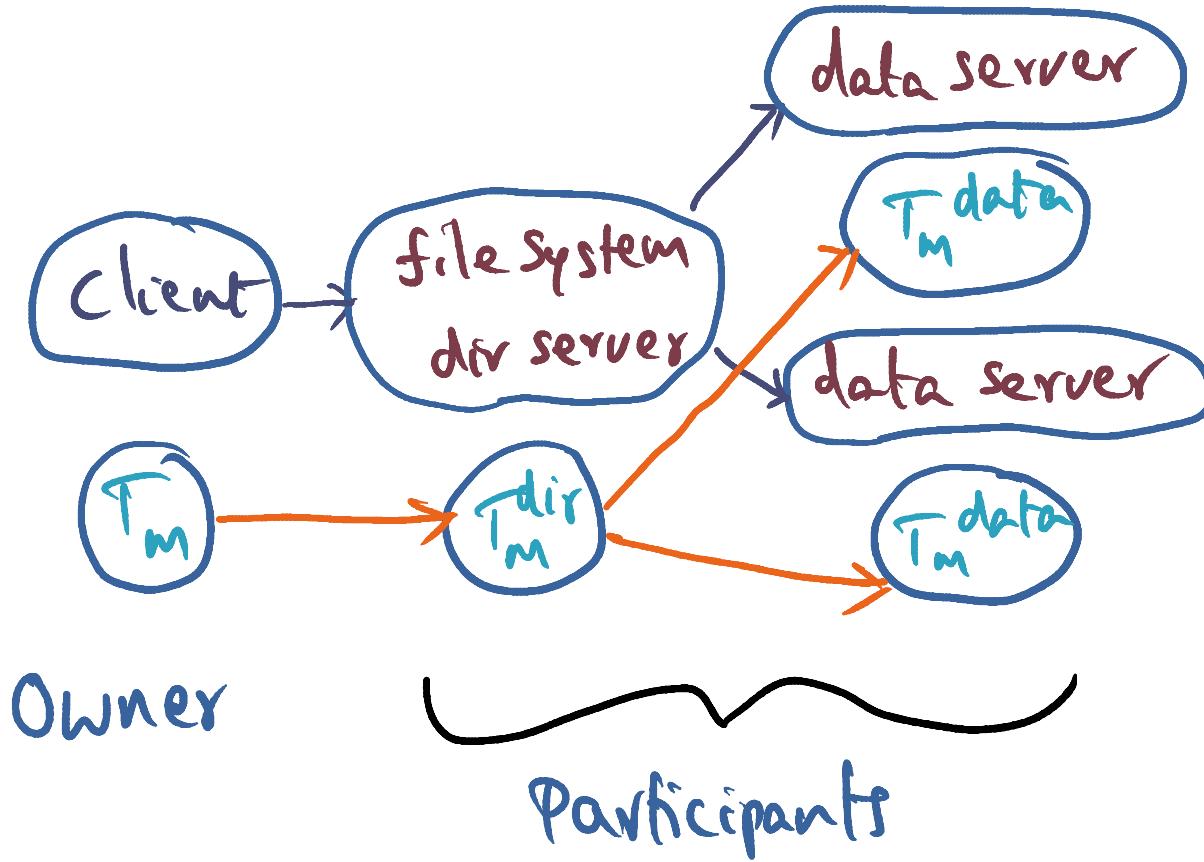


Examples:

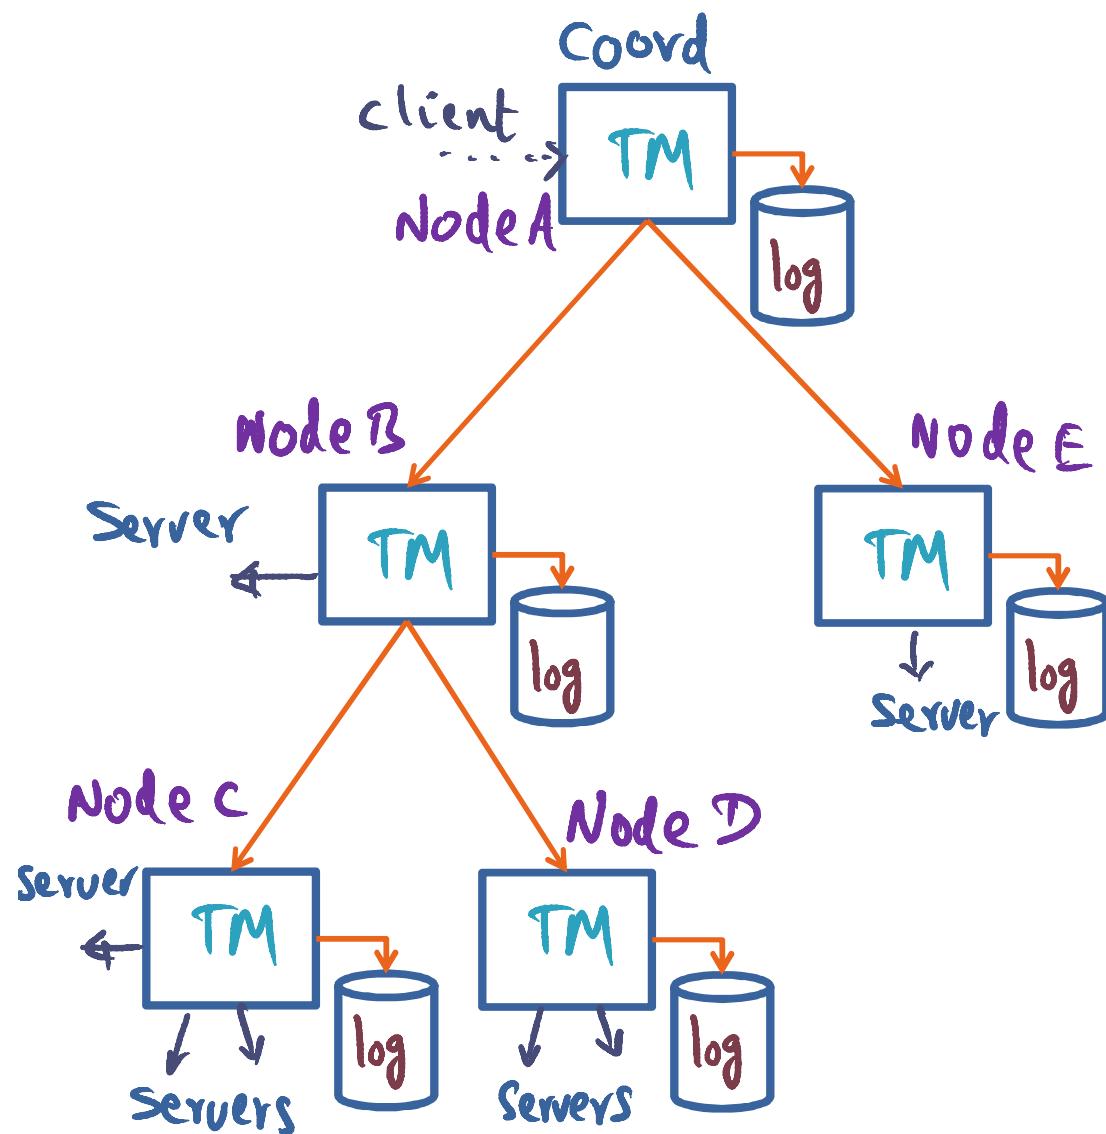
memory, file handles,  
Communication handles,  
orphan Windows, ...

## Transaction management

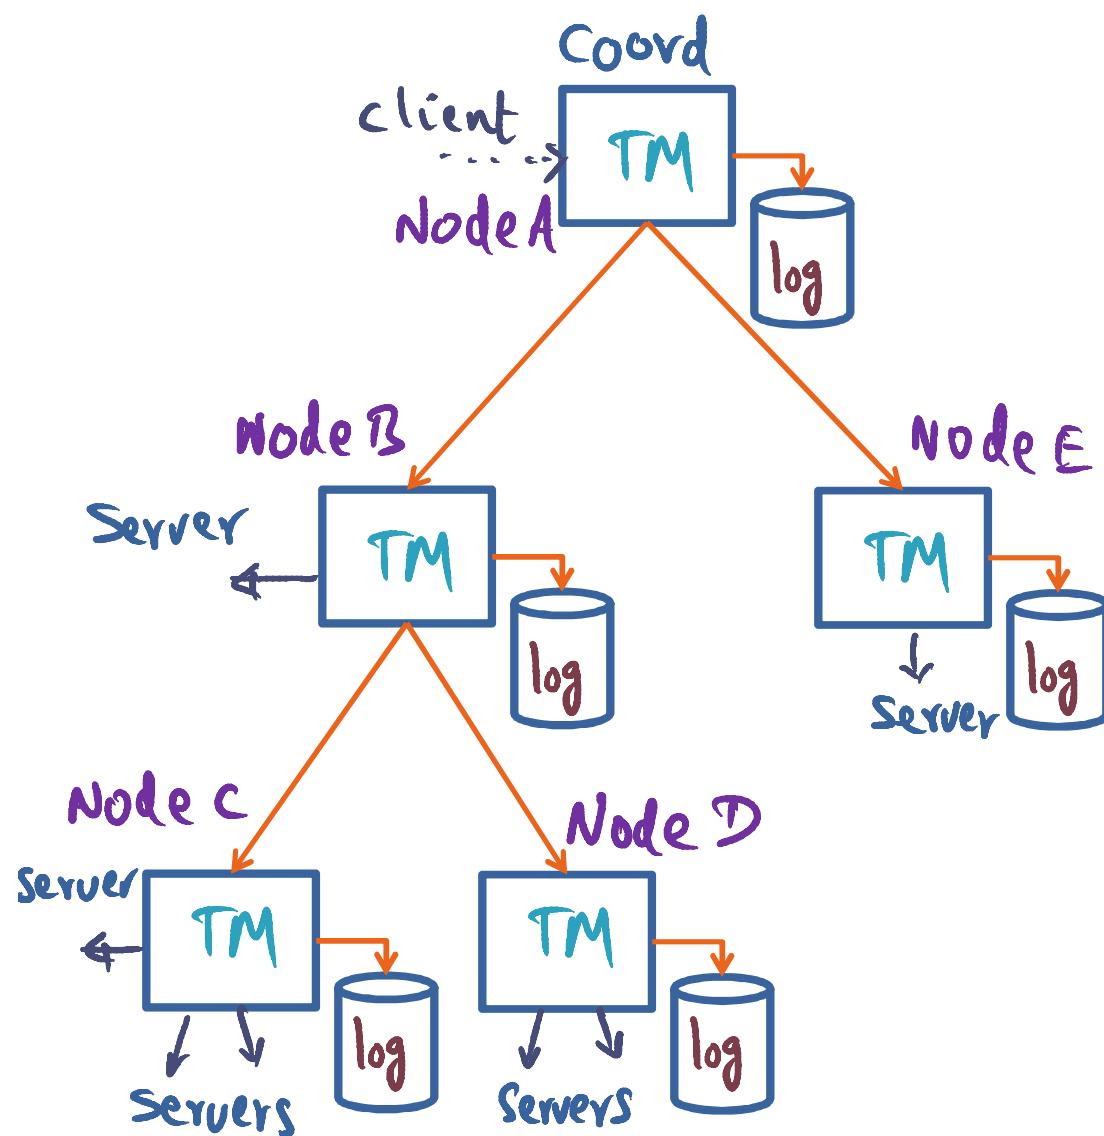
Coordinator can be different from owner



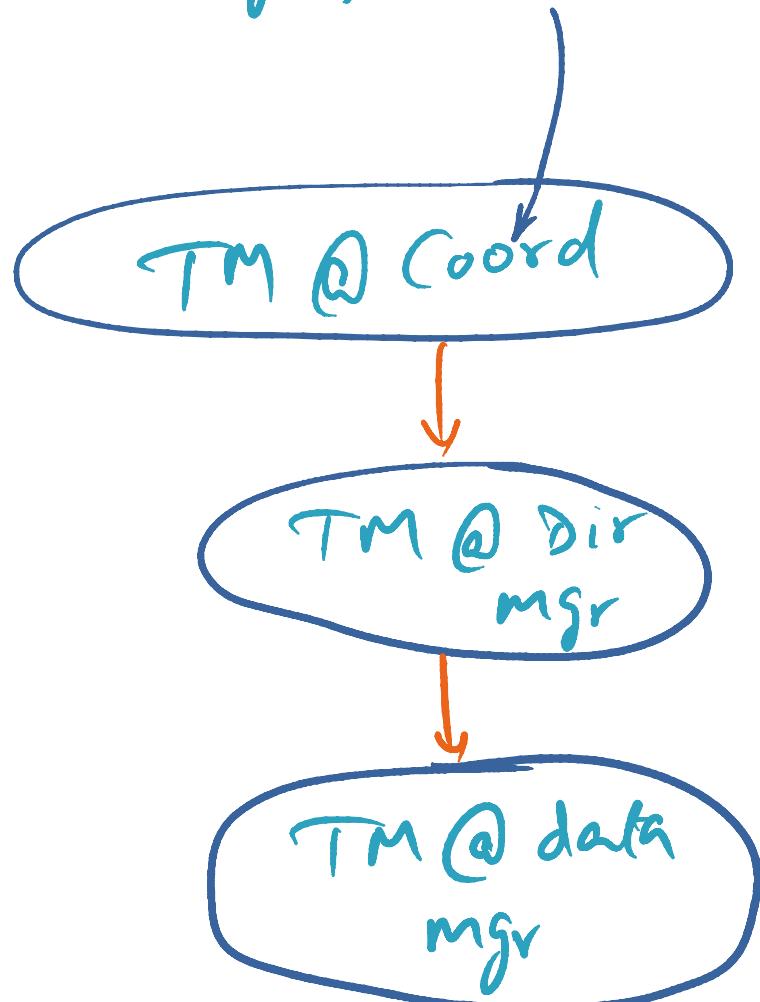
# Distributed Transaction



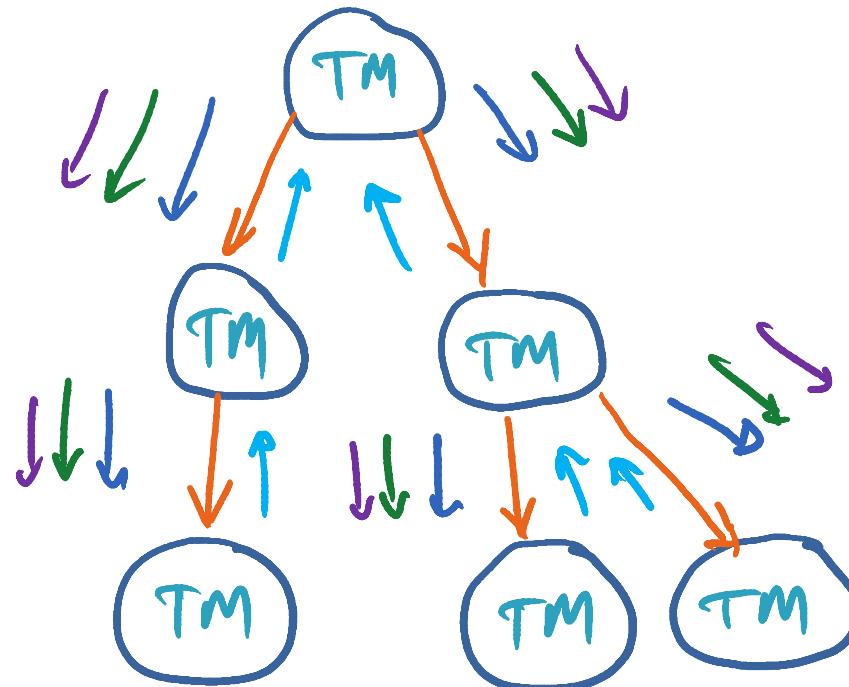
# Distributed Transaction



All TMs not equal  
- e.g. file server



# Commit initiated by Coordinator



Down the tree

→ Vote request

→ Abort request

→ End Commit/abort

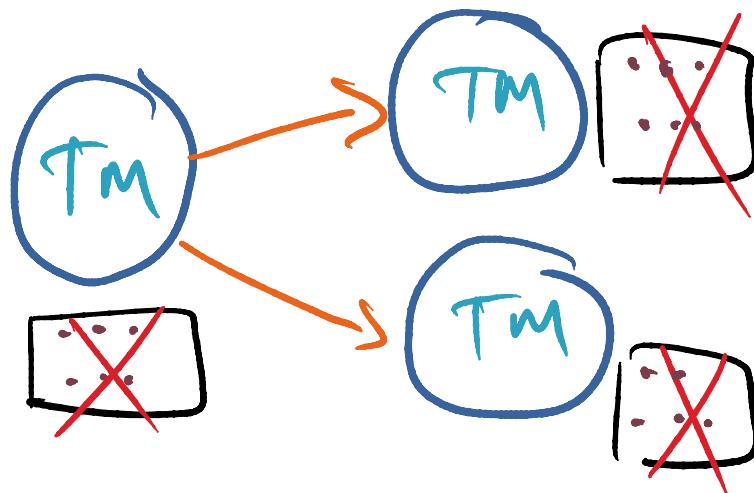
Up the tree

→ response commensurate  
with request

## upshot of bundling IPC + Recovery

### Reclaim Resources

- breadcrumbs left behind by failed clients/servers



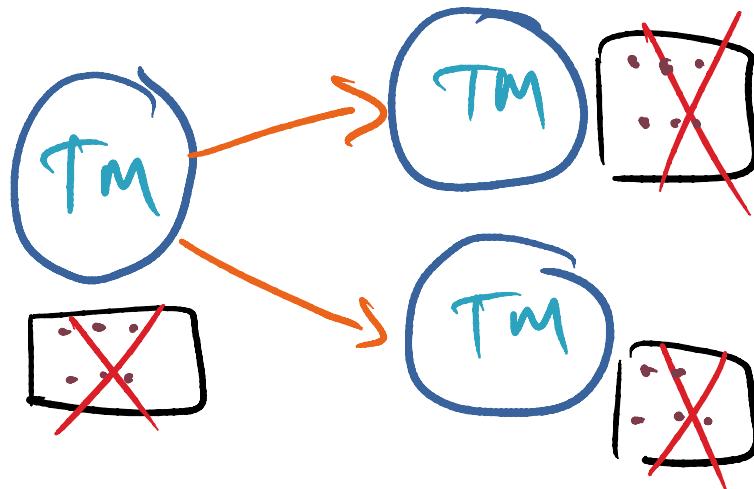
Examples:

memory, file handles,  
Communication handles,  
orphan Windows, ...

## upshot of bundling IPC + Recovery

### Reclaim Resources

- breadcrumbs left behind by failed clients/servers



Examples:

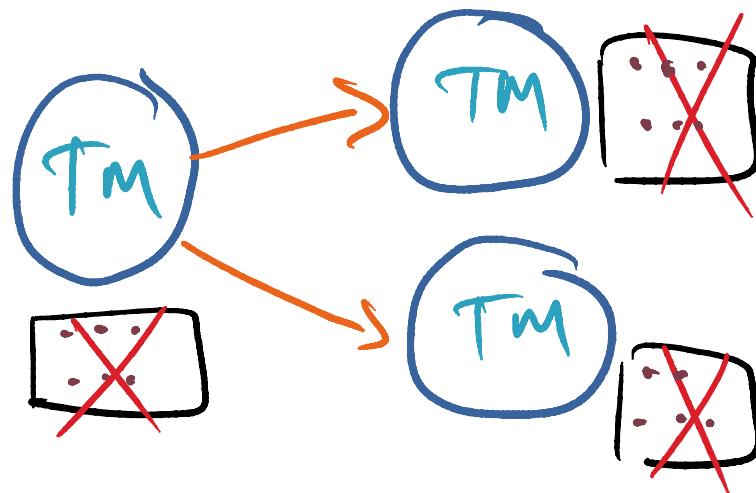
memory, file handles,  
Communication handles,  
orphan Windows, ...

NO Extra communication for recovery

## upshot of bundling IPC + Recovery

### Reclaim Resources

- breadcrumbs left behind by failed clients/servers



Examples:

memory, file handles,  
Communication handles,  
orphan Windows, ...

NO Extra Communication for recovery

only mechanism in OS, policy up to each service

- low overhead mechanisms for simple services
- weighty mechanisms for services such as FS

## Implementation Notes

### Log maintenance

- TMs write log records for recovering persistent state
- Frequency of "log force" impacts performance

⇒ Services have to choose mechanisms commensurate with their recovery requirements

# Key Takeaways

- Transaction as a fundamental OS mechanism to bundle in state recovery of OS services...resurgence in the 90's in LRVM for providing persistence.
- Again in 2010, resurgence in the form of providing safeguard against system vulnerability to malicious attacks in another research OS called TxOS...a later module on system security.
- Computer industry still focused on performance...reliability always takes a backseat...nasty things go on under cover in OS to gain performance!
  - You write a file and you think it is on the hard disk...think again! It is a while before your file write actually gets persisted on the disk...what if the system crashes in the meanwhile?...too bad!

Future?

- New kinds of memories called “storage class memories” that have latency properties similar to DRAM and yet non-volatile.

Will this new technology result in a resurgence of exploring transactions in OS again? Only time will tell...