

Today

Case Studies of parallel OS

* #



* Corey
* cellular Disc. } Partial reading

Today

Case Studies of parallel OS

*



midterm :
monday

* Corey
* Cellular Disc. } Partial reading

Today

Case Studies of parallel OS

* #



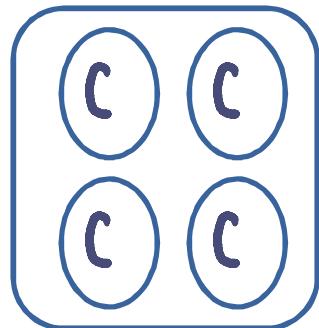
midterm : Oct 5
monday

* Corey
* Cellular Disco } Partial reading

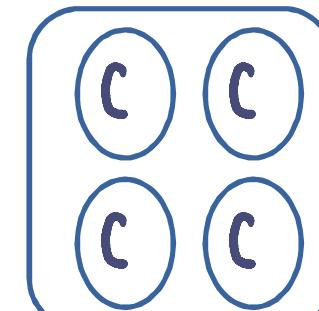
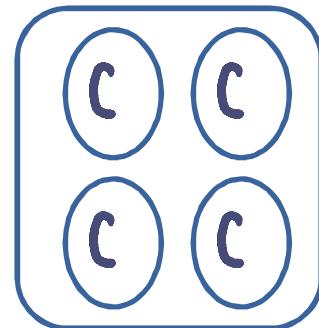
Next week

Distributed systems

* watch Lamport clock video
(First video in Lesson 5 on Nodacity)

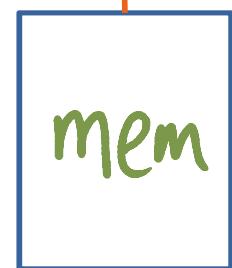
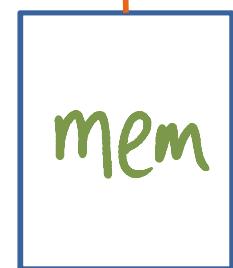
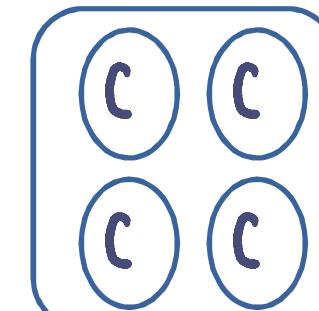


OS for Parallel machines



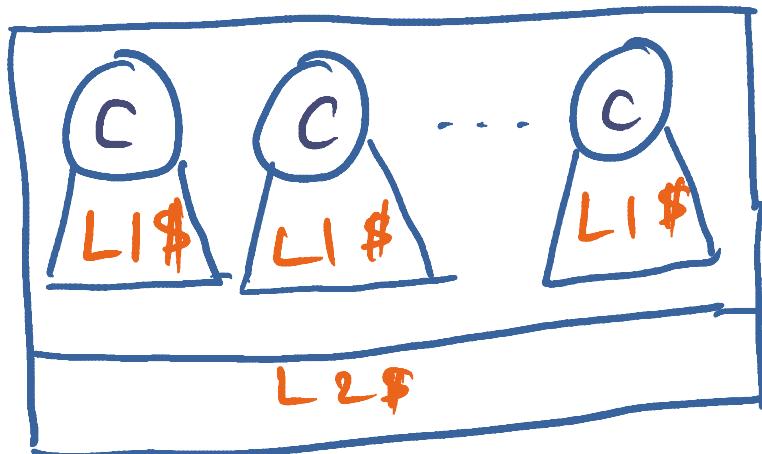
challenges

- NUMA effects
- Deep mem hierarchy
- False sharing



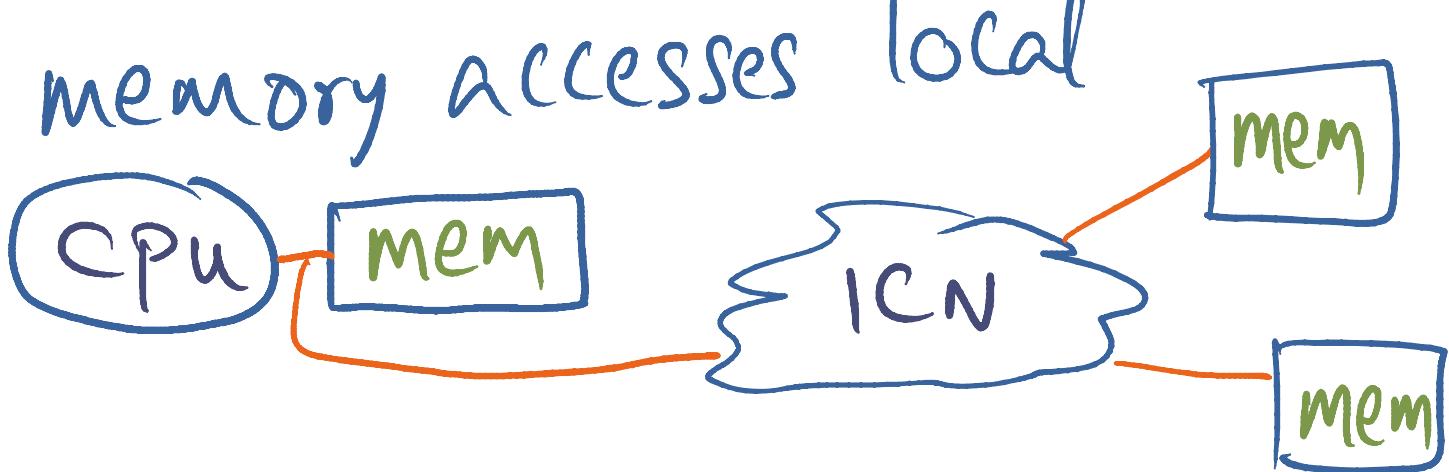
Principles

Cache Conscious decisions



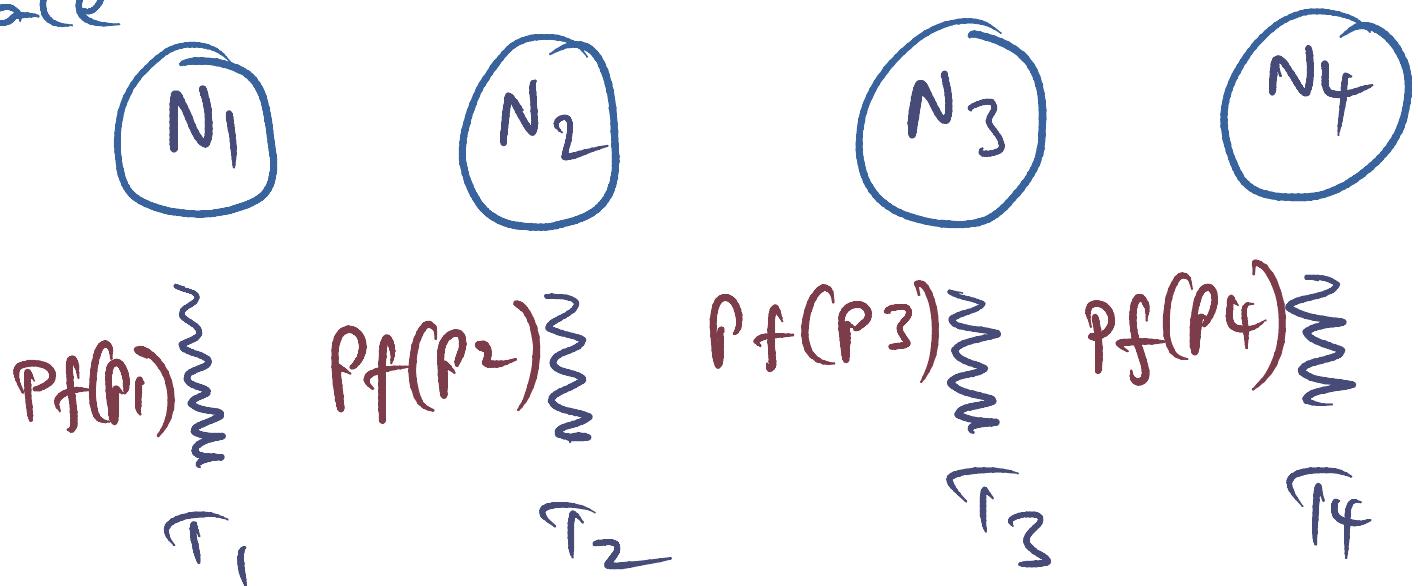
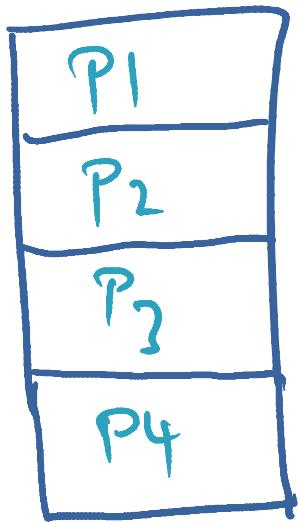
Limit shared System Data Structures

Keep memory accesses local



Concrete Example - P.F. Service

Address Space



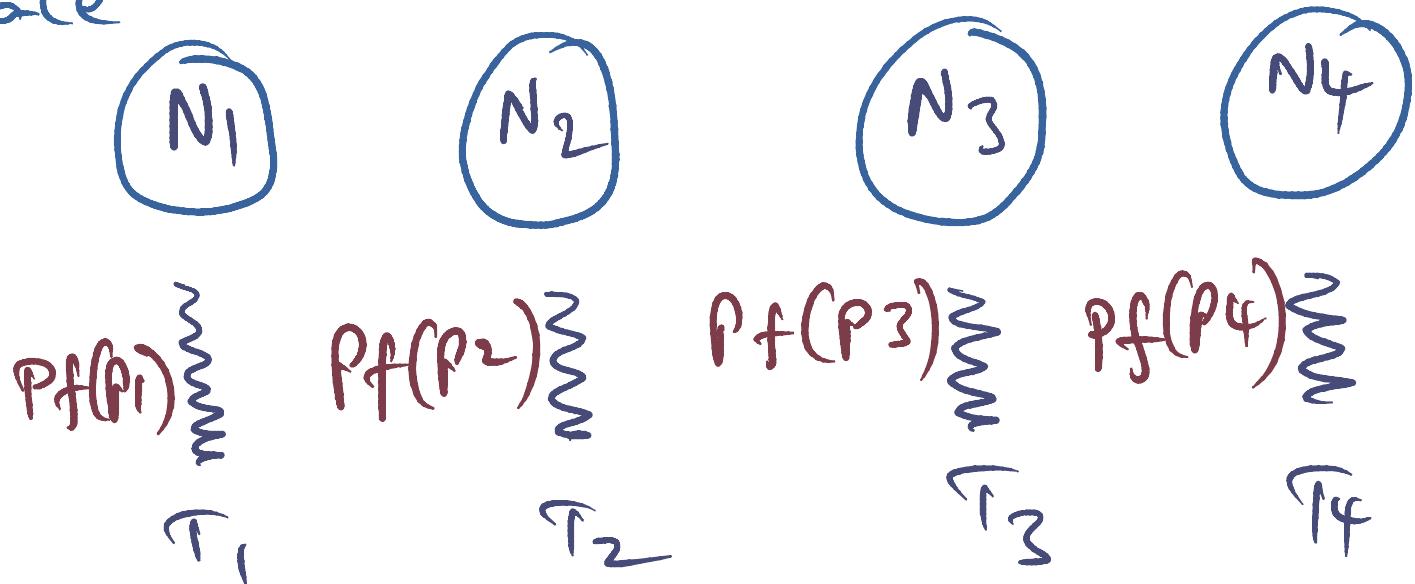
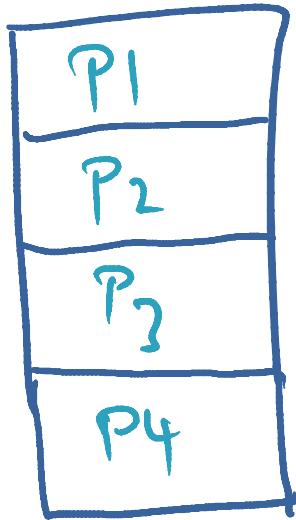
Page table Common

Page fault handler on each node

Page table data structure shared by
all handlers

Concrete Example - P.F. Service

Address Space

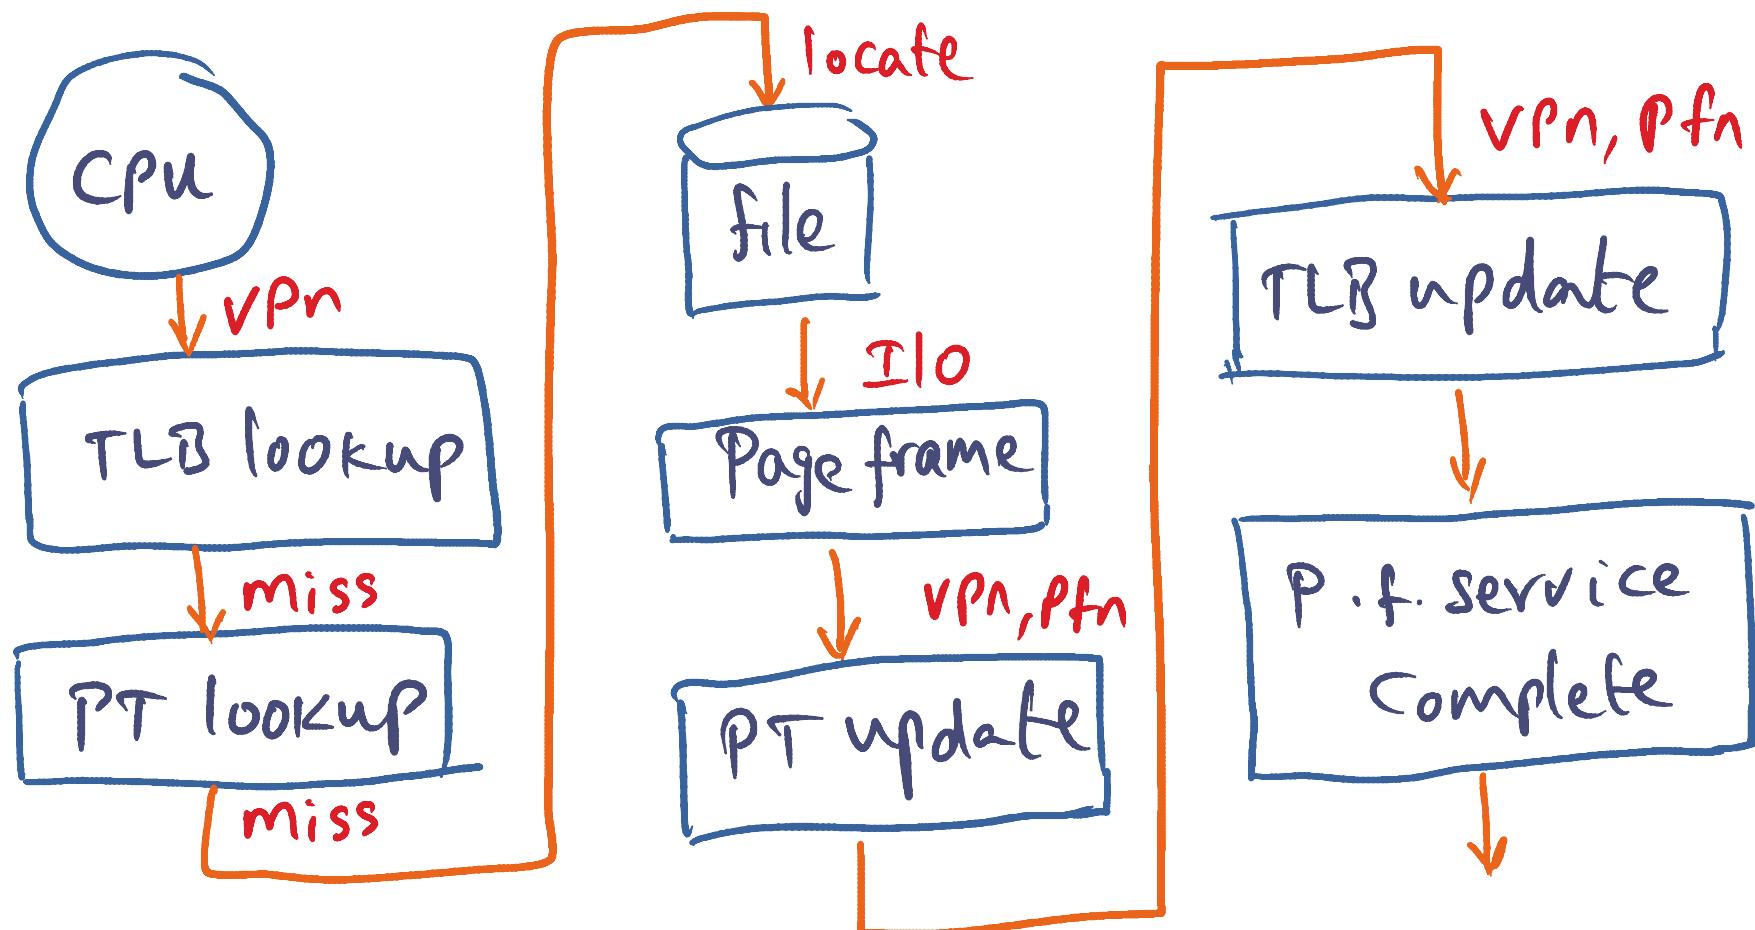


Page table Common

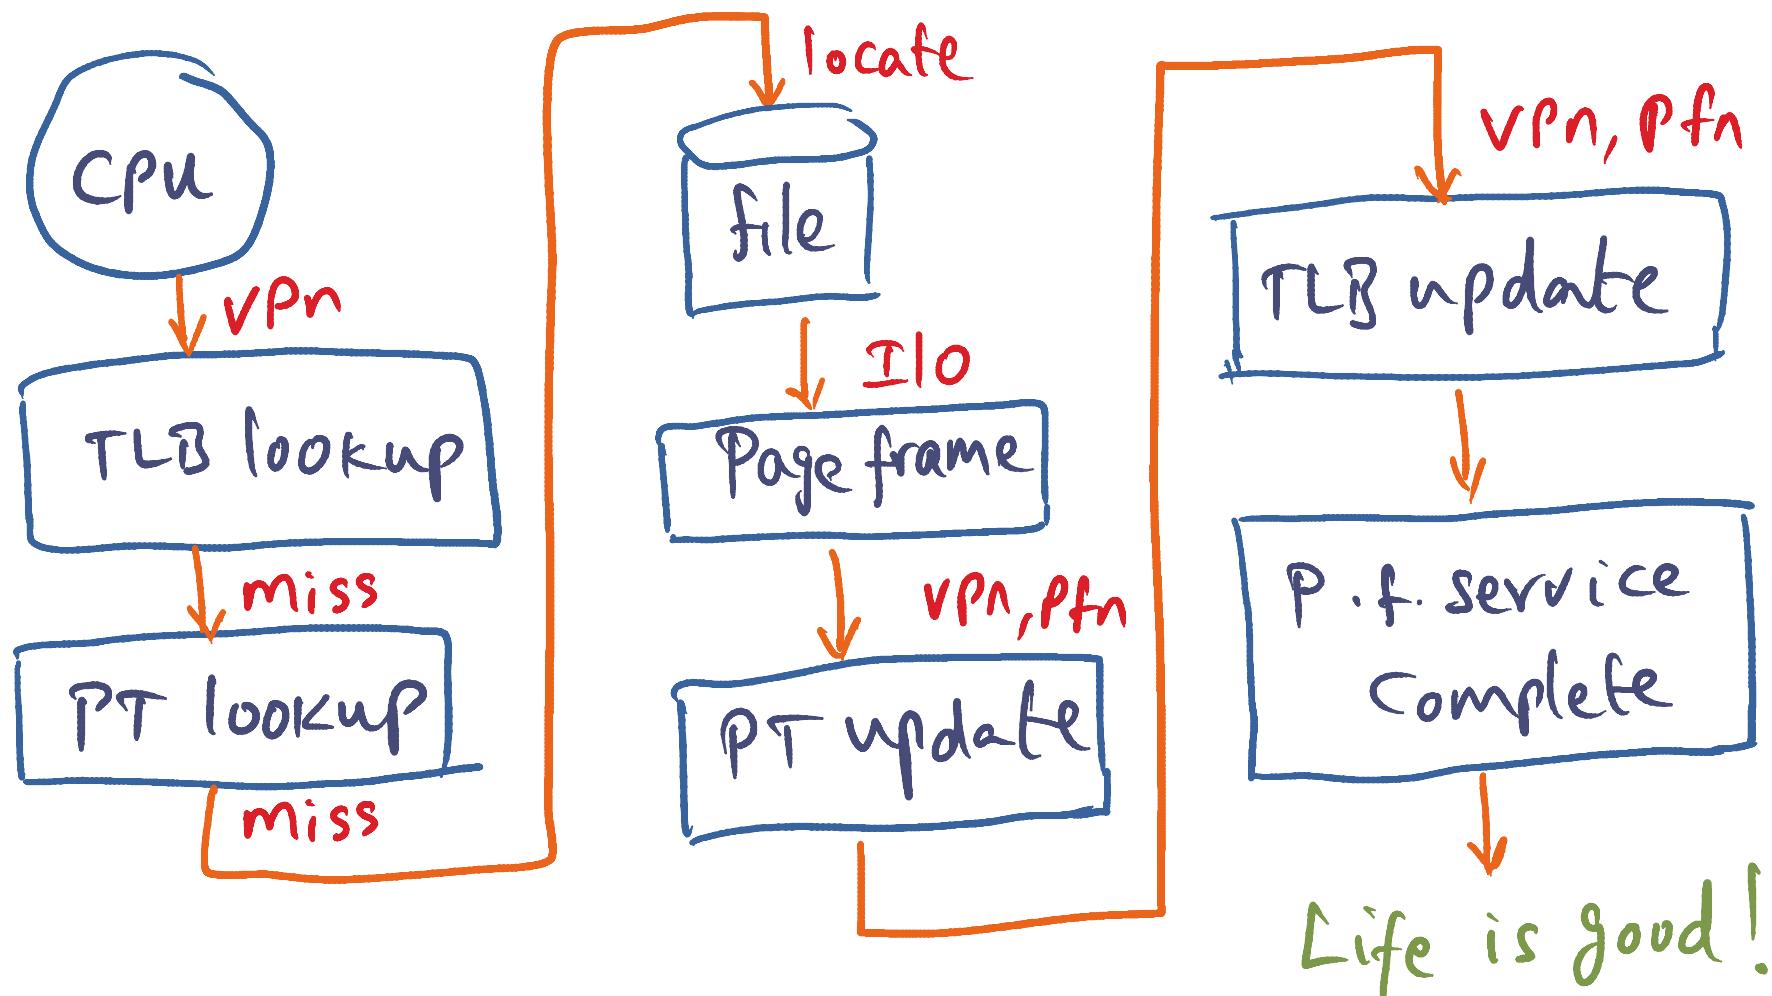
Page fault handler on each node

Page table data structure shared by
all handlers \Rightarrow serialization 😞

Refresher on page fault service

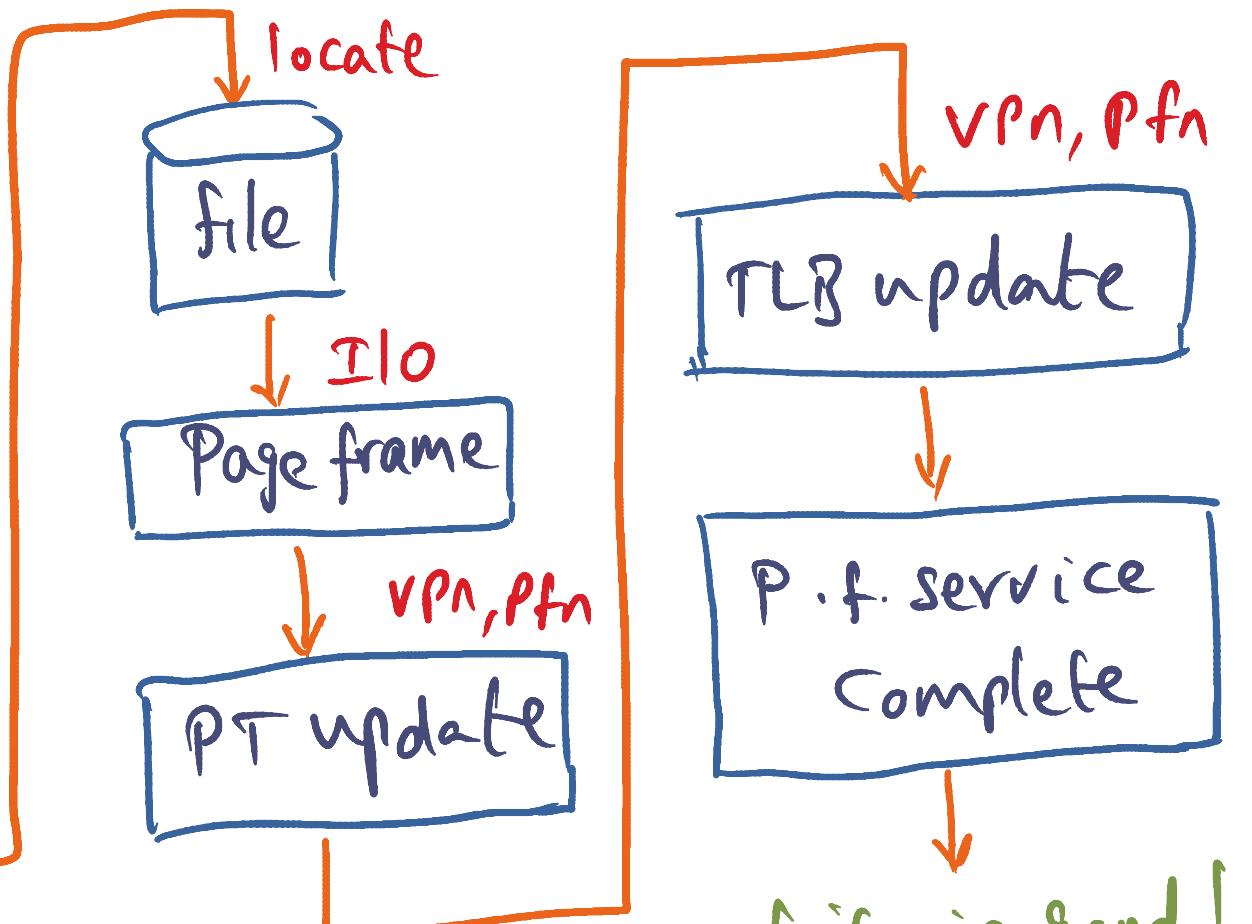
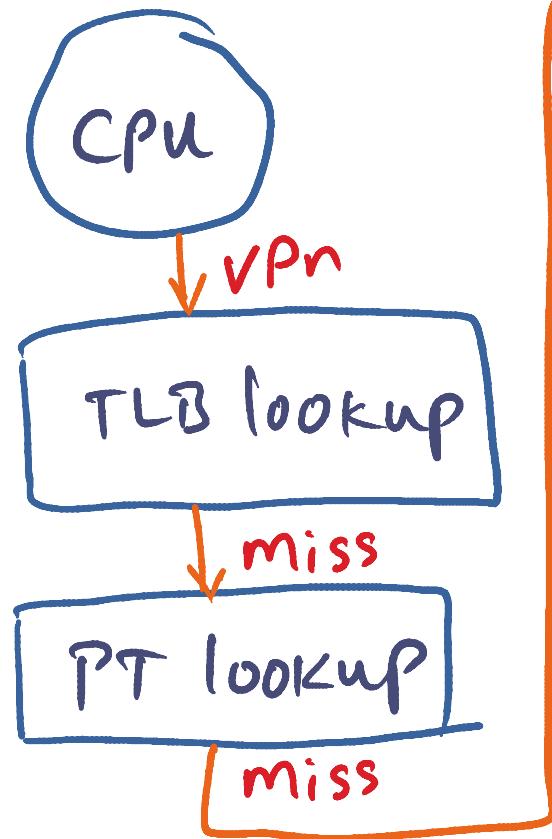


Refresher on page fault service



Refresher on page fault service

Thread
Specific



Life is good!

Refresher on page fault service

Thread
Specific



↓ VPN



↓ miss



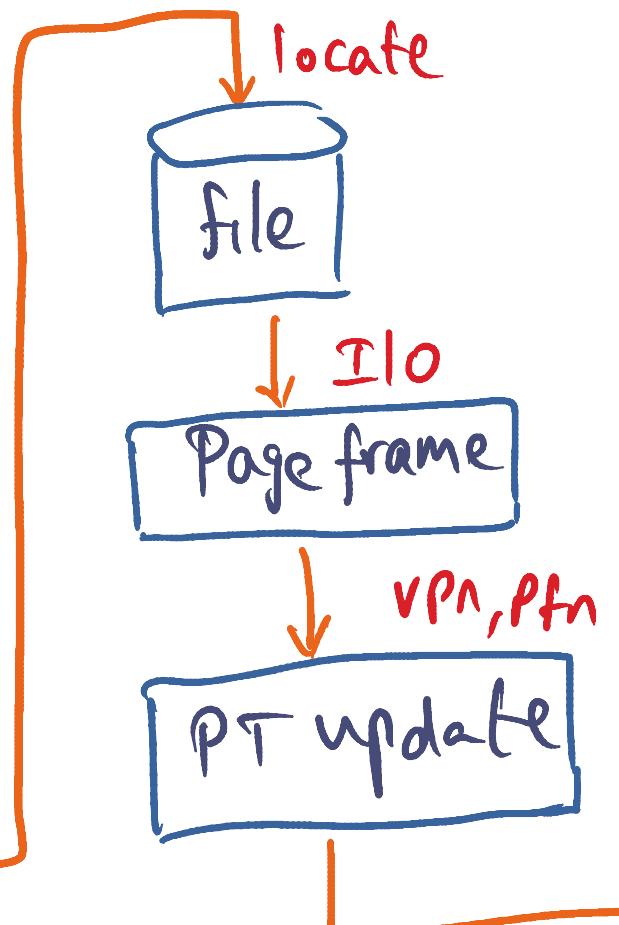
miss

Processor
Specific

↓ VPN, Pfn

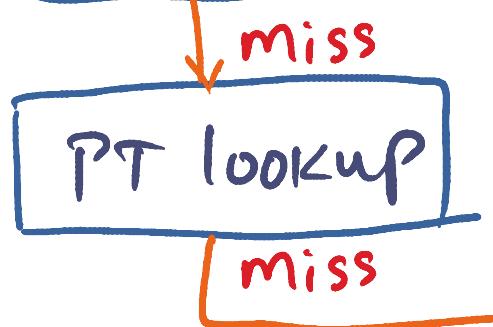
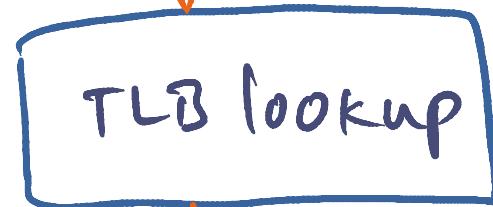
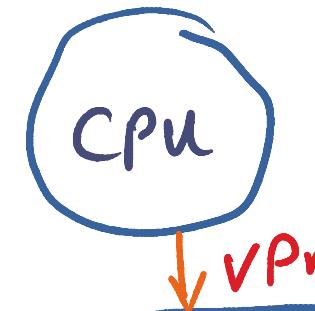


↓
Life is good!

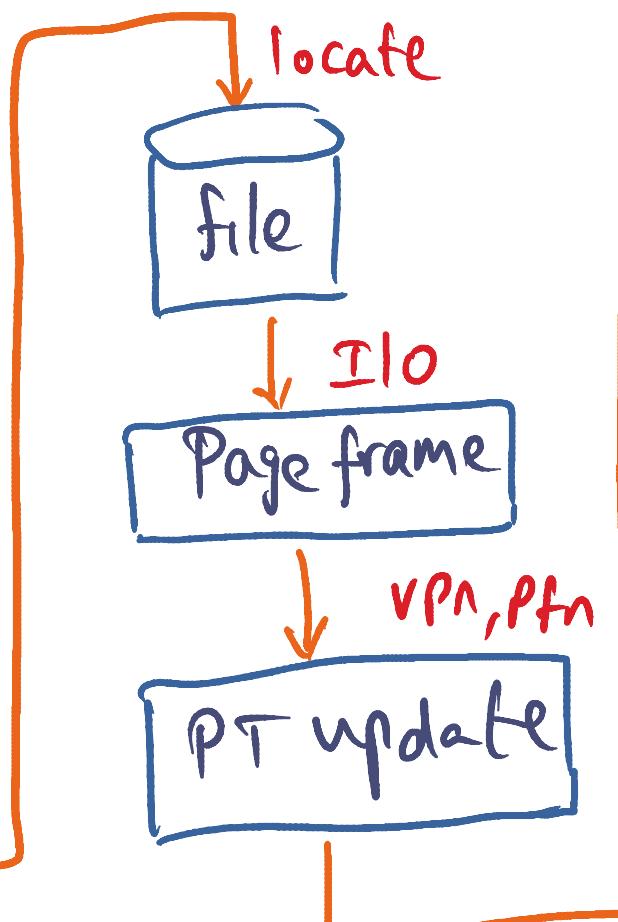


Refresher on page fault service

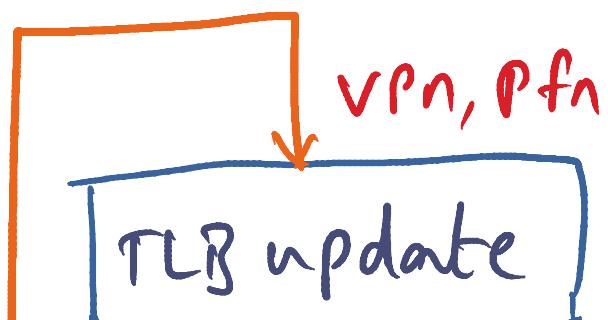
Thread
Specific



Avoid
Serialization



Processor
Specific

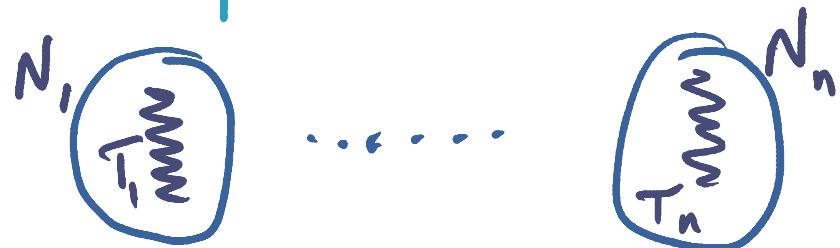


Life is good!

Parallel OS + page fault service

Easy scenario

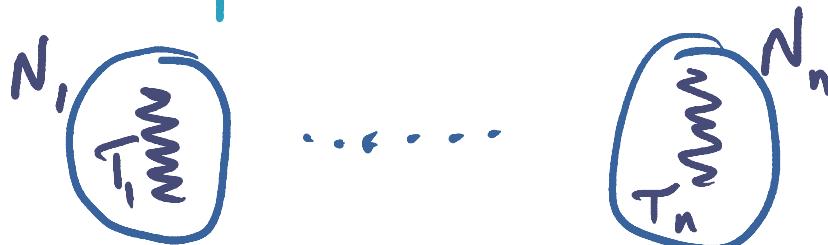
- multiprocess workload



Parallel OS + page fault service

Easy scenario

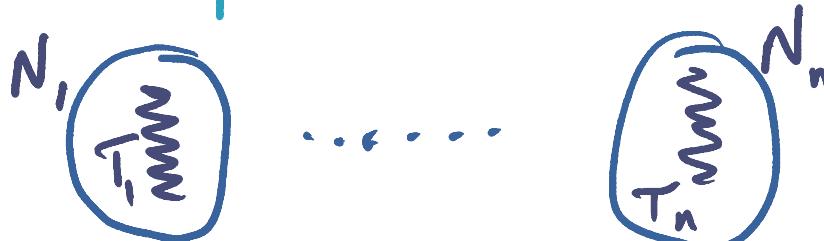
- multiprocess workload
 - * Threads independent
 - * page tables distinct
 - * no serialization



Parallel OS + page fault service

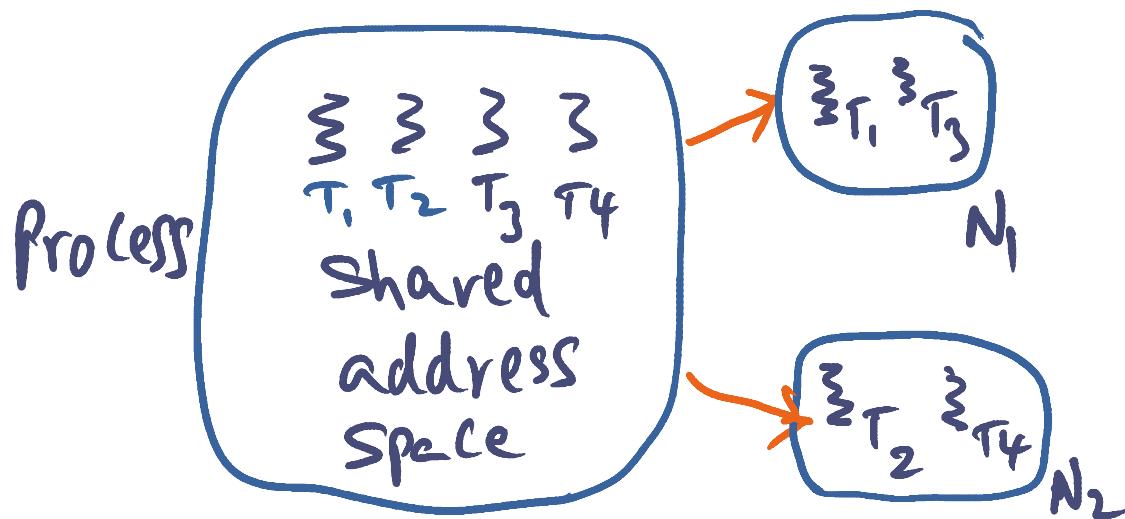
Easy scenario

- multiprocess workload
 - * Threads independent
 - * page tables distinct
 - * no serialization



Hard scenario

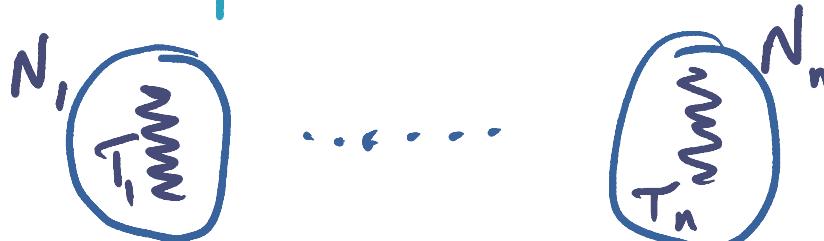
- multi-threaded workload



Parallel OS + page fault service

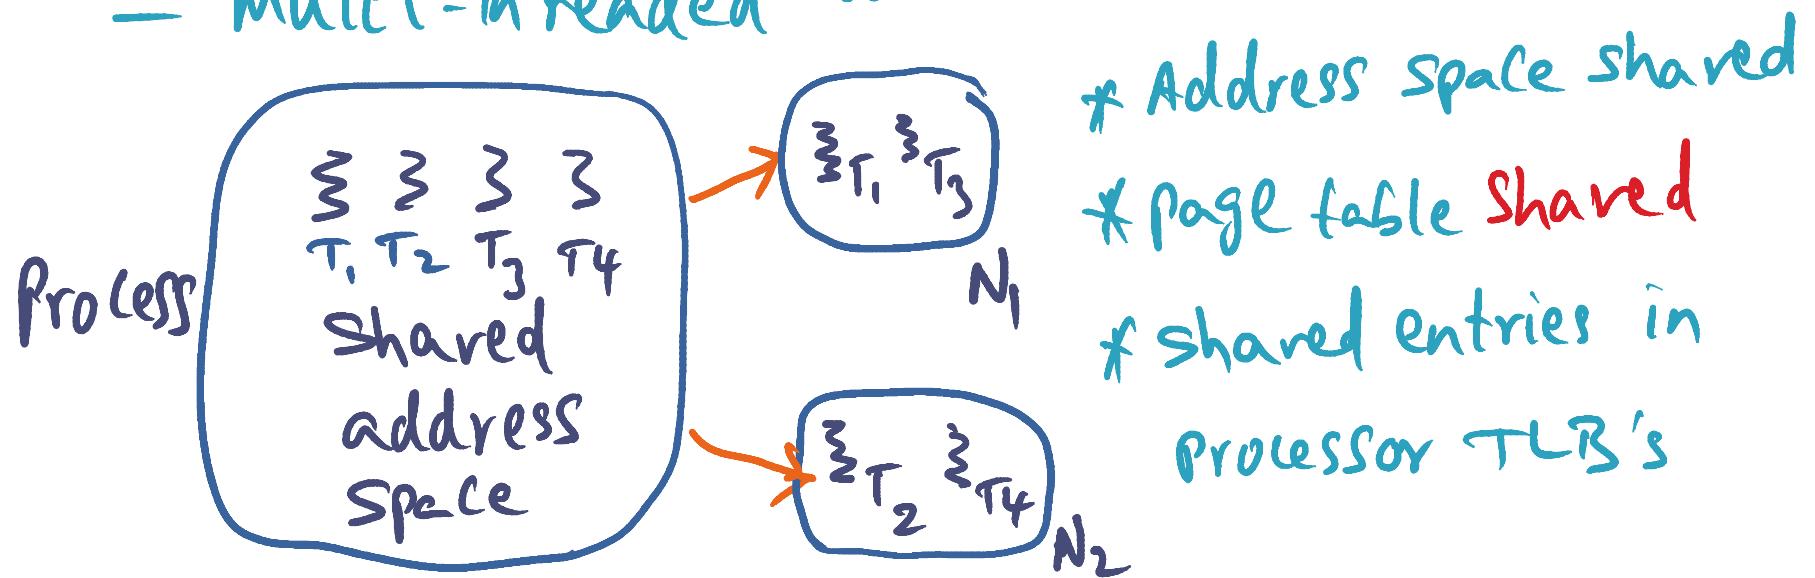
Easy scenario

- multiprocess workload
 - * Threads independent
 - * page tables distinct
 - * no serialization



Hard scenario

- multi-threaded workload



Recipe for Scalable Structure in Parallel OS

For Every Subsystem

- determine functionally needs of that service
- To ensure concurrent execution of service
 - * Minimize shared data structures

Less Sharing \Rightarrow more scalable

- where possible replicate/partition system data structures
 - \Rightarrow less locking
 - \Rightarrow more concurrency

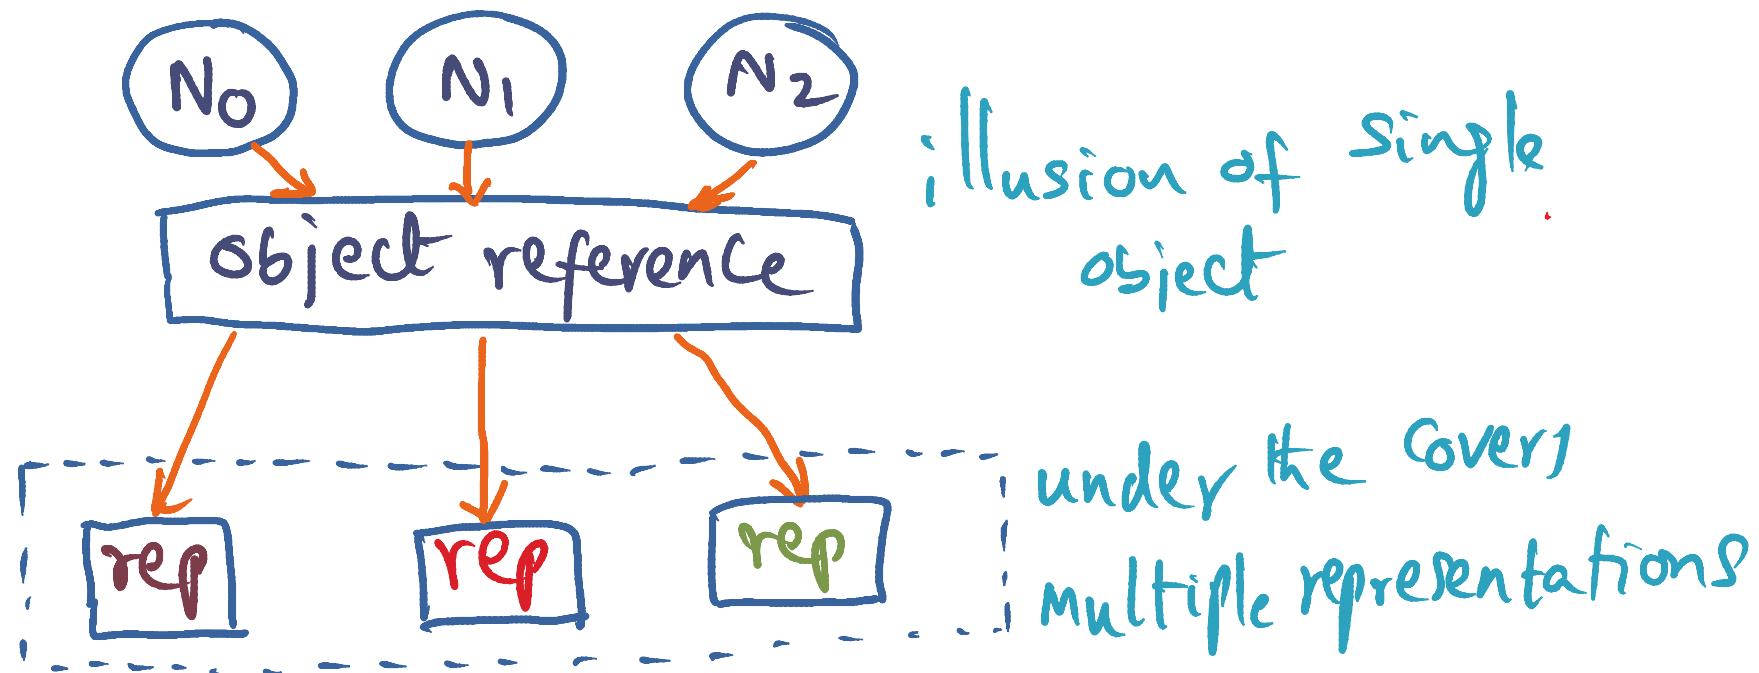


Watch !!



Watch !!
Did you ??

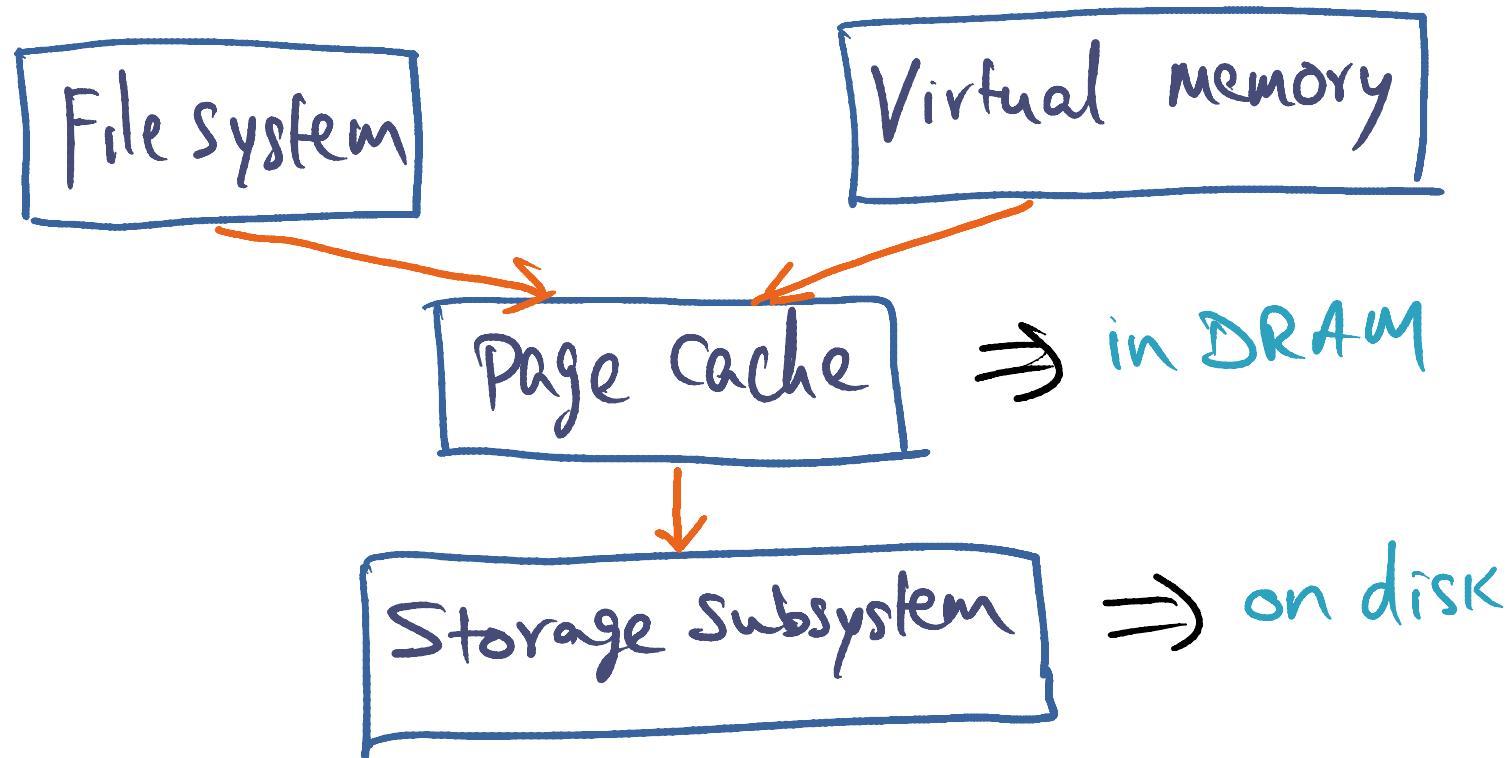
Tornado's secret sauce: clustered object



Degree of clustering?

- choice of implementor of service
 - * Singleton rep, one per core, one per CPU, one per group of CPUs, ...
- PPC for consistency

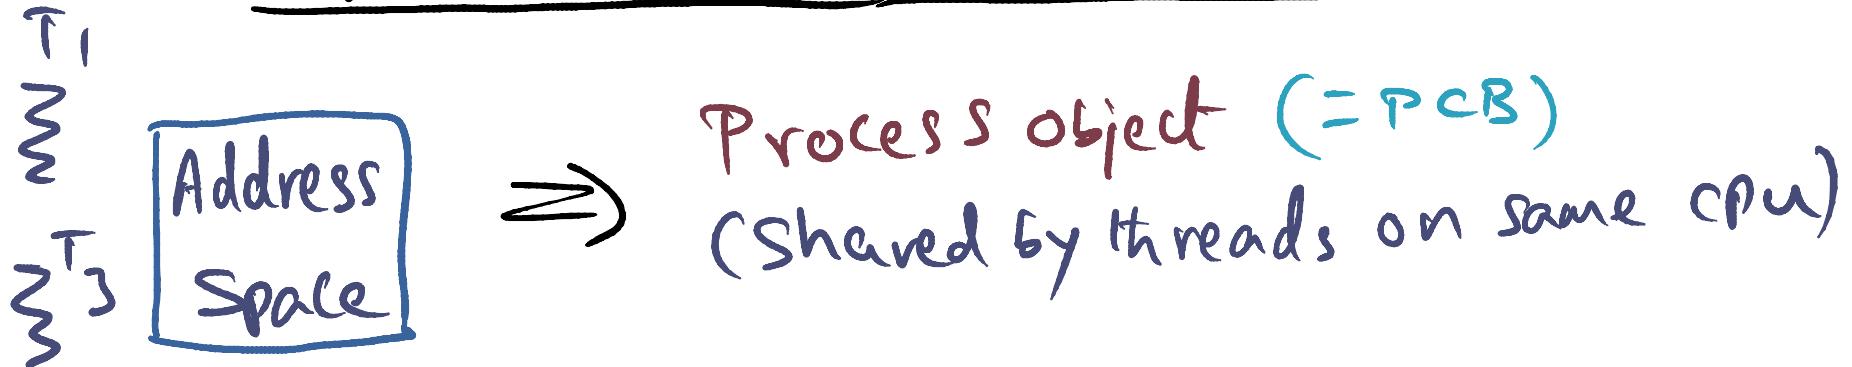
Traditional Structure



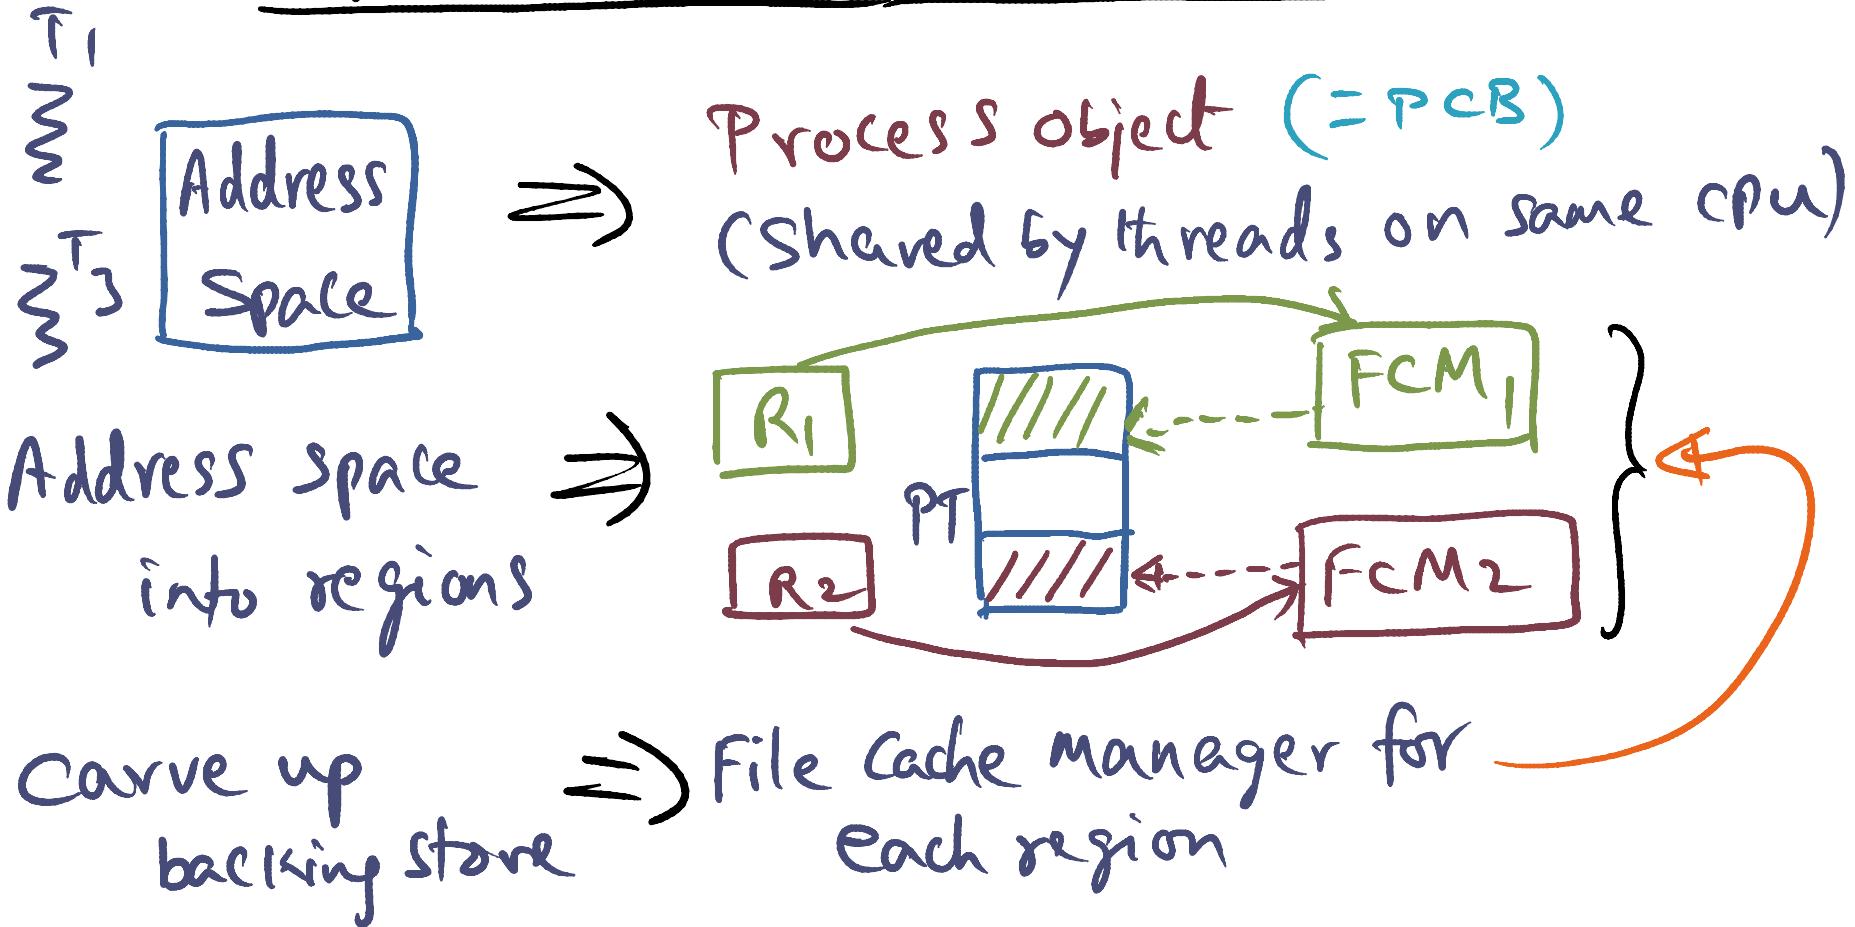
VM data structures

- PCB, TLB, PT
 - virtual pages on disk
- } => per process

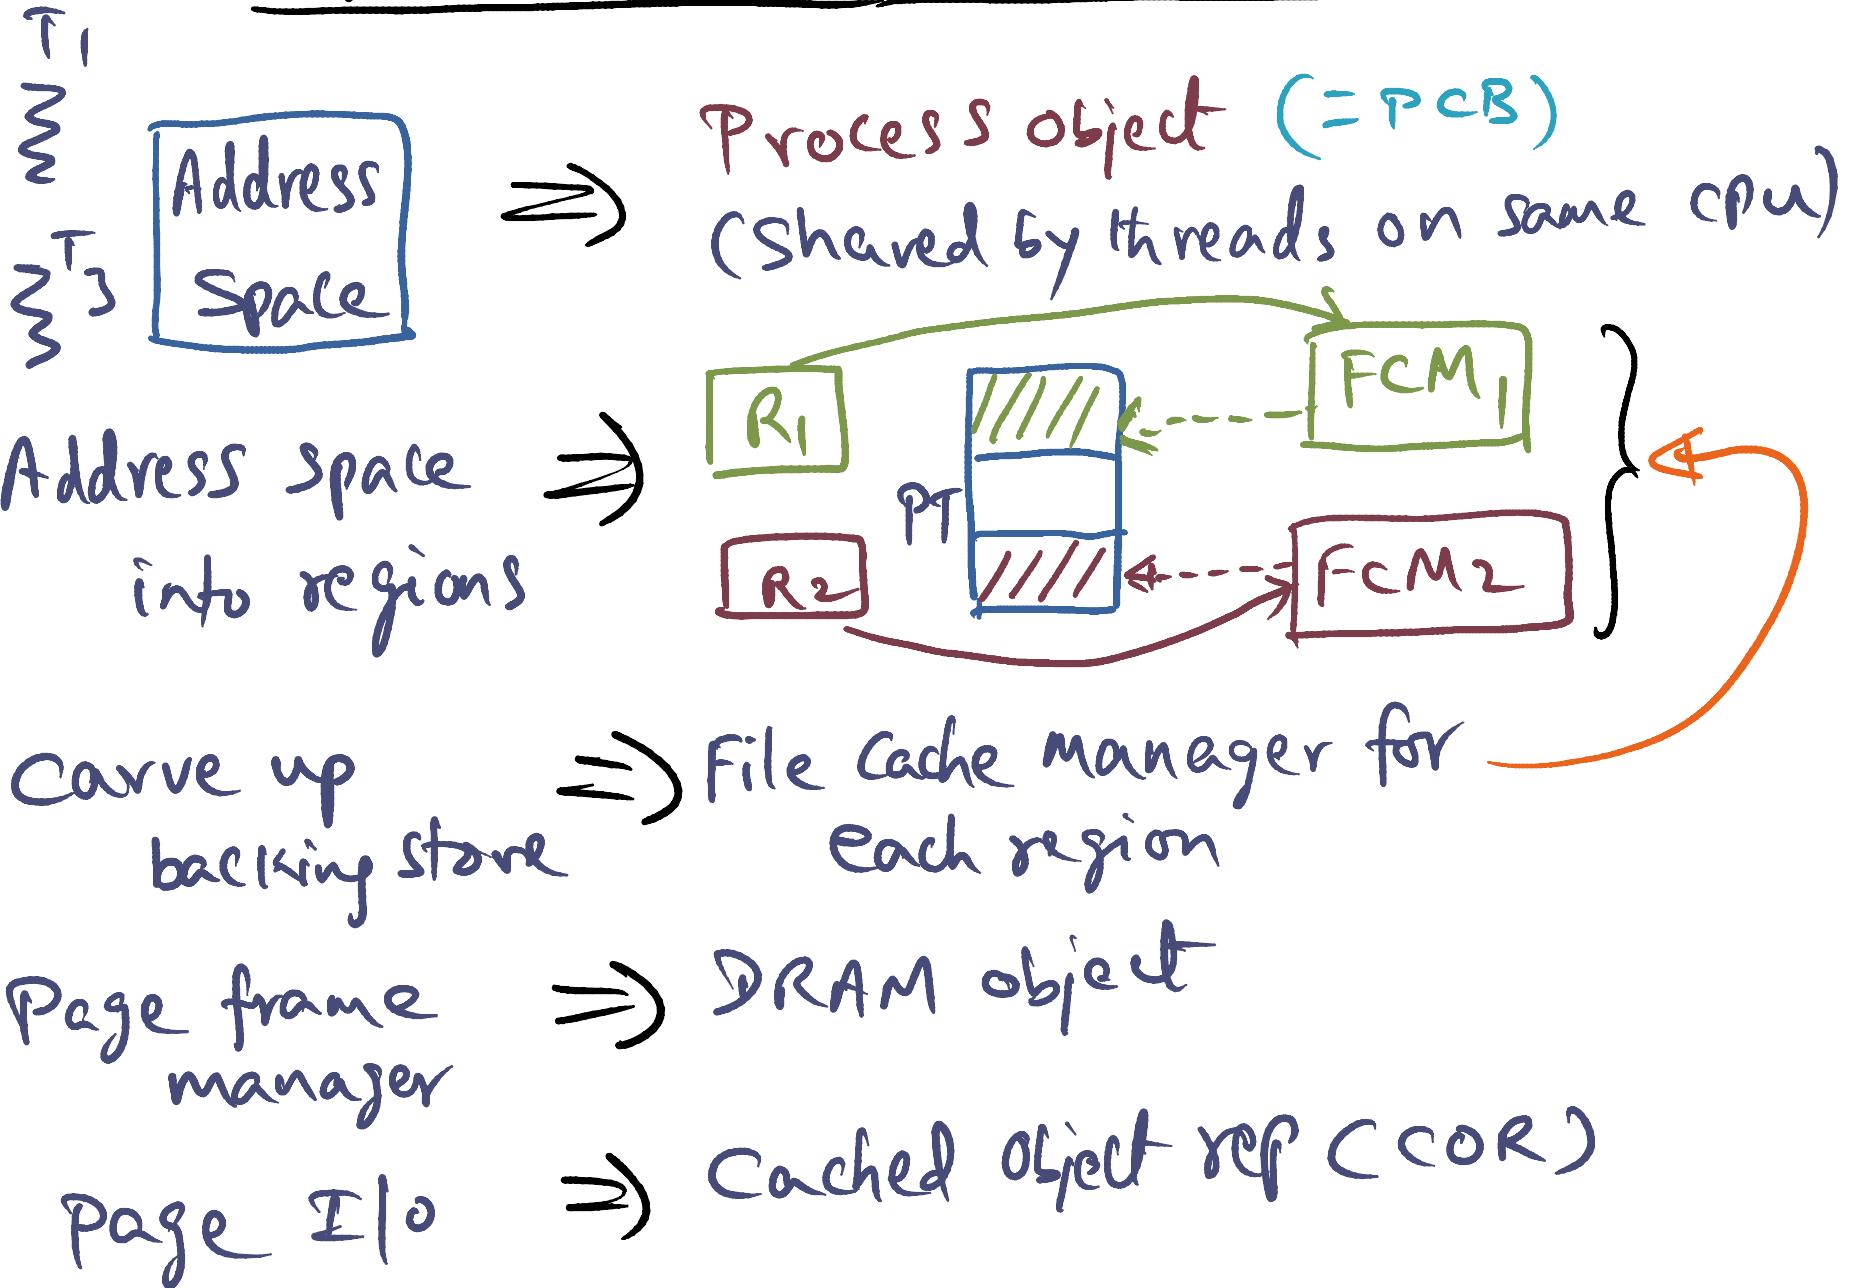
Objectification of memory management



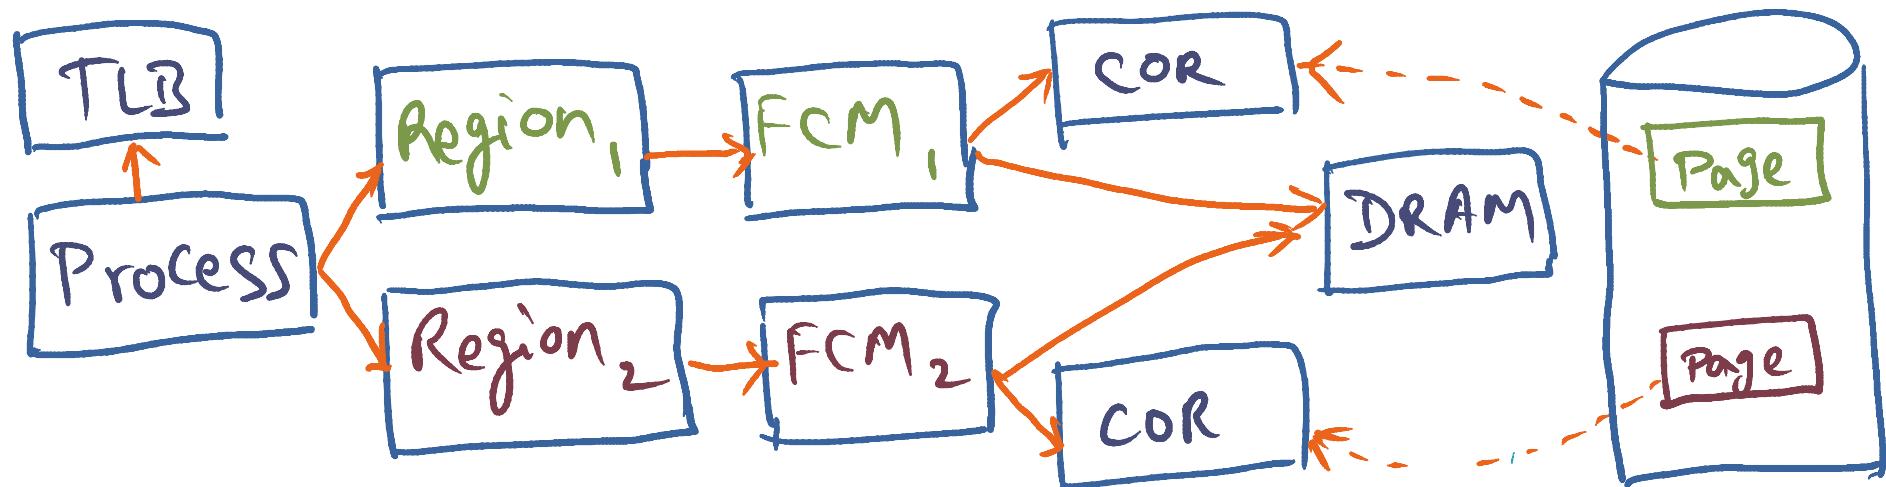
Objectification of memory management



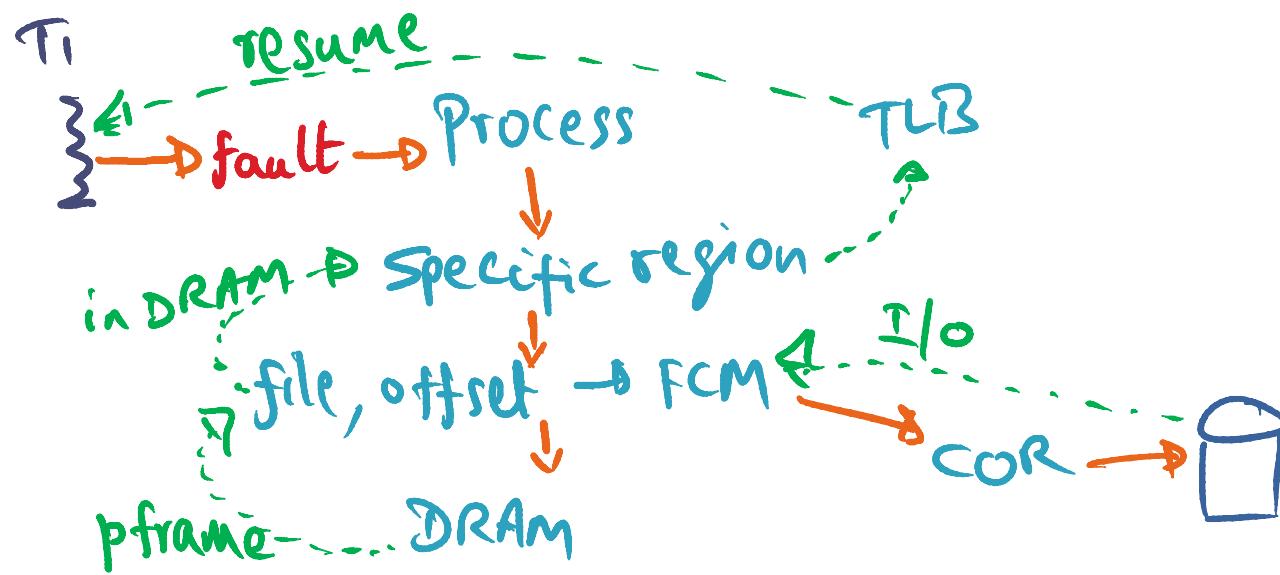
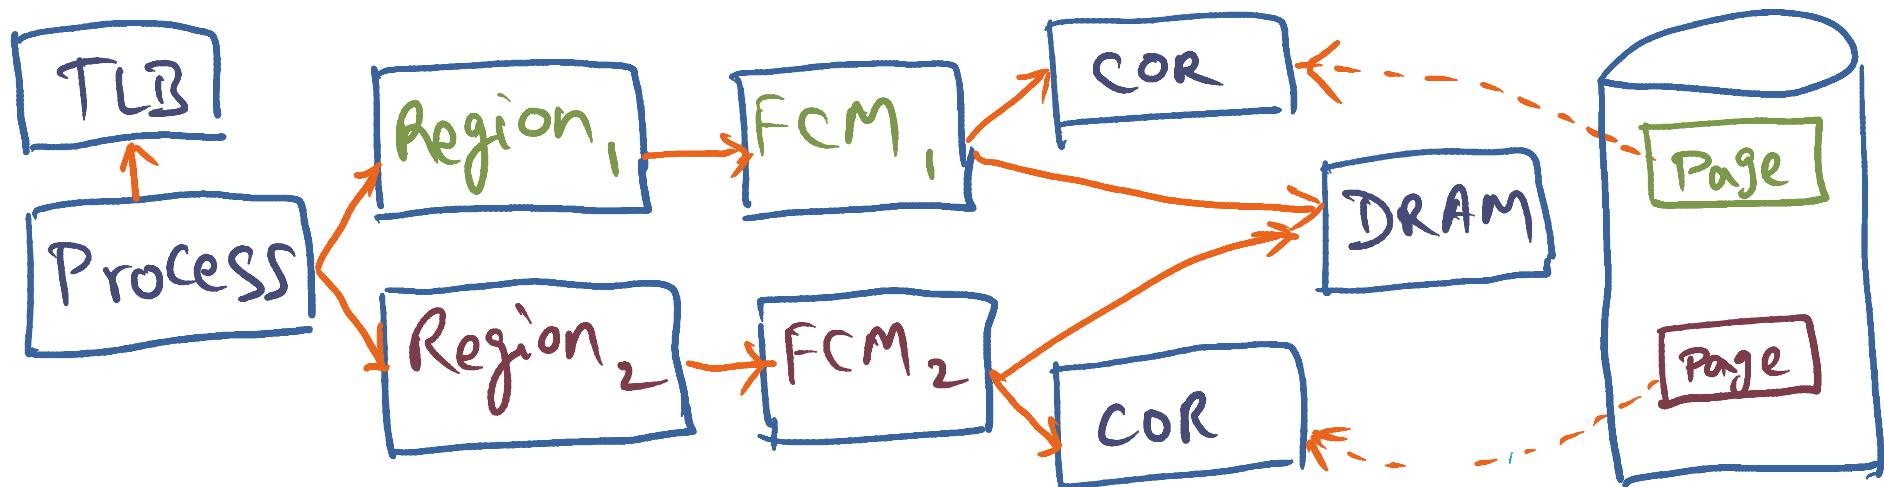
Objectification of memory management



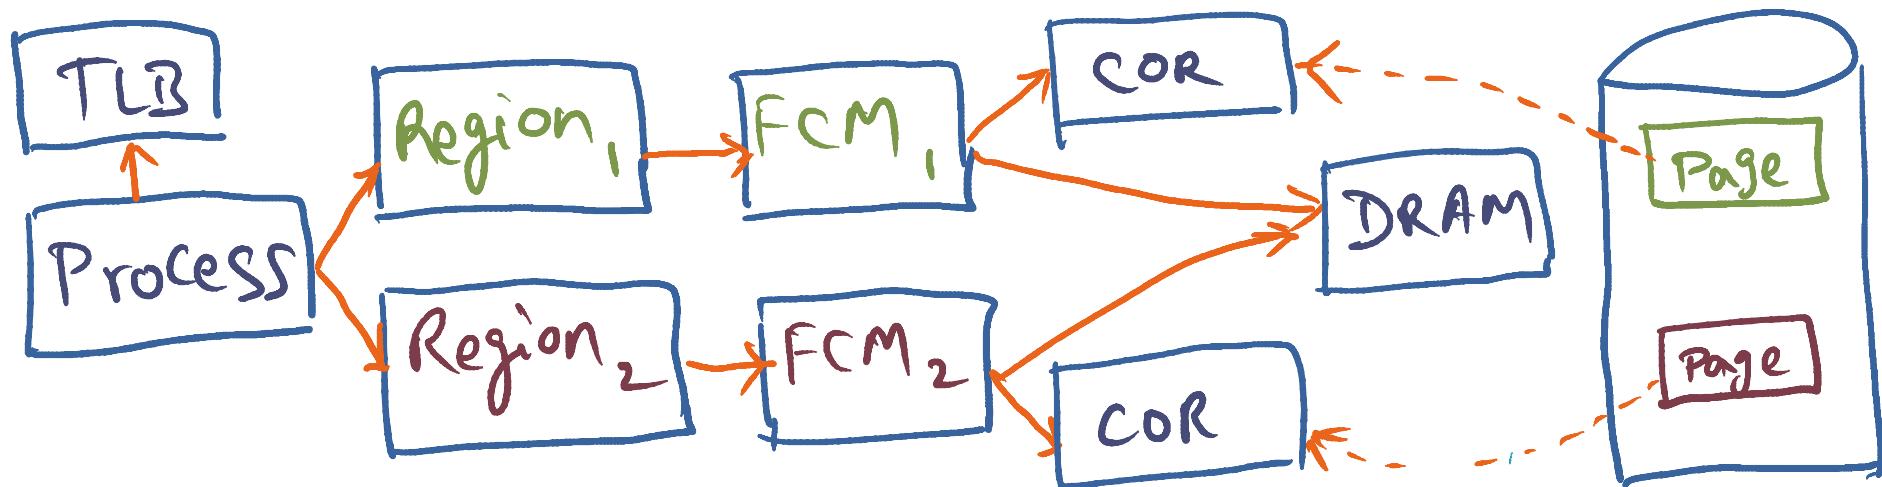
Objectified structure of VM manager



Objectified structure of VM manager

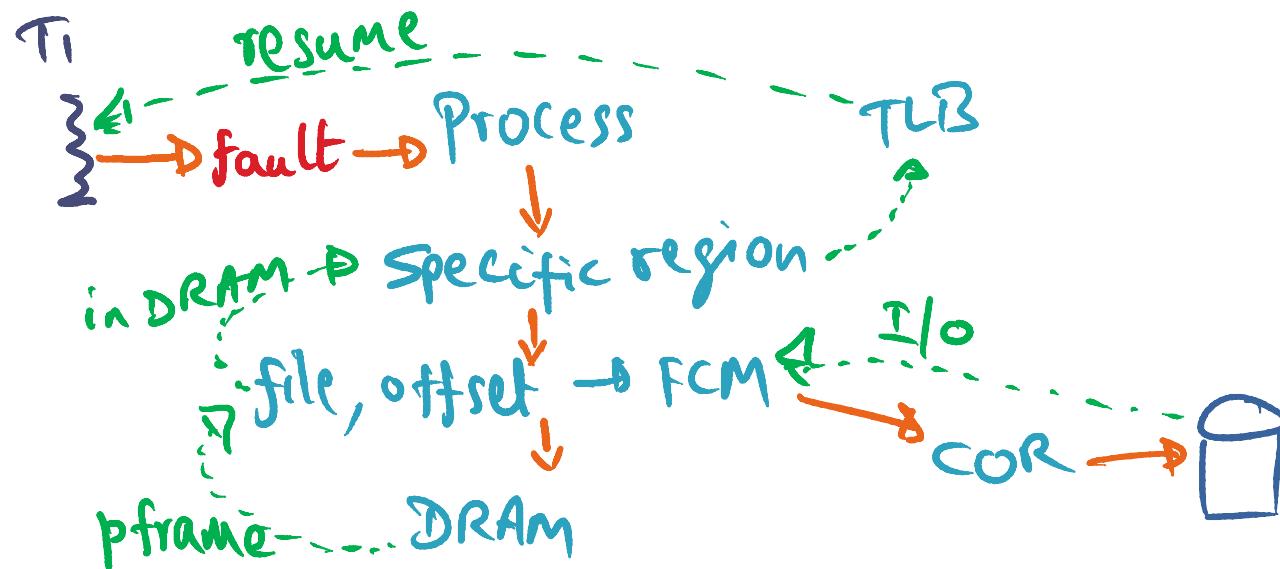
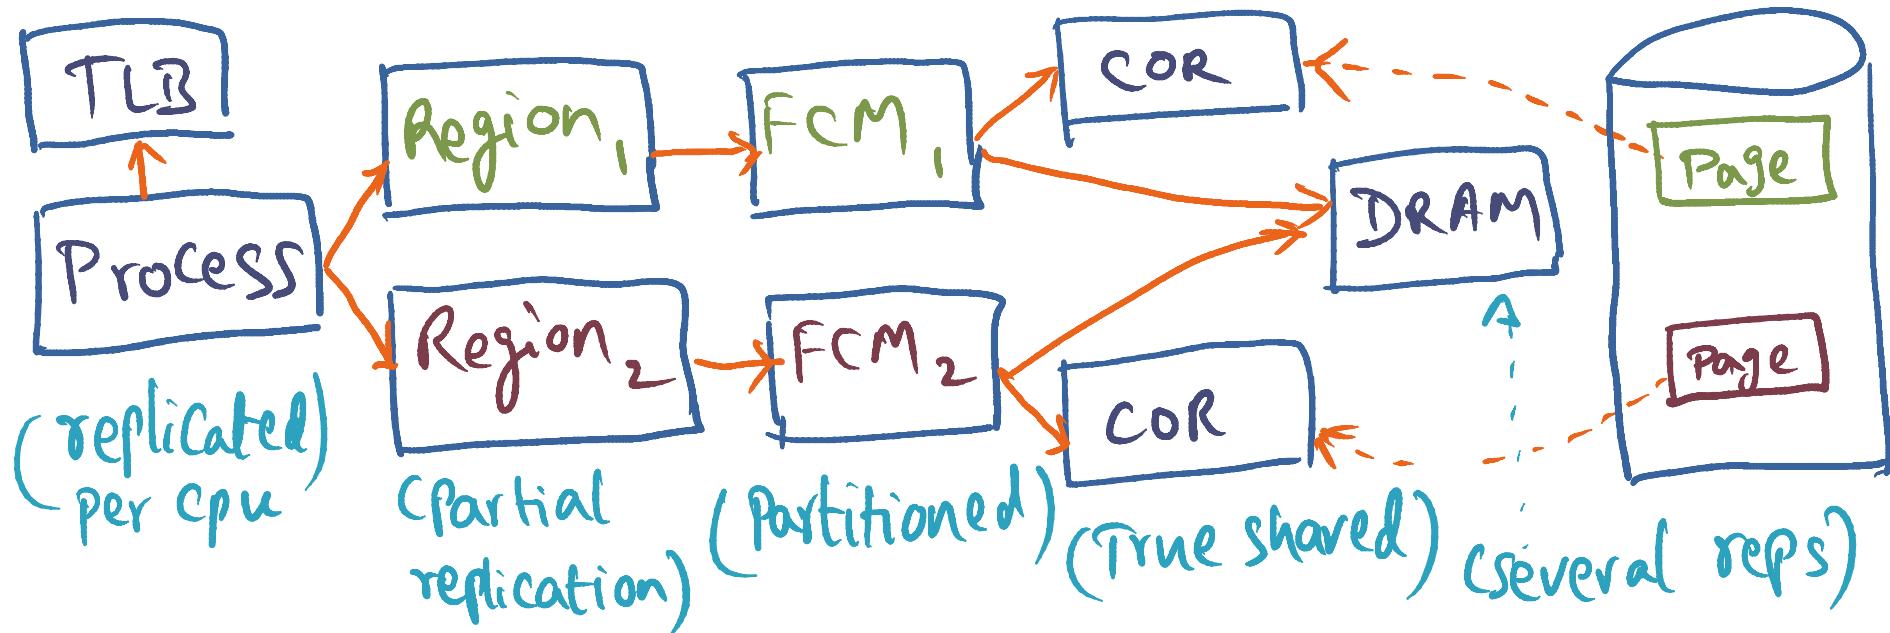


Group activity

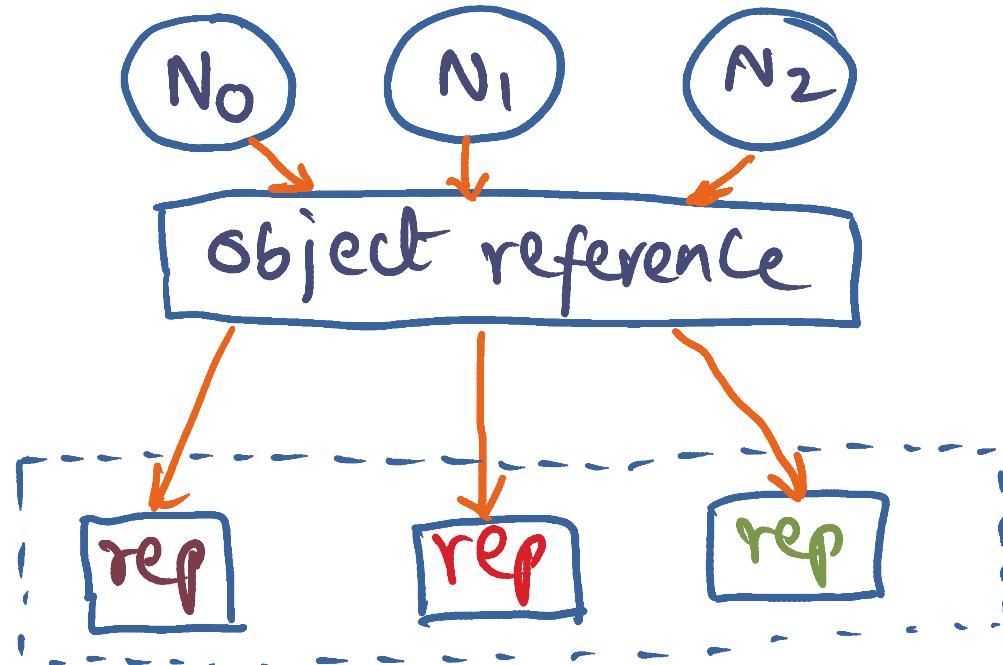


Starting with the above figure, discuss for each object in the path of the memory manager which clustering strategy (a single shared object, partial replication, full replication, partitioned reps) is most appropriate.

Objectified structure of VM manager



Advantages of clustered object



same object reference on all nodes

Allows incremental optimization

* usage pattern determines level of replication