

Today

Midterm: Oct 5

Synch + Comm. in Parallel Systems

\* MCS paper

✓ \* Parallel Systems Basics

✓ \* Spin lock algorithms

✓ \* Barrier algorithms

✓ - Counting, Sense, tree, MCS

✓ - tournament, dissemination

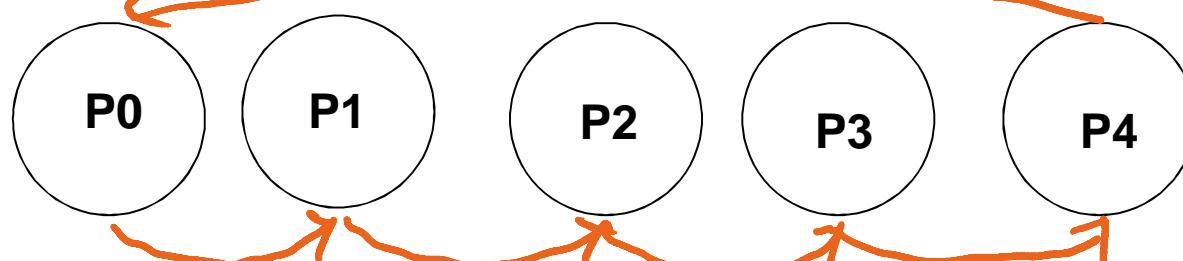
⇒ \* Cross domain communication (C�pc)

Wednesday

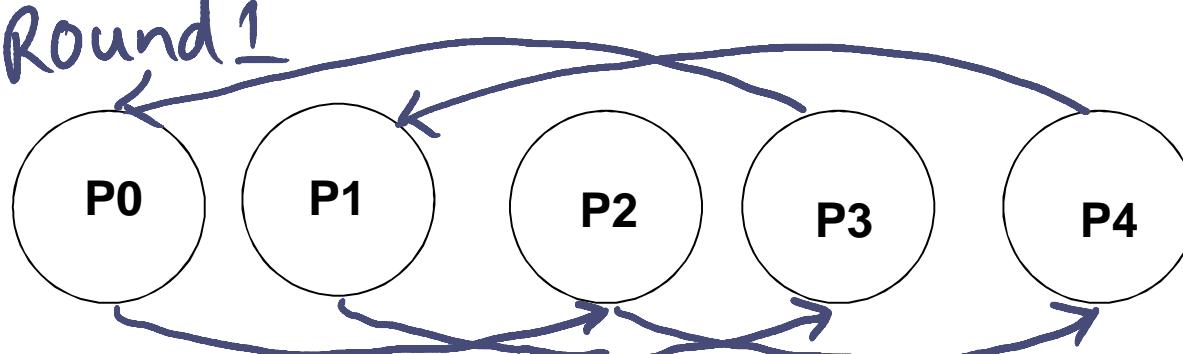
\* Sched in SM systems

# Dissemination Barrier

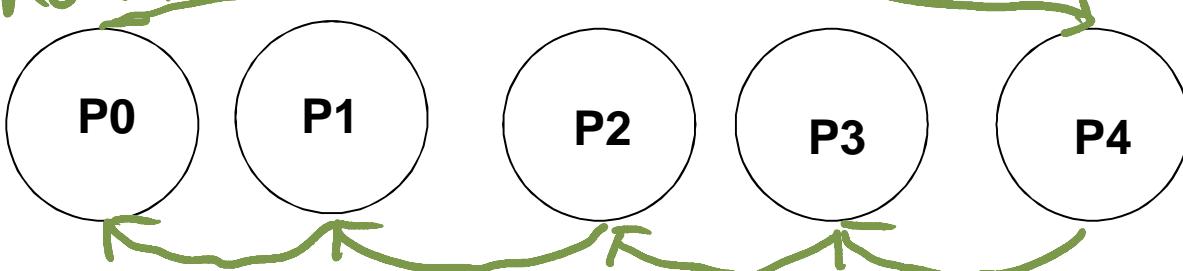
Round 0



Round 1



Round 2



N need not be power of 2

Each round(K):

$$P_i \rightarrow P_{(i+2^k) \bmod N}$$

$O(N)$  comm. events  
per round

-  $\lceil \log_2 N \rceil$  rounds

All Awake!

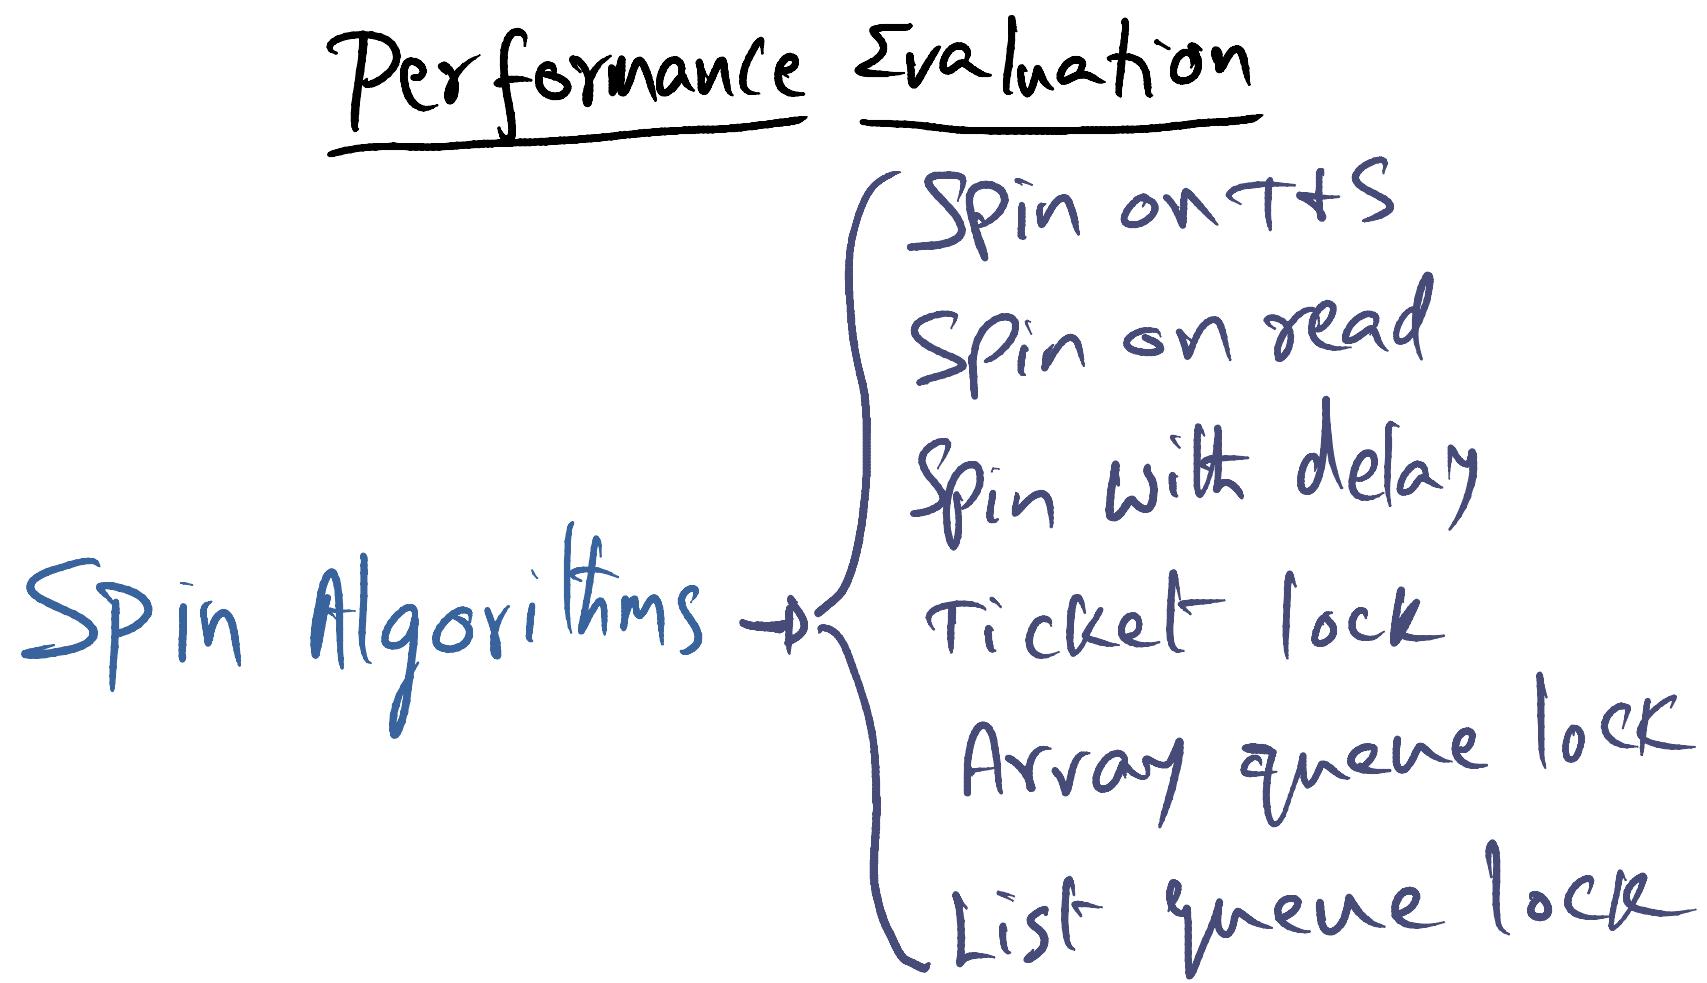
- No hierarchy

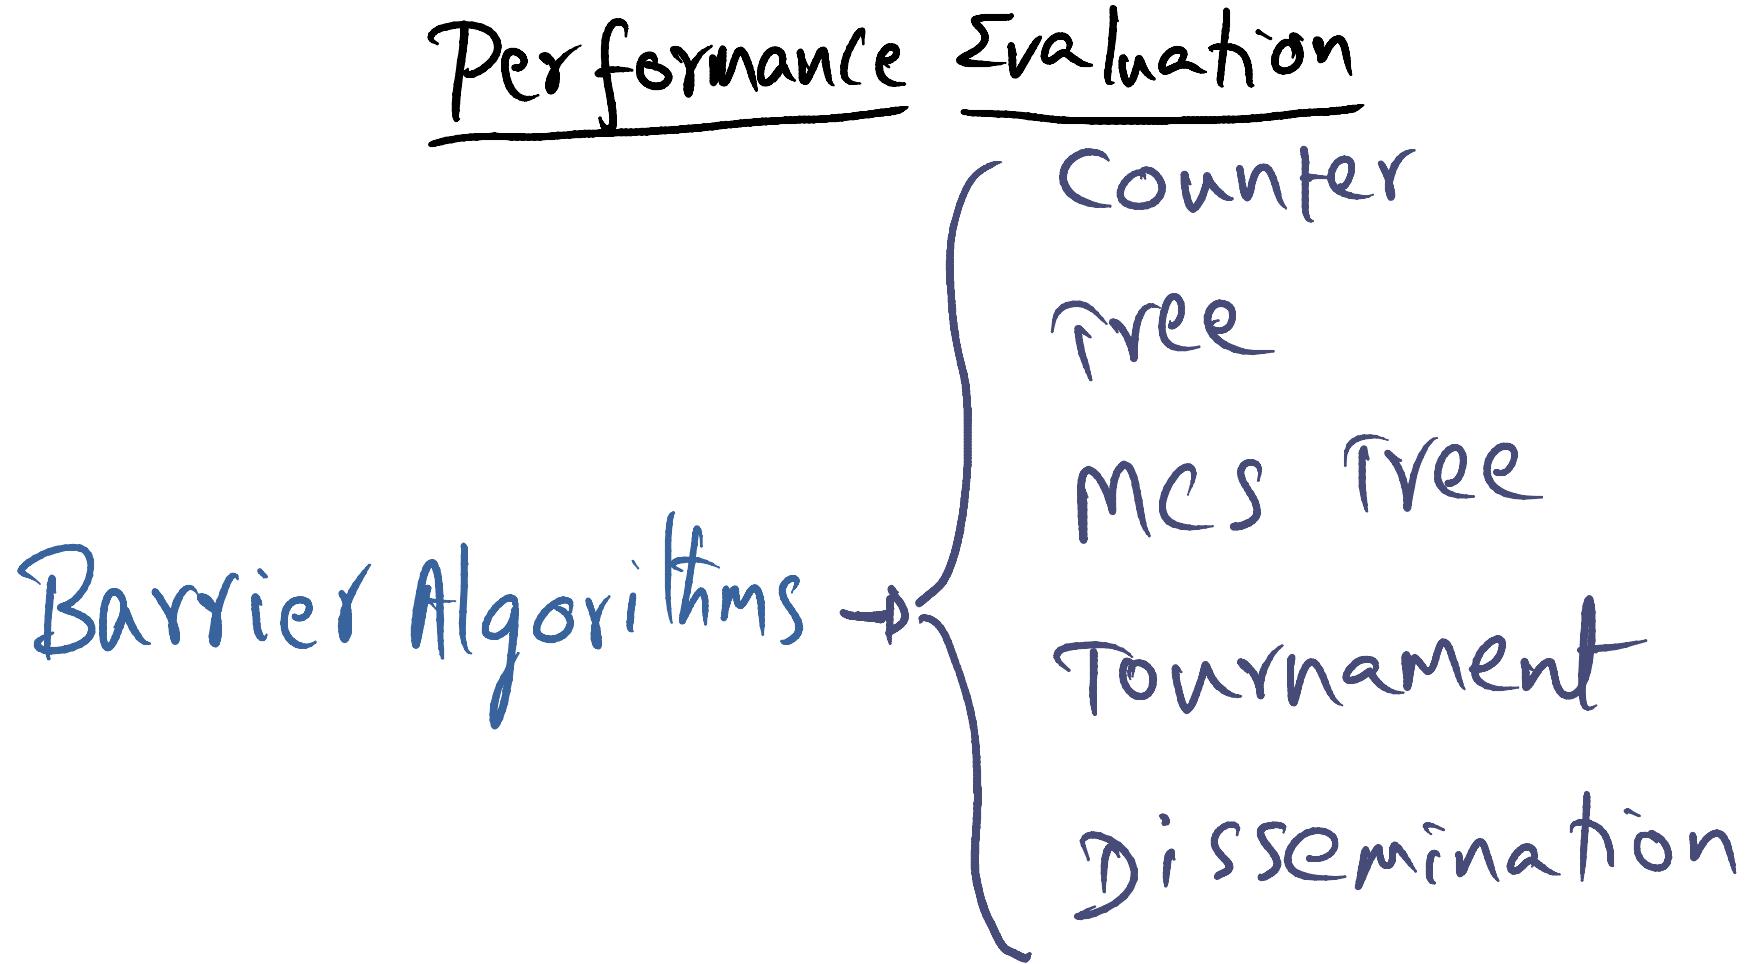
- works for NCC  
and clusters

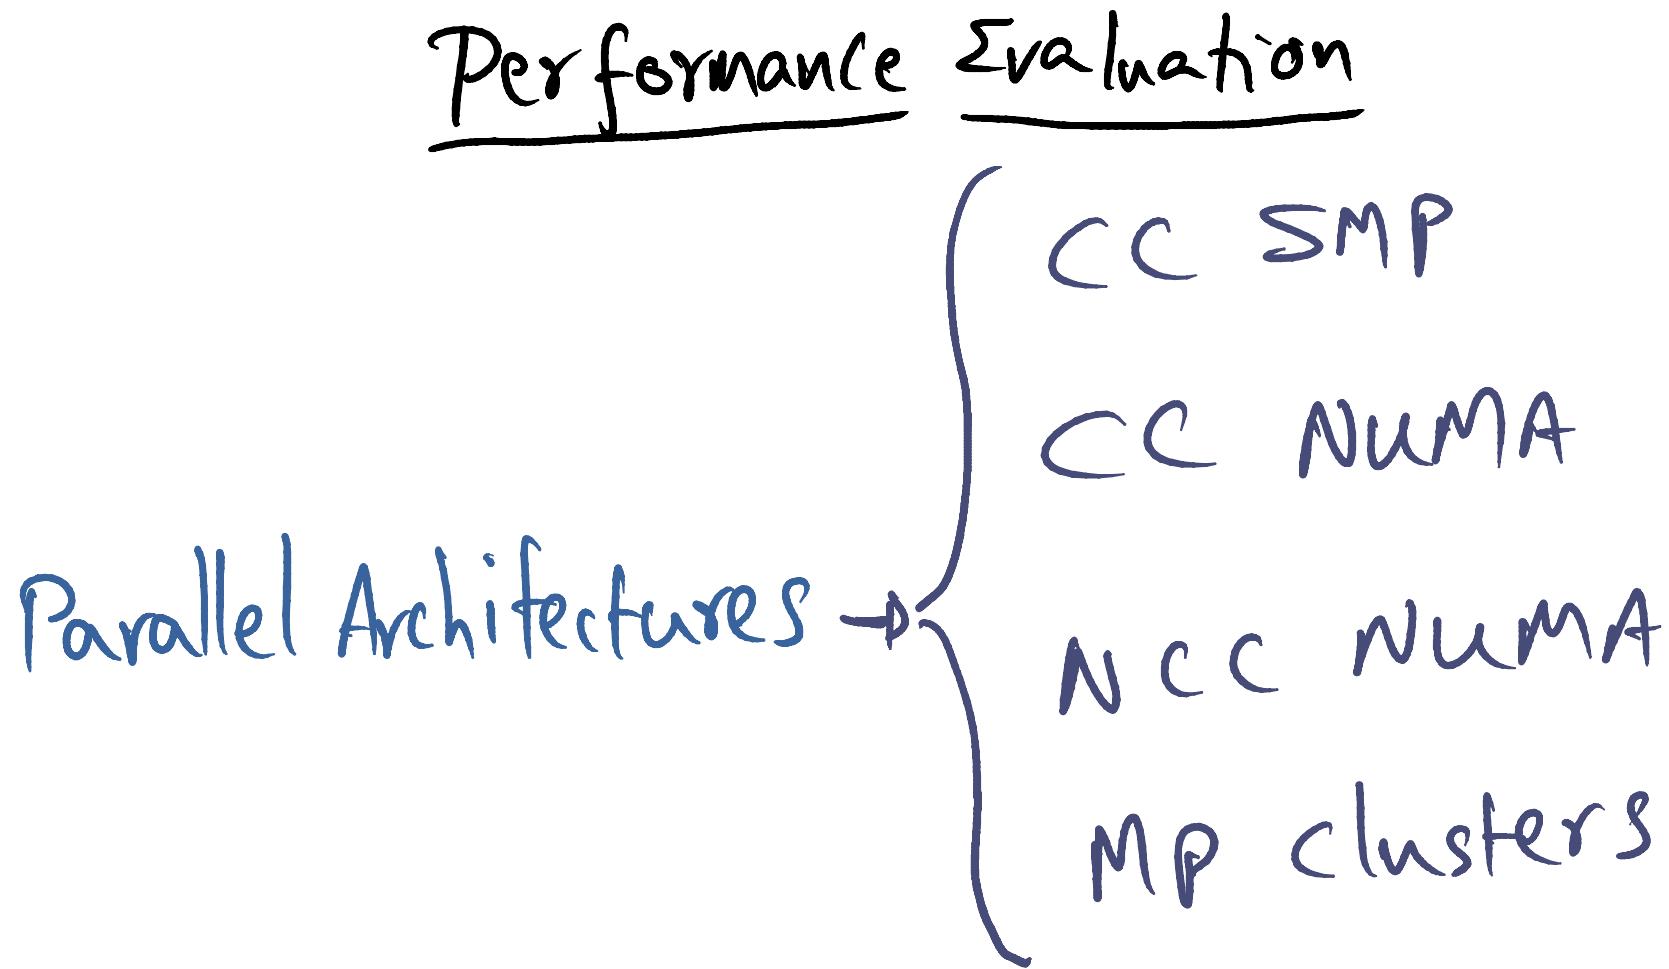
# Virtues of Dissemination

- No hierarchy
- No pair wise sync
- Each node works independently to send out the messages per protocol
- All will realize the barrier is complete when they have received the  $\text{ceil}(\log_2 N)$  messages from their peers
- Total amount of comm:  $O(N \log_2 N)$
- In comparison, tournament incurs  $O(\log_2 N)$  comm since in that set up the comm decreases as we go up the tree; since dissemination is based on gossip, the amount of comm in each round is constant and equal to  $N$
- Works for NCC and MP machines in addition to SM

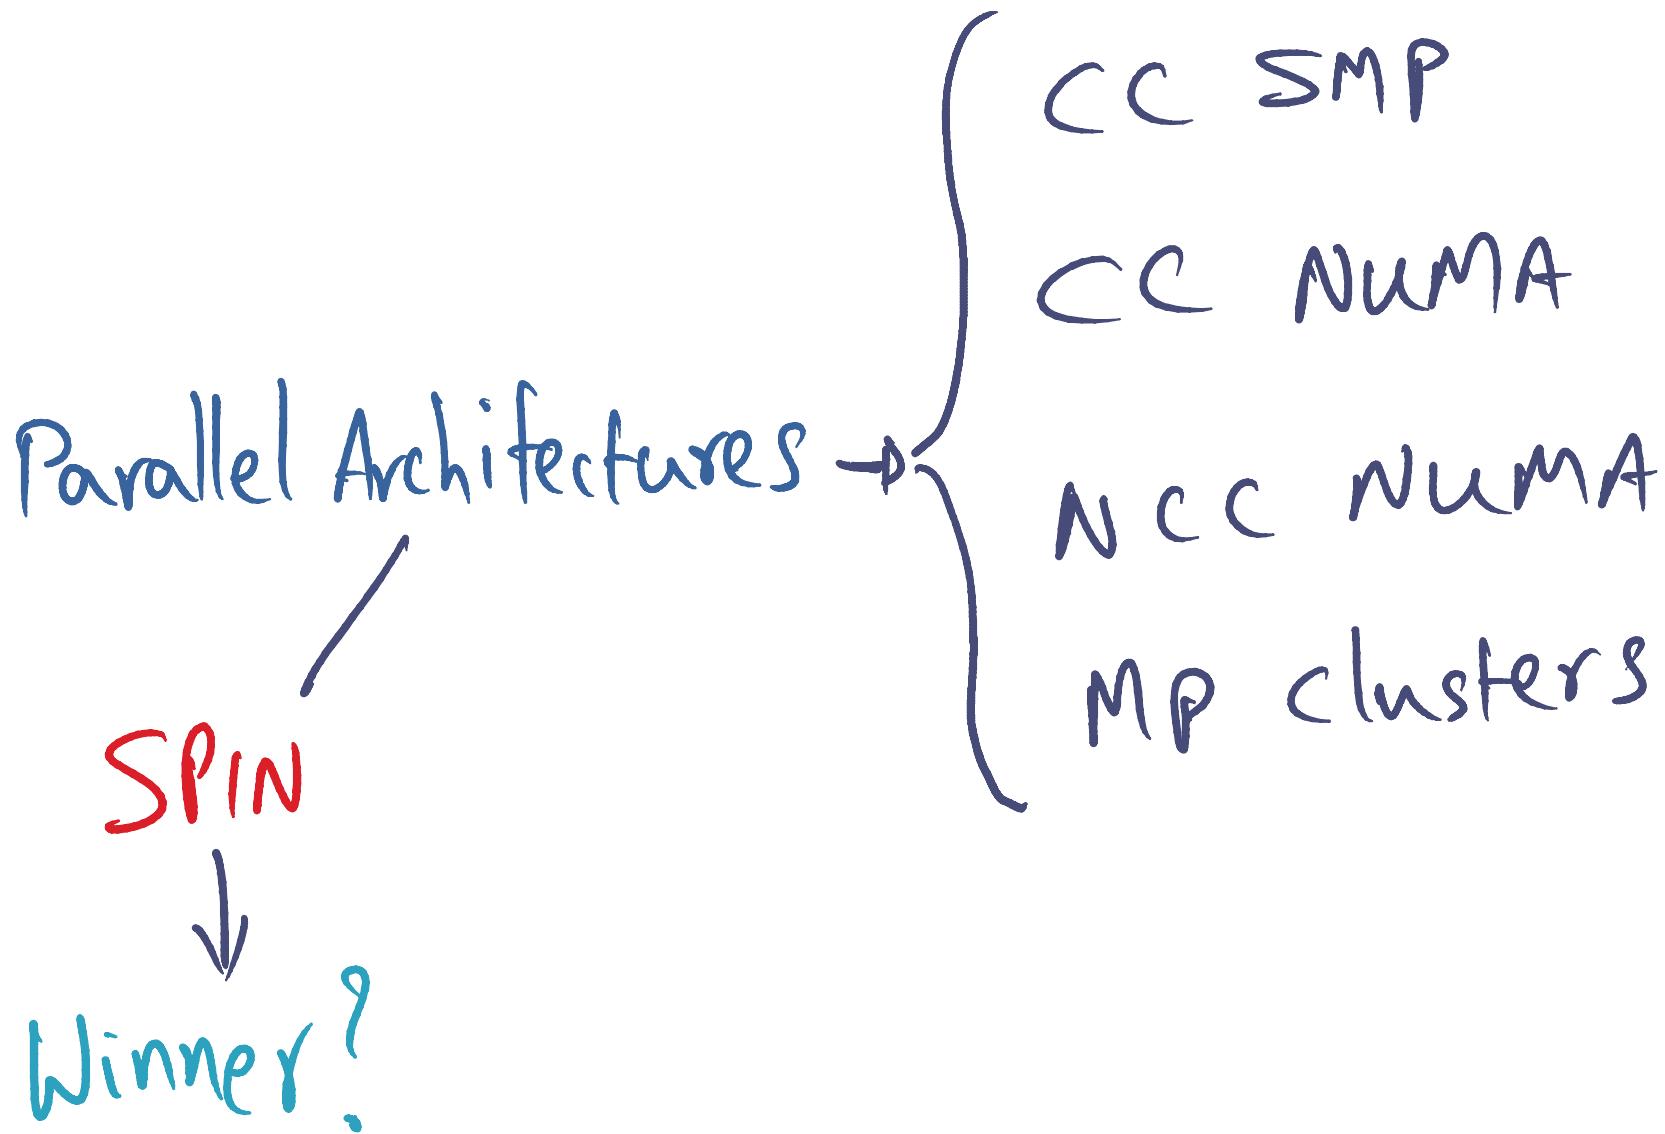
# Performance Evaluation

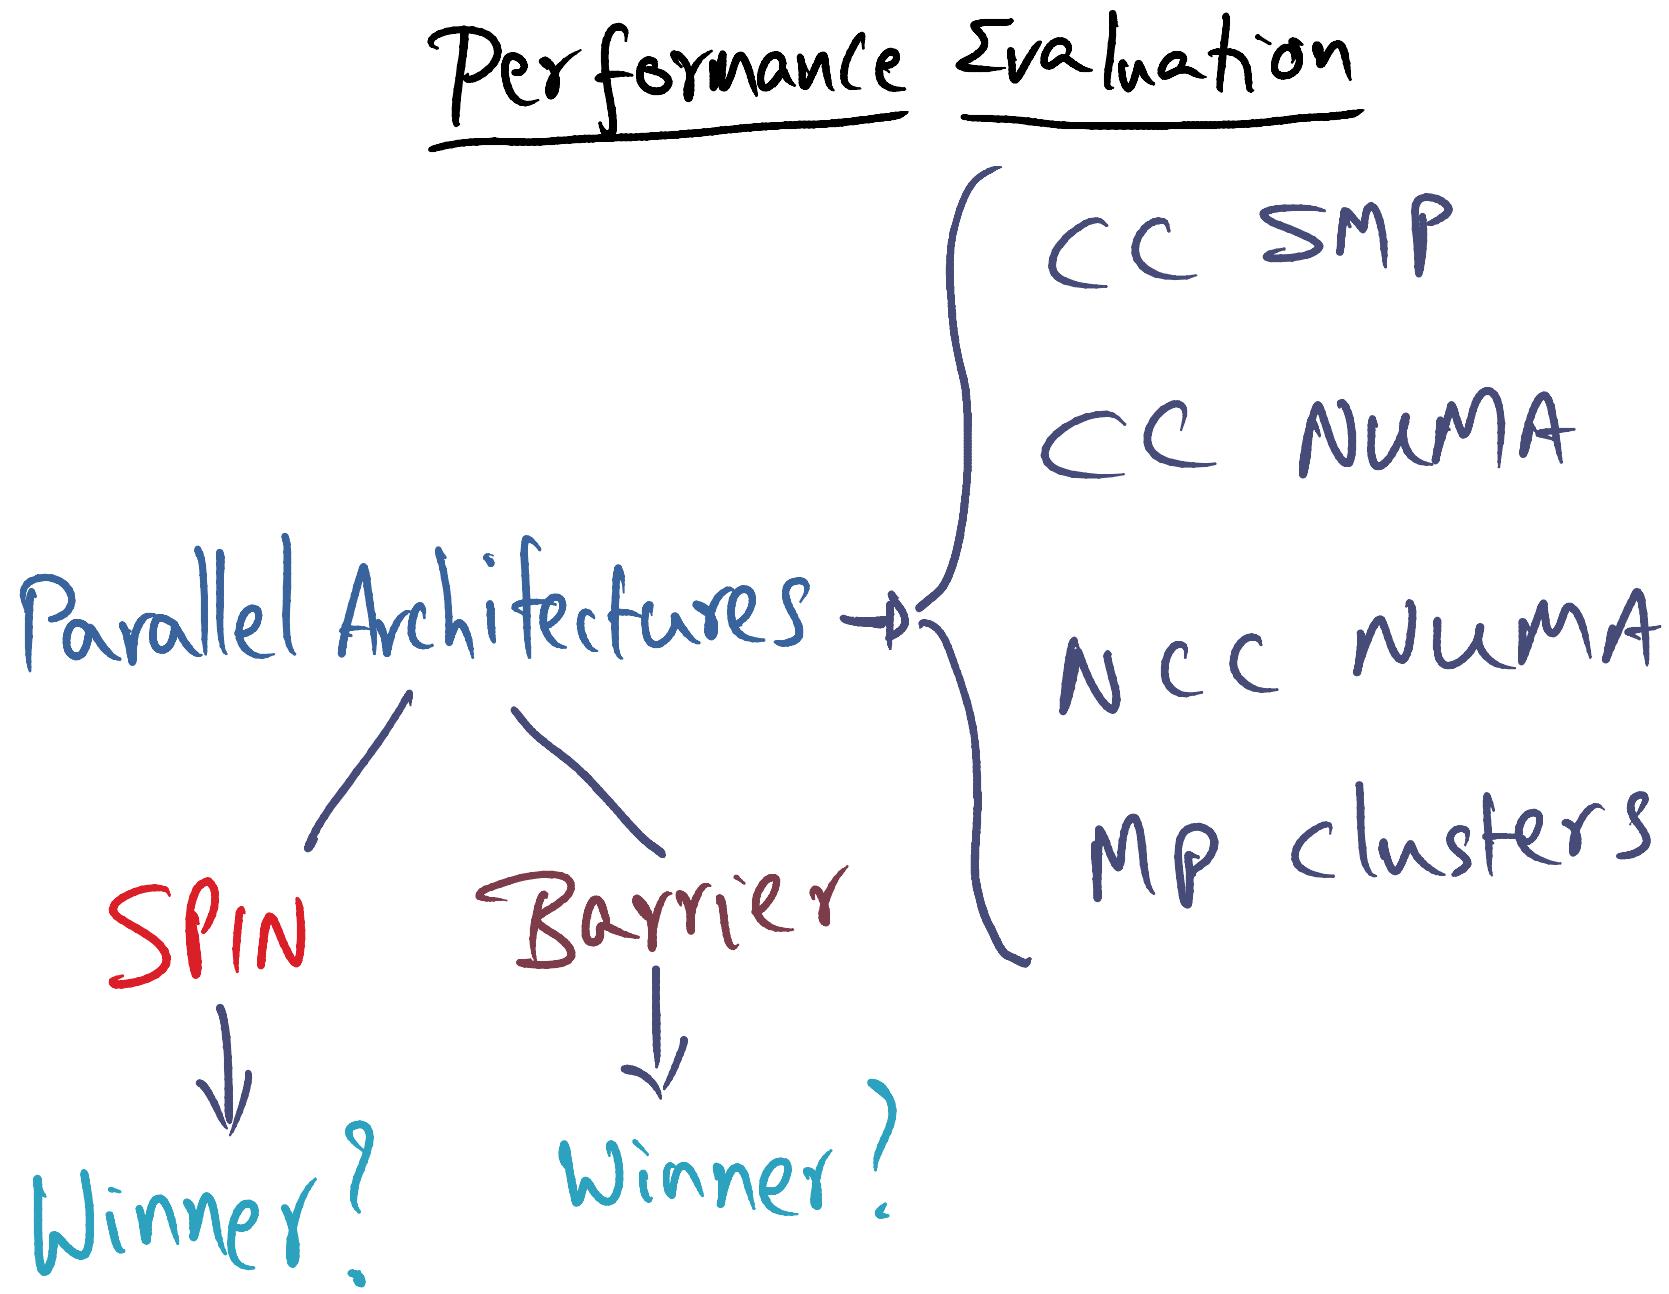






## Performance    Evaluation





# BBN Butterfly

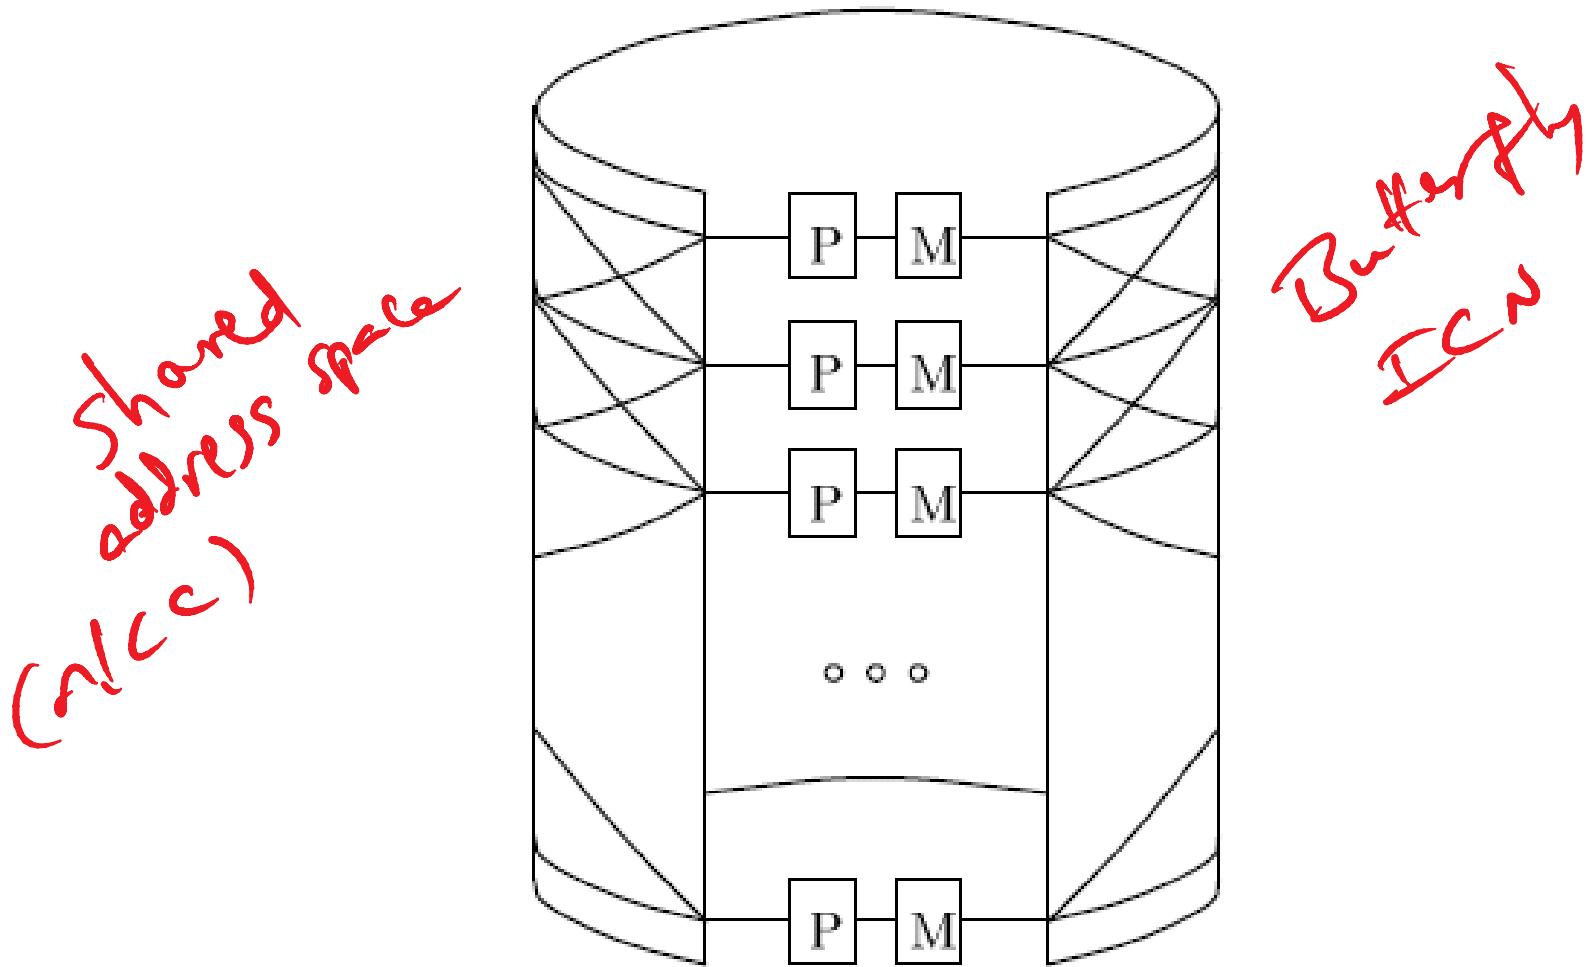


Figure 2: The BBN Butterfly 1.

# Sequent Symmetry

CC-SMP  
with shared bus or  
ICP

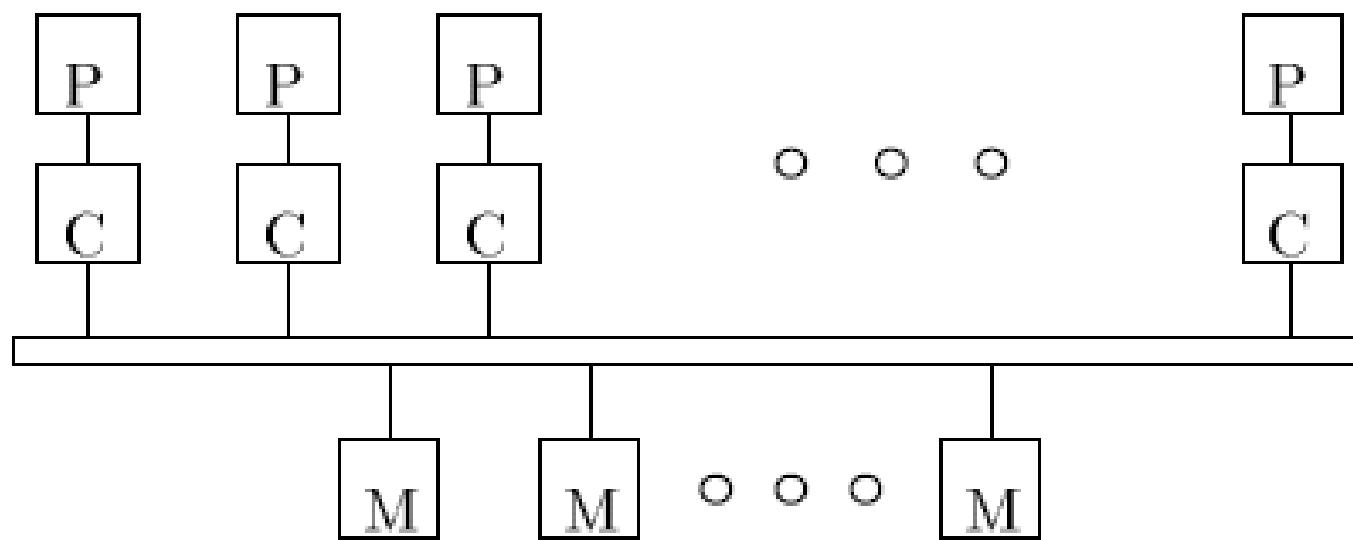


Figure 3: The Sequent Symmetry Model B.

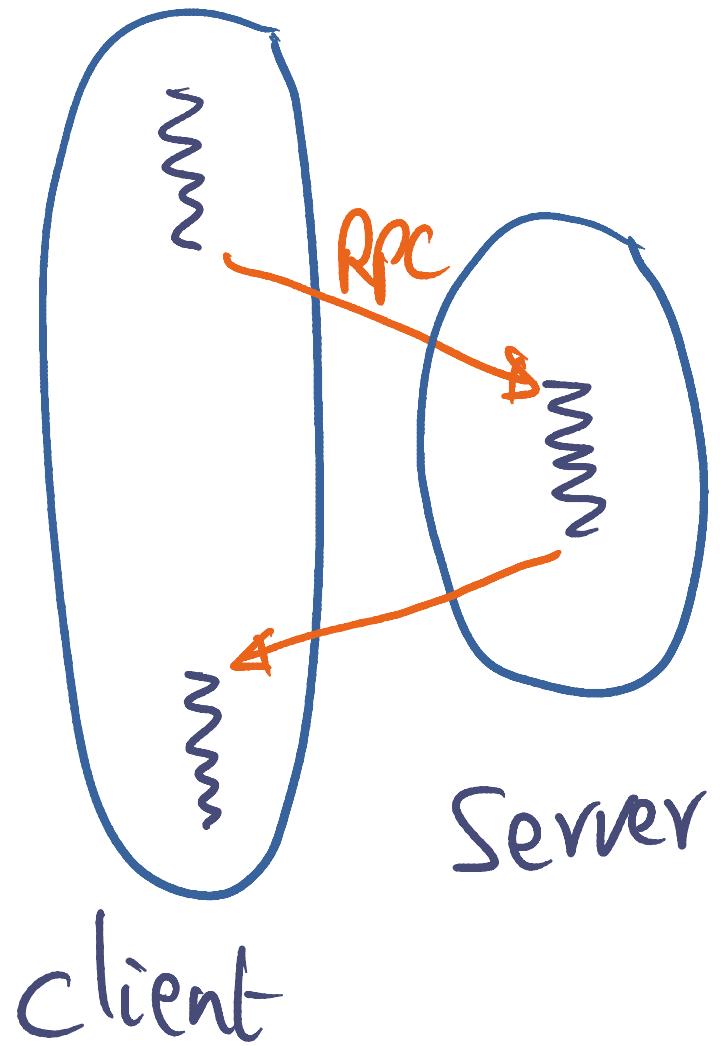
# Performance

- Spinlock
  - Which one do you think will win on symmetry?
  - Which one do you think will win on BBN?
- Barrier
  - Which one do you think will win on symmetry?
  - Which one do you think will win on BBN?

*- form an intuition  
- verify whether reported  
result agree with  
your opinion*

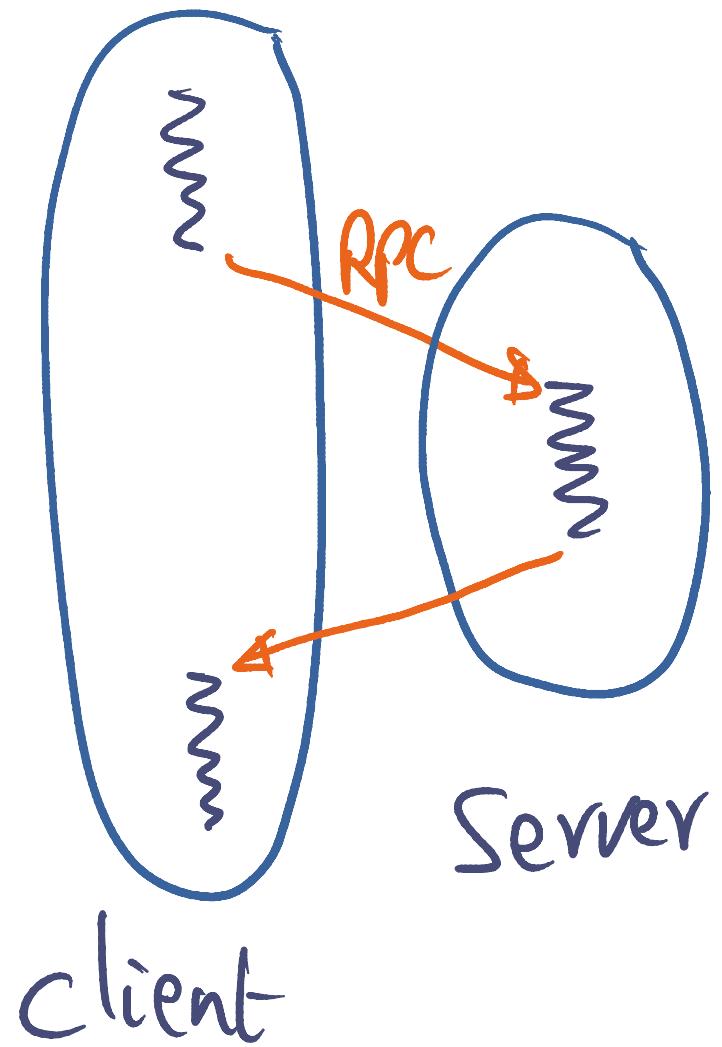
**Compare their barrier and lock results on symmetry and BBN against the results reported for the same algorithms on a different parallel machine, namely, KSR-1 in the paper “Scalability study of KSR-1”**  
**(See T-Square Modules: Section 10.10)**

# RPC and client - server Systems



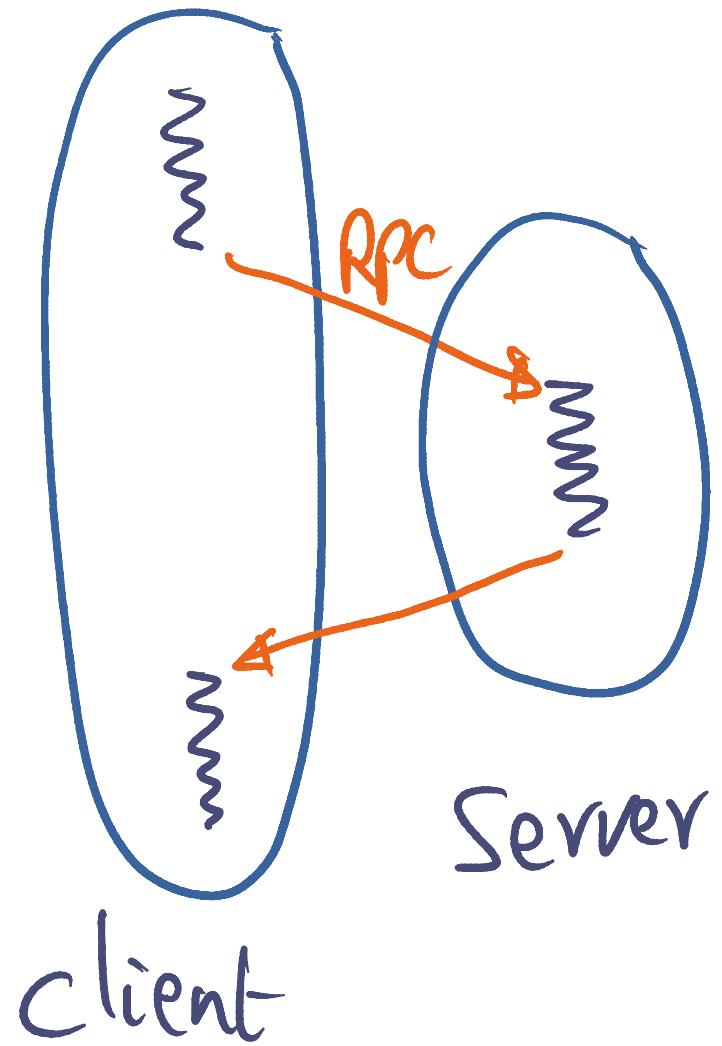
# RPC and client - server Systems

E.g. File System



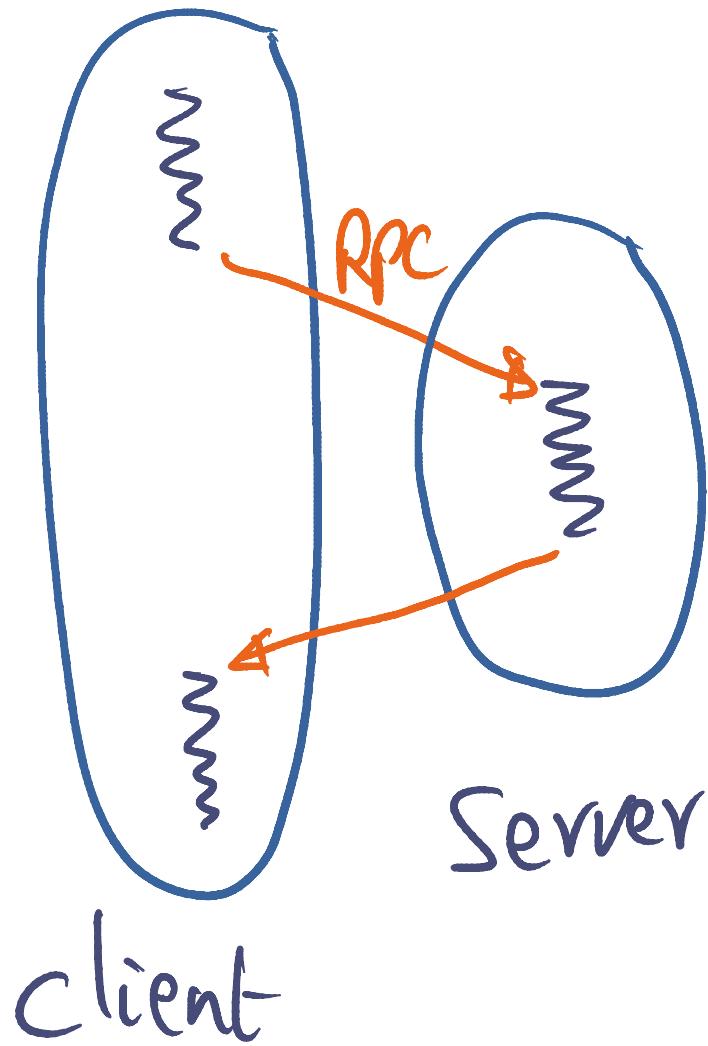
# RPC and client - server Systems

E.g. File System



RPC: usually remote

# RPC and client - server Systems

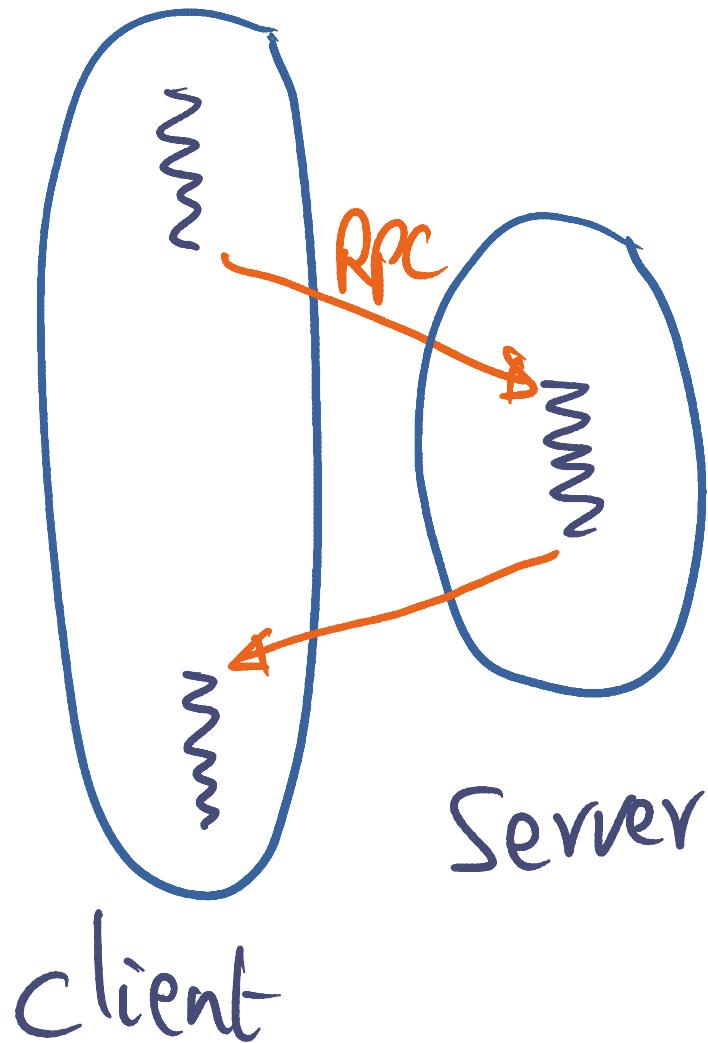


E.g. File System

RPC: usually remote  
client - Server on  
same machine?

# RPC and client - server Systems

---



E.g. File System

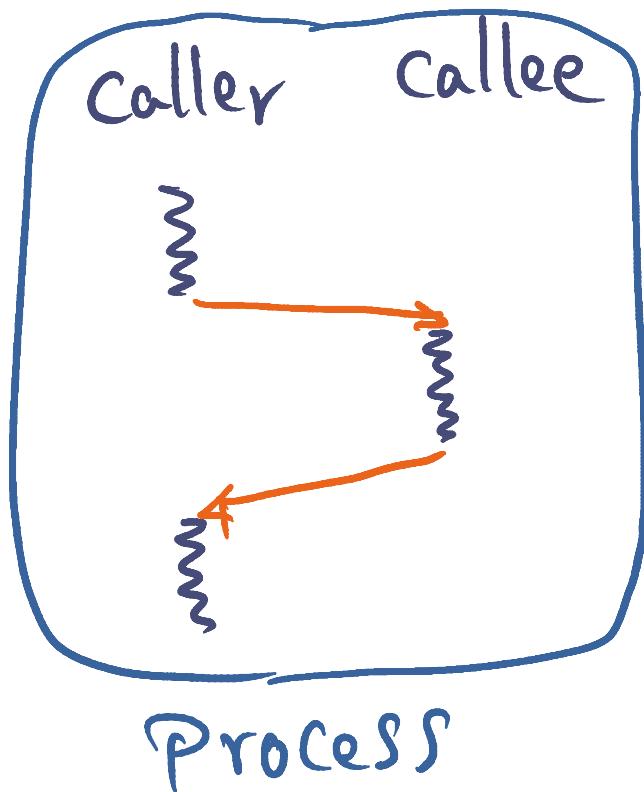
RPC: usually remote

client - Server on  
same machine?

— Performance  
Vs.  
Safety

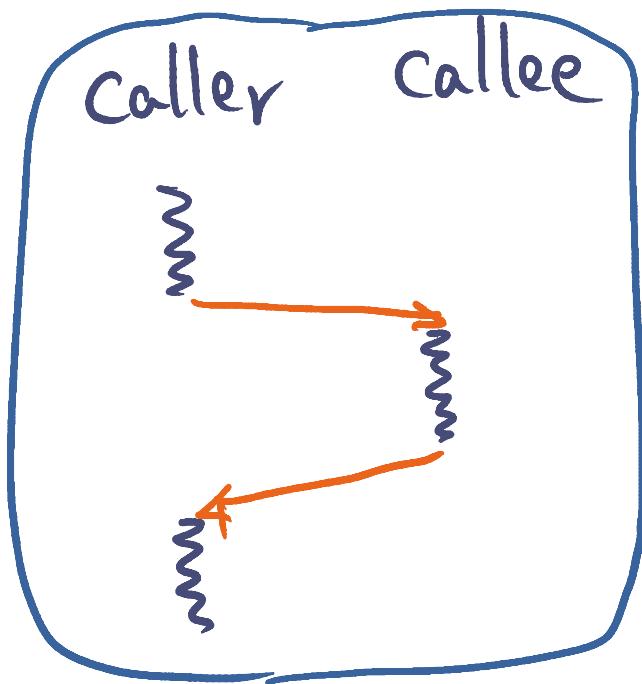
# RPC Vs. Simple Procedure Call

## Procedure Call



# RPC Vs. Simple Procedure Call

## Procedure Call

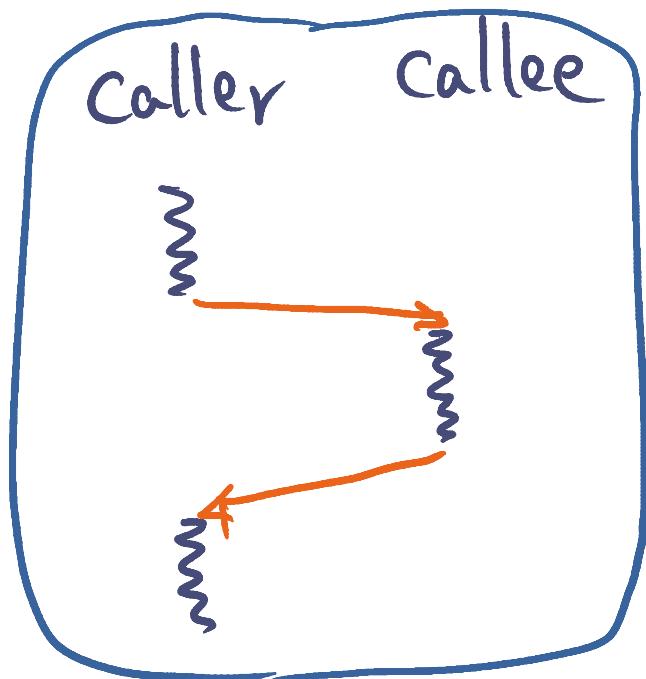


Process

- All at  
Completion

# RPC Vs. Simple Procedure Call

Procedure Call

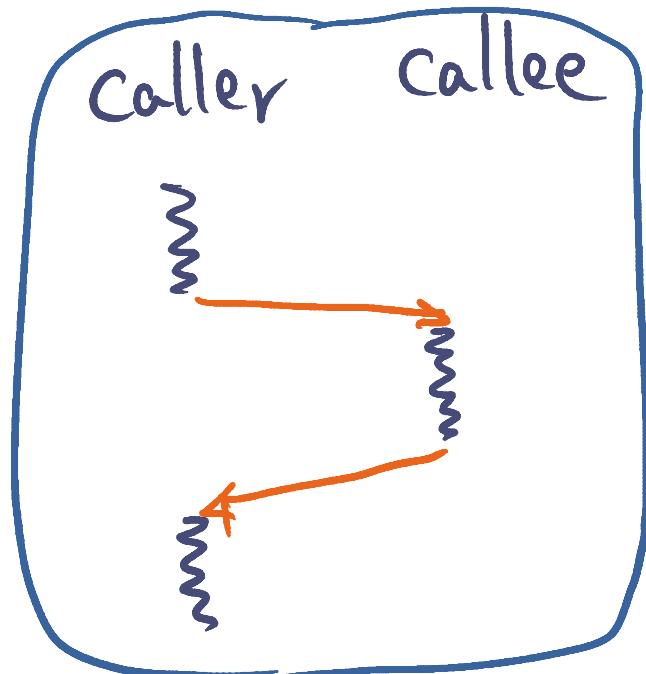


Process

- All at  
compiletime

# RPC Vs. Simple Procedure Call

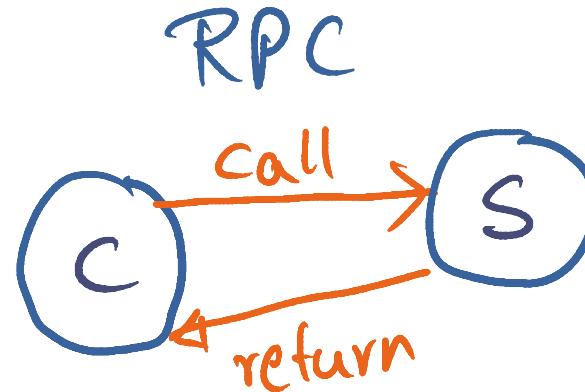
Procedure Call



Process

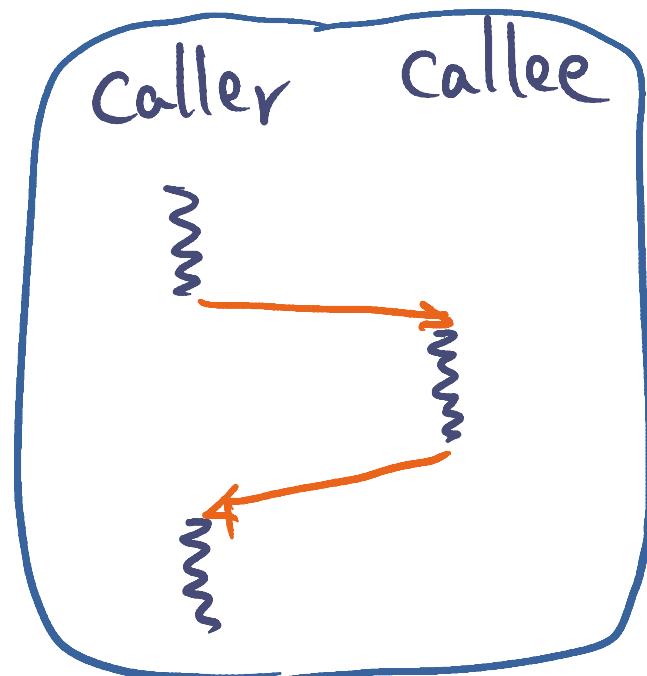
- All at  
Completion

RPC



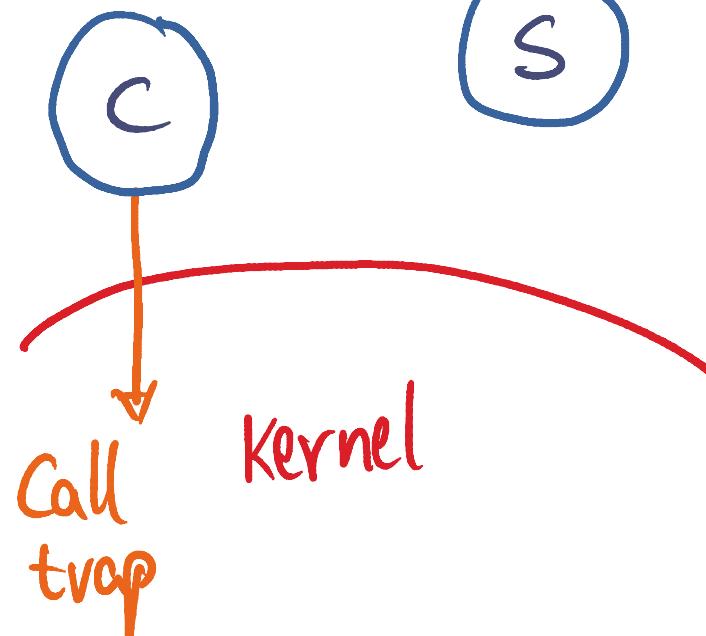
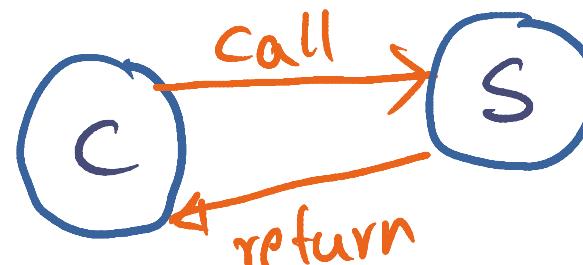
# RPC Vs. Simple Procedure Call

## Procedure Call



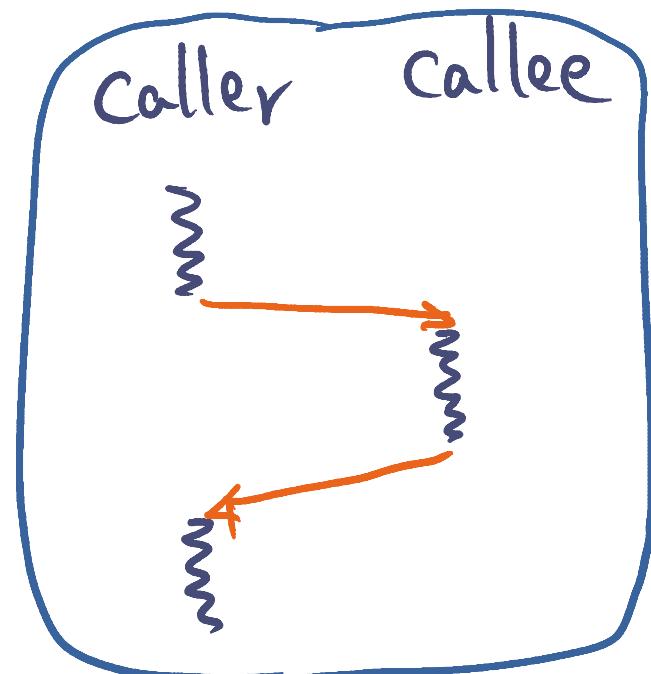
Process  
- All at  
Completion

## RPC



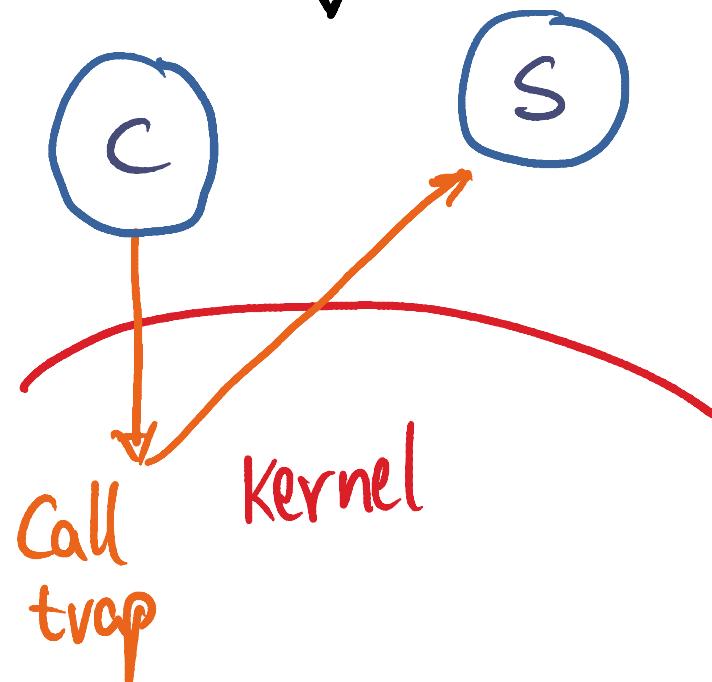
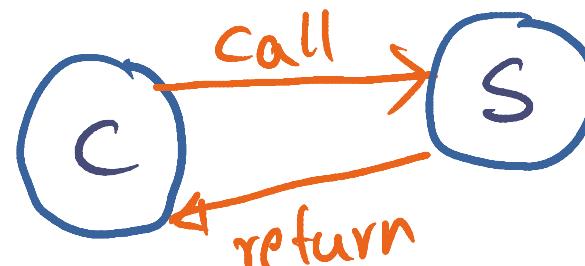
# RPC Vs. Simple Procedure Call

## Procedure Call



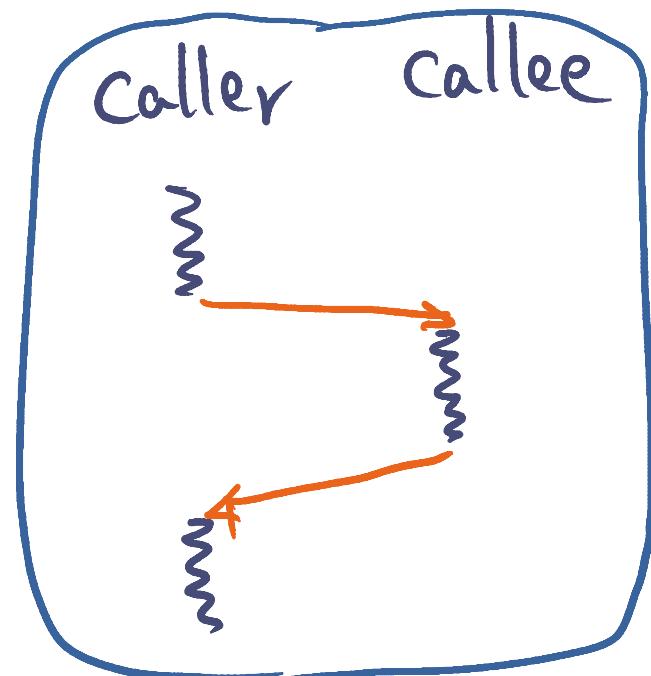
Process  
- All at  
Completion

## RPC



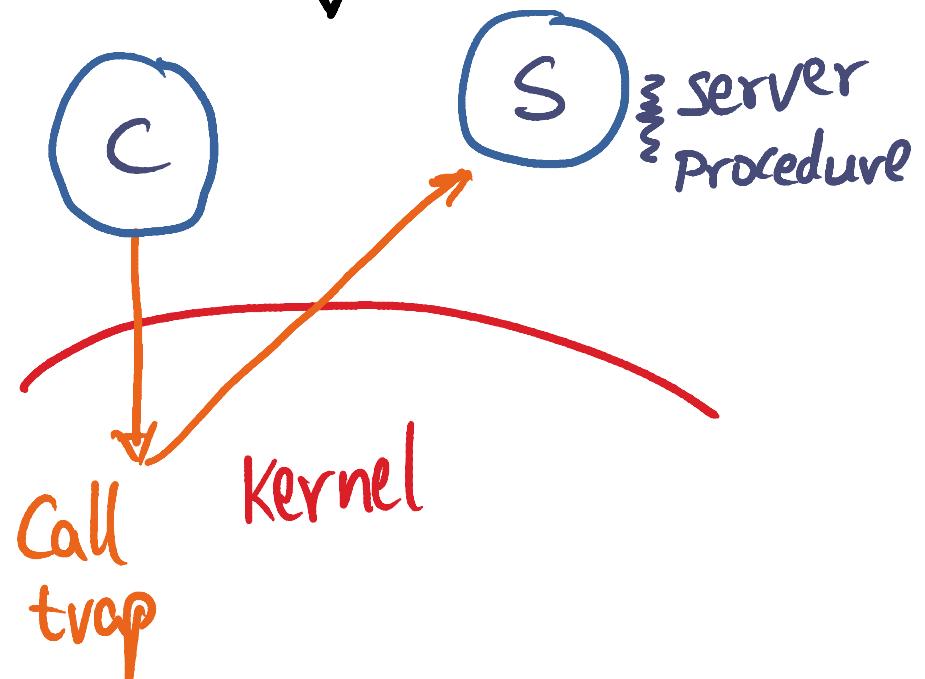
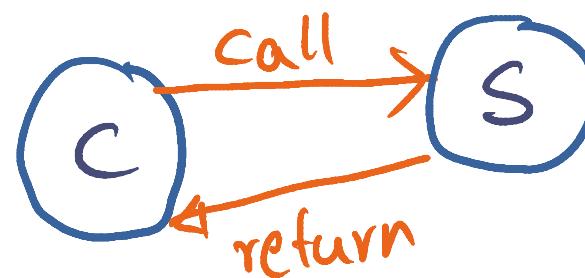
# RPC Vs. Simple Procedure Call

## Procedure Call



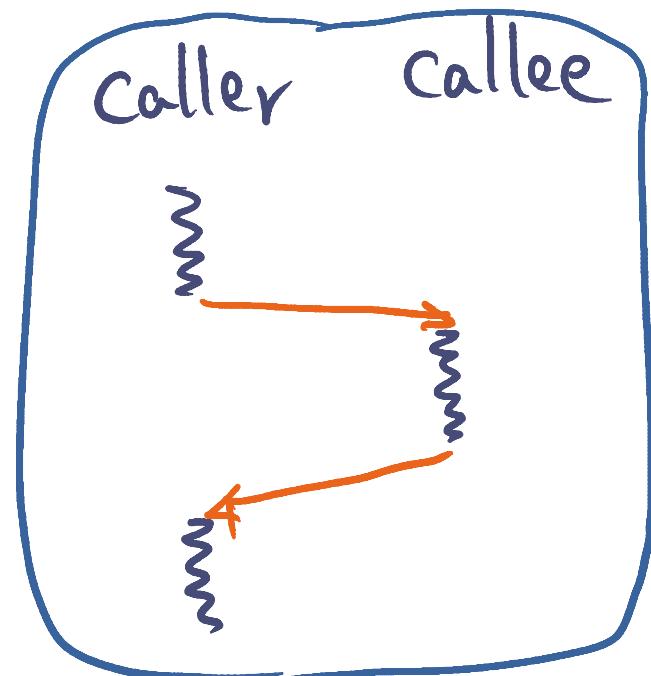
Process  
- All at  
Completion

## RPC



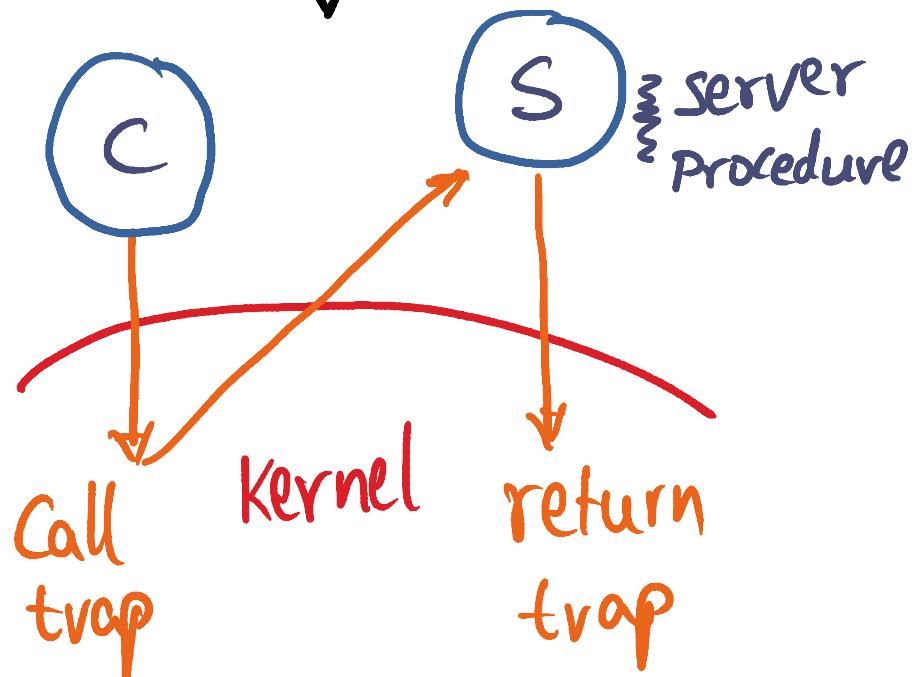
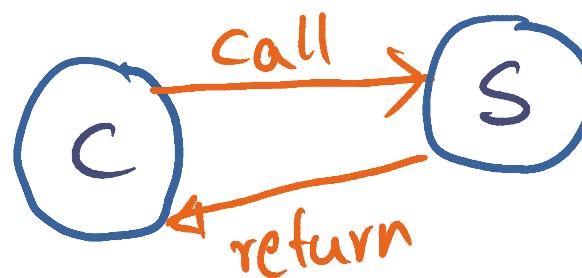
# RPC Vs. Simple Procedure Call

## Procedure Call



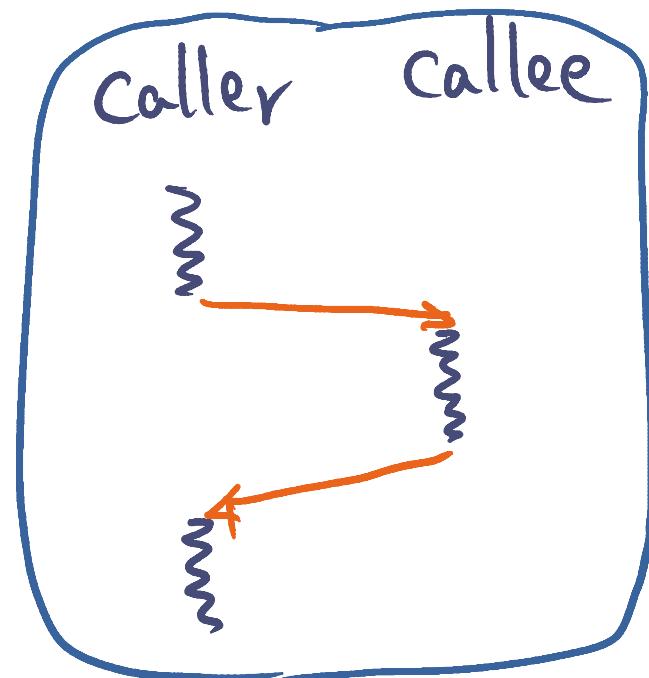
Process  
- All at  
Completion

## RPC



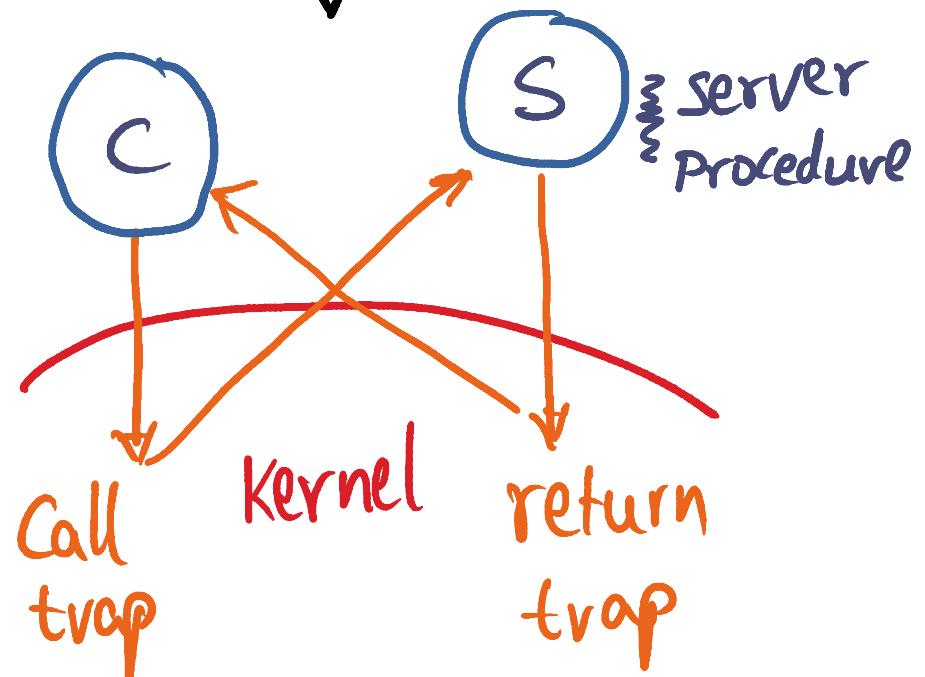
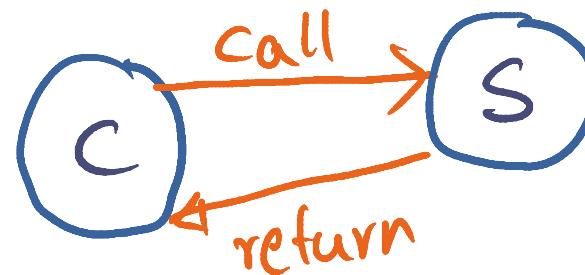
# RPC Vs. Simple Procedure Call

## Procedure Call



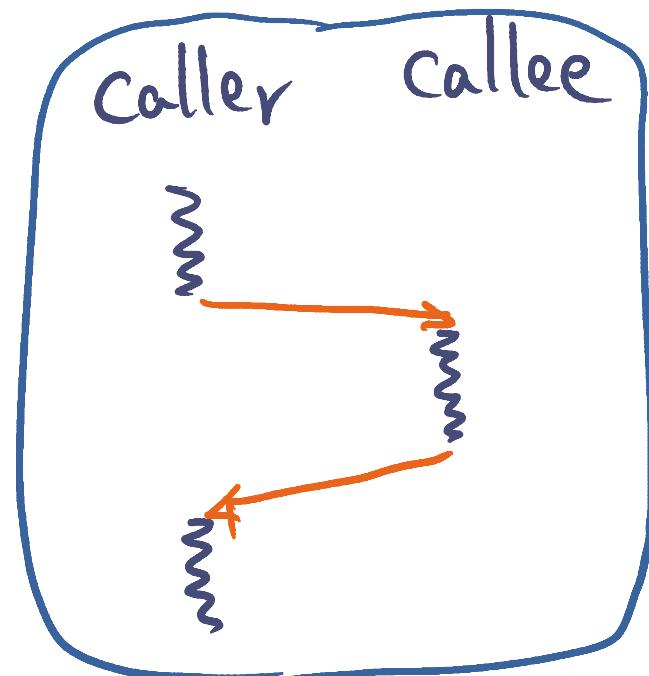
Process  
- All at  
Completion

## RPC



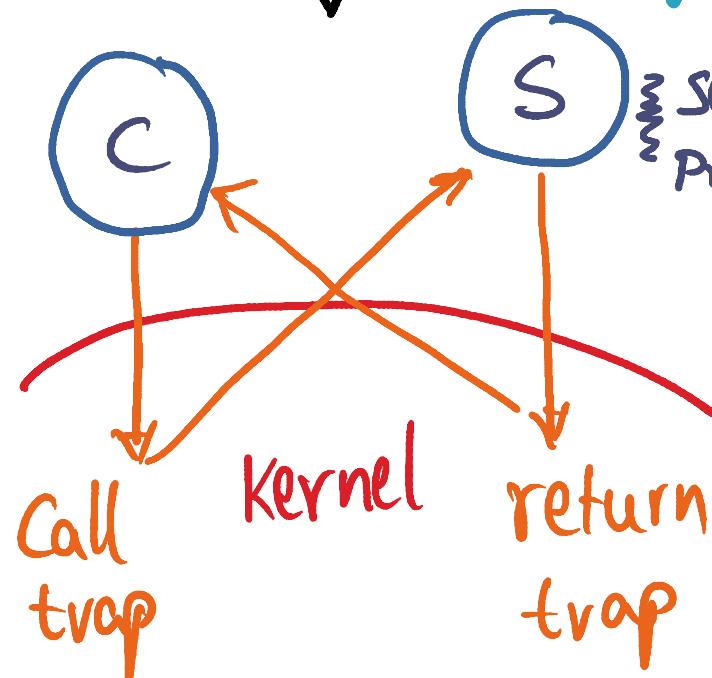
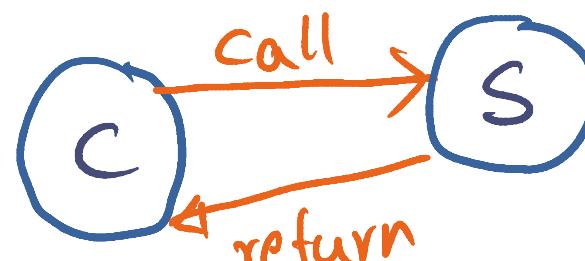
# RPC Vs. Simple Procedure Call

## Procedure Call



Process  
— All at  
Compiletime

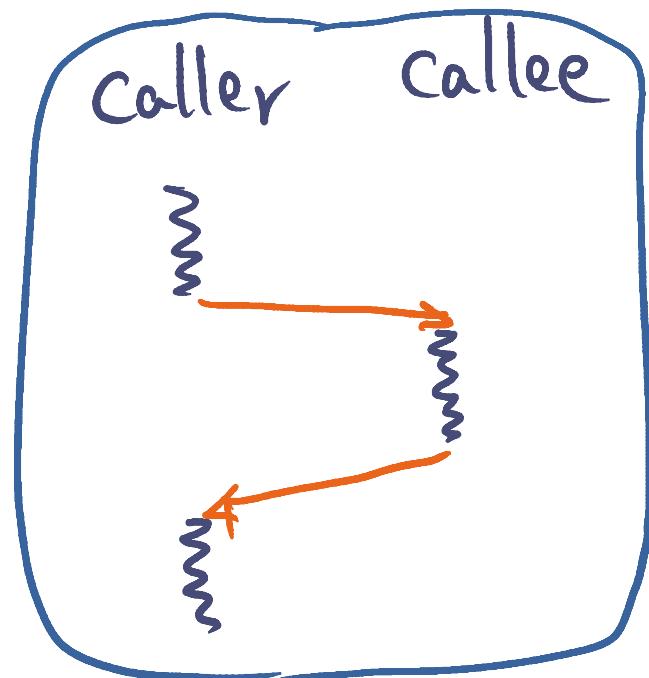
## RPC



All at  
runtime  
server  
procedure

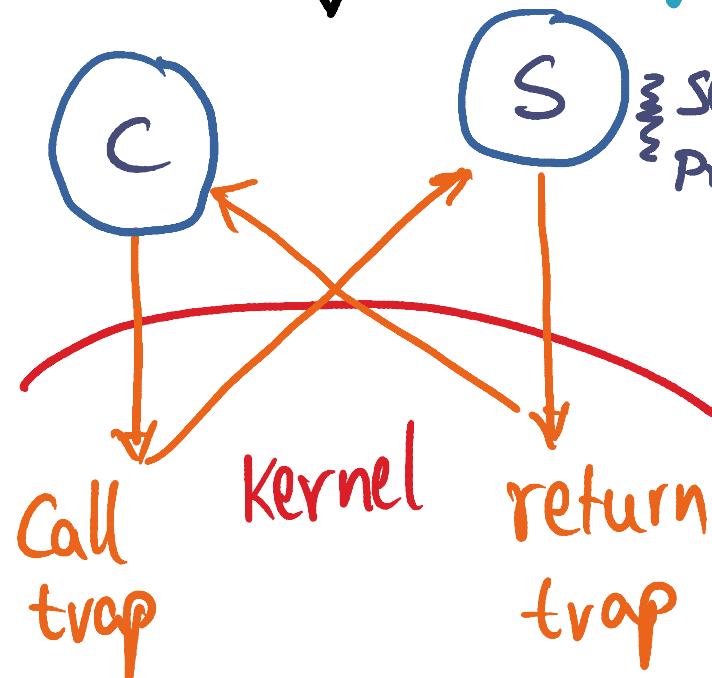
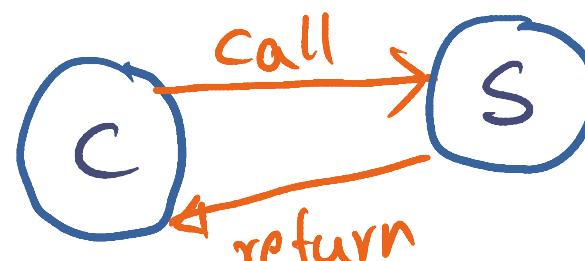
# RPC Vs. Simple Procedure Call

## Procedure Call



Process  
- All at  
Compiletime

## RPC

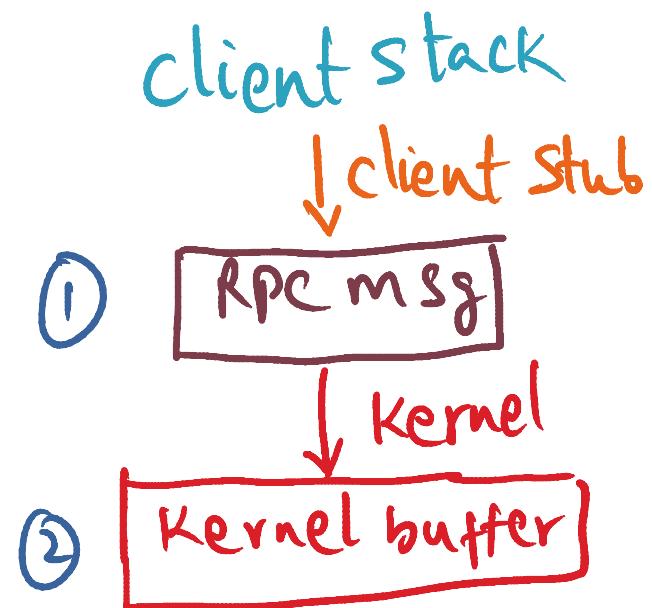


All at  
runtime  
server  
procedure

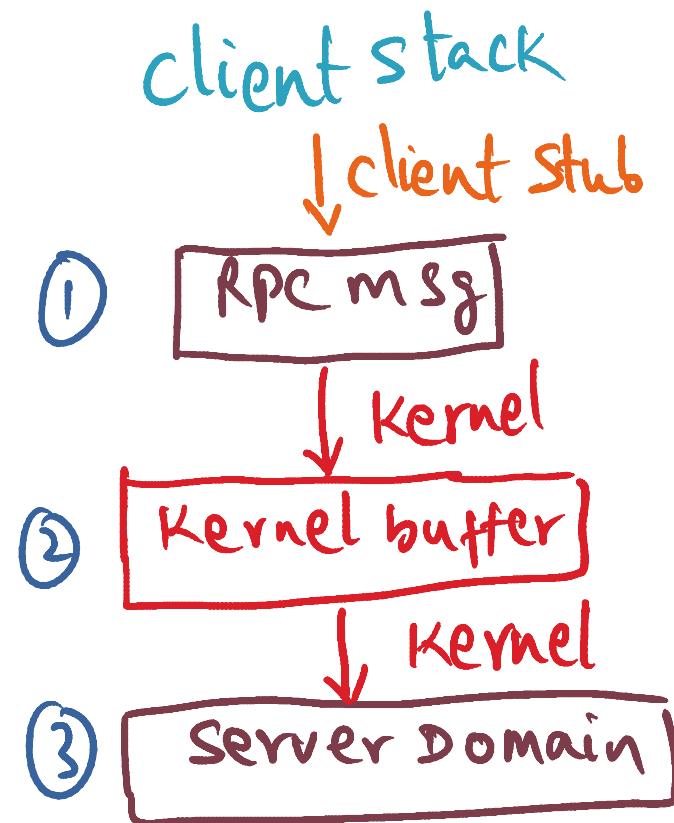
# Copying overhead



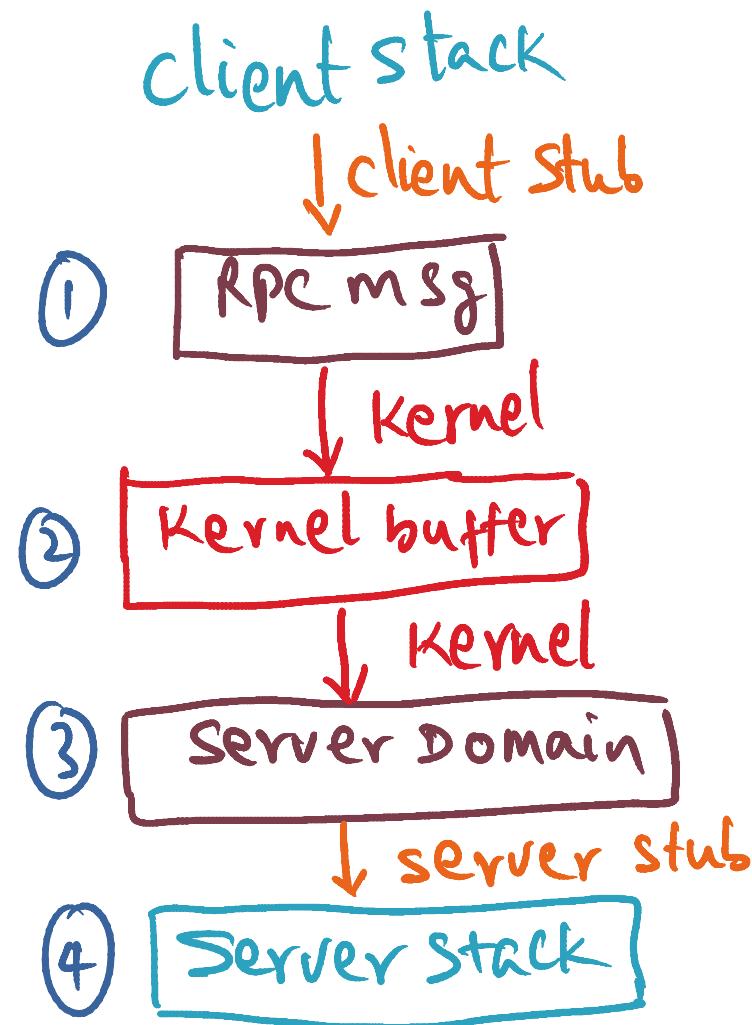
# Copying overhead



# Copying overhead



# Copying overhead



Making RPC cheap

Making RPC cheap  
How to remove overheads?

## Making RPC cheap (Binding)

How to remove overheads?

## Making RPC cheap (Binding)

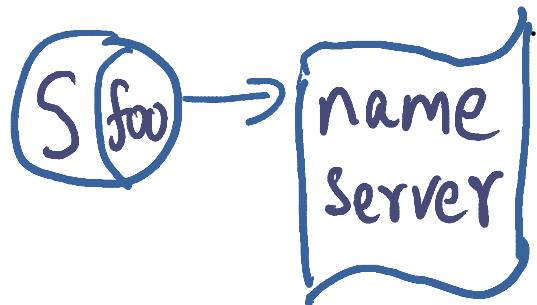
How to remove overheads?

- Set up (Binding)  $\Rightarrow$  one time cost

## Making RPC cheap (Binding)

How to remove overheads?

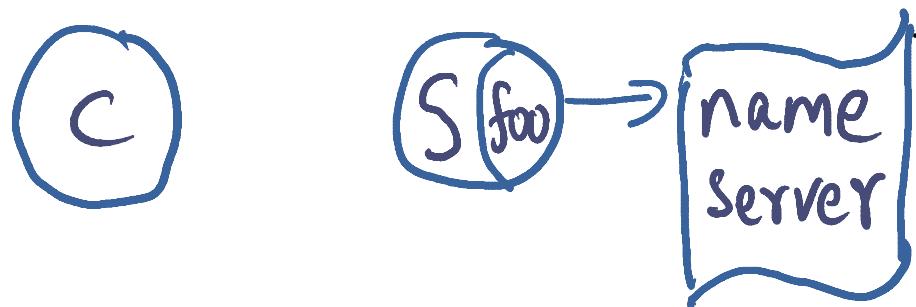
- Set up (Binding)  $\Rightarrow$  one time cost



## Making RPC cheap (Binding)

How to remove overheads?

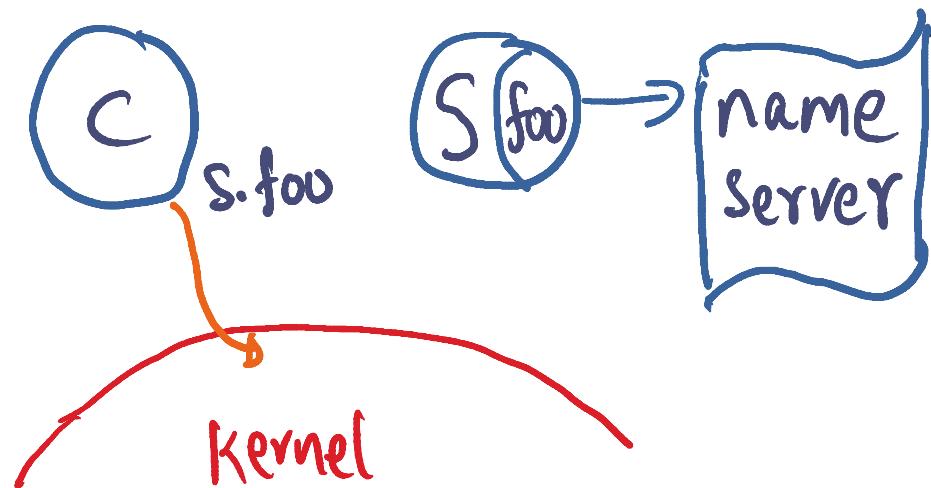
- Set up (Binding)  $\Rightarrow$  one time cost



## Making RPC cheap (Binding)

How to remove overheads?

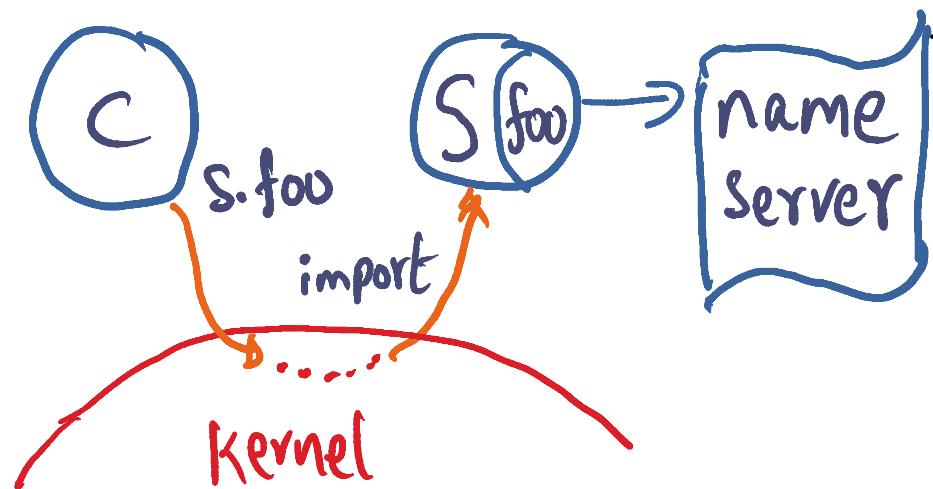
- Set up (Binding)  $\Rightarrow$  one time cost



## Making RPC cheap (Binding)

How to remove overheads?

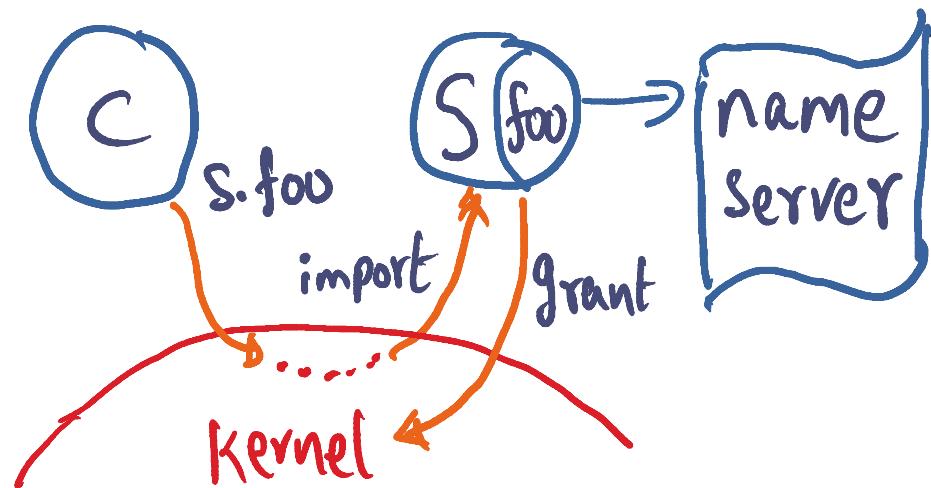
- Set up (Binding)  $\Rightarrow$  one time cost



# Making RPC cheap (Binding)

How to remove overheads?

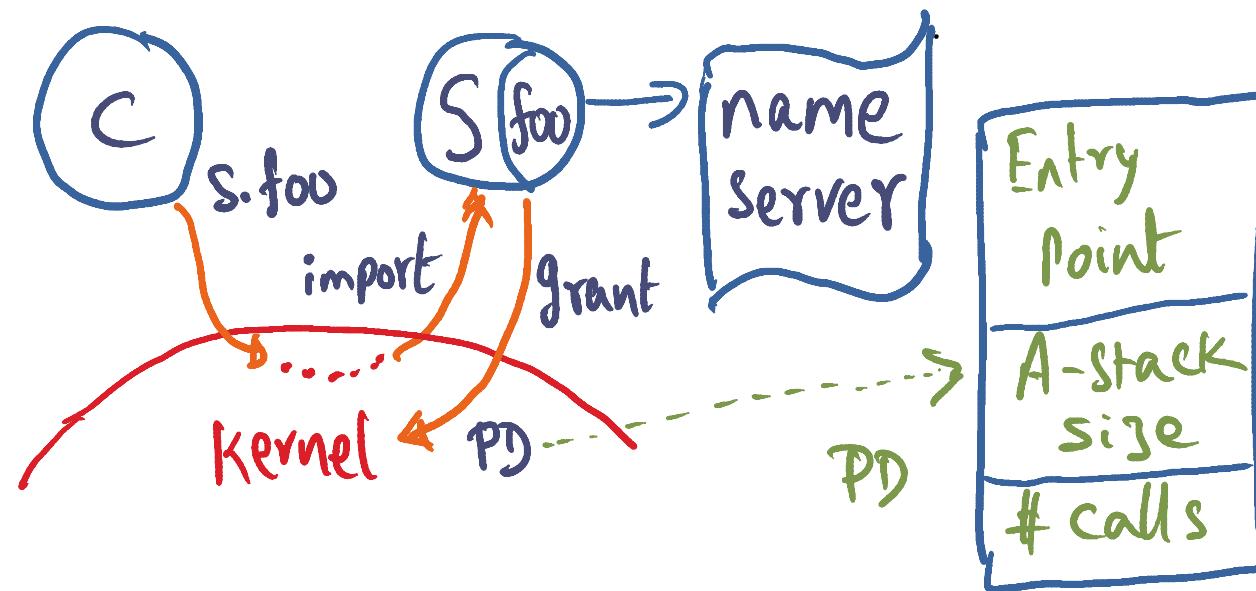
- Set up (Binding)  $\Rightarrow$  one time cost



# Making RPC cheap (Binding)

How to remove overheads?

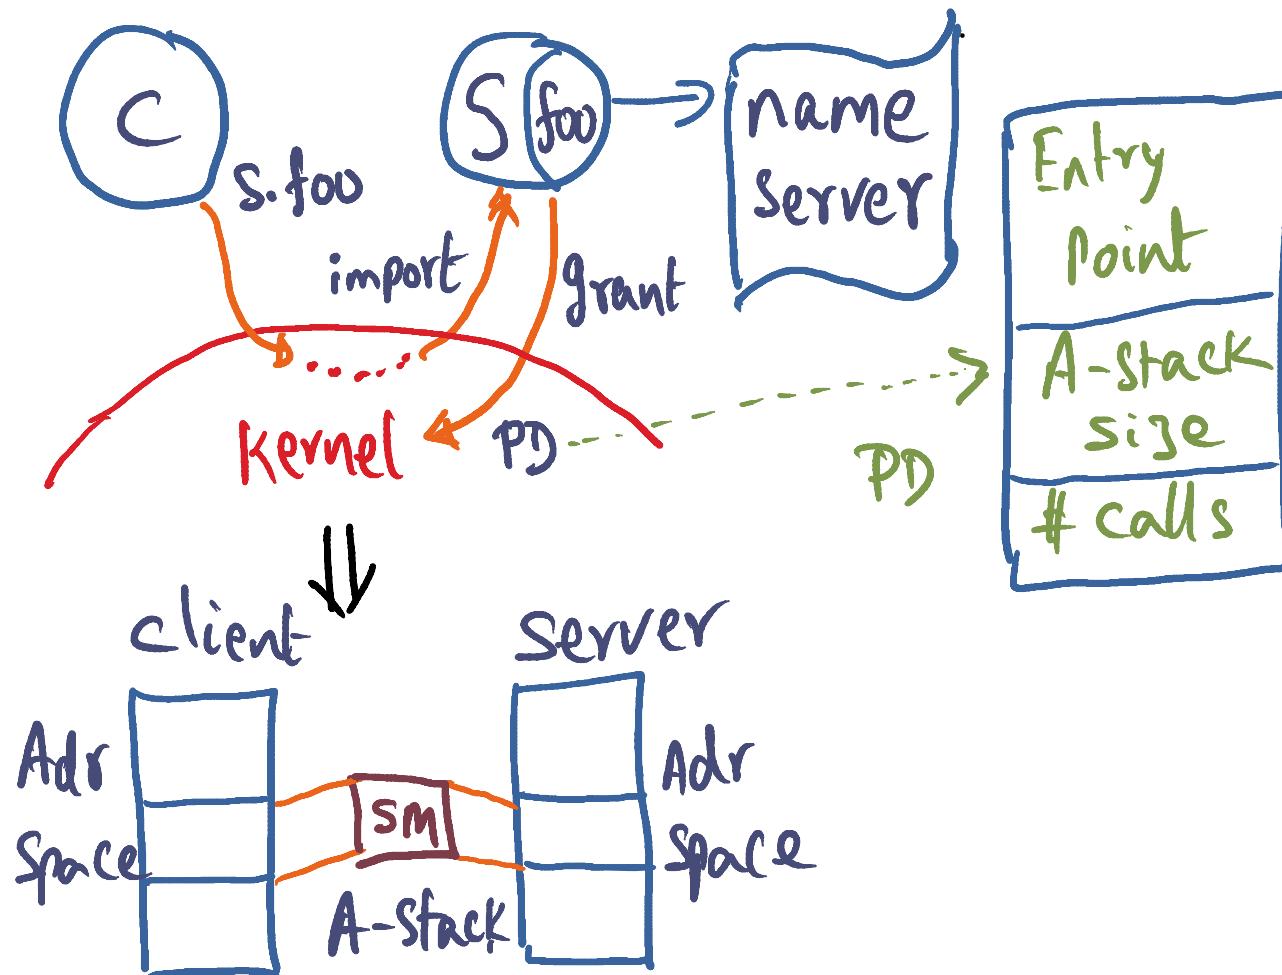
- Set up (Binding)  $\Rightarrow$  one time cost



# Making RPC cheap (Binding)

How to remove overheads?

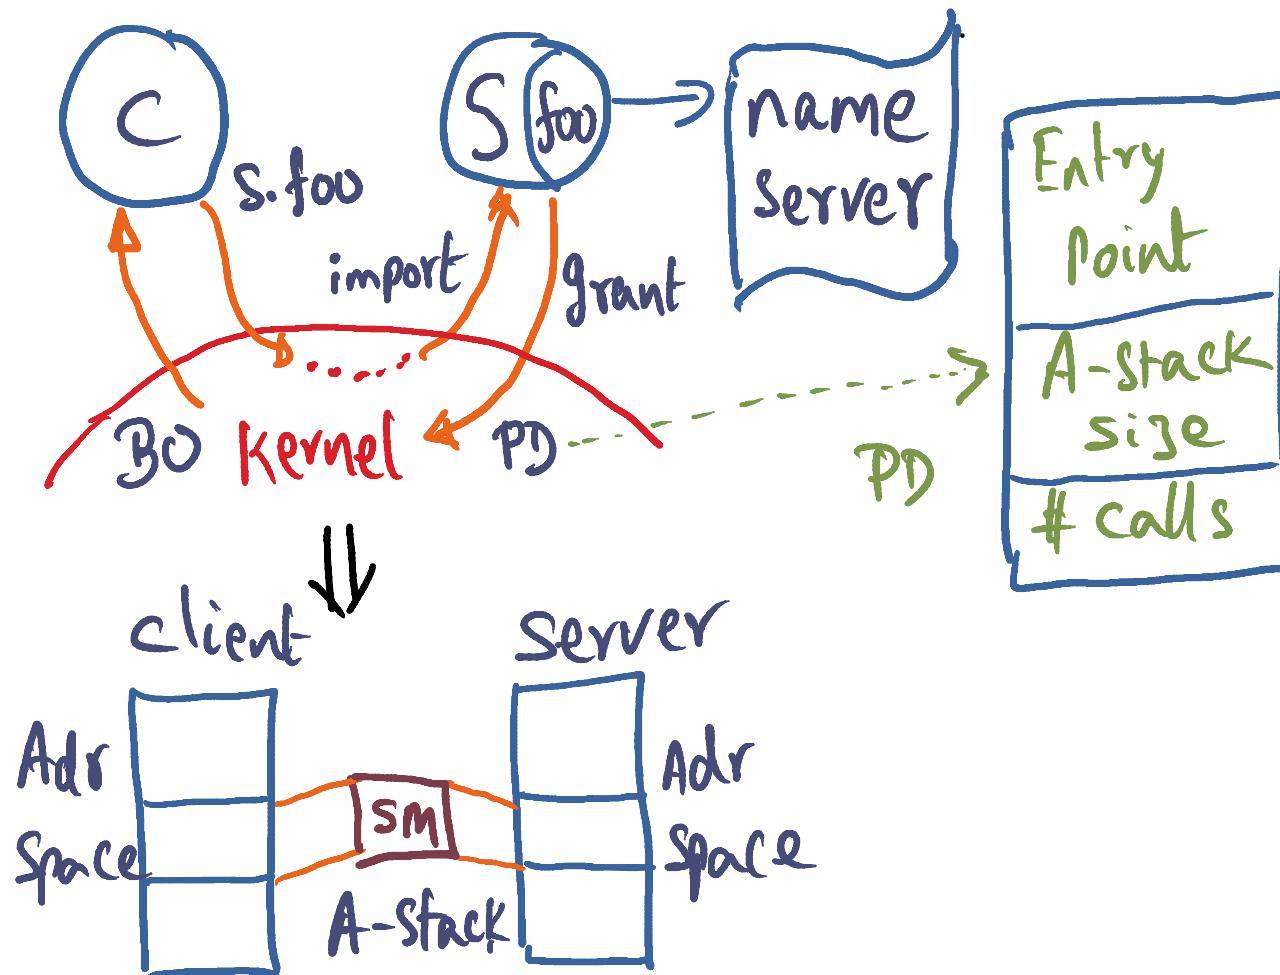
- Set up (Binding)  $\Rightarrow$  one time cost



# Making RPC cheap (Binding)

How to remove overheads?

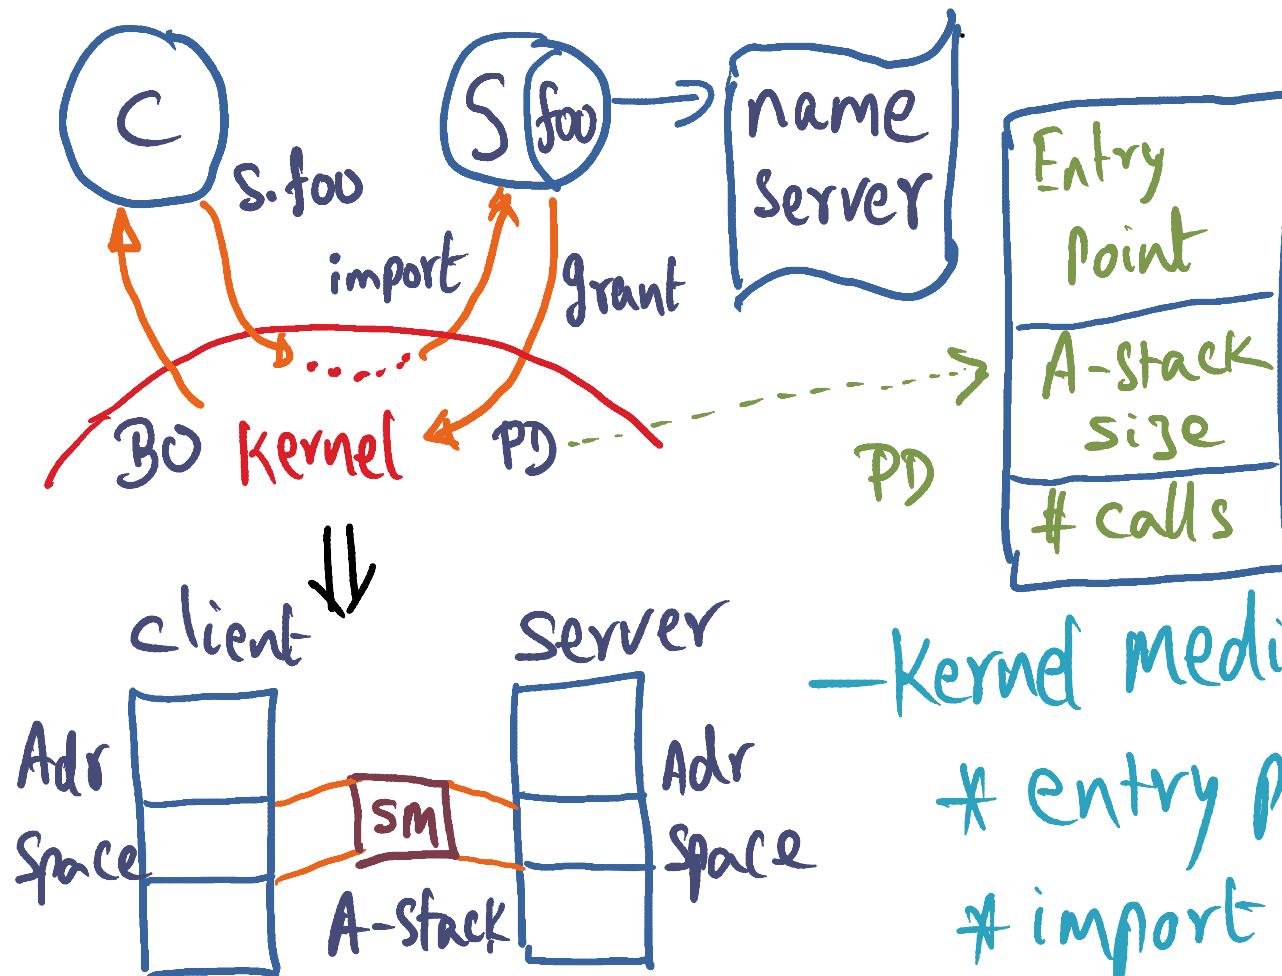
- Set up (Binding)  $\Rightarrow$  one time cost



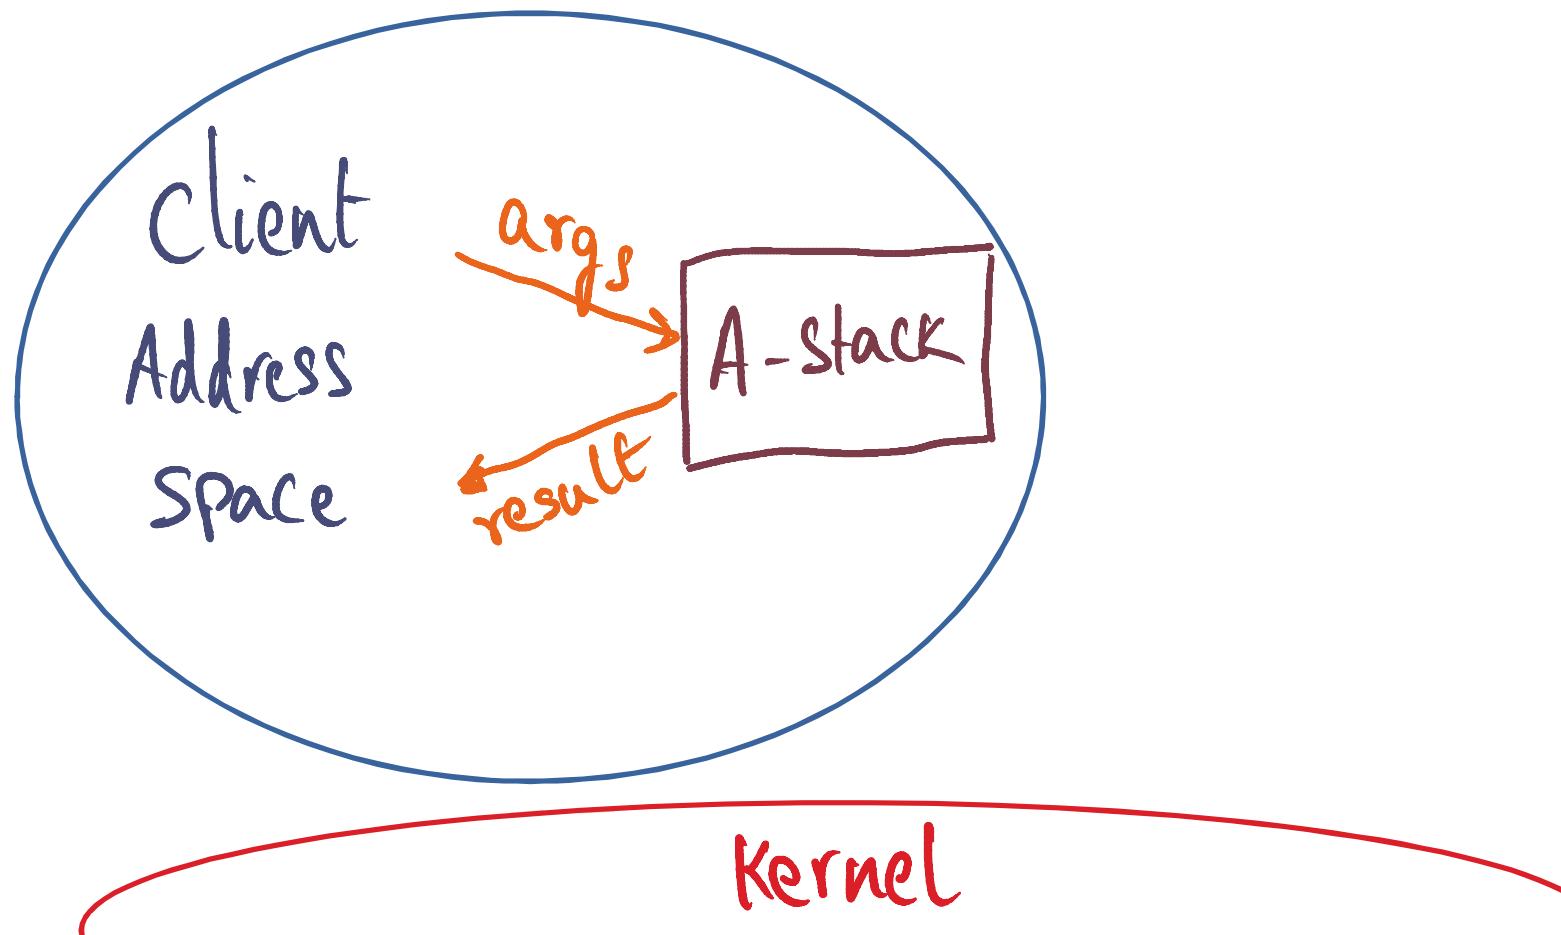
# Making RPC cheap (Binding)

How to remove overheads?

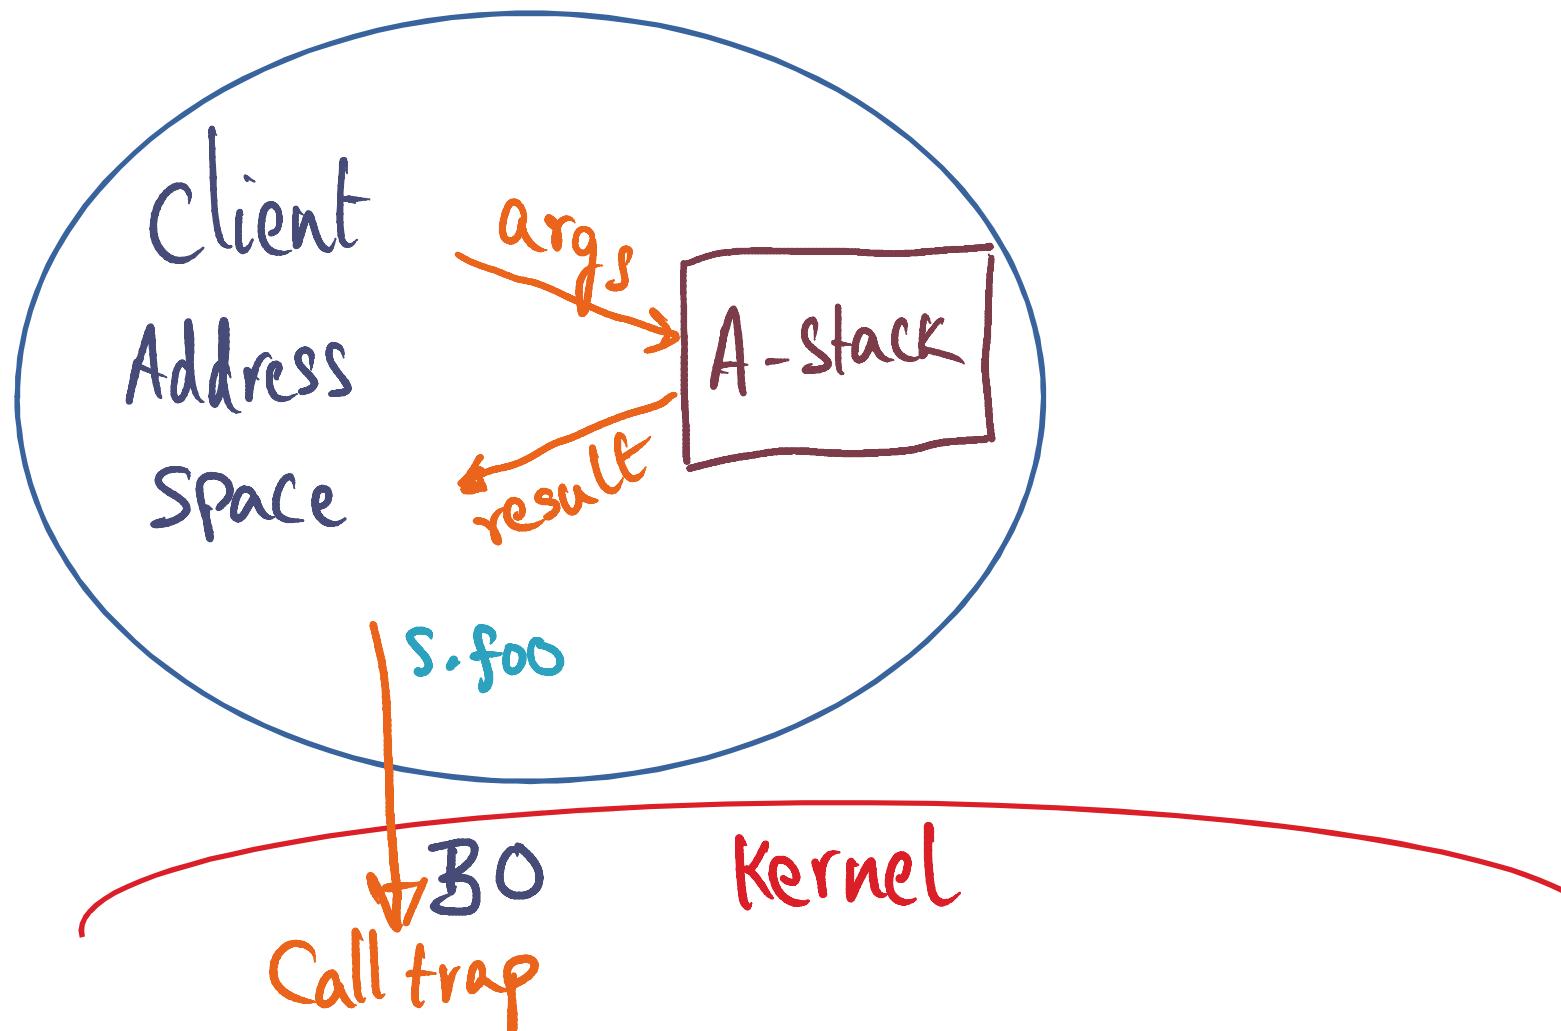
- Set up (Binding)  $\Rightarrow$  one time cost



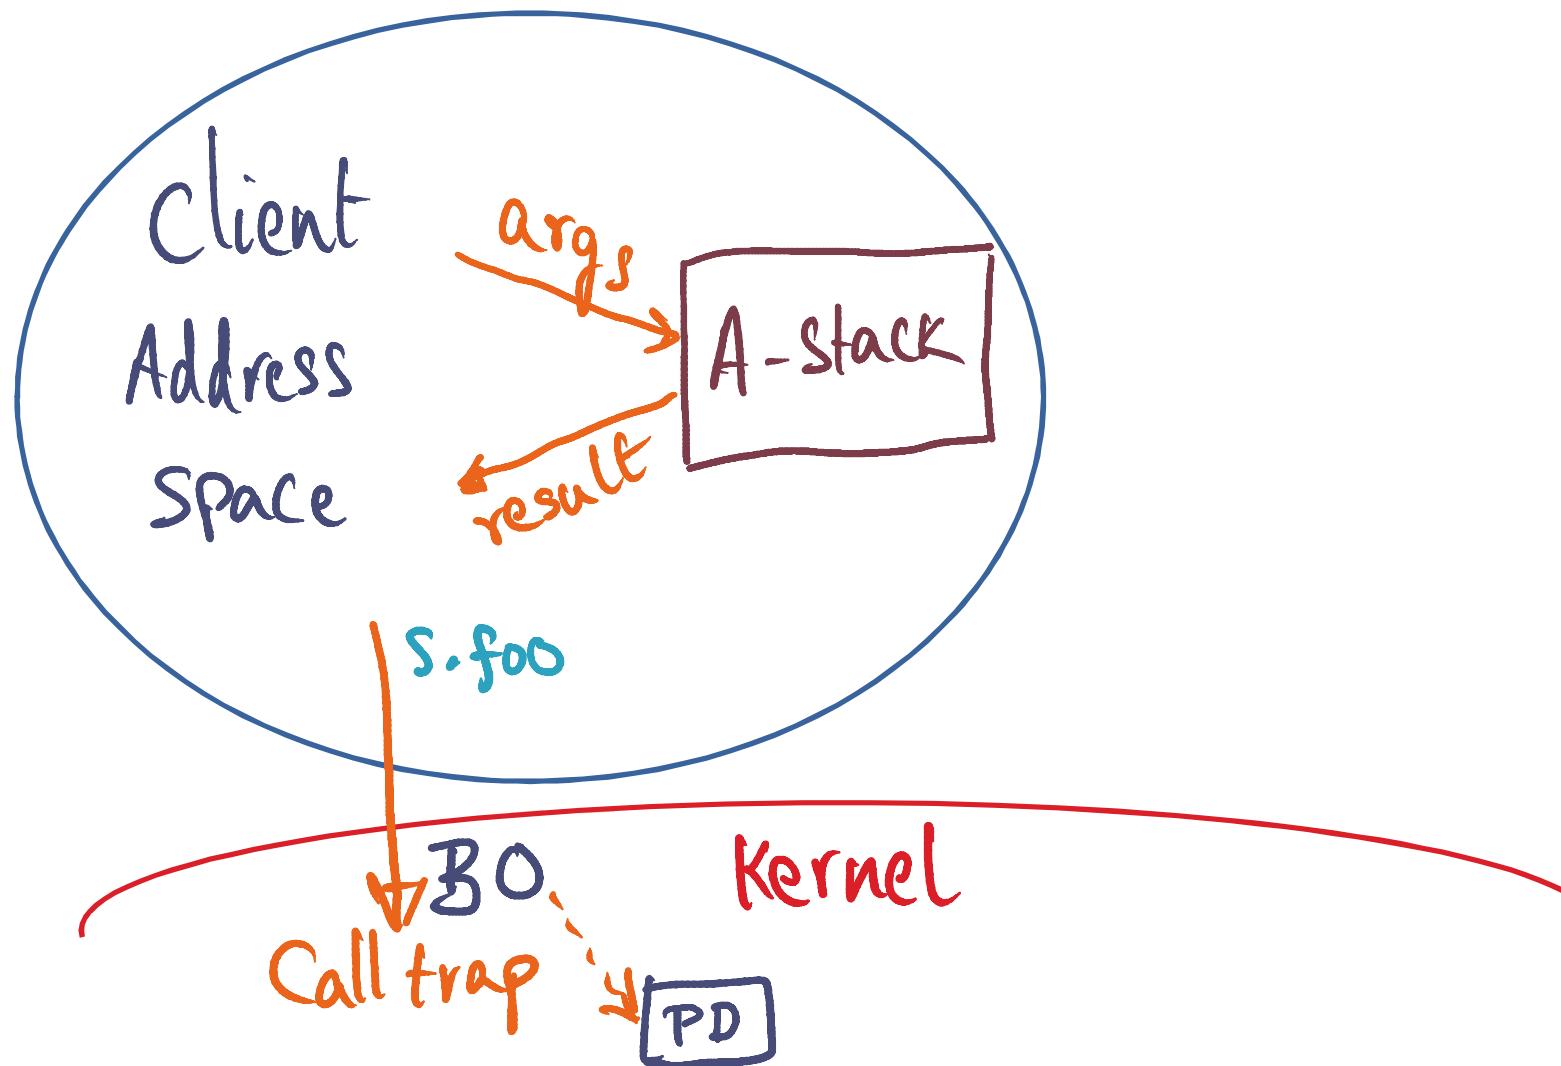
## Making RPC cheap (Actual Calls)



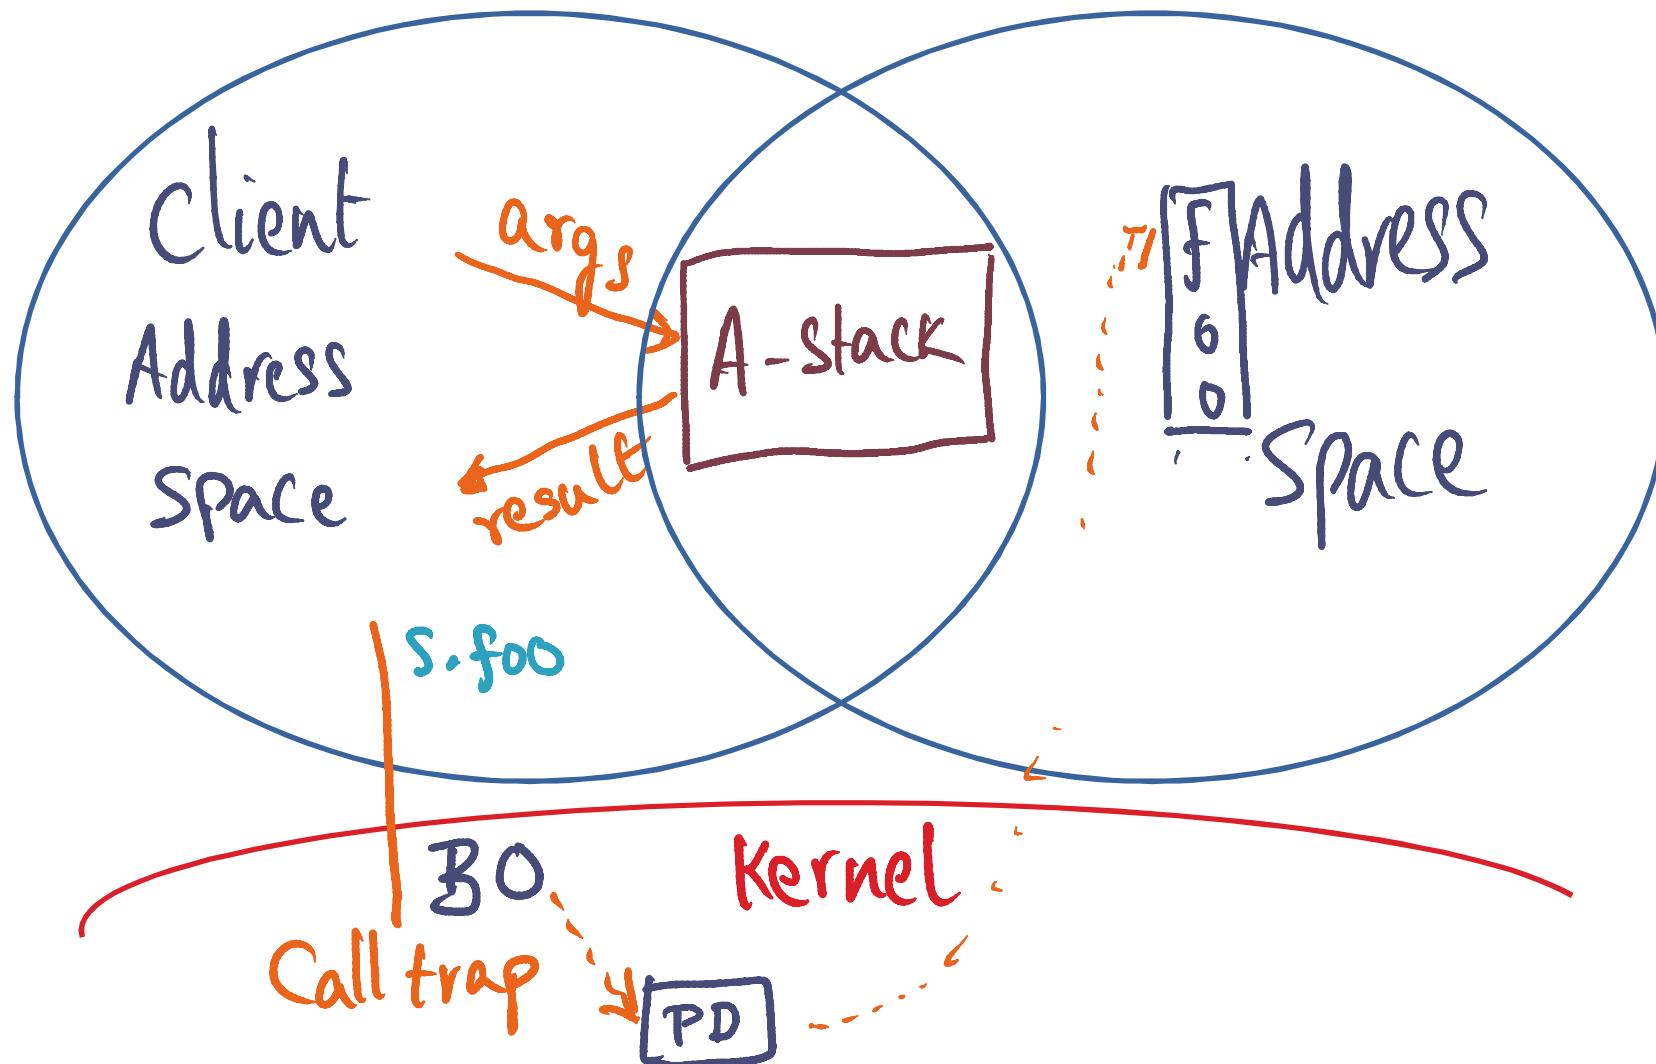
## Making RPC cheap (Actual Calls)



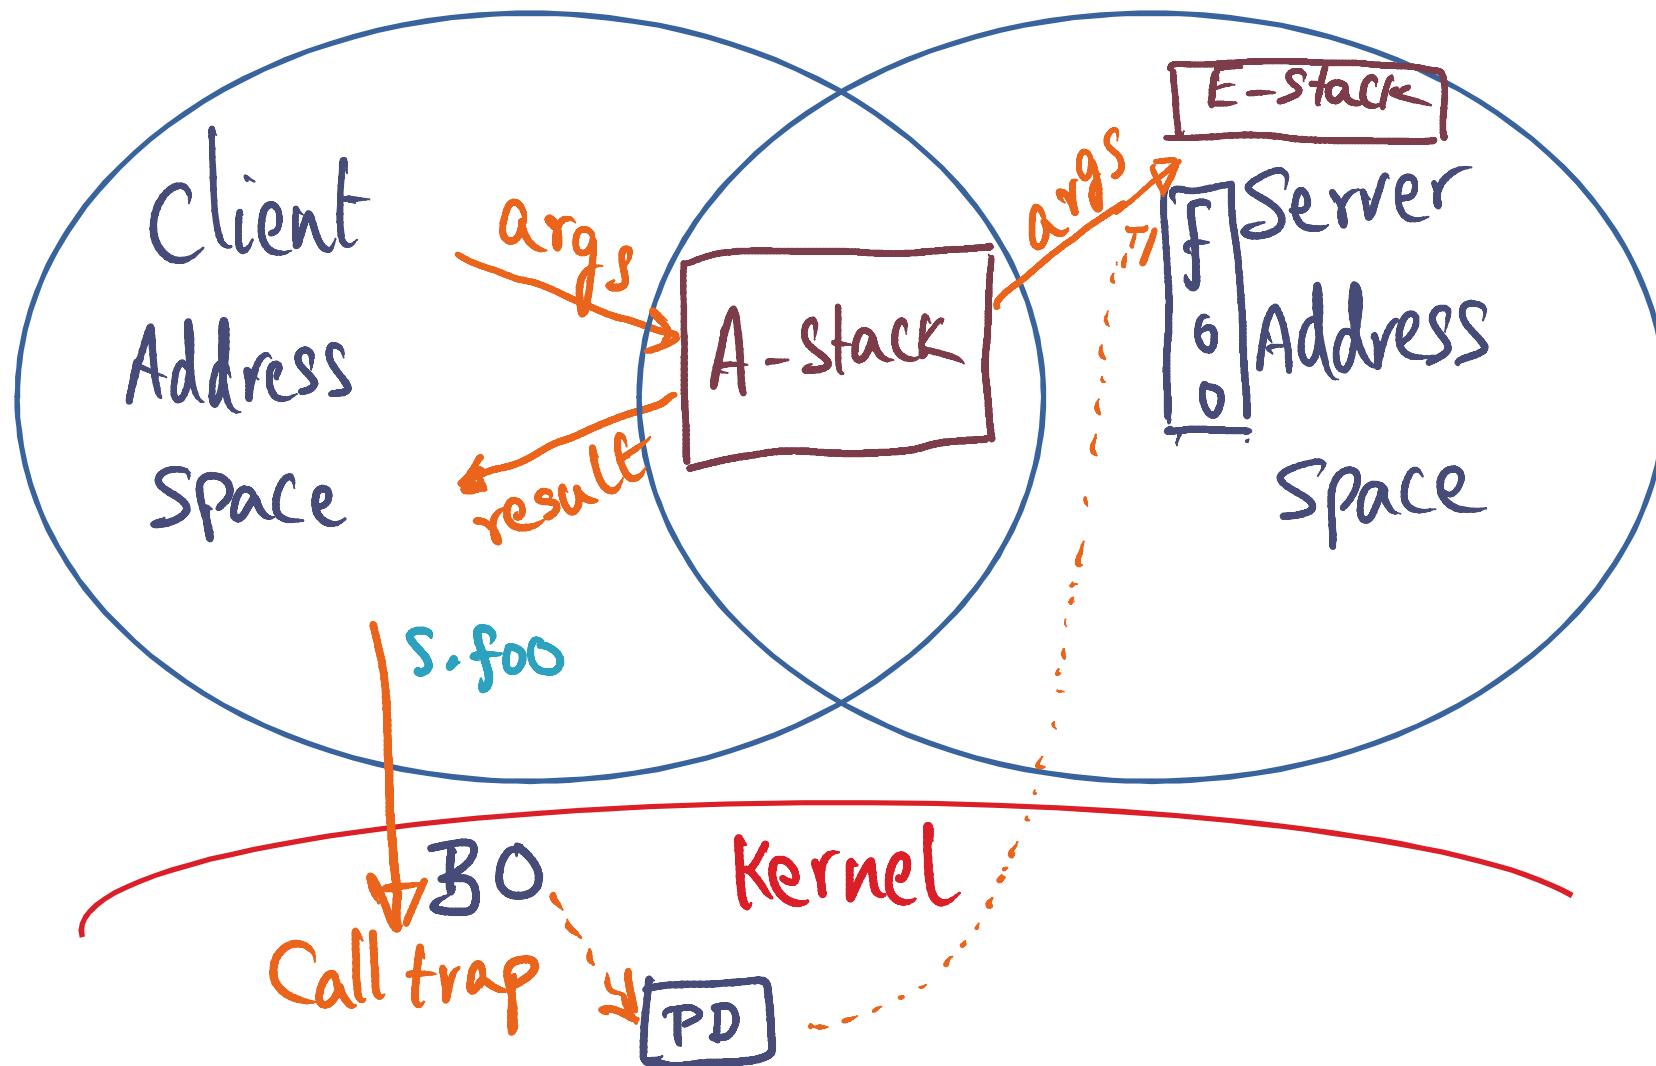
## Making RPC cheap (Actual Calls)



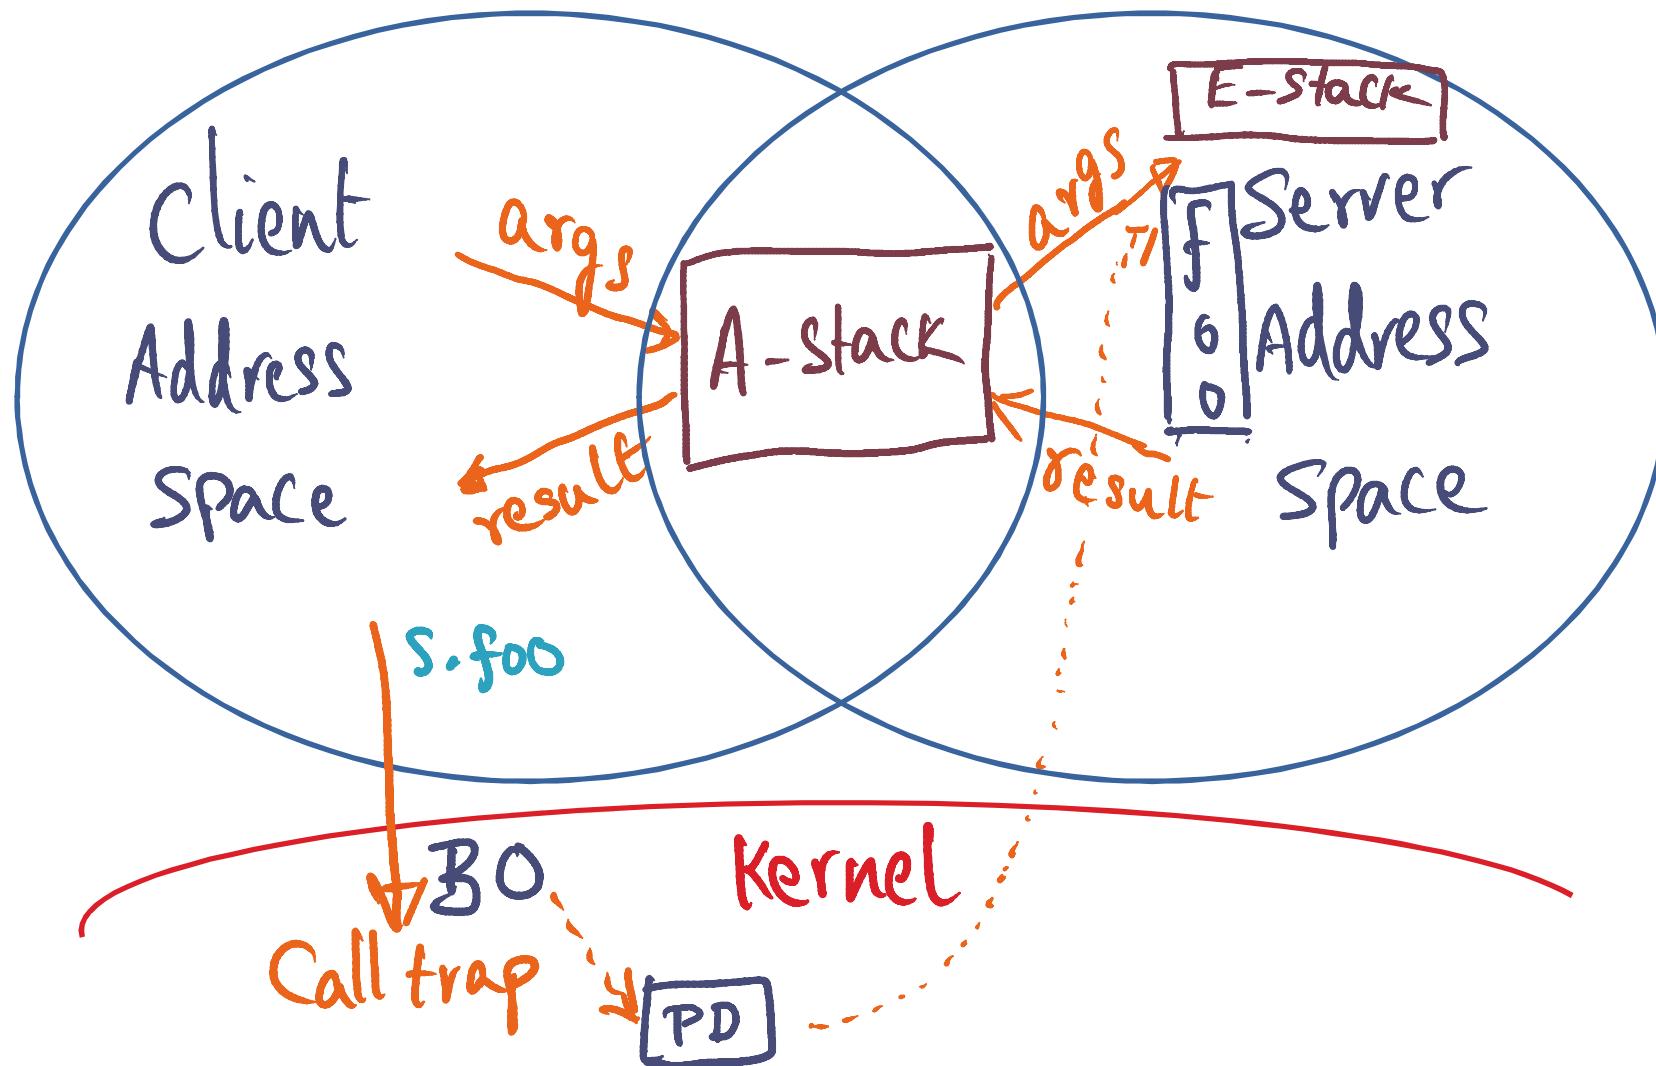
## Making RPC cheap (Actual Calls)



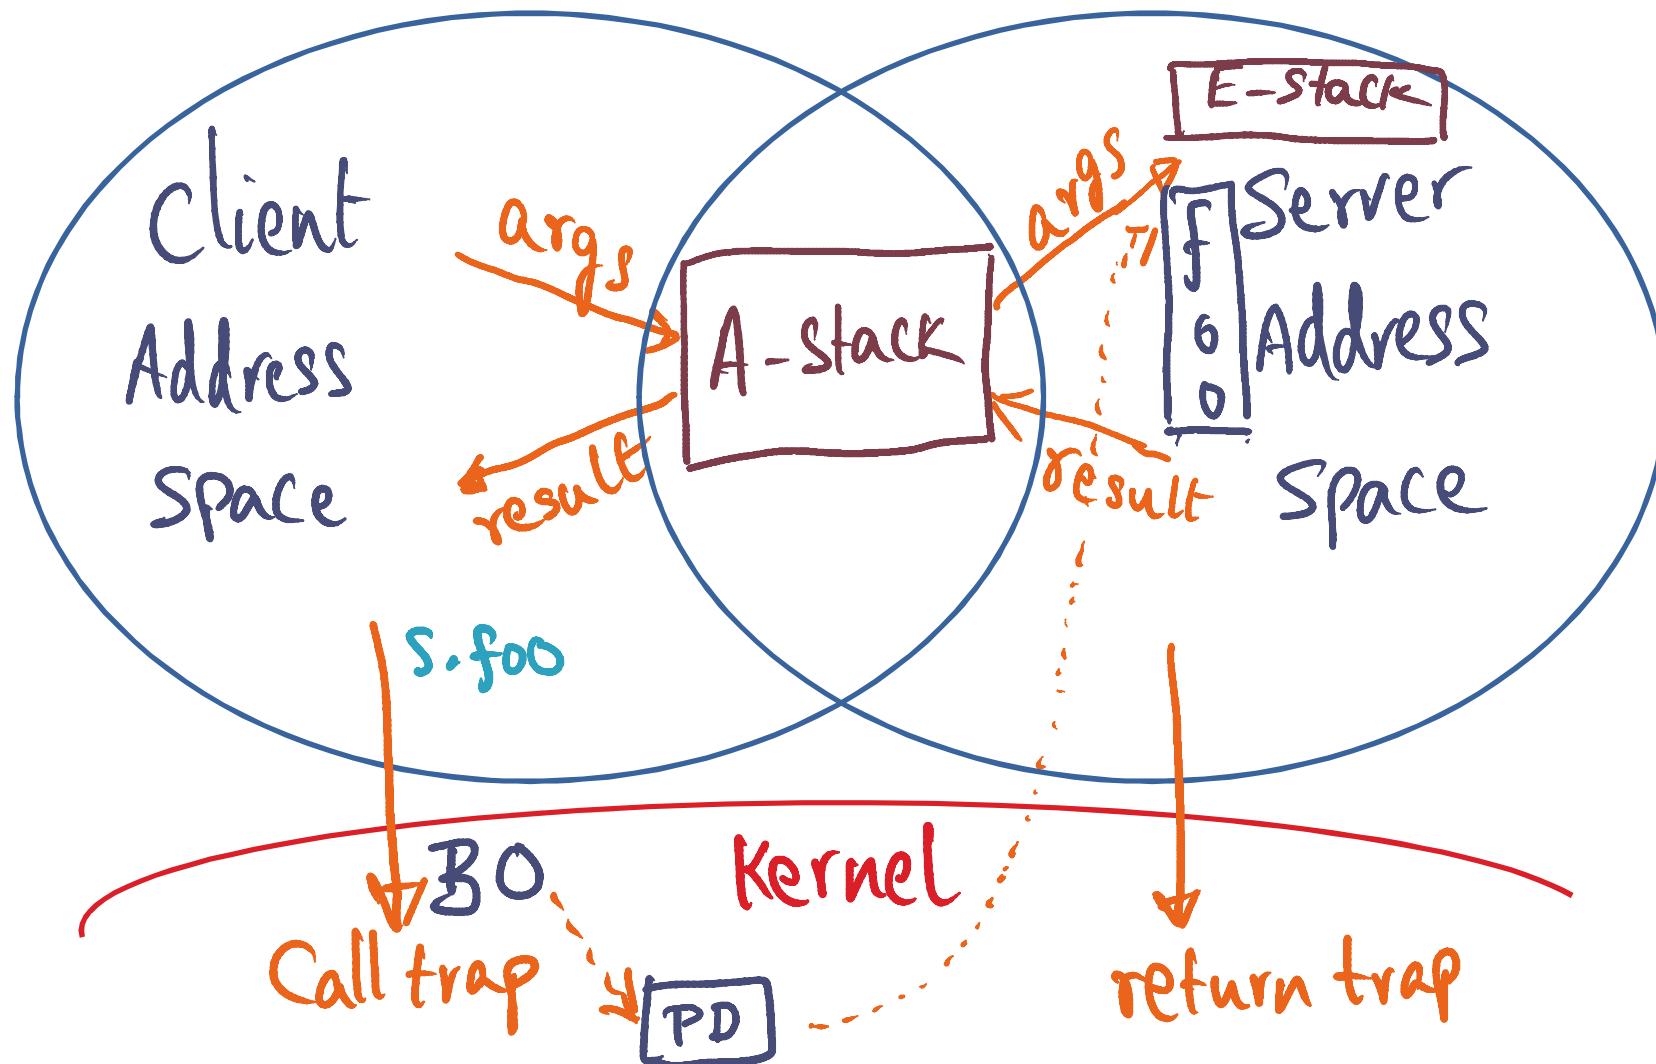
# Making RPC cheap (Actual Calls)



# Making RPC cheap (Actual Calls)



## Making RPC cheap (Actual Calls)



## Making RPC cheap (Actual Calls)

Original

client stack

↓ client stub

① RPC msg

## Making RPC cheap (Actual Calls)

Original

client stack

↓ client stub

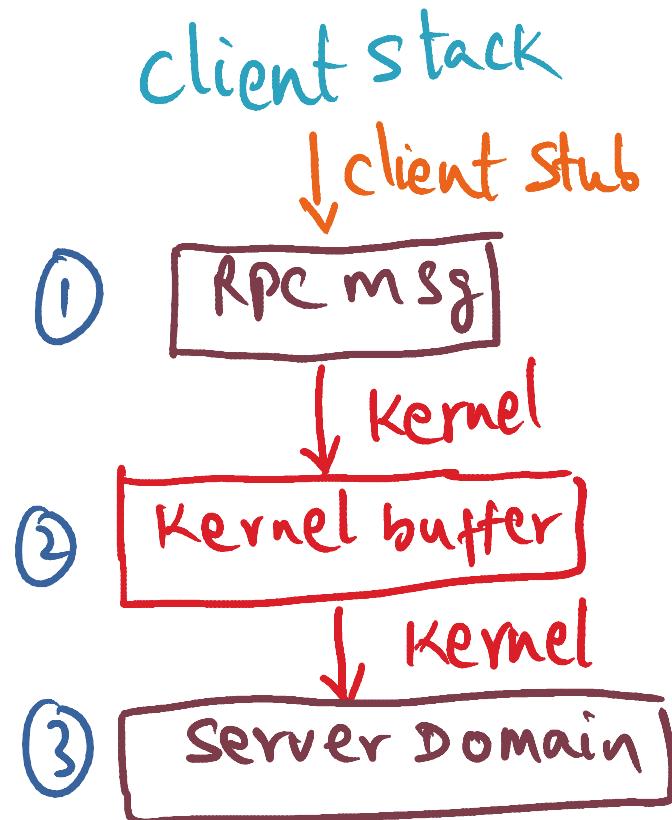
① RPC msg

↓ Kernel

② Kernel buffer

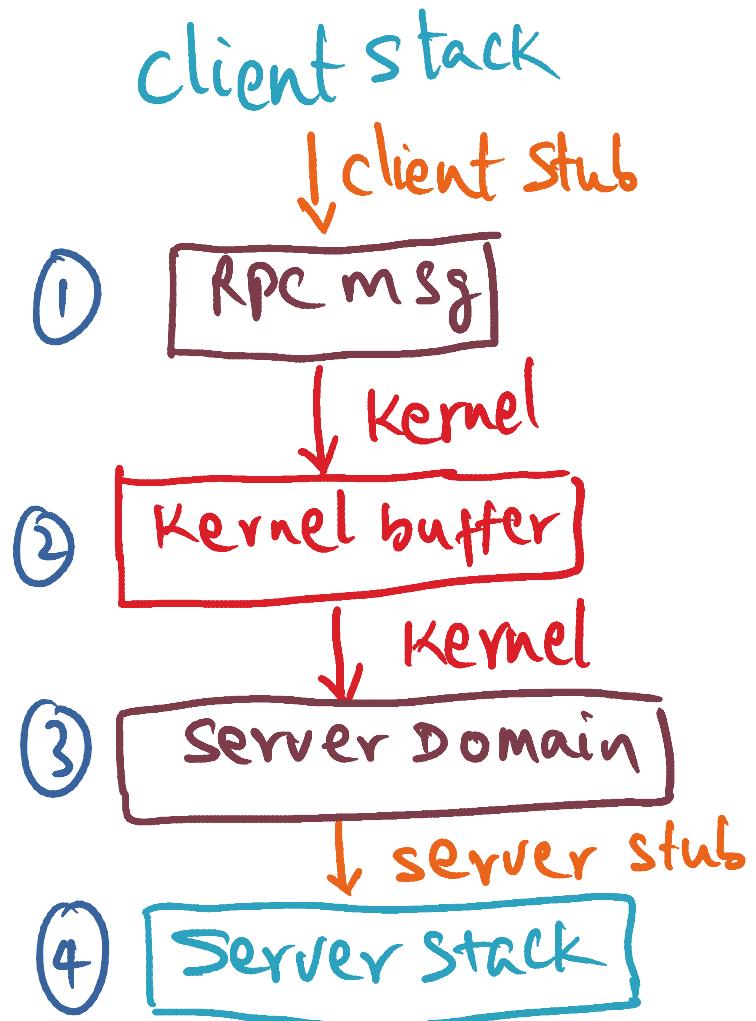
## Making RPC cheap (Actual Calls)

Original



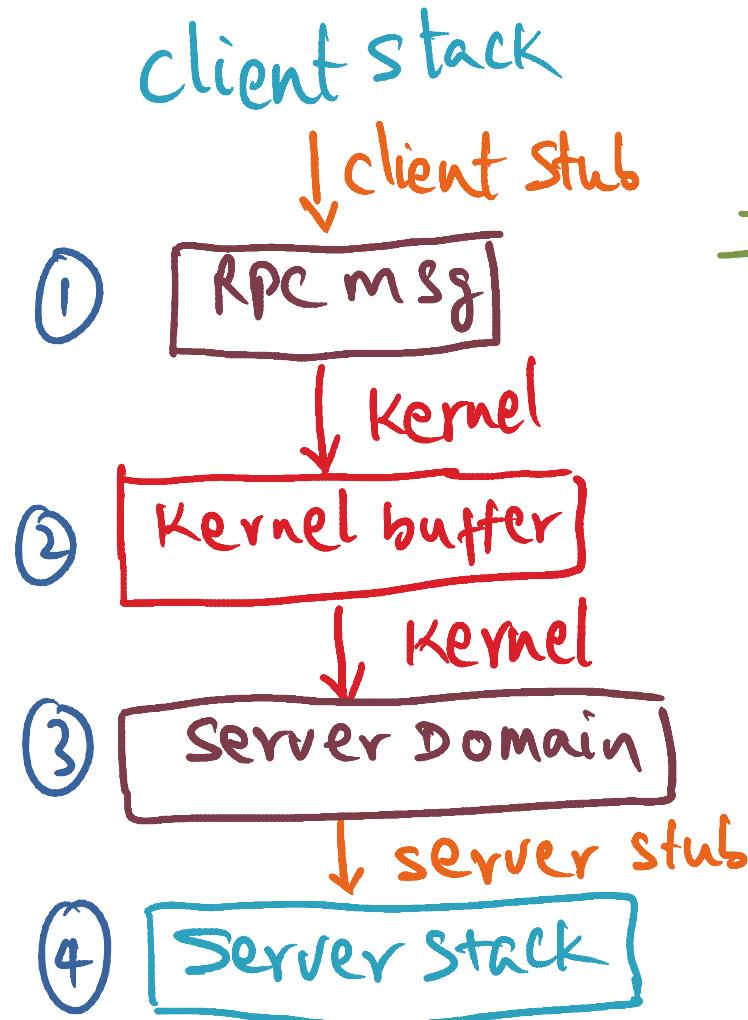
# Making RPC cheap (Actual Calls)

Original

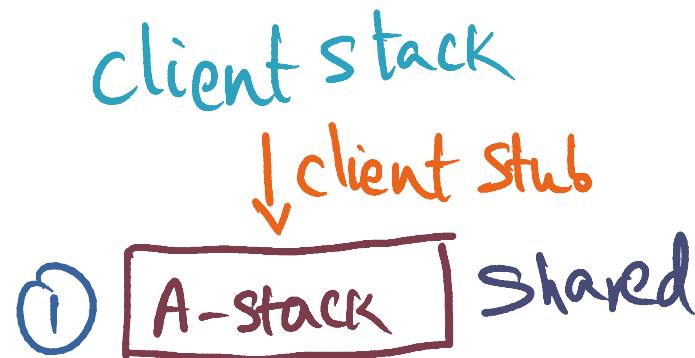


# Making RPC cheap (Actual Calls)

Original

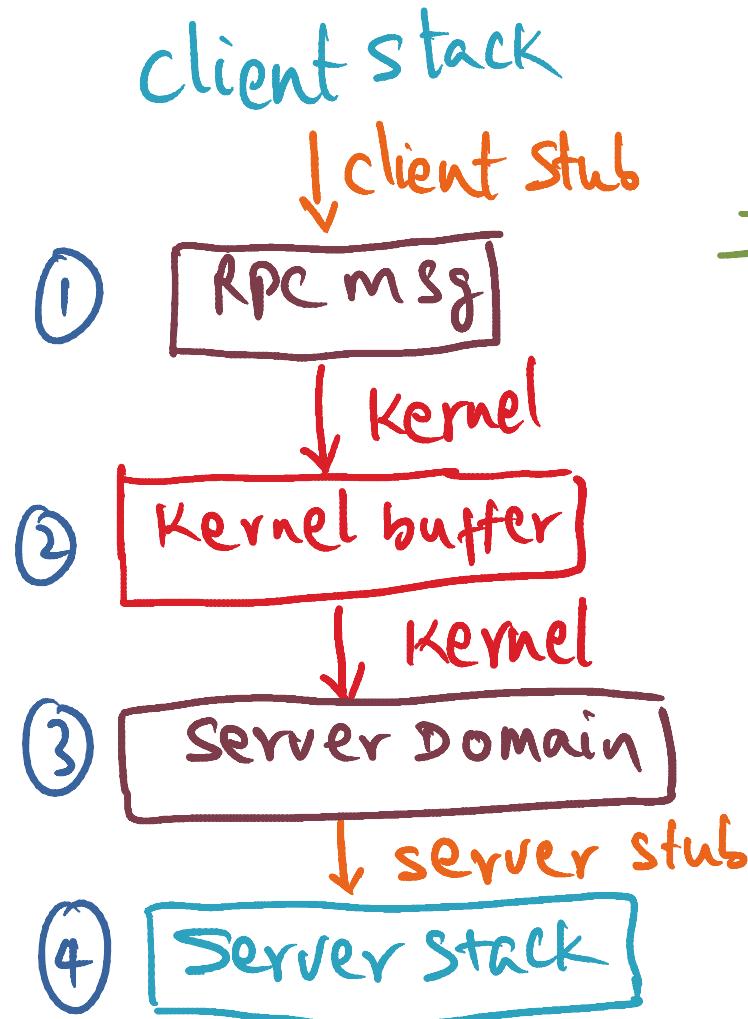


NOW

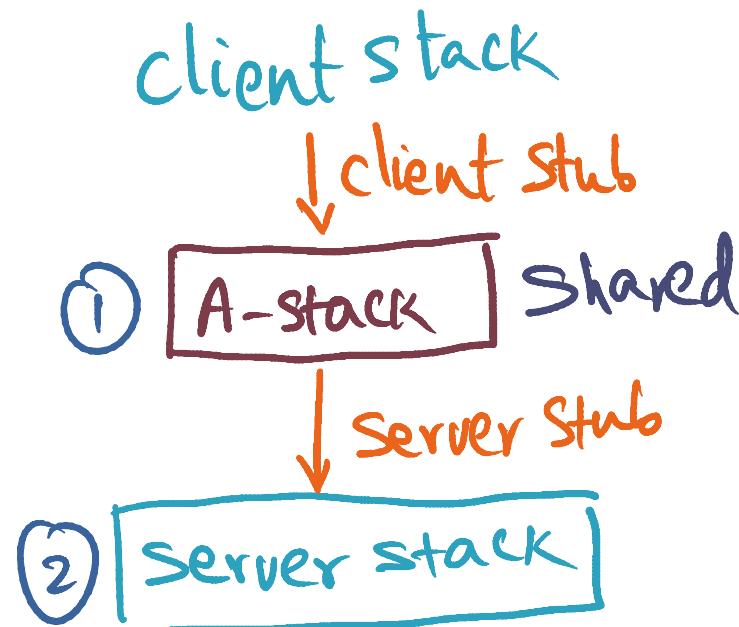


# Making RPC cheap (Actual Calls)

Original

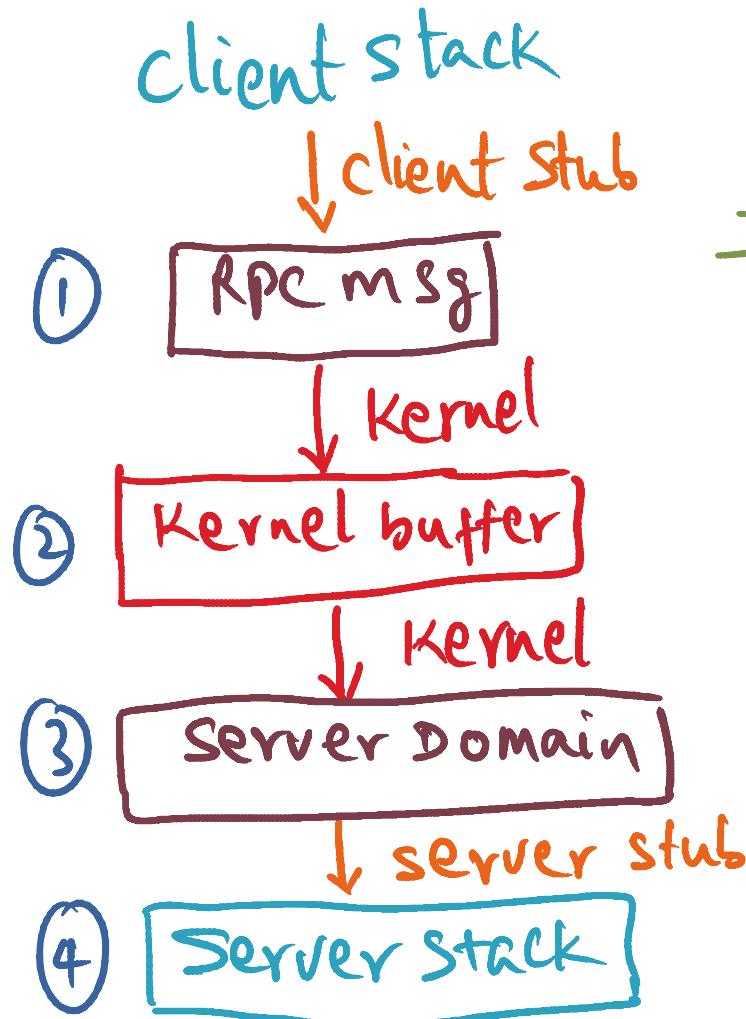


NOW

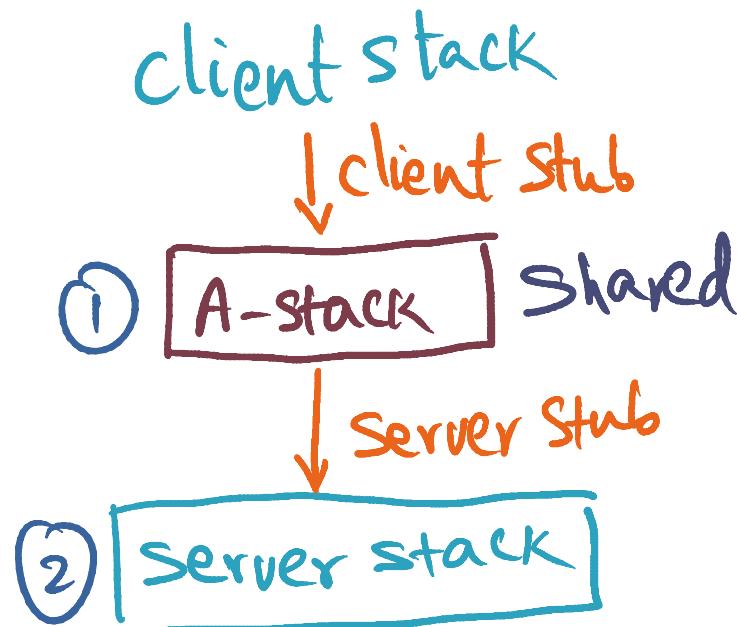


# Making RPC cheap (Actual Calls)

Original



NOW

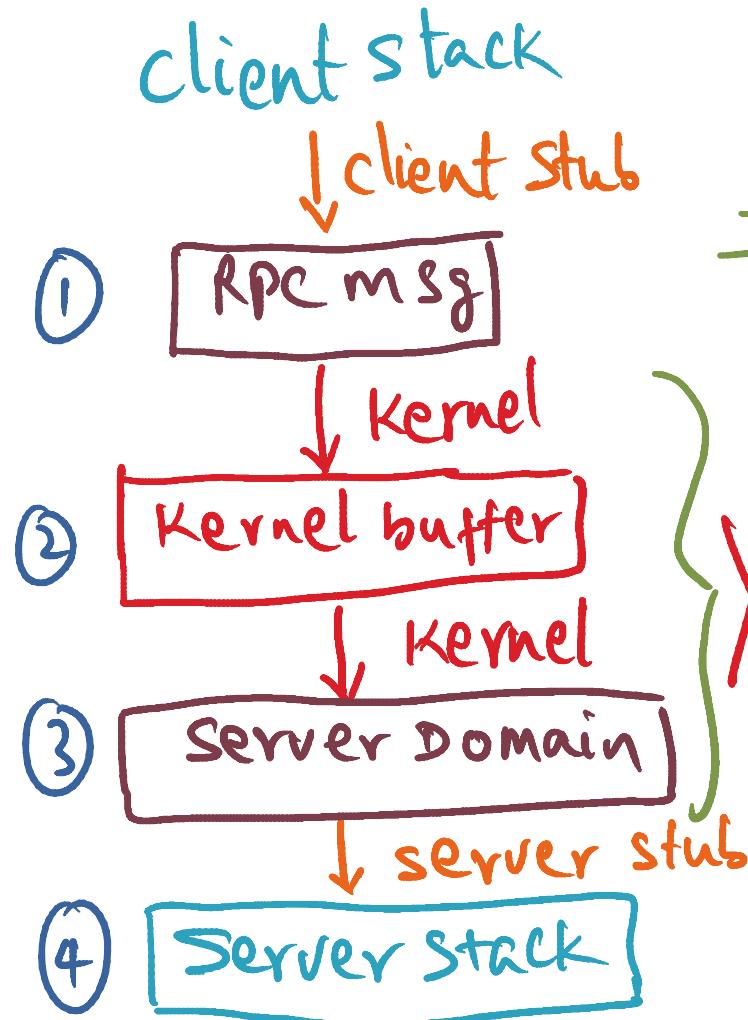


TWO Copies

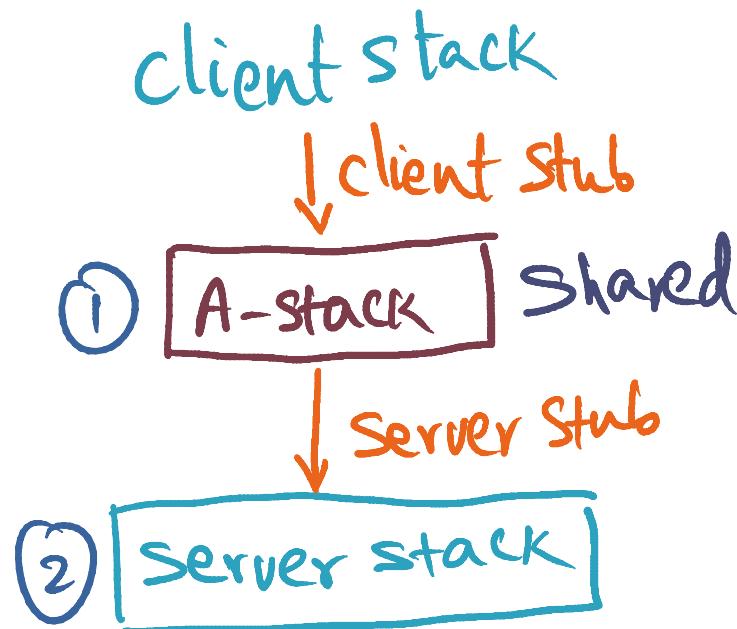
- marshal
- unmarshal

# Making RPC cheap (Actual Calls)

Original



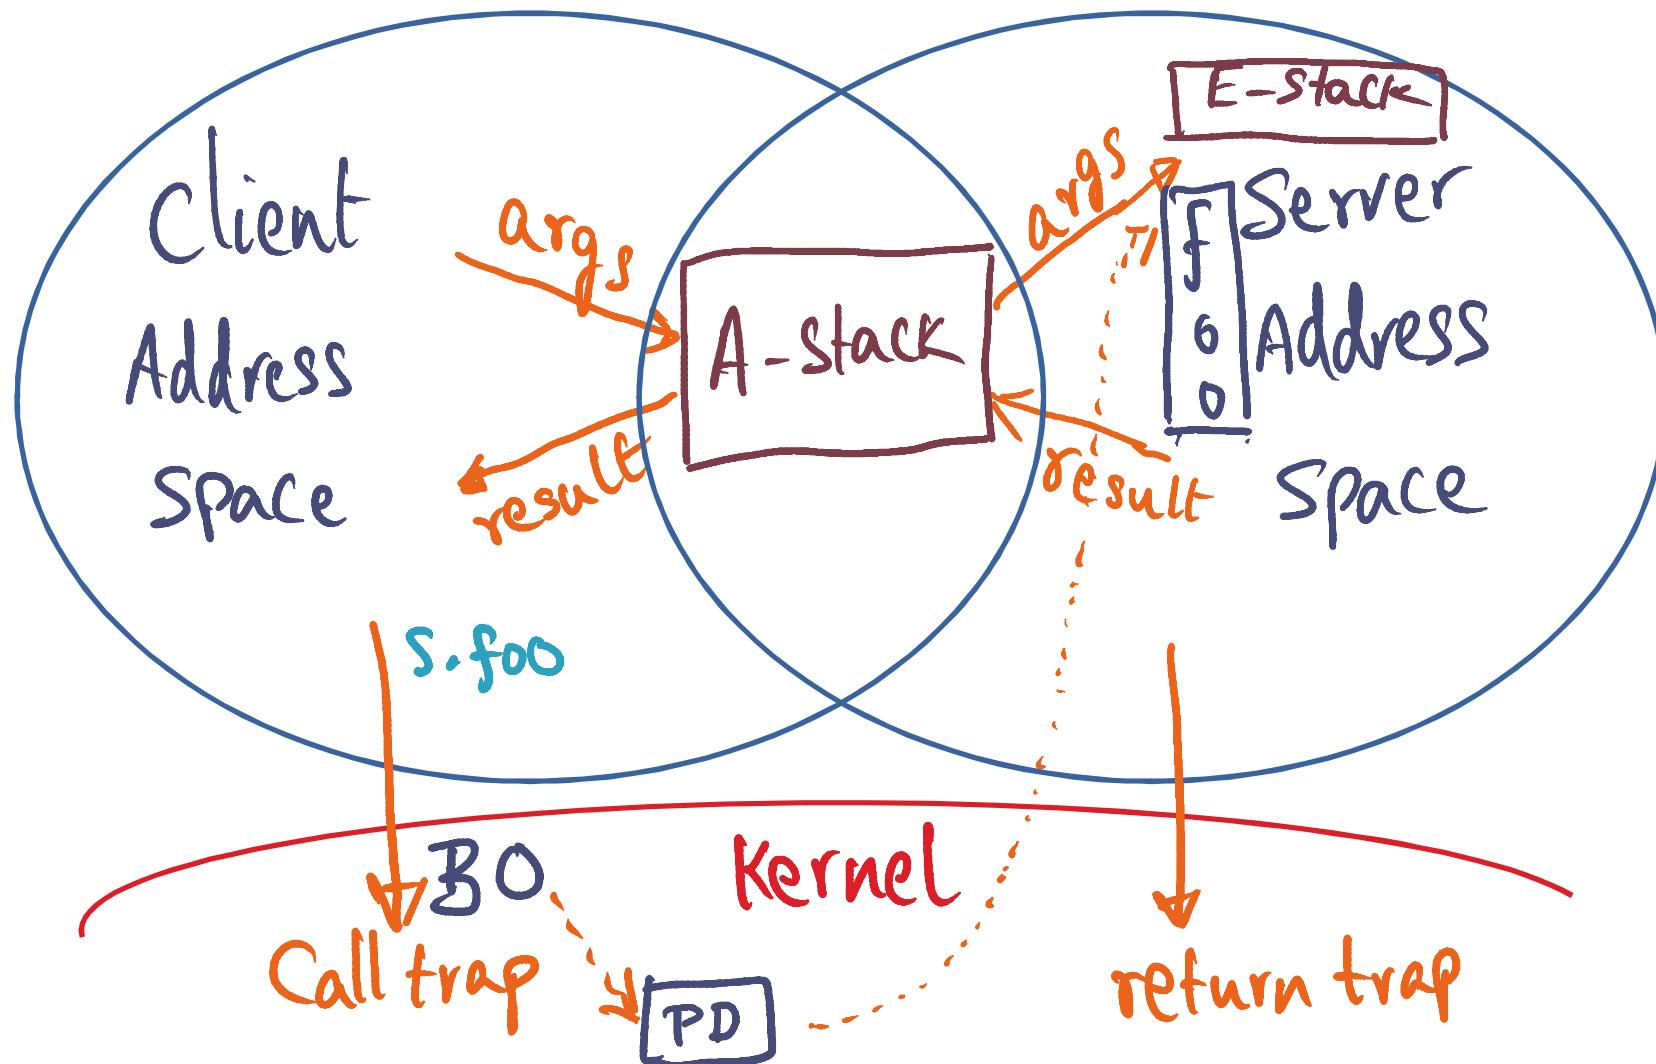
NOW



TWO Copies

- marshal
- unmarshal

# Making RPC cheap (Actual Calls)

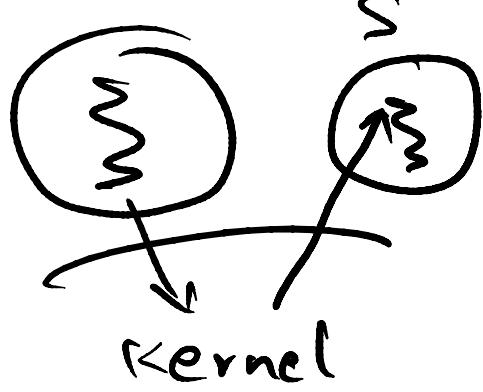


## Optimizations

Thread switching from client to server

\* Doctor client thread to run in  
Server domain

C S  $\Rightarrow$  Save some explicit cost  
in thread switching



Implicit cost ??

$\Rightarrow$  use Liedtke's domain  
packing tricks

# Summary

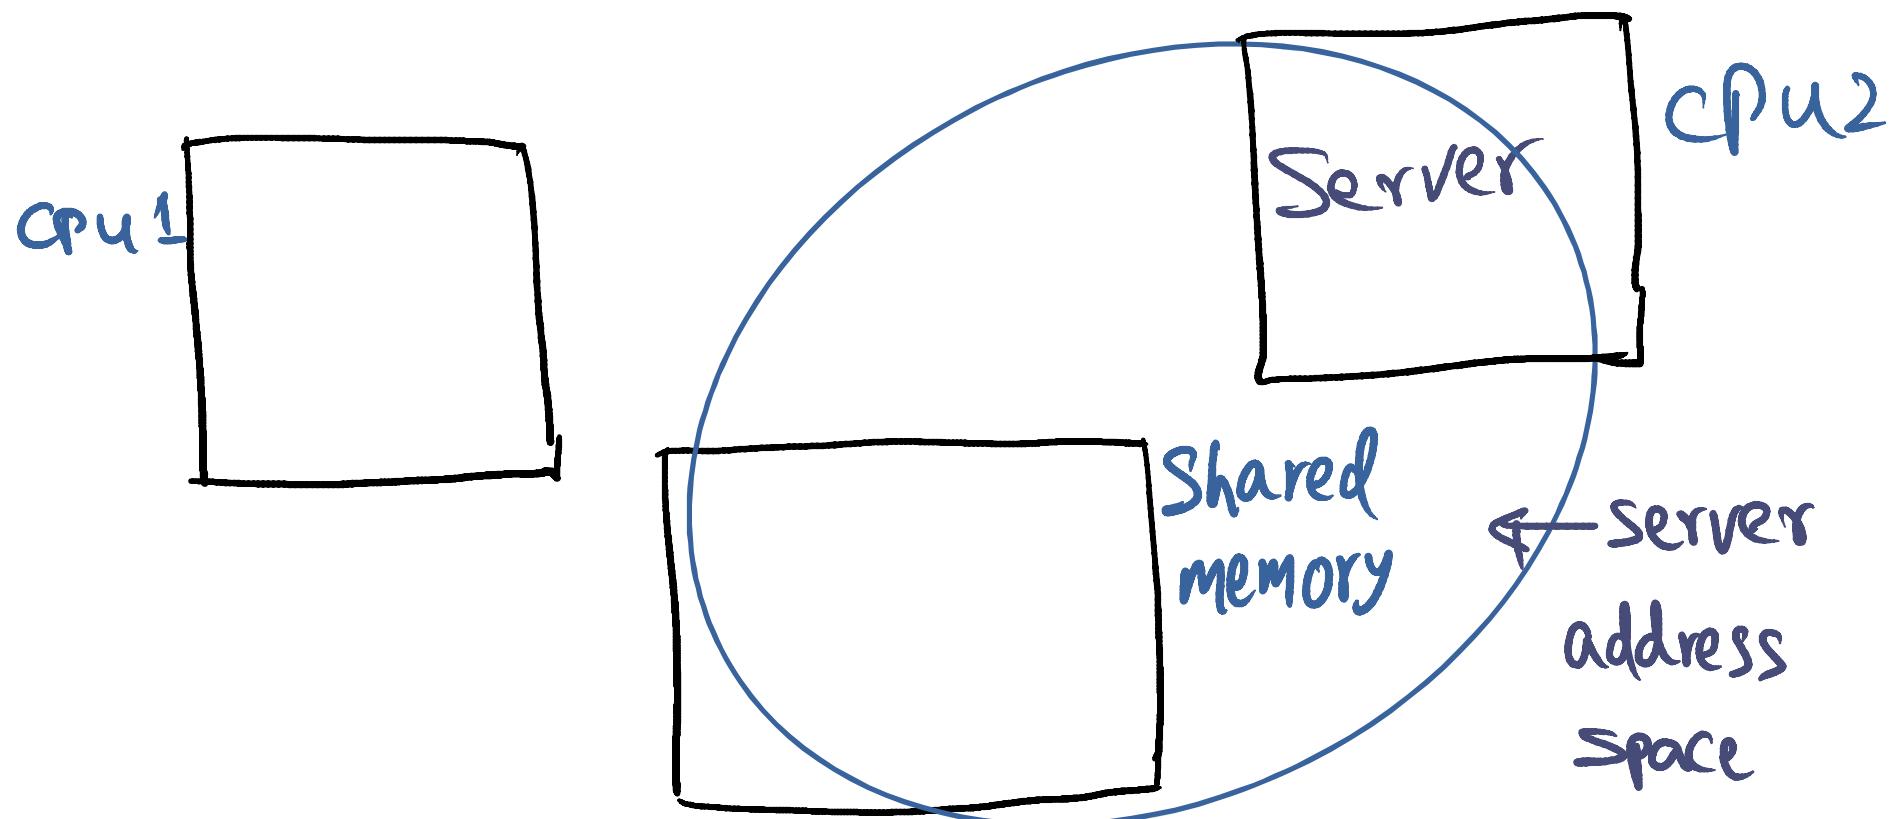
- One time binding cost
- Actual call eliminate kernel copies
- We are left with
  - Client/server traps (unavoidable)
  - Switching domains (optimizations to reduce explicit costs)
  - Loss of locality due to domain switching (small domains use Liedtke's tricks to reduce implicit costs)

## RPC on SMP

Exploit multiple processors

\* Pre-load server Domain

\* Keep caches warm

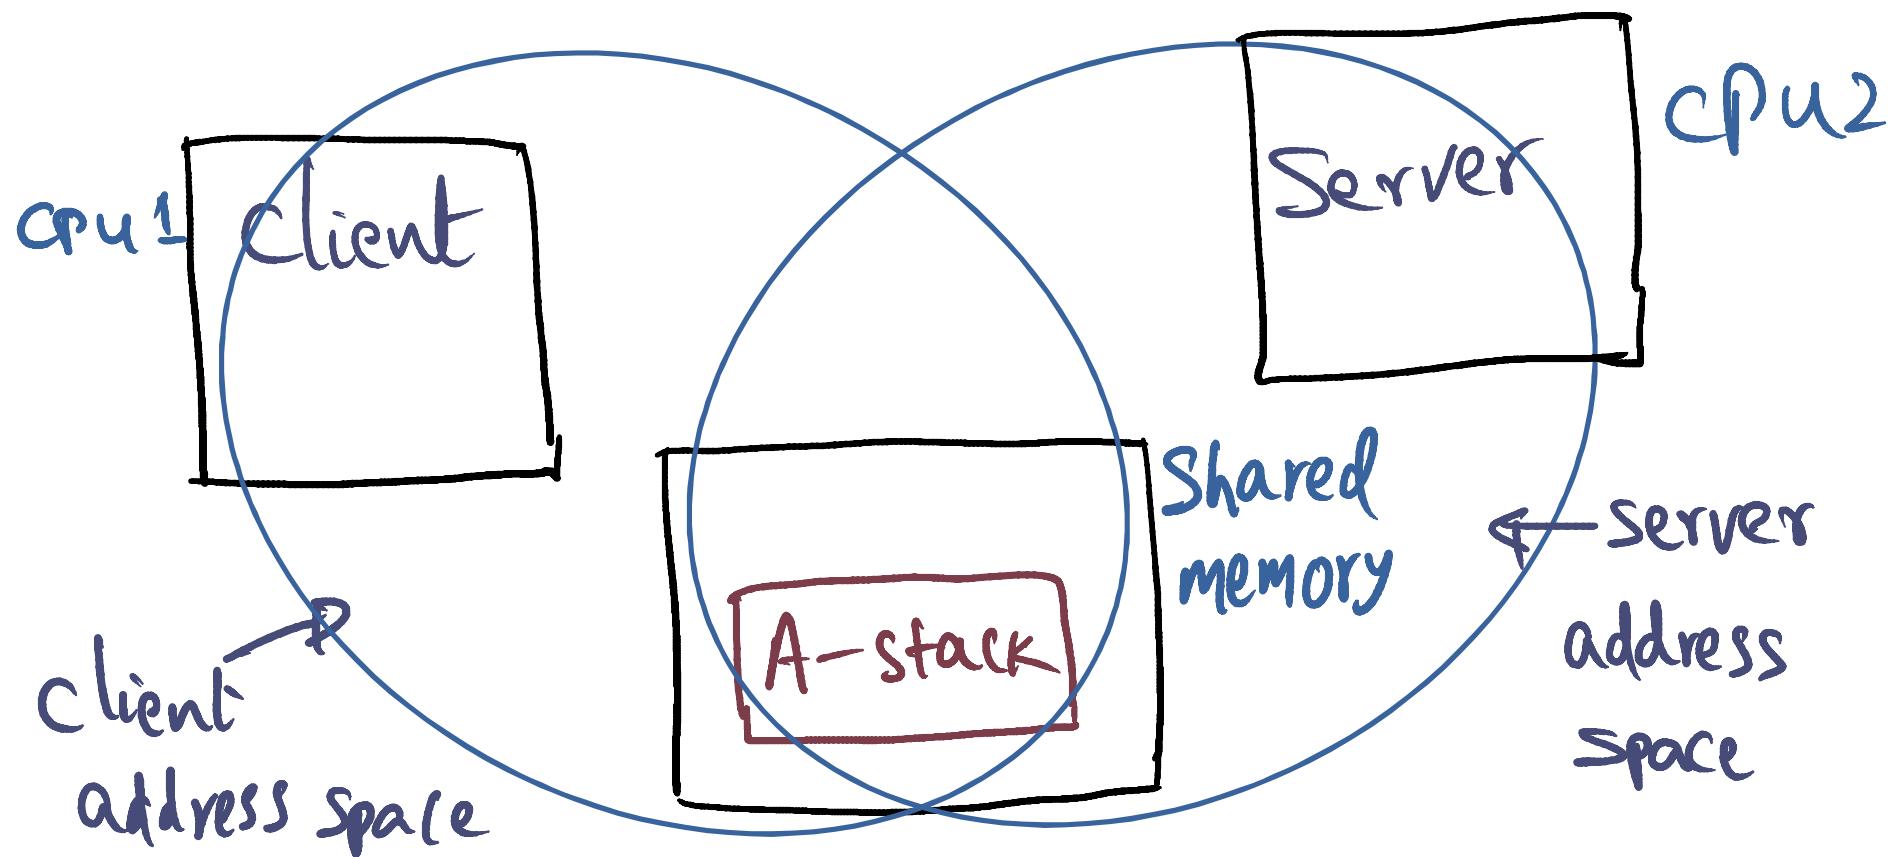


## RPC on SMP

Exploit multiple processors

\* Pre-load server Domain

\* Keep caches warm



# RPC ON SMP

Exploit multiple processors

\* Pre-load server Domain

\* Keep caches warm

