

Today

✓ * Sync (MCS)

✓ * Comm (LRPC)

⇒ * Sched in SM Systems

PI due

Squillante

Fedorova

Friday

Case Studies of parallel OS

* 

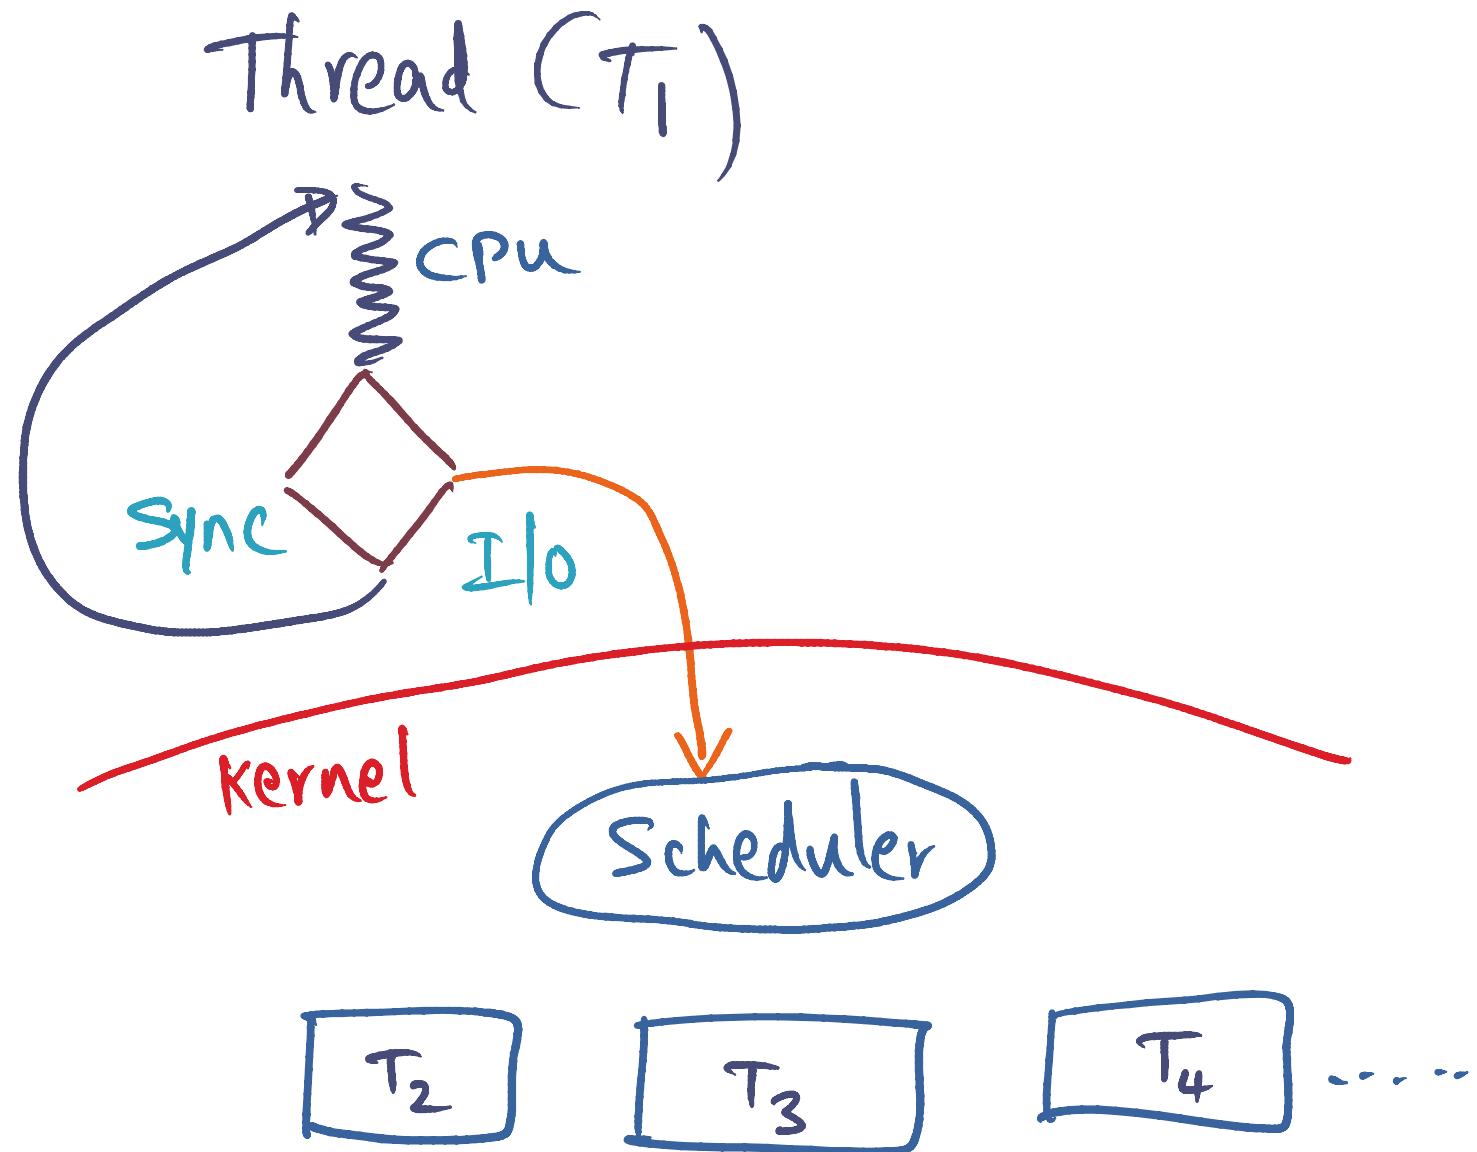
* Corey

* Cellular Disc.

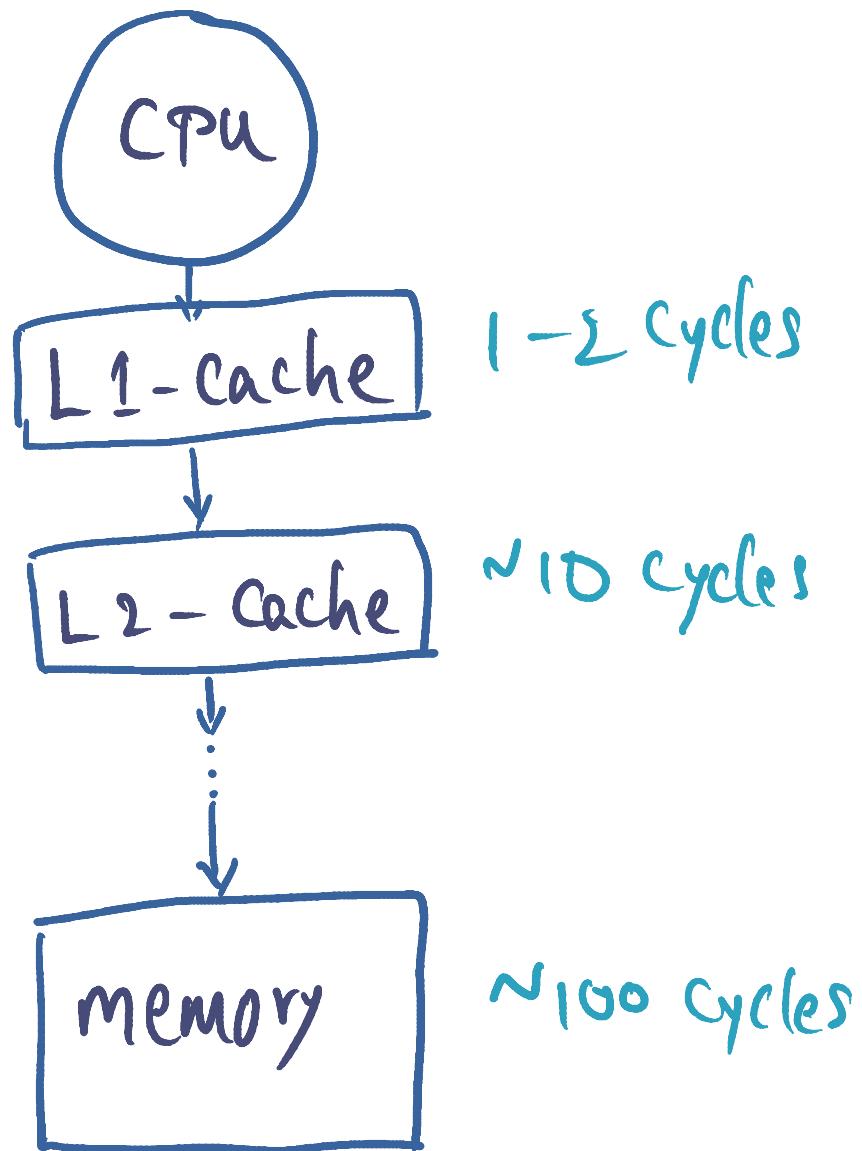
} Partial reading

Please
watch
video
before
class

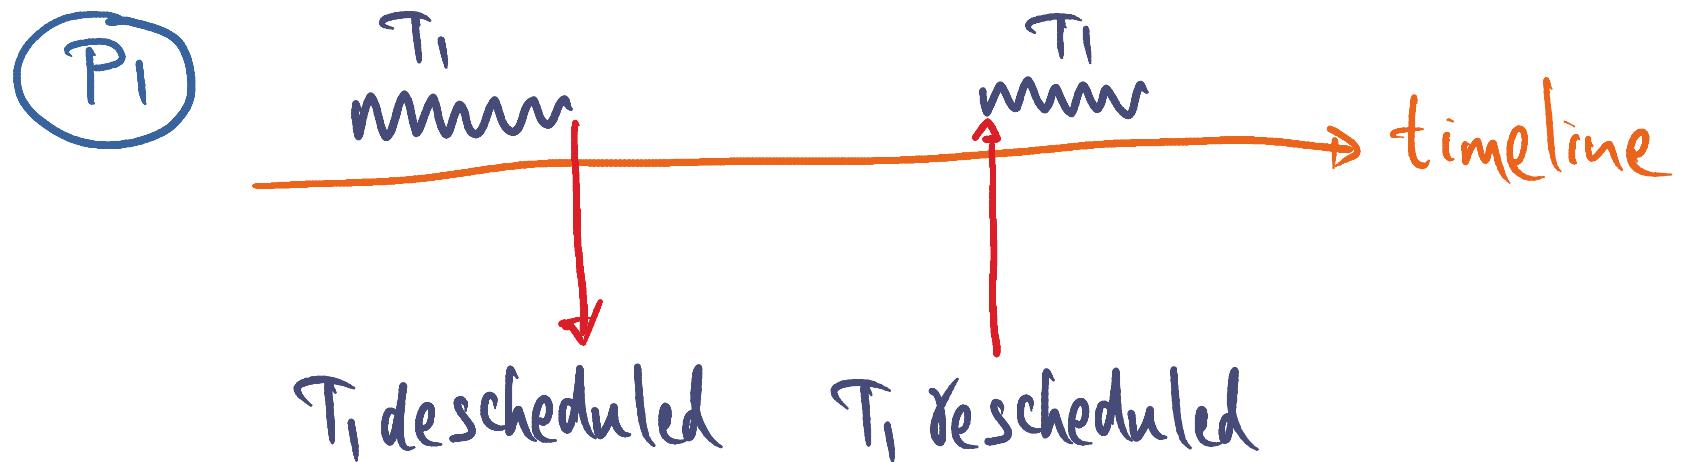
Scheduling - First Principles



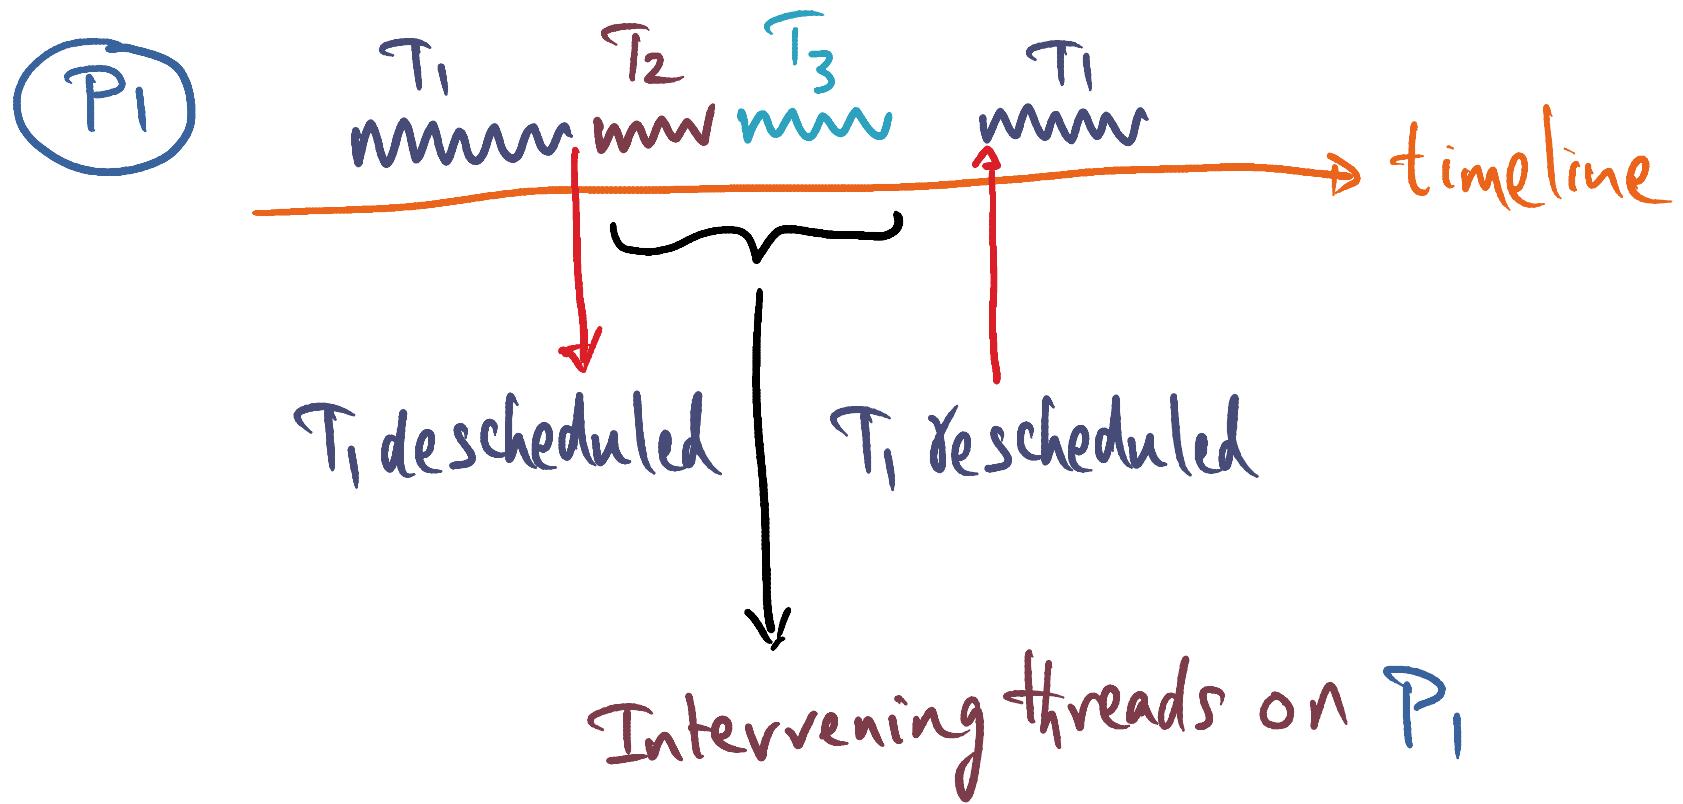
Memory hierarchy refresher



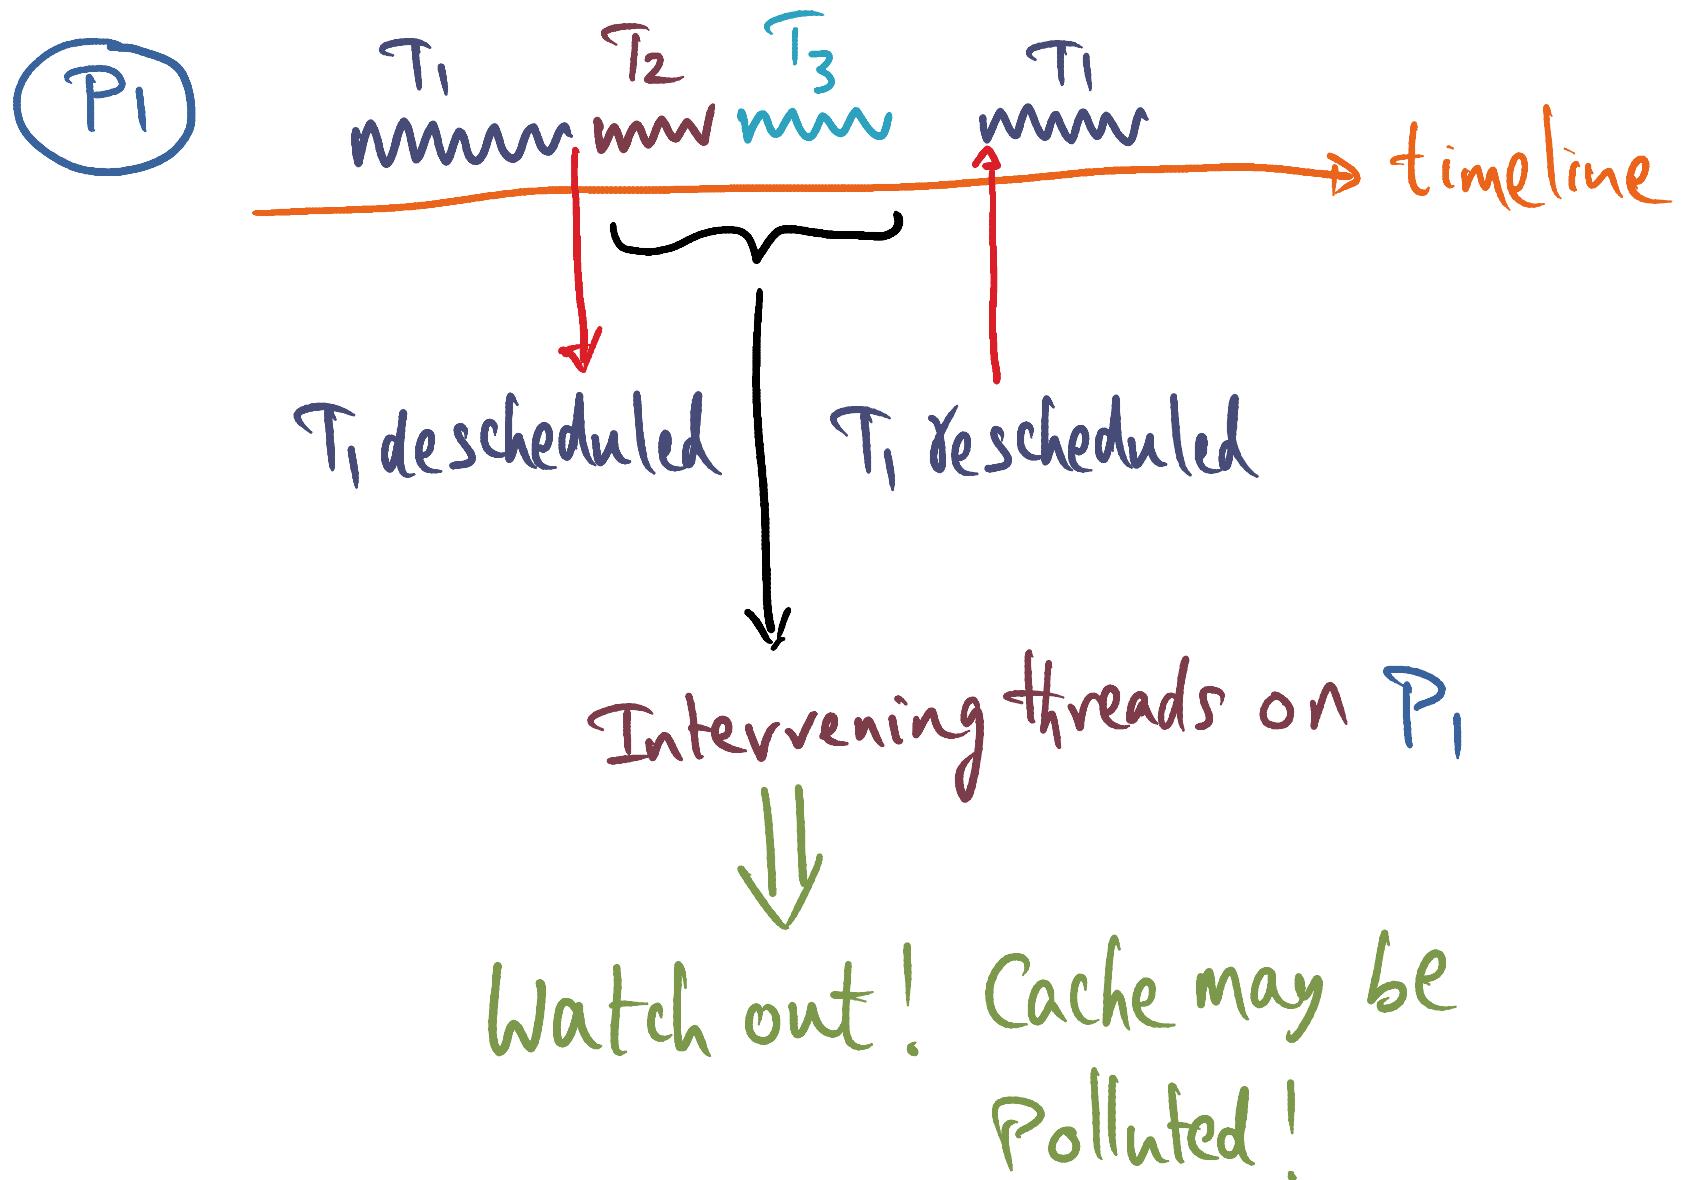
Cache affinity scheduling



Cache affinity Scheduling



Cache affinity Scheduling



Scheduling Policies

FCFS

Scheduling Policies

FCFS : Ignores affinity for fairness

Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

Scheduling Policies

FCFS : Ignores affinity for fairness

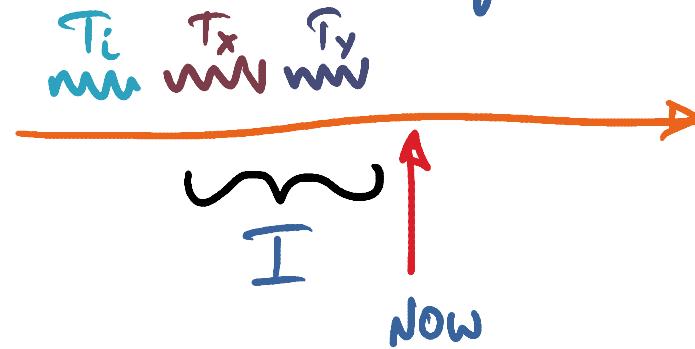
Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

Minimum Intervening: $T_i \rightarrow P_j^{I_{min}}$

Minimum Intervening: $T_i \rightarrow P_j^{I_{\min}}$

P_j 's timeline



Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

Minimum Intervening: $T_i \rightarrow P_j^{I_{min}}$

Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{Last}

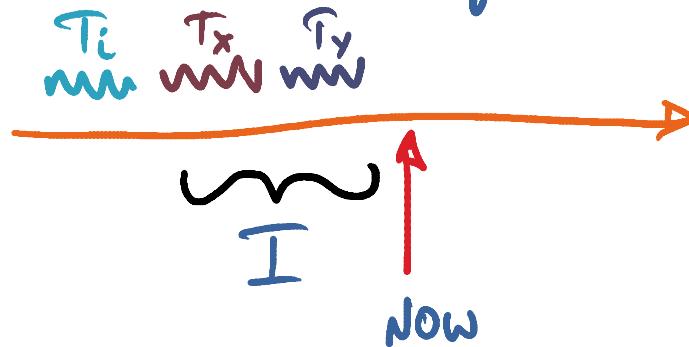
Minimum Intervening: $T_i \rightarrow P_j^{T_{\min}}$

Minimum Intervening plus queue:

$T_i \rightarrow P_j^{(T+Q)_{\min}}$

Minimum Intervening: $T_i \rightarrow P_j^{I_{\min}}$

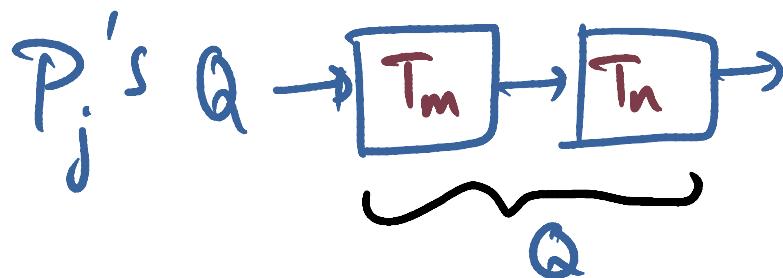
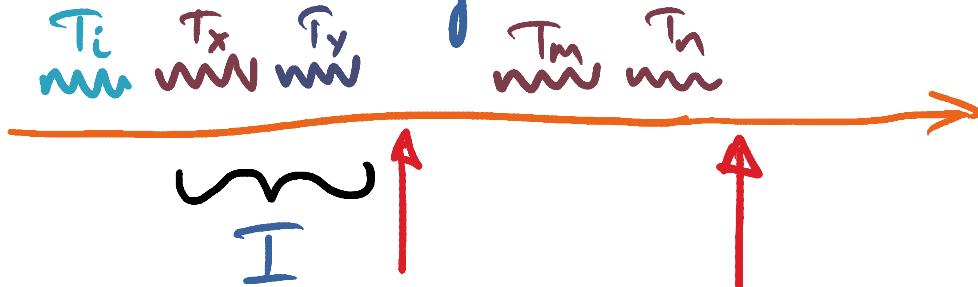
P_j 's timeline



Minimum Intervening plus queue:

$T_i \rightarrow P_j^{(I+Q)_{\min}}$

P_j 's timeline



Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

Minimum Intervening: $T_i \rightarrow P_j^{I_{min}}$

Minimum Intervening plus queue:

$T_i \rightarrow P_j^{(I+Q)_{min}}$

Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

} focus on
Cache
affinity of
 T_i

Minimum Intervening: $T_i \rightarrow P_j^{I_{min}}$

Minimum Intervening plus queue:

$T_i \rightarrow P_j^{(I+Q)_{min}}$

Scheduling Policies

FCFS : Ignores affinity for fairness

Fixed processor: T_i always on P_{fixed}

Last processor: T_i on P_{last}

Minimum Intervening: $T_i \rightarrow P_j^{T_{min}}$

Minimum Intervening plus queue:

$$T_i \rightarrow P_j^{(T+Q)_{min}}$$

} focus on
Cache
affinity of
 T_i

} focus on
Cache
Pollution
when T_i
gets to run

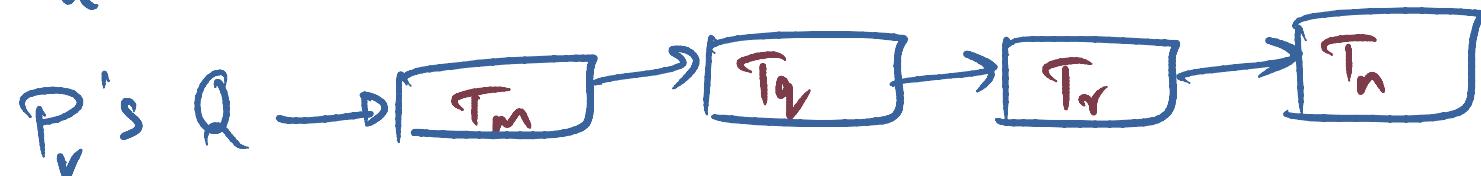
Scheduling Policies

FCFS : Ignores affinity for fairness

thread centric { Fixed processor: T_i always on P_{fixed} } focus on Cache affinity of T_i
Last processor: T_i on P_{last}

Processor centric { Minimum Intervening: $T_i \rightarrow P_j^{T_{min}}$ } focus on Cache
Minimum Intervening plus queue: $T_i \rightarrow P_j^{(T+Q)_{min}}$ } Pollution when T_i gets to run

Question



T_y 's affinity: $P_u^T = 2; P_v^T = 1 \}$ Num Intervening tasks

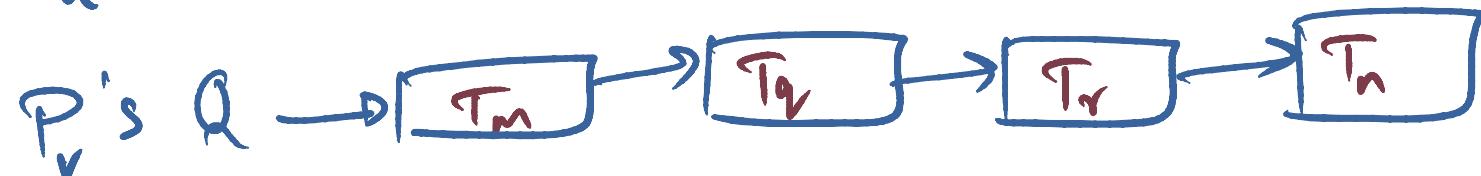
Scheduling Policy: Min Inter. plus Q

Which Q for T_y ?

P_u 's Q

P_v 's Q

Question / Solution



T_y 's affinity: $P_u^T = 2$; $P_v^T = 1$ } Num Intervening tasks

Scheduling Policy: Min Inter. plus Q

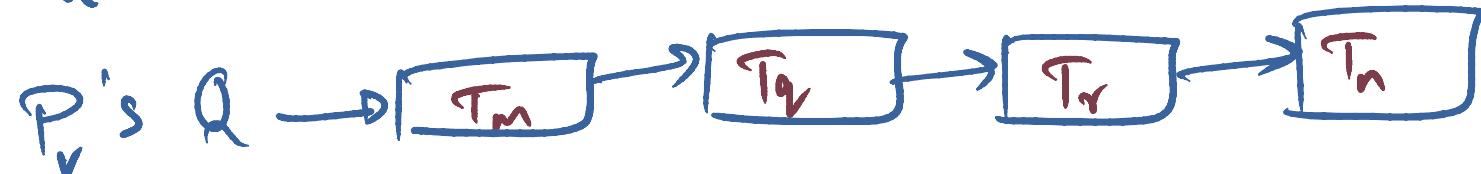
Which Q for T_y ?

P_u 's Q

P_v 's Q

$$\text{Min}_{P_u}^{T_y} = 2 + 1 = 3$$

Question / Solution



T_y 's affinity: $P_u^T = 2; P_v^T = 1 \} \text{Num Intervening tasks}$

Scheduling Policy: Min Inter. plus Q

Which Q for T_y ?

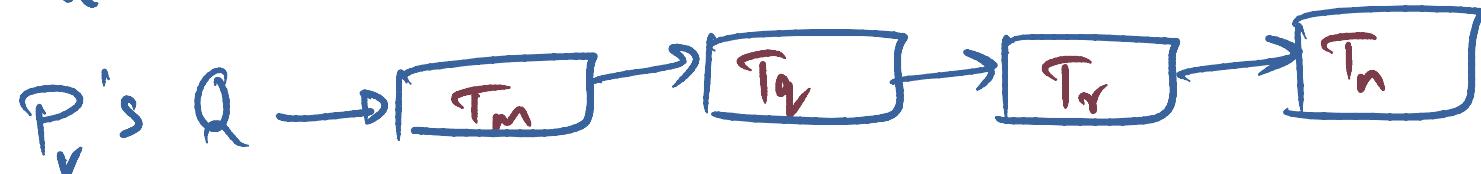
P_u 's Q

P_v 's Q

$$\text{Min}_{P_u}^{T_y} = 2 + 1 = 3$$

$$\text{Min}_{P_v}^{T_y} = 1 + 4 = 5$$

Question / Solution



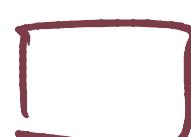
T_y 's affinity: $P_u^T = 2$; $P_v^T = 1$ } Num Intervening tasks

Scheduling Policy: Min Inter. plus Q

Which Q for T_y ?



P_u 's Q



P_v 's Q

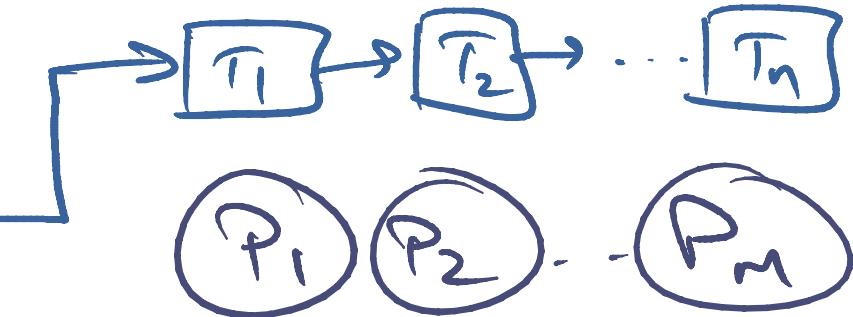
$$\text{Min}_{P_u}^{T_y} = 2 + 1 = 3$$

$$\text{Min}_{P_v}^{T_y} = 1 + 4 = 5$$

Implementation Issues

Queue Based

- Global Que

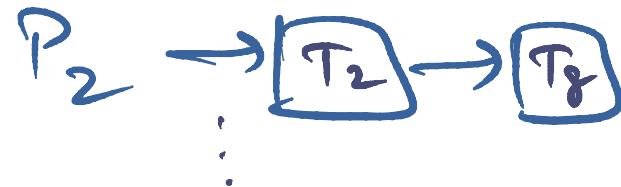
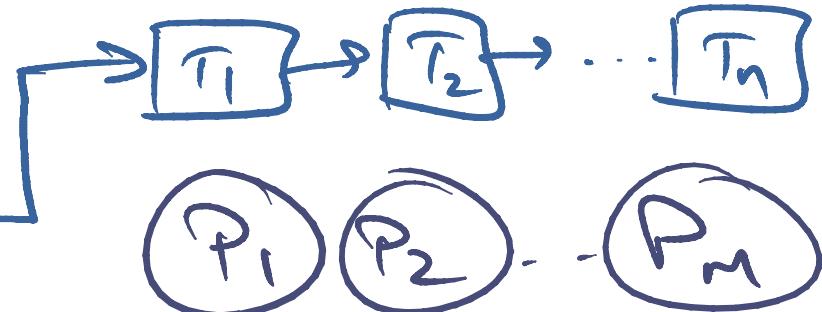


Implementation Issues

Queue Based

- Global Que

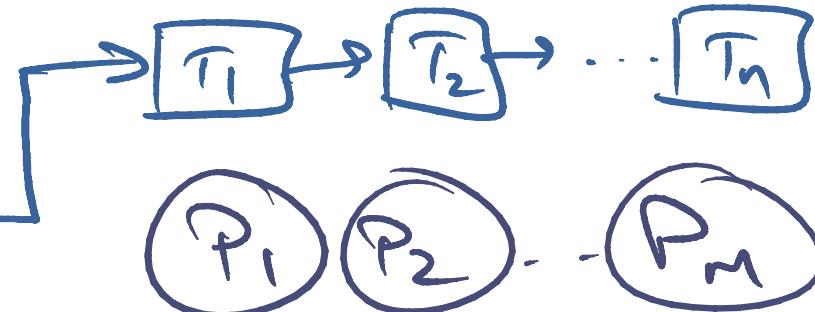
- Affinity-based local queues



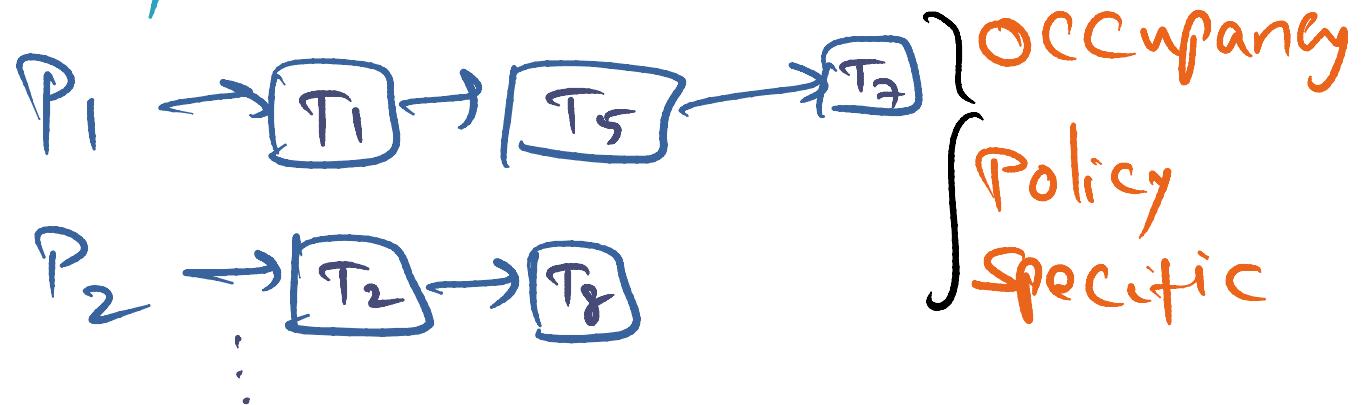
Implementation Issues

Queue Based

- Global Que



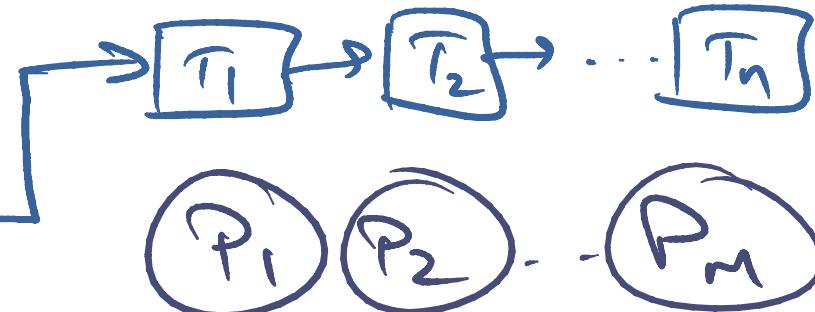
- Affinity-based local queues



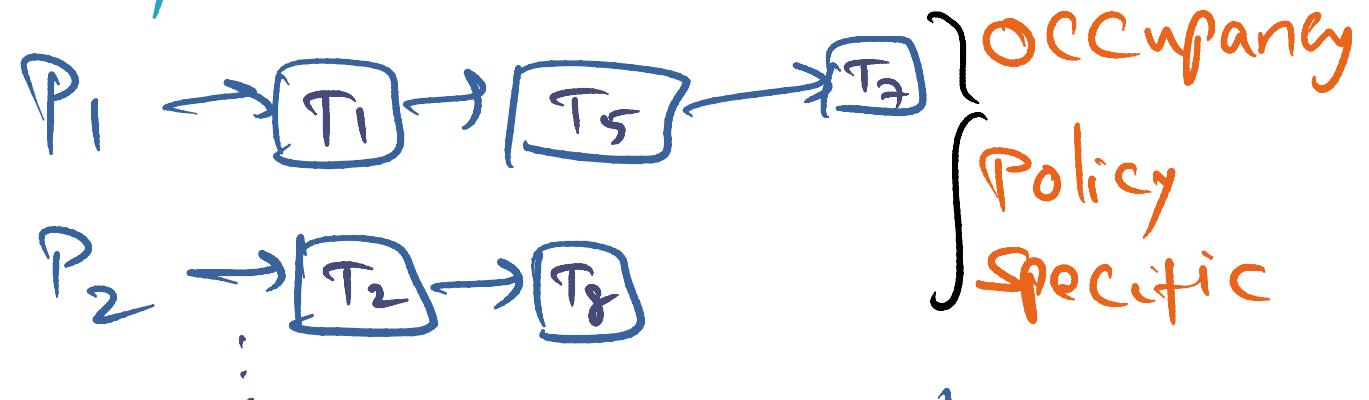
Implementation Issues

Queue Based

- Global Que



- Affinity-based local queues



$$T_i \text{ 's Priority} = BP_i + Age_i + \text{Affinity}_i$$

- Determines position in the queue

Performance

Figures of merit

- Throughput
- Response time
- Variance

Performance

Figures of merit

- Throughput \rightarrow System Centric
- Response time
- Variance

Performance

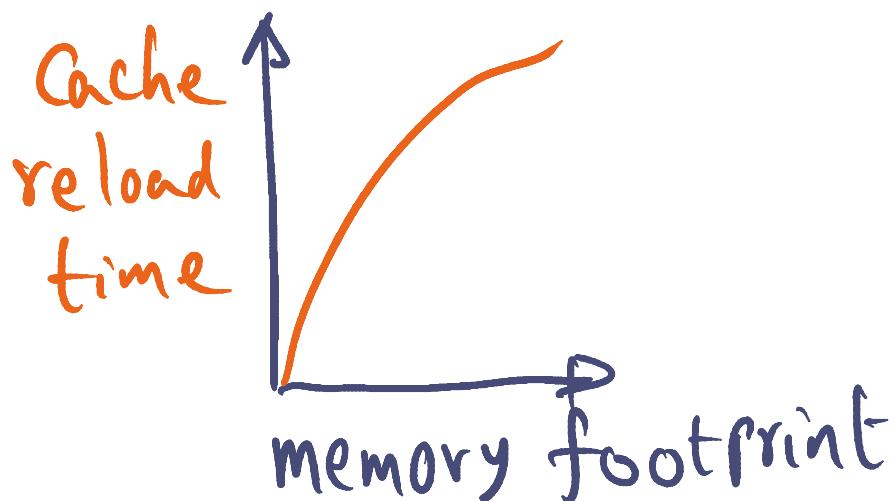
Figures of merit

- Throughput \rightarrow System Centric
 - Response time
 - Variance
- } \rightarrow User Centric

Performance

Figures of merit

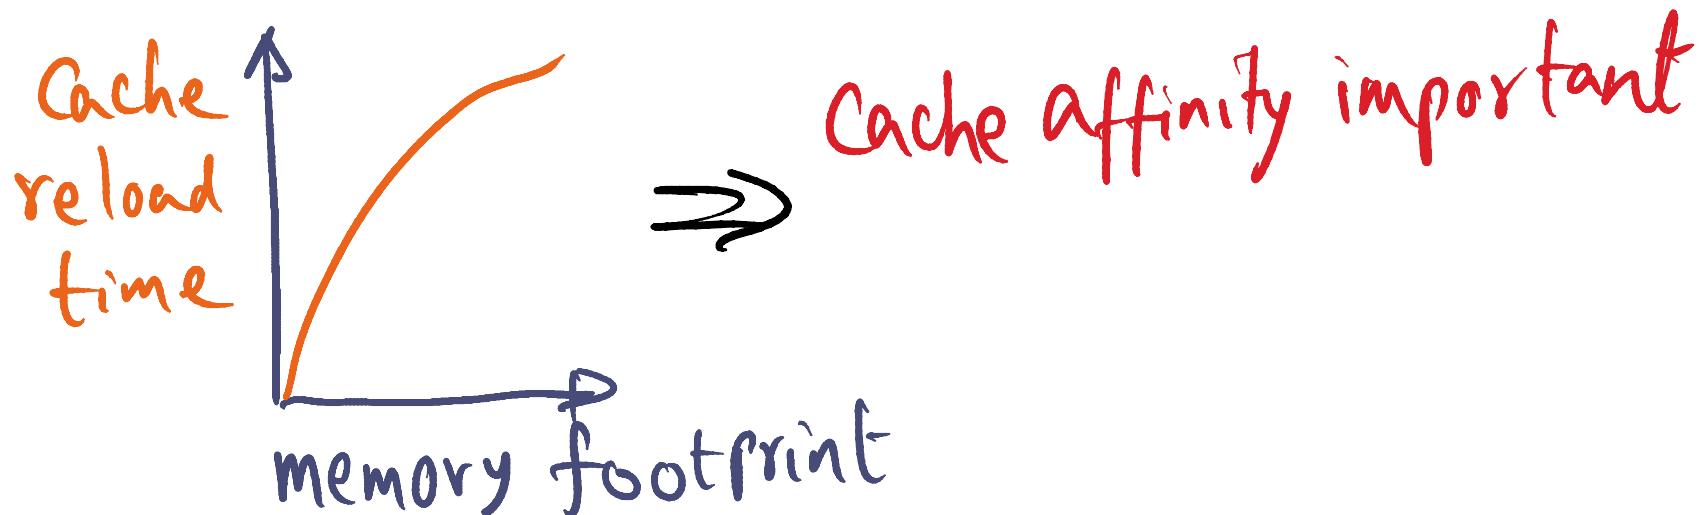
- Throughput \rightarrow System Centric
 - Response time
 - Variance
- } \rightarrow User Centric



Performance

Figures of merit

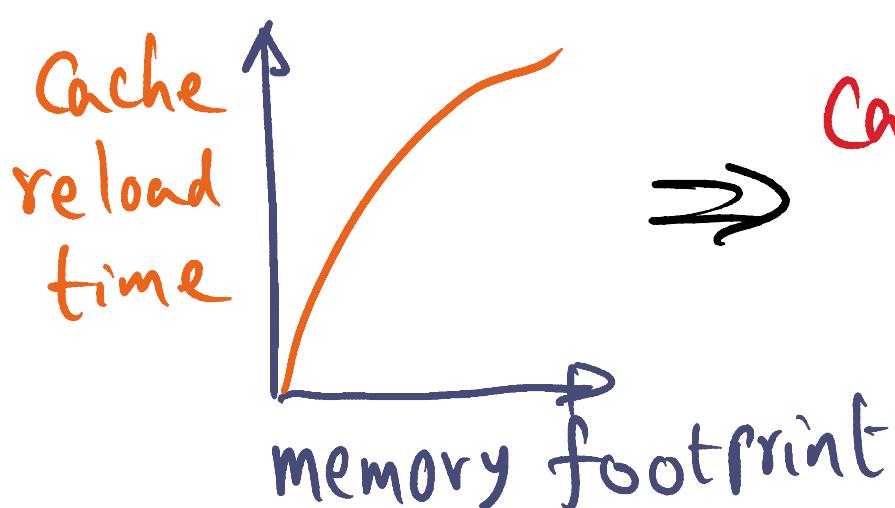
- Throughput \rightarrow System Centric
 - Response time
 - Variance
- } \rightarrow User Centric



Performance

Figures of merit

- Throughput \rightarrow System Centric
 - Response time
 - Variance
- } \rightarrow User Centric



\Rightarrow Cache affinity important
heavy load
Fixed Processor

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.
- Moral of the story
 - you have to design your scheduler commensurate with the workload you expect

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.
- Moral of the story
 - you have to design your scheduler commensurate with the workload you expect
- Inserting idle loop in processor scheduling
 - Good idea?

Summary

- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.
- Moral of the story
 - you have to design your scheduler commensurate with the workload you expect
- Inserting idle loop in processor scheduling
 - Good idea?
 - Waiting for “high affinity” thread

Summary

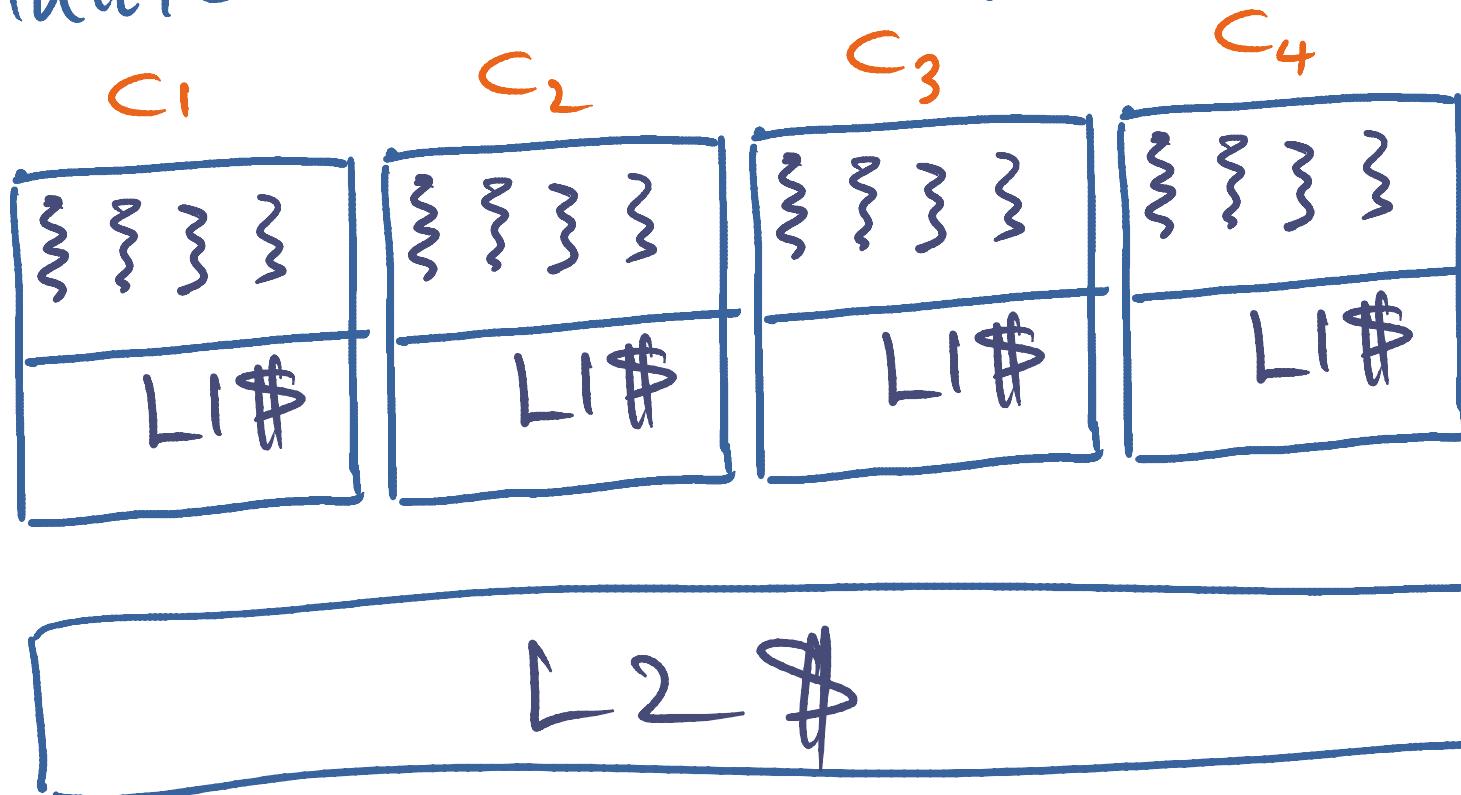
- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.
- Moral of the story
 - you have to design your scheduler commensurate with the workload you expect
- Inserting idle loop in processor scheduling
 - Good idea?
 - Waiting for “high affinity” thread
 - Another instance where “procrastination” may help improve system performance!

Summary

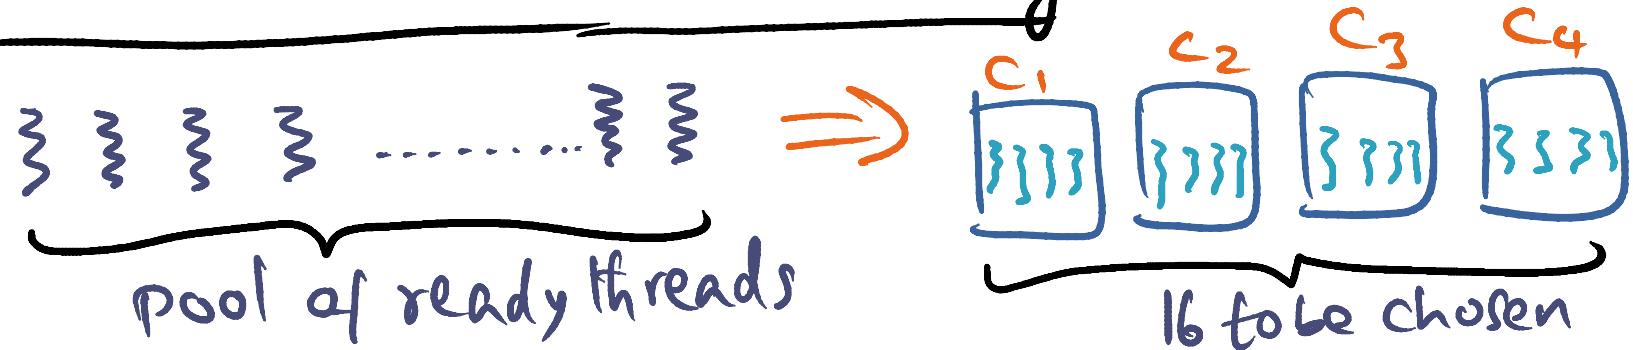
- FCFS is fair but can lead to high variance in response time (small job stuck behind a large job)
- MI and variants good at light load
- Heavy load cache is polluted anyway => fixed processor may in fact lead to better perf.
- Moral of the story
 - you have to design your scheduler commensurate with the workload you expect
- Inserting idle loop in processor scheduling
 - Good idea?
 - Waiting for “high affinity” thread
 - Another instance where “procrastination” may help improve system performance!
 - Have we seen this before?

Cache Affinity and Multicore

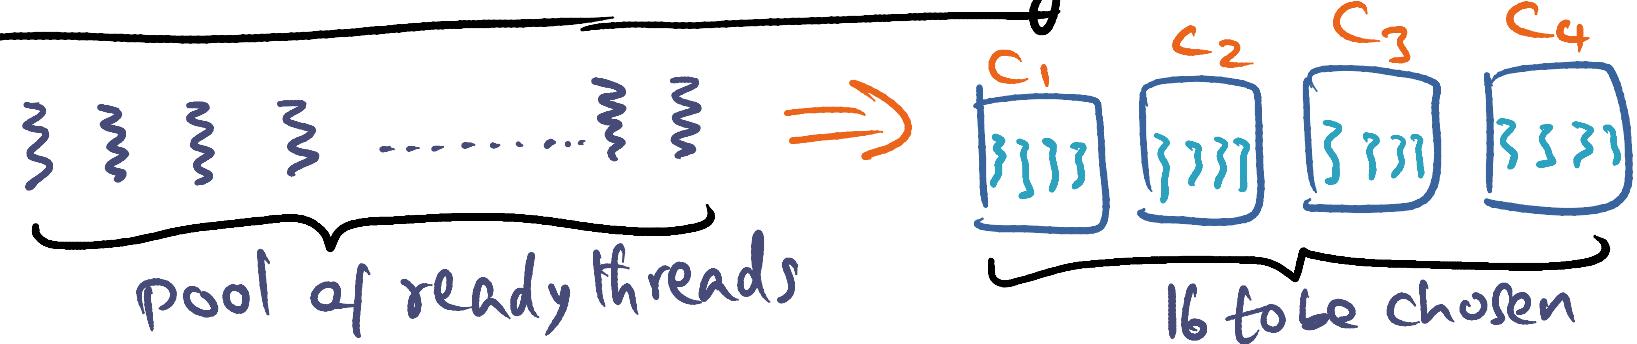
Multicore Multithreaded Processors



Cache Aware Scheduling

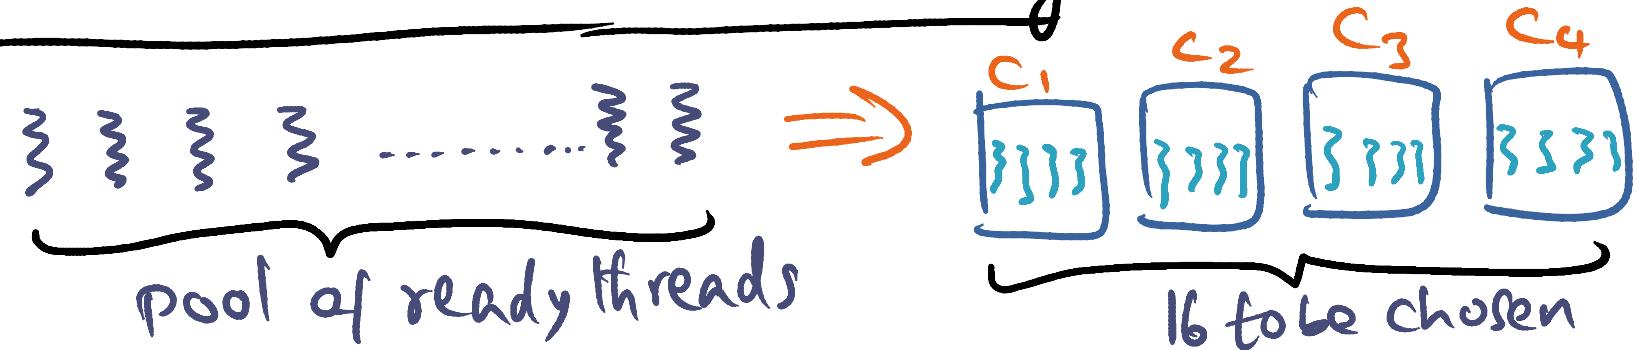


Cache Aware Scheduling



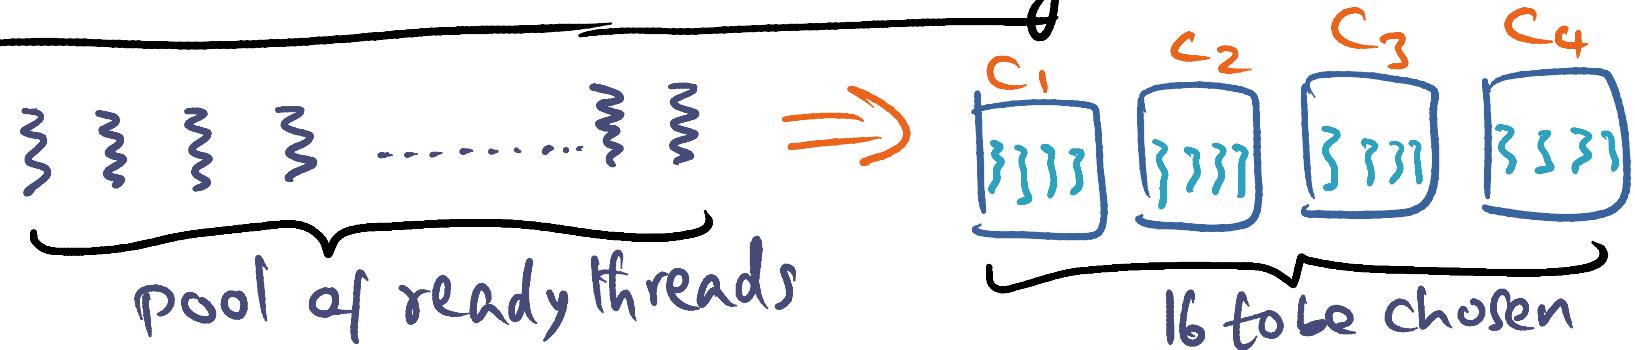
— Cache frugal threads — C_{ft}

Cache Aware Scheduling



- Cache frugal threads - C_{ft}
- Cache hungry threads - C_{ht}

Cache Aware Scheduling



- Cache frugal threads - C_{ft}

- Cache hungry threads - C_{ht}

$$\sum_1^n C_{ft} + \sum_1^m C_{ht} < \text{Size}(L^2 \#)$$

$$m+n = 16$$

Friday

Case Studies of parallel OS

* #



- * Corey
- * Cellular Disco } Partial reading

Please

Watch

video

before
class

Next week

Distributed systems

* watch Lamport clock video

(First video in Lesson 5 on Nodacity)